

Ripasso concetti base: PIPELINE BASE

Andrea Bartolini <a.bartolini@unibo.it>

(Architettura dei) Calcolatori Elettronici, 2023/2024

Instruction Execution

- PC \rightarrow instruction memory, fetch instruction
- Register numbers \rightarrow register file, read registers
- Depending on instruction class
 - Use ALU to calculate
 - Arithmetic result
 - Memory address for load/store
 - Branch comparison
 - Access data memory for load/store
 - PC \leftarrow target address or PC + 4

RISC-V Pipeline

- Five stages, one step per stage
 1. IF: Instruction fetch from memory
 2. ID: Instruction decode & register read
 3. EX: Execute operation or calculate address
 4. MEM: Access memory operand
 5. WB: Write result back to register

RISC Pipeline – Stadi

IF: Instruction Fetch

- $IR \leftarrow Mem[PC]$ Invia il program counter (PC) alla memoria e riceve (fetch) l'istruzione corrente dalla memoria.
- $NPC \leftarrow PC + 4$ Incrementa il PC per indirizzare l'istruzione successiva (istruzioni da 4 bytes e memoria indirizzabile al singolo byte.)

ID: Instruction Decode

- $Opcode, funct, Imm, Rd, Rs1, Rs2, .. \leftarrow IR$ Decodifica l'istruzione, siccome pochi formati e registri sempre nelle stesse posizioni. Decodifica semplice => In decode si può fare altro.
- $A, B \leftarrow RegisterFile[Rs1, Rs2]$ Leggi il contenuto del RegisterFile a indirizzo Rs1, Rs2
- $Imm_se \leftarrow SignExt.[Imm]$ Estendi il bit di segno del campo immediato: è sempre nella stessa posizione per istruz. di load e alu.
- $Cond = ? (A == B)$ Test di uguaglianza tra i dati contenuti nei due registri sorgente per possibili istruzioni branch.
- $NPC = NPC + Imm_se$ Calcola il possibile *branch target address* sommando al PC il campo immediato esteso di segno

RISC Pipeline – Stadi

EX: Execution

- $ALU_output = A + Imm_se$
- $ALU_output = alu_op(A, B)$
- $ALU_output = alu_op(A + Imm_se)$
- $Cond = ? (A == B)$
- $NPC = NPC + Imm_se$

Load: calcola l'*indirizzo effettivo* da cui leggere il dato in memoria sommando il campo immediato al dato letto dal Register File (RF).

ALU reg-reg: Esegue l'operazione ALU sui due valori letti dal Register File

ALU reg-imm: Esegue l'operazione ALU tra immediato e dato letto dal Register File

Branch: Test di ugualinza tra i dati contenuti nei due registri sorgente per possibili istruzioni branch.

Branch: Calcola il possibile *branch target address* sommando al PC il campo immediato esteso di segno

MEM: Memory access / branch completion:

- $PC \leftarrow NPC$
- $LMD \leftarrow Mem[ALU_output]$
- $Mem[ALU_output] \leftarrow B$

Branch: Aggiorna il program counter per puntare alla prossima istruzione.

Load: Legge il dato contenuto in memoria all'indirizzo effettivo contenuto in ALU_output e salva il valore nel *load memory data (LMD)*.

Store: Scrive in memoria all'*indirizzo effettivo* contenuto in ALU_output il dato letto dal RegisterFile.

RISC Pipeline – Stadi

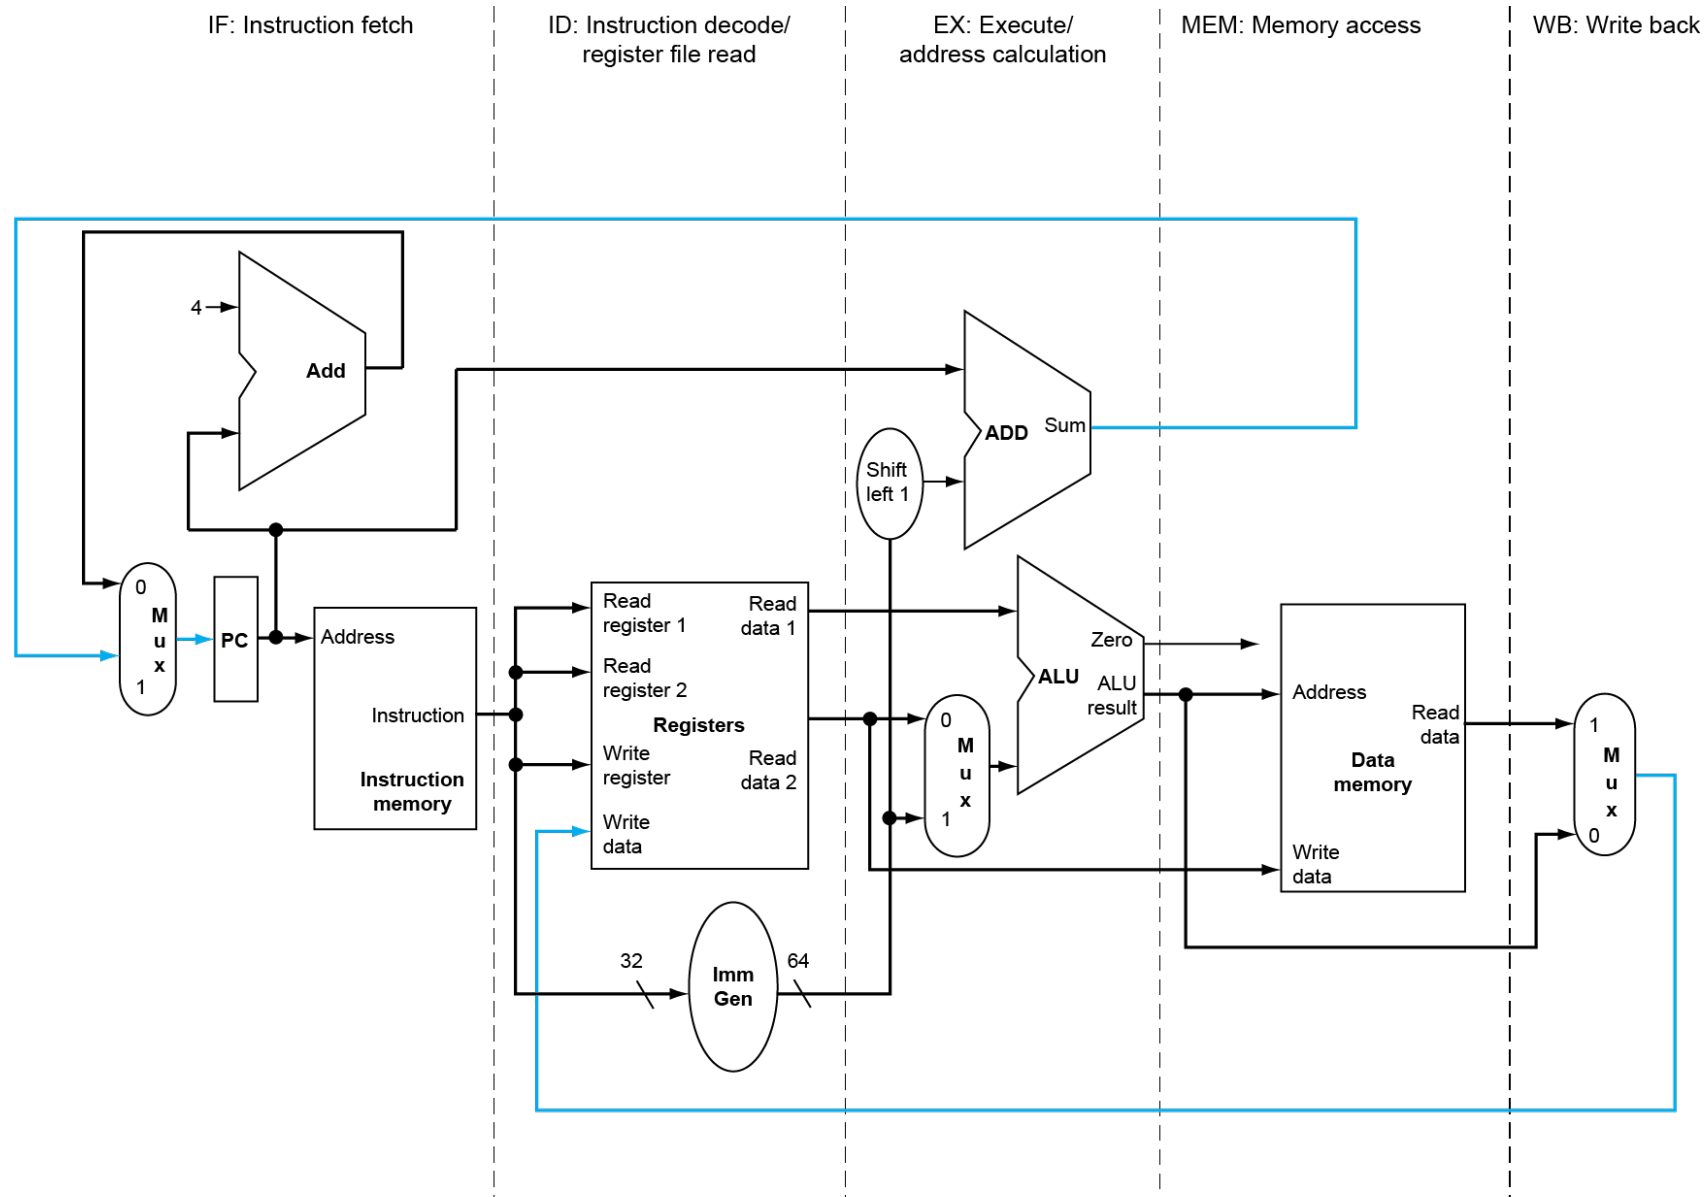
WB: Write Back

- RegisterFile[Rd] \leftarrow ALU_output
- RegisterFile[Rd] \leftarrow LMD

ALU reg-reg: memorizza il risultato dell'operazione nel Register File

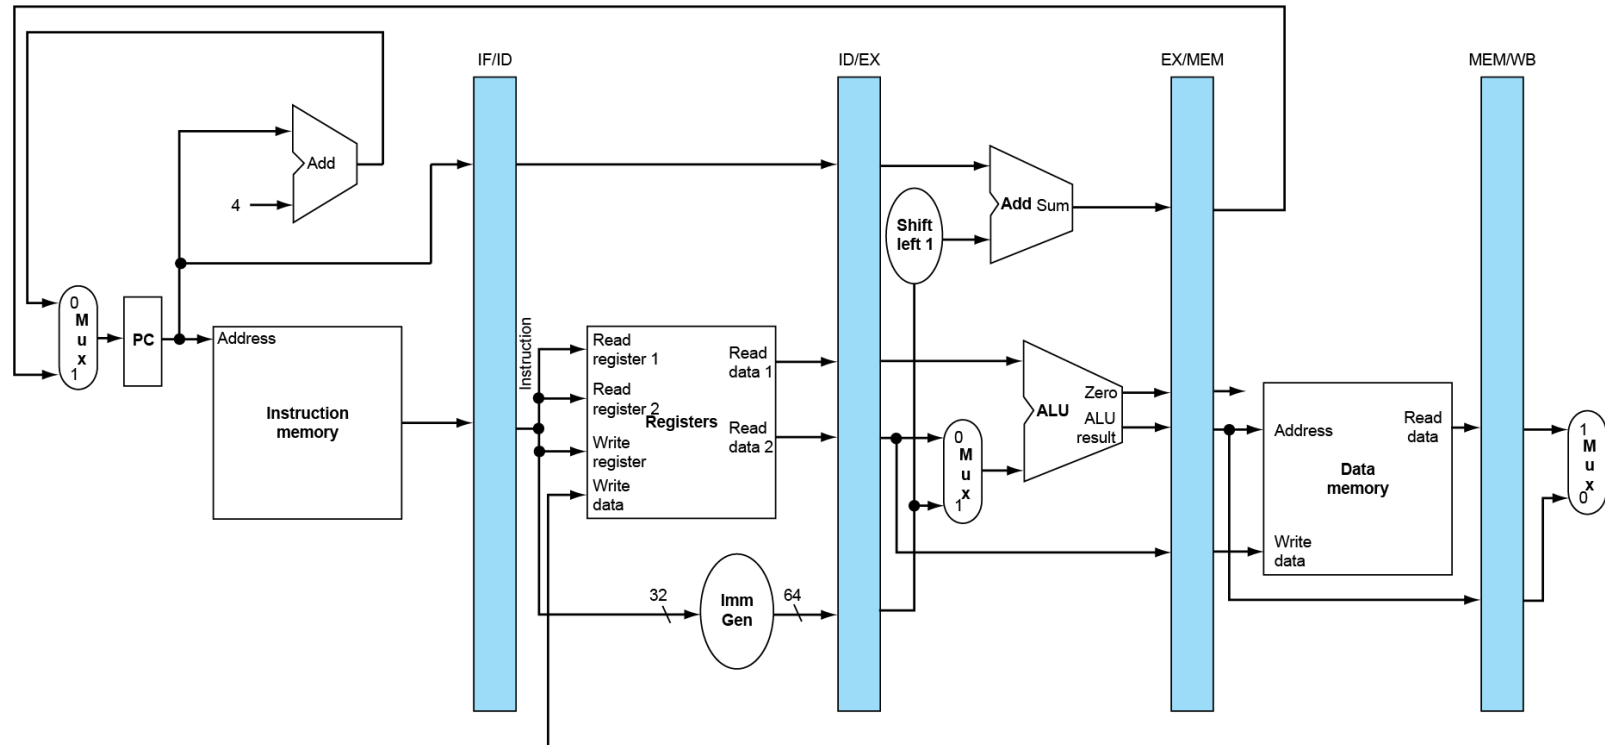
Load: Salva il dato letto nella memoria nel Register File

RISC-V Pipelined Datapath



Pipeline registers

- Need registers between stages
 - To hold information produced in previous cycle



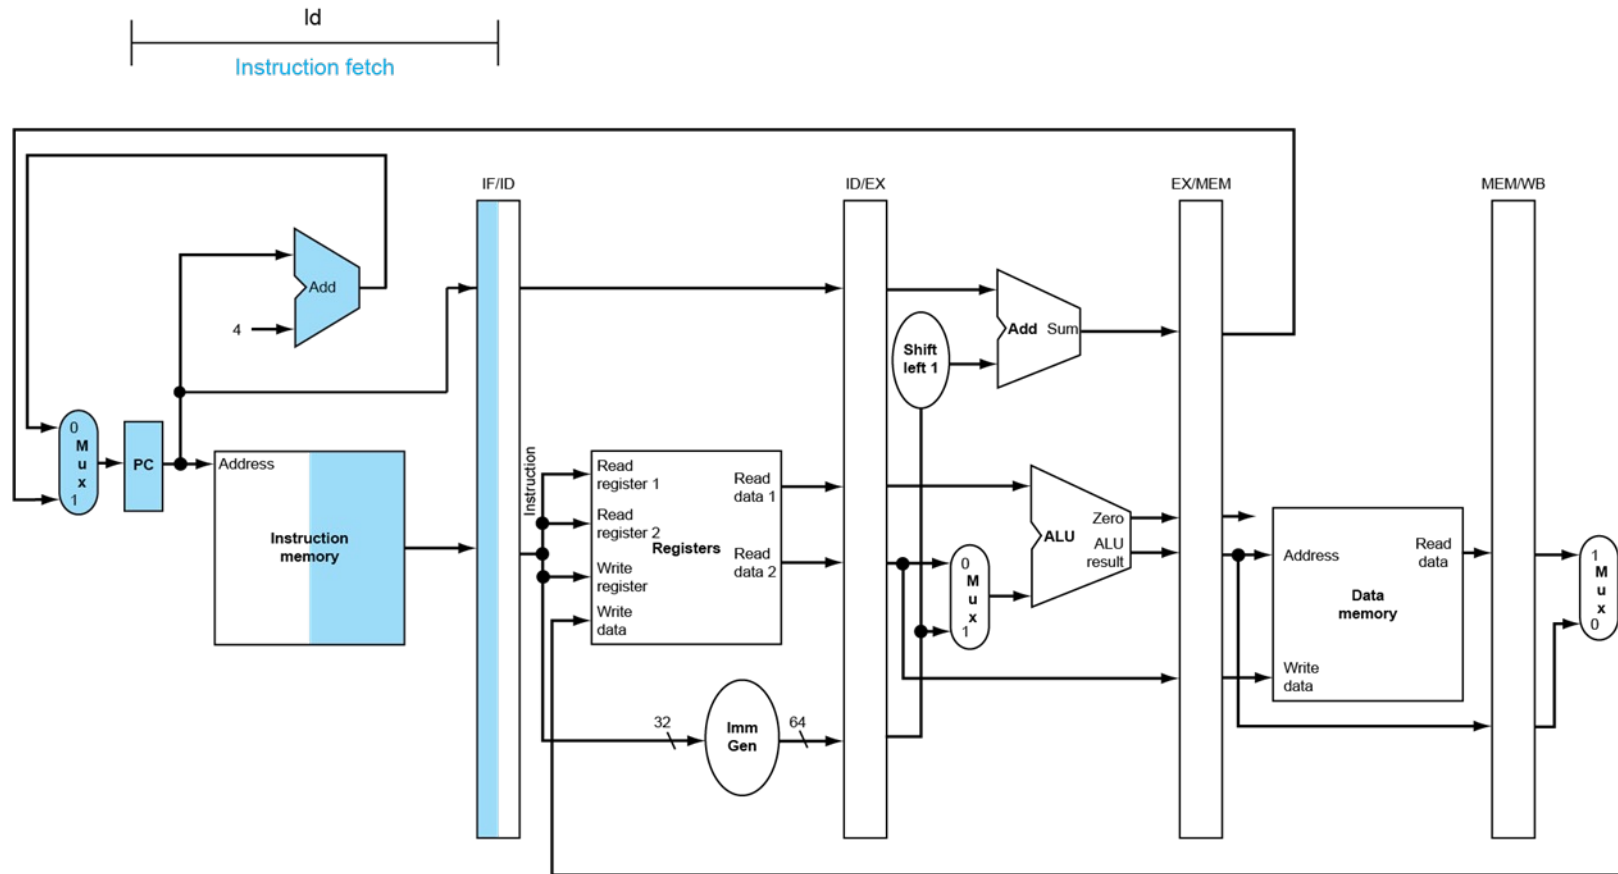
Pipeline Speedup

- If all stages are balanced
 - i.e., all take the same time
 - $\text{Time between instructions}_{\text{pipelined}} = \frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$
- If not balanced, speedup is less
- Speedup due to increased throughput
 - Latency (time for each instruction) does not decrease

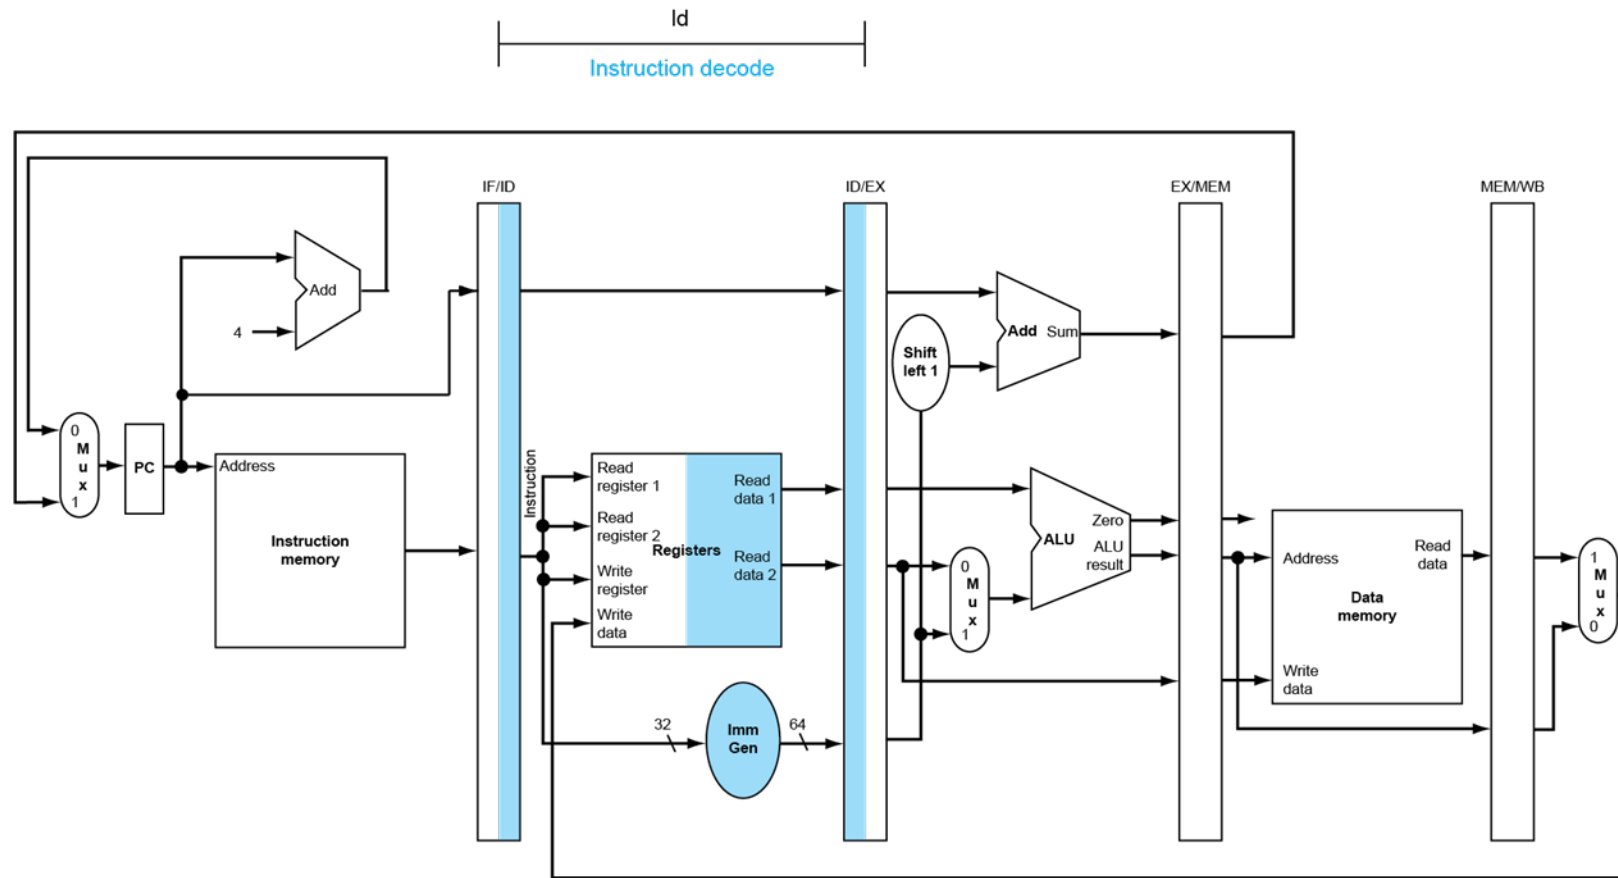
Pipelining and ISA Design

- RISC-V ISA designed for pipelining
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage

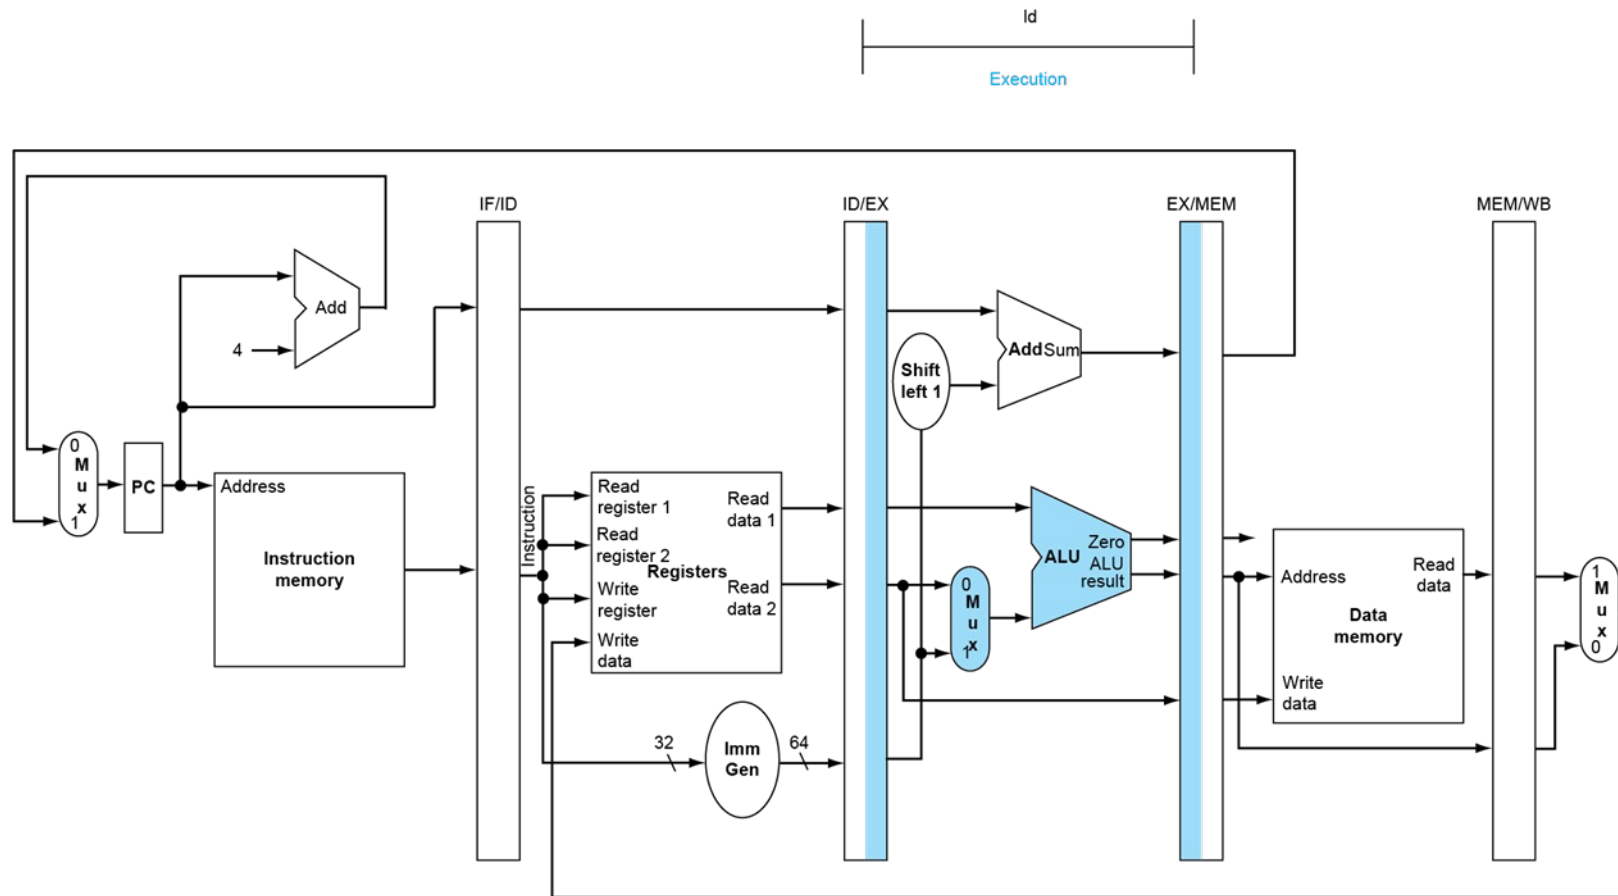
IF for Load, Store, ...



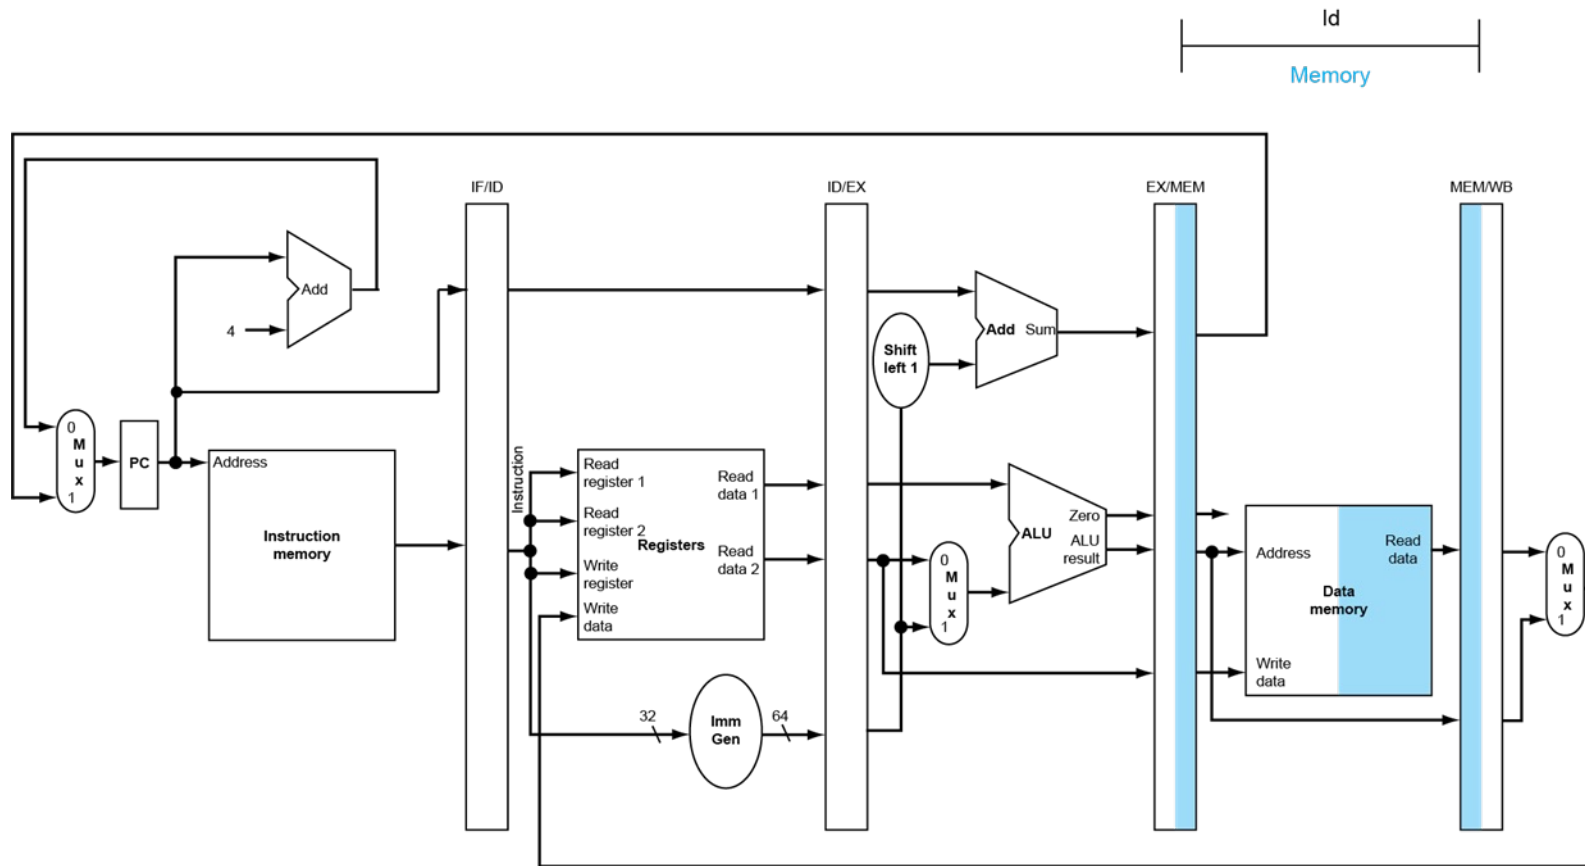
ID for Load, Store, ...



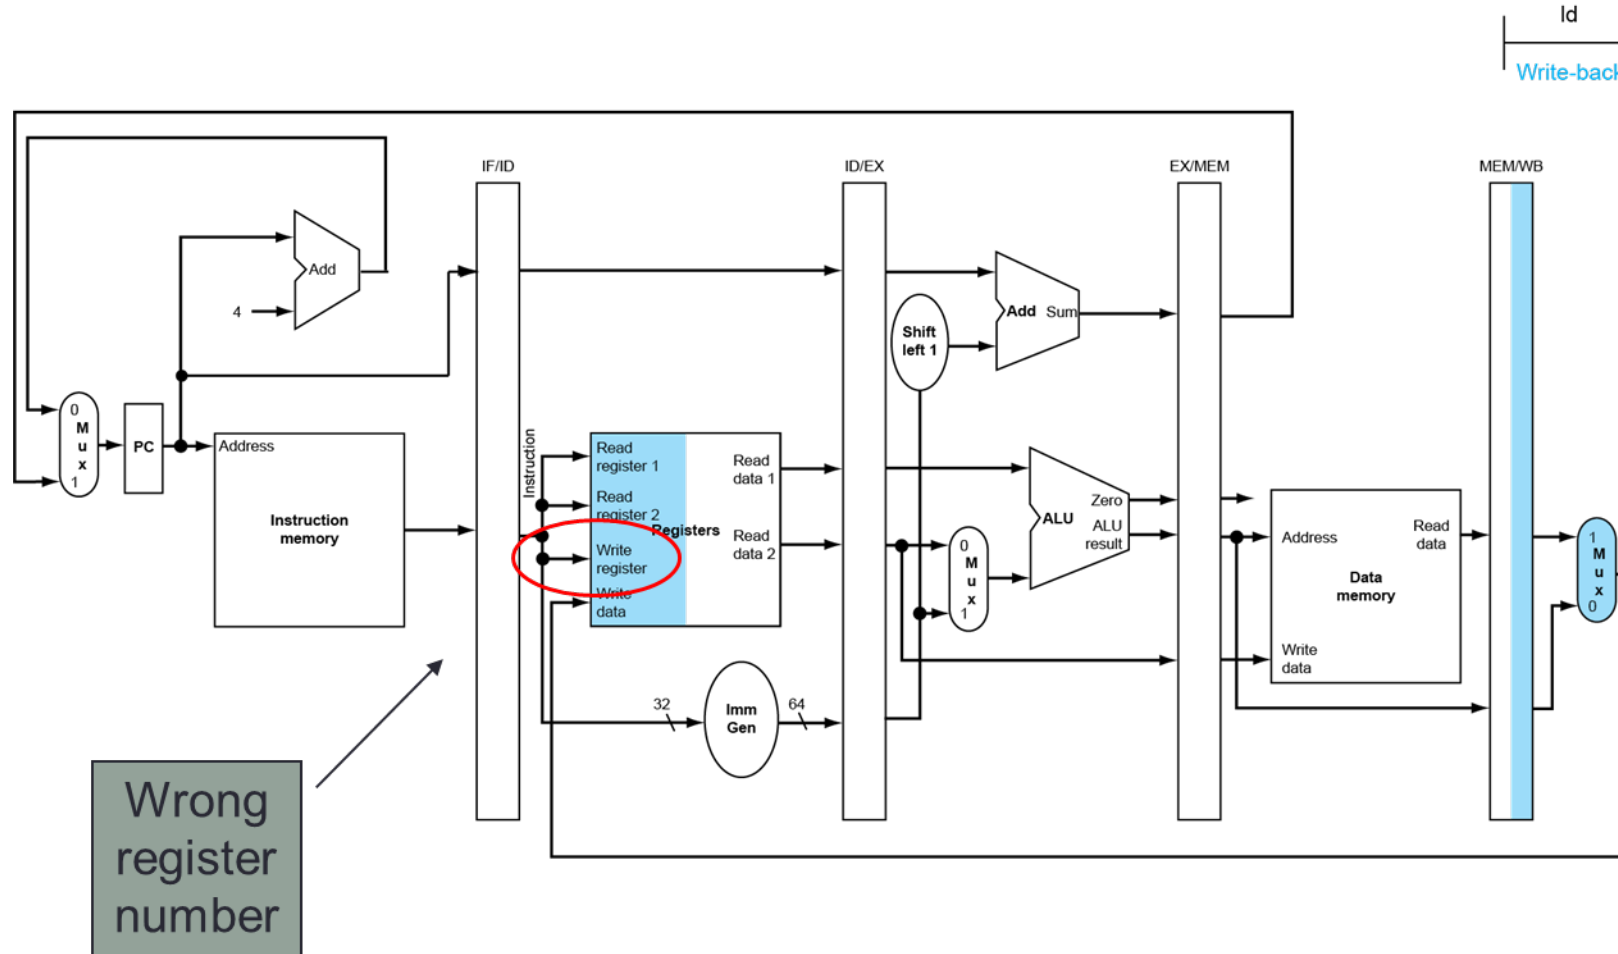
EX for Load



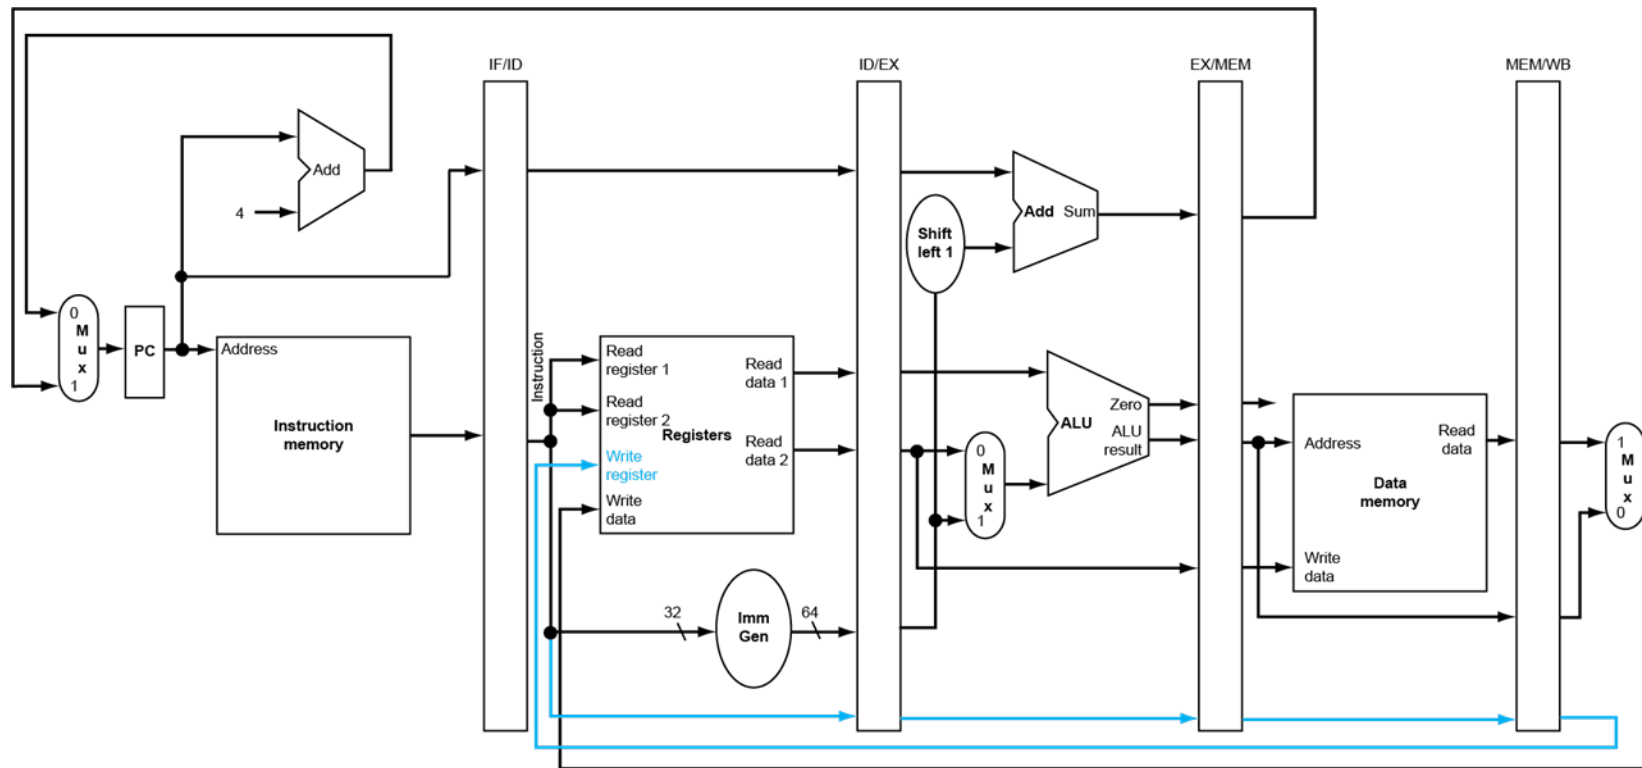
MEM for Load



WB for Load

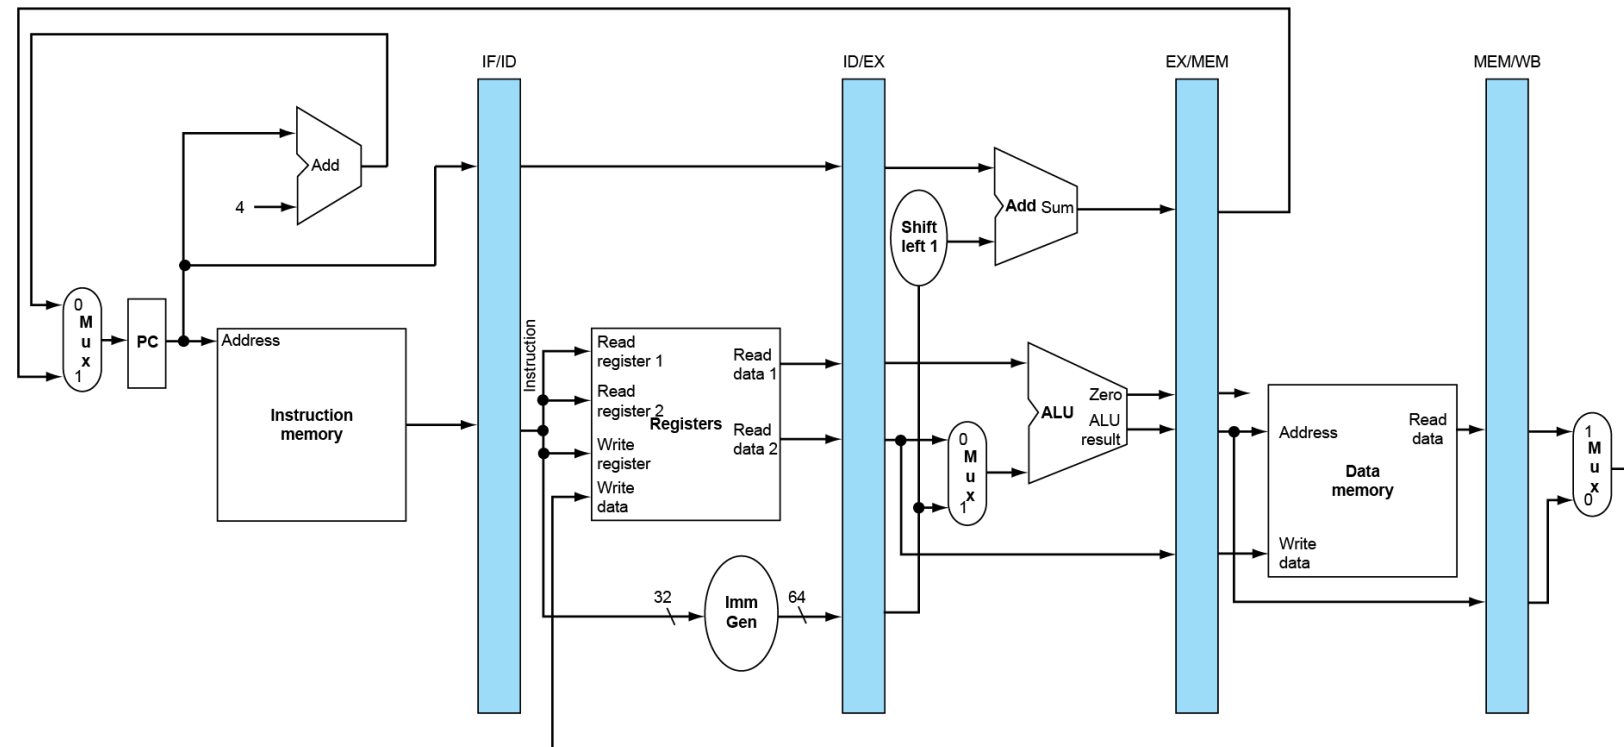


Corrected Datapath for Load



Pipeline registers

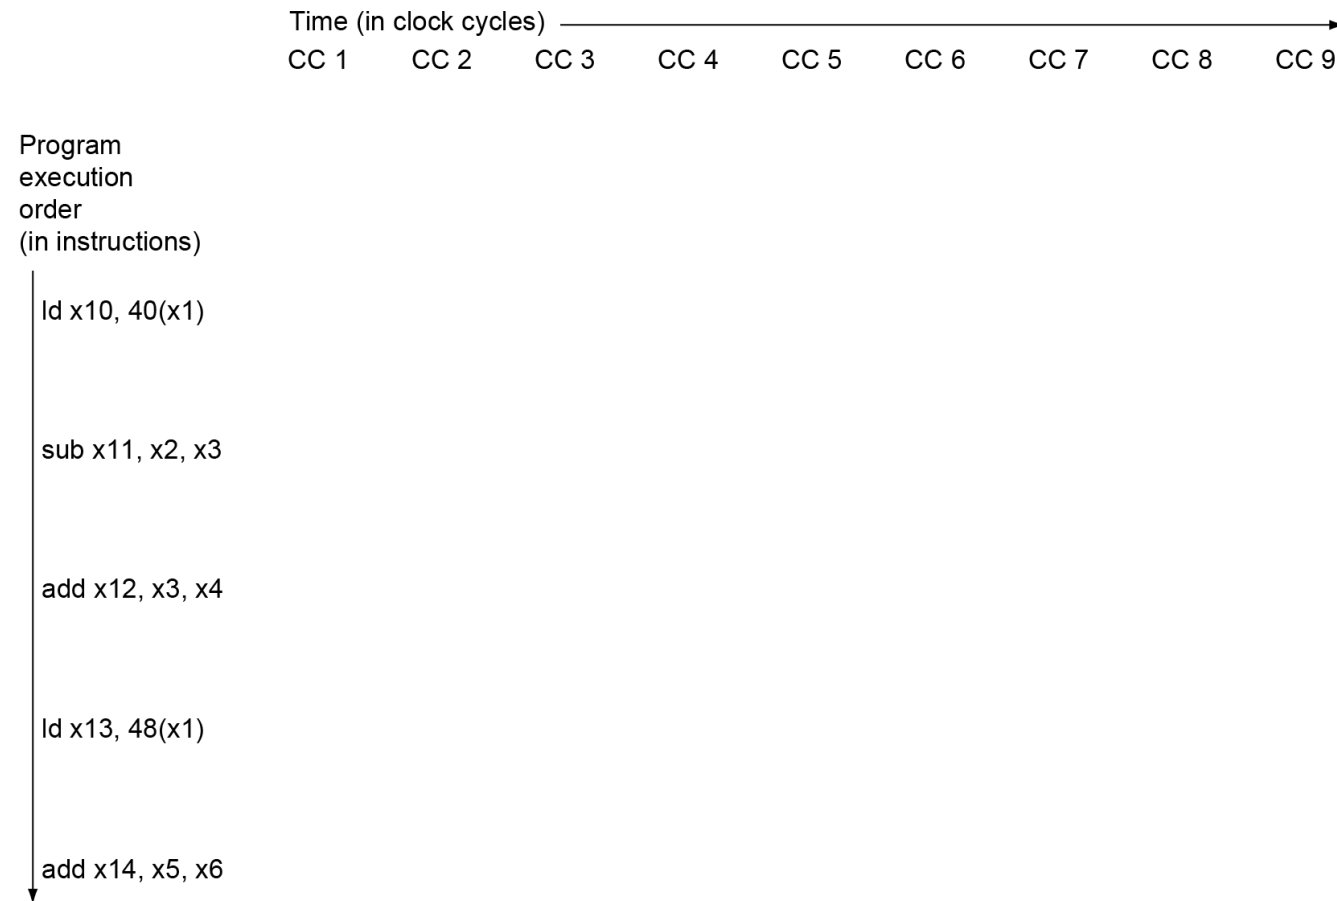
- Need registers between stages
 - To hold information produced in previous cycle
 - Progress the instruction fields with the pipeline



Stage	Any instruction		
IF	$IF/ID.IR \leftarrow Mem[PC]$ $IF/ID.NPC, PC \leftarrow (if ((EX/MEM.opcode = branch) \& EX/MEM.cond) \{EX/MEM.ALUOutput\} else \{PC+4\});$		
ID	$ID/EX.A \leftarrow Regs[IF/ID.IR[rs1]]$; $ID/EX.B \leftarrow Regs[IF/ID.IR[rs2]]$; $ID/EX.NPC \leftarrow IF/ID.NPC$; $ID/EX.IR \leftarrow IF/ID.IR$; $ID/EX.Imm \leftarrow sign-extend(IF/ID.IR[immediate\ field]);$		
	ALU instruction	Load instruction	Branch instruction
EX	$EX/MEM.IR \leftarrow ID/EX.IR$; $EX/MEM.ALUOutput \leftarrow$ $ID/EX.A\ func\ ID/EX.B$; or $EX/MEM.ALUOutput \leftarrow$ $ID/EX.A\ op\ ID/EX.Imm$;	$EX/MEM.IR\ to\ ID/EX.IR$ $EX/MEM.ALUOutput \leftarrow$ $ID/EX.A + ID/EX.Imm$; $EX/MEM.B \leftarrow ID/EX.B$;	$EX/MEM.ALUOutput \leftarrow$ $ID/EX.NPC +$ $(ID/EX.Imm < < 2)$; $EX/MEM.cond \leftarrow$ $(ID/EX.A == ID/EX.B)$;
MEM	$MEM/WB.IR \leftarrow EX/MEM.IR$; $MEM/WB.ALUOutput \leftarrow$ $EX/MEM.ALUOutput$;	$MEM/WB.IR \leftarrow EX/MEM.IR$; $MEM/WB.LMD \leftarrow$ $Mem[EX/MEM.ALUOutput]$; or $Mem[EX/MEM.ALUOutput] \leftarrow$ $EX/MEM.B$;	
WB	$Regs[MEM/WB.IR[rd]] \leftarrow$ $MEM/WB.ALUOutput$;	For load only: $Regs[MEM/WB.IR[rd]] \leftarrow$ $MEM/WB.LMD$;	

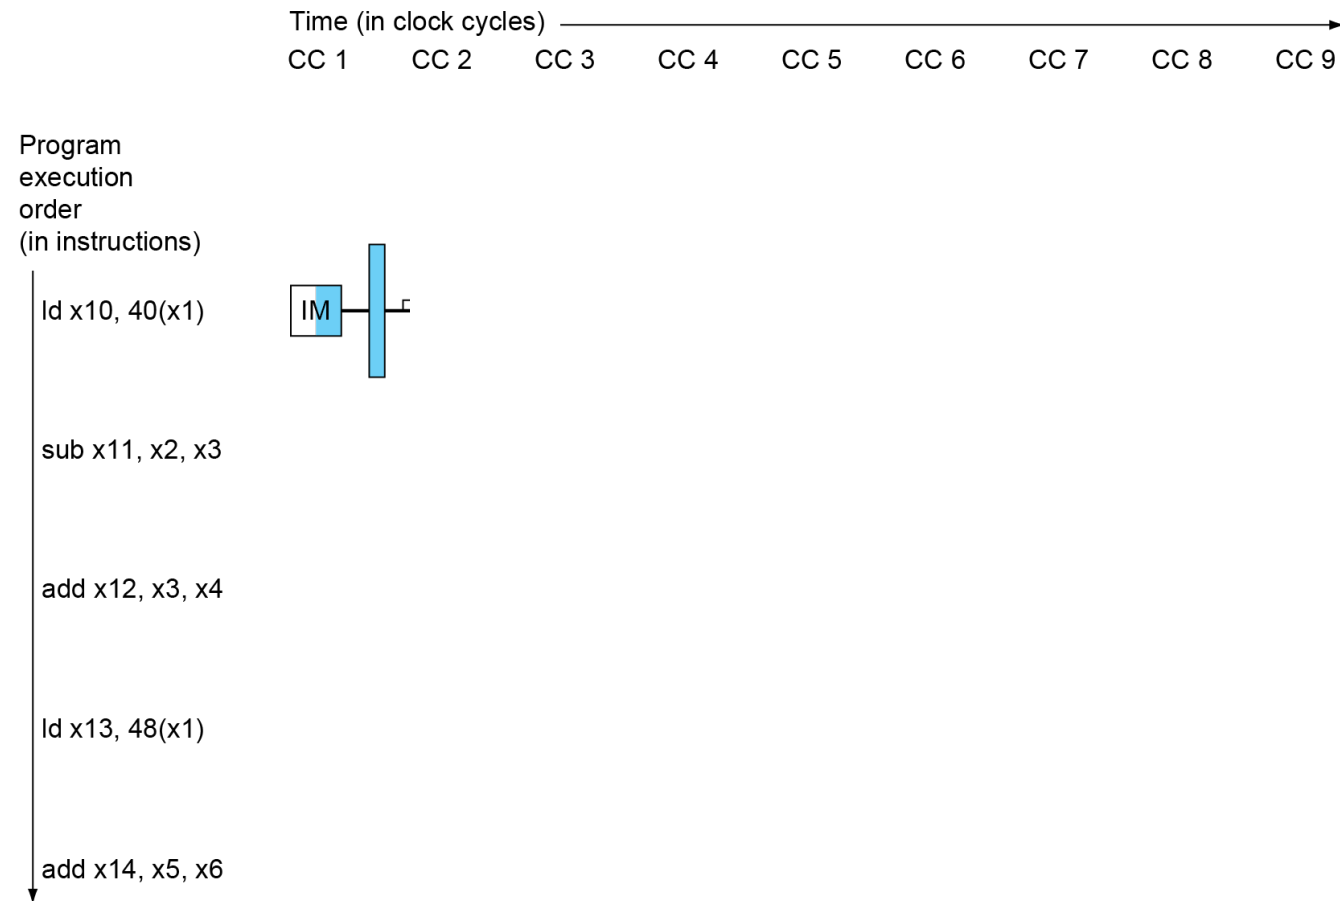
Multi-Cycle Pipeline Diagram

- Form showing resource usage



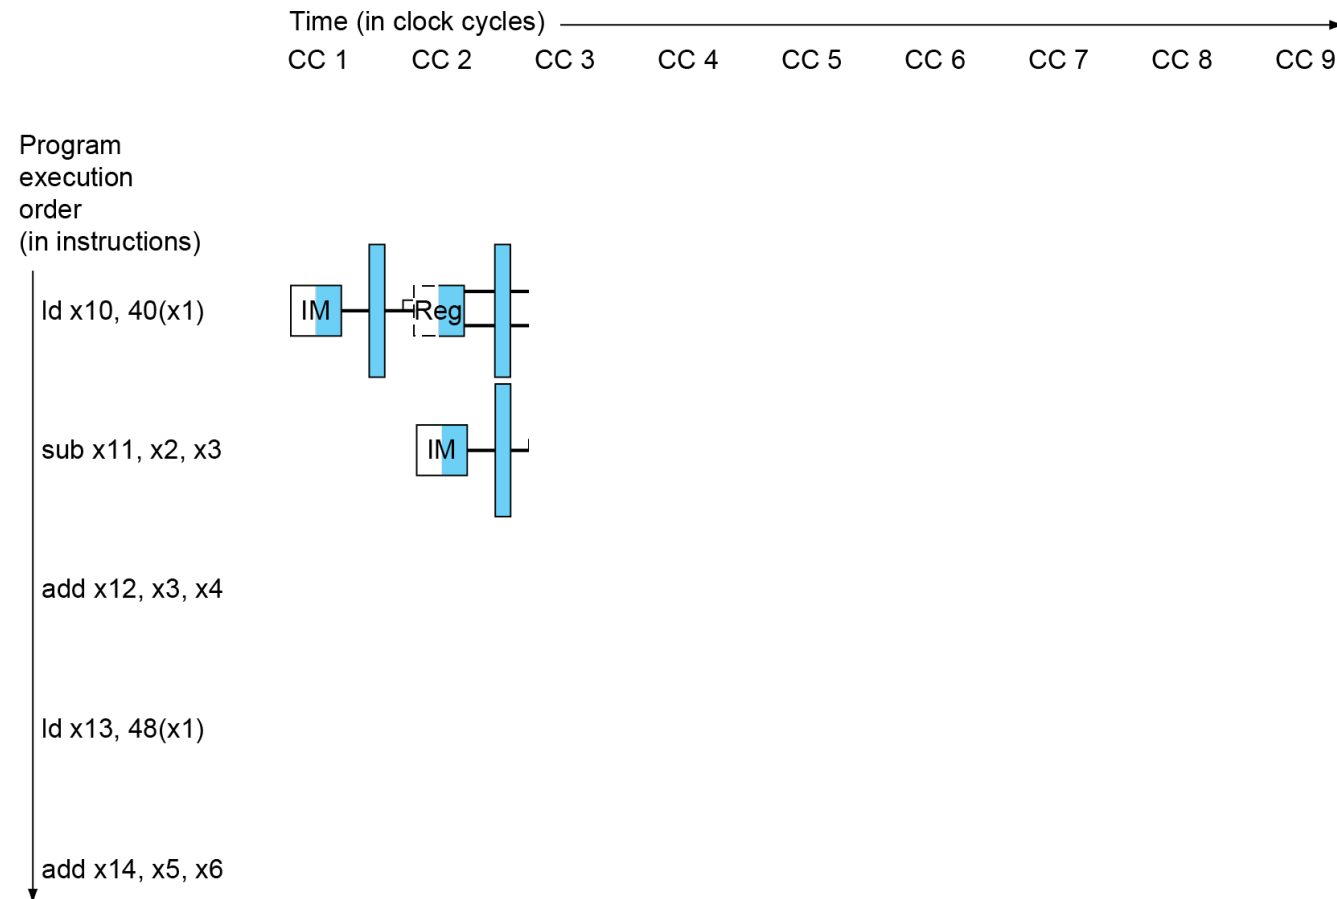
Multi-Cycle Pipeline Diagram

- Form showing resource usage



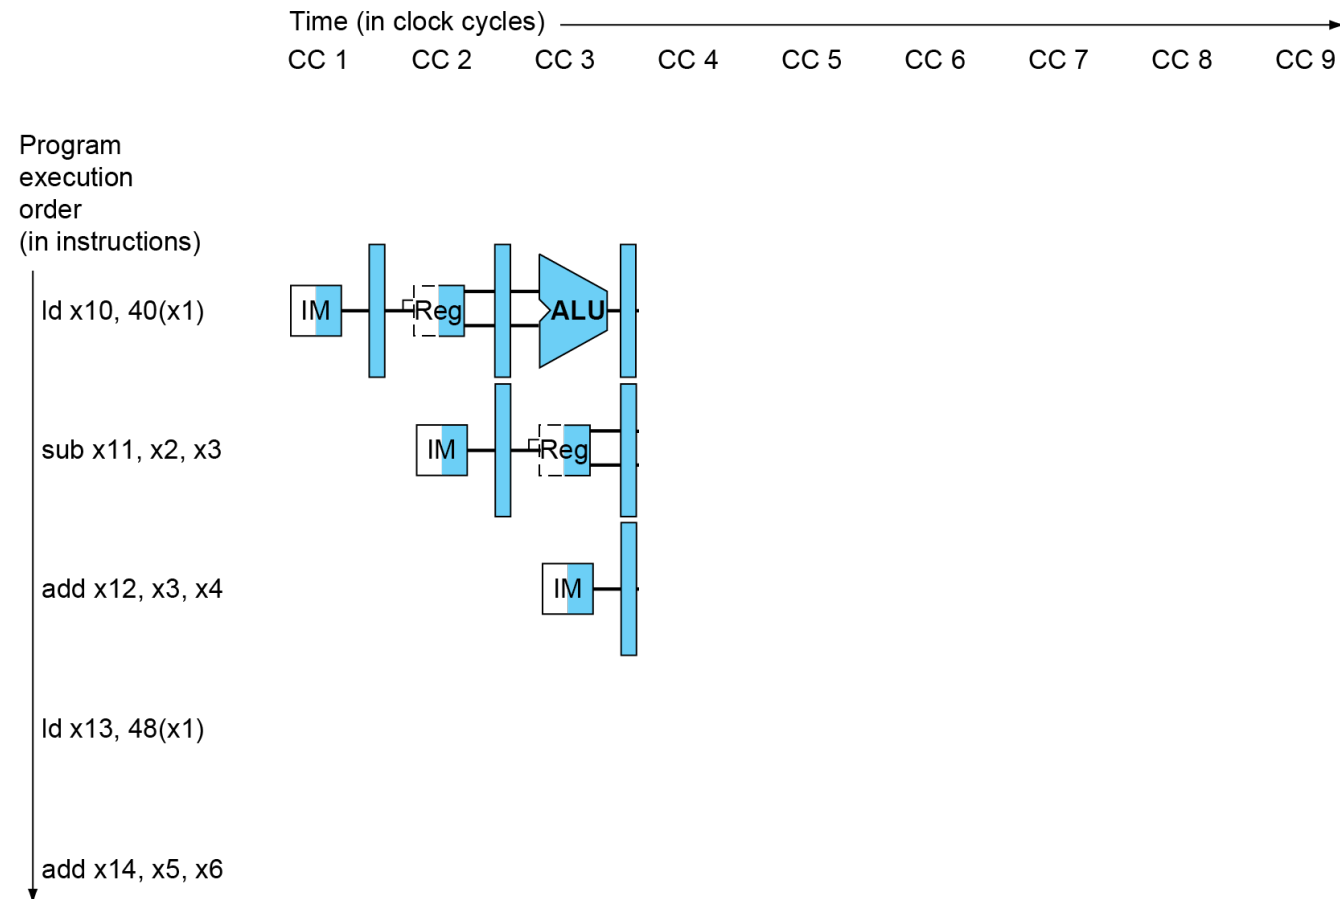
Multi-Cycle Pipeline Diagram

- Form showing resource usage



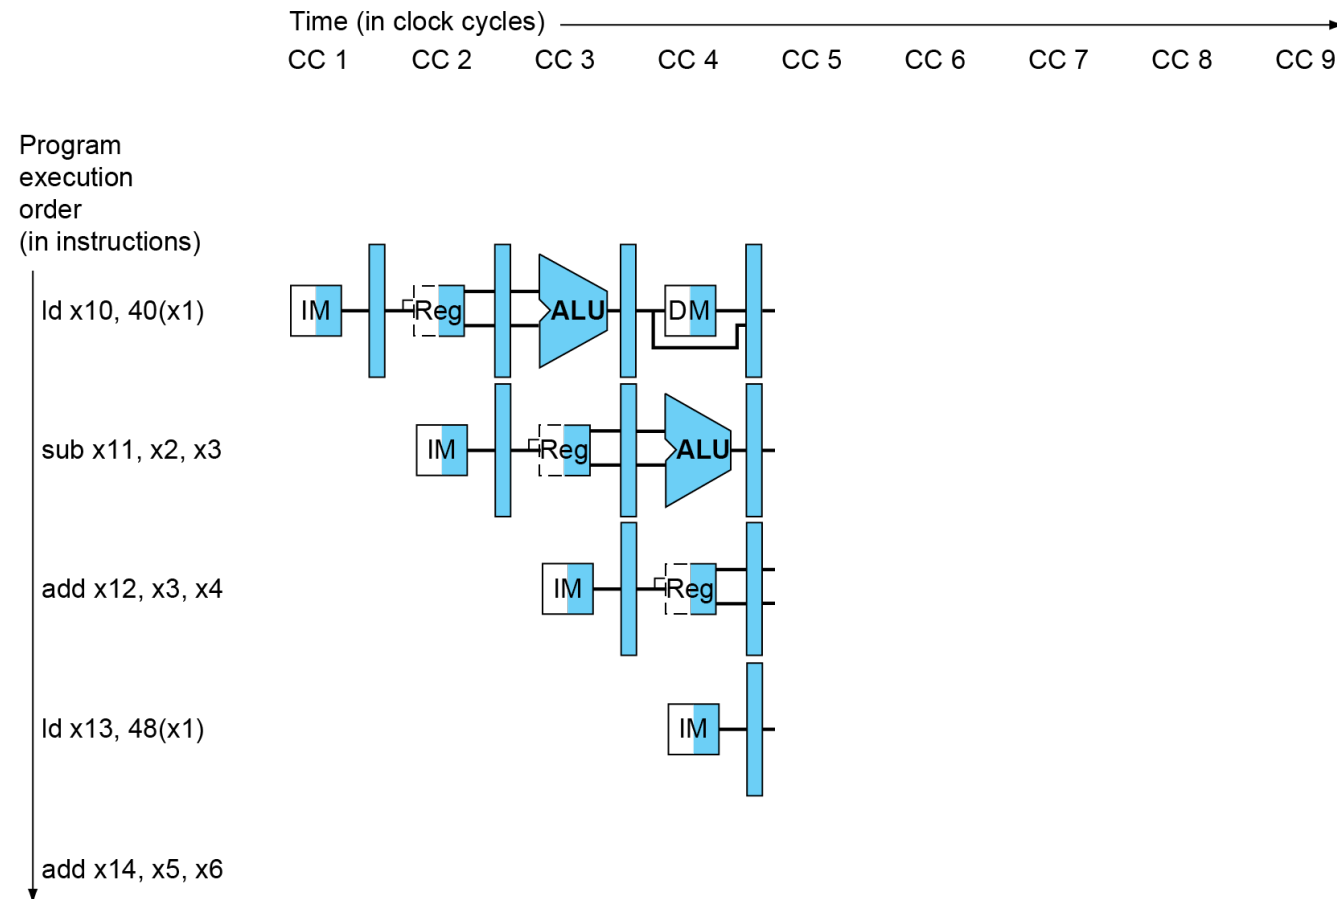
Multi-Cycle Pipeline Diagram

- Form showing resource usage



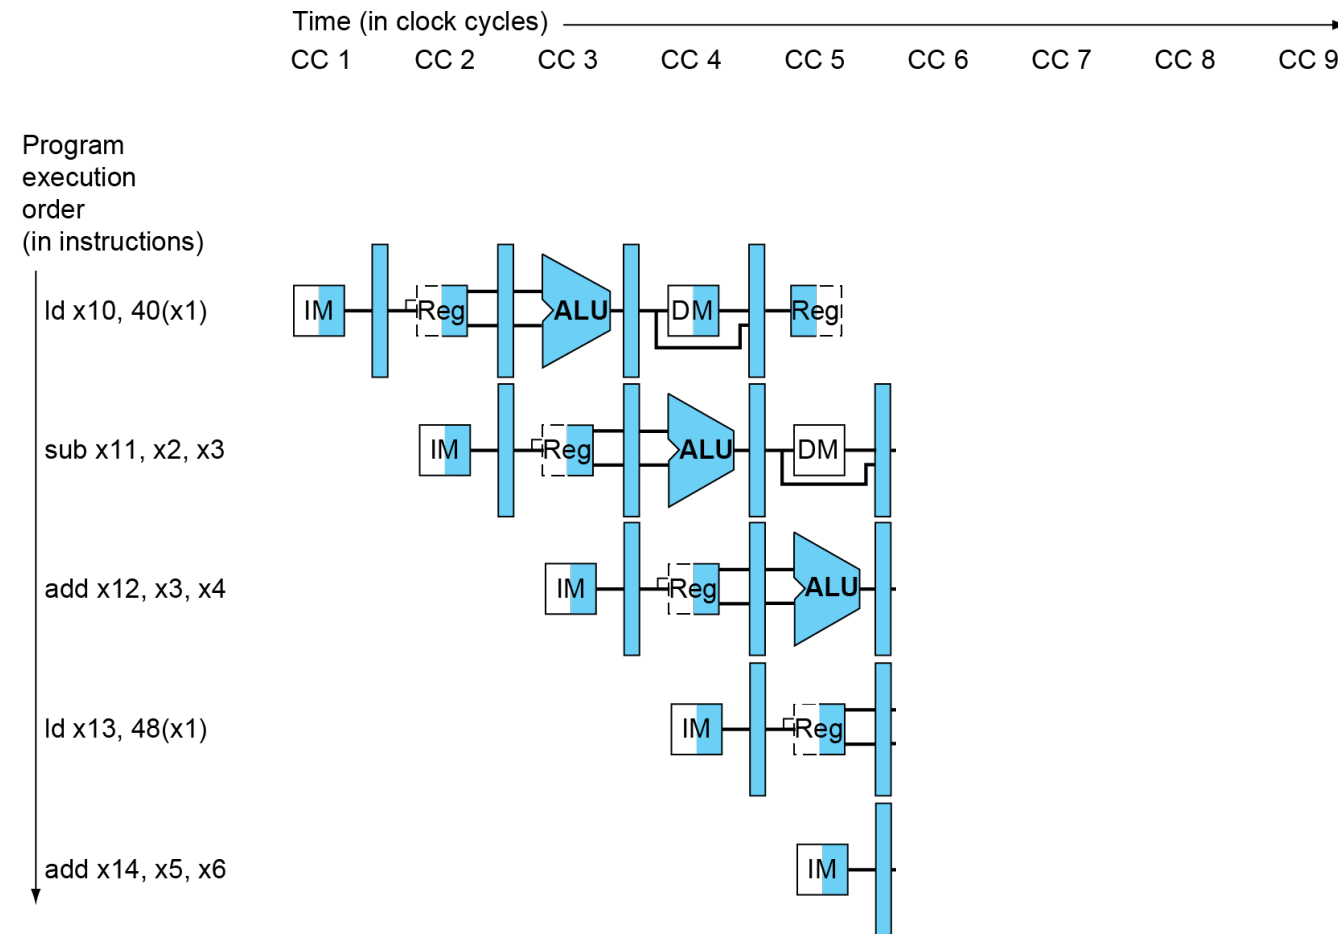
Multi-Cycle Pipeline Diagram

- Form showing resource usage



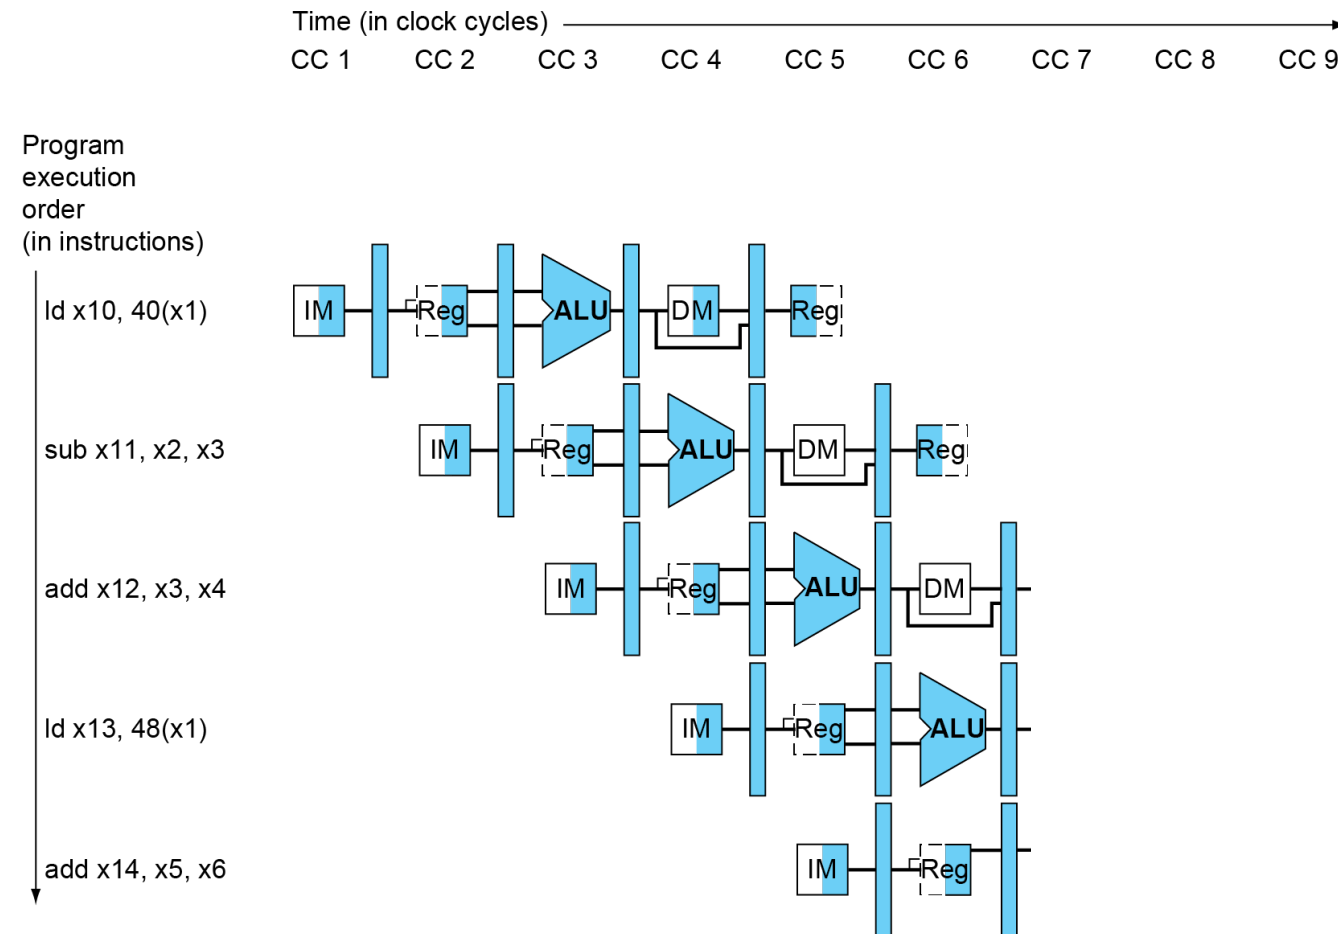
Multi-Cycle Pipeline Diagram

- Form showing resource usage



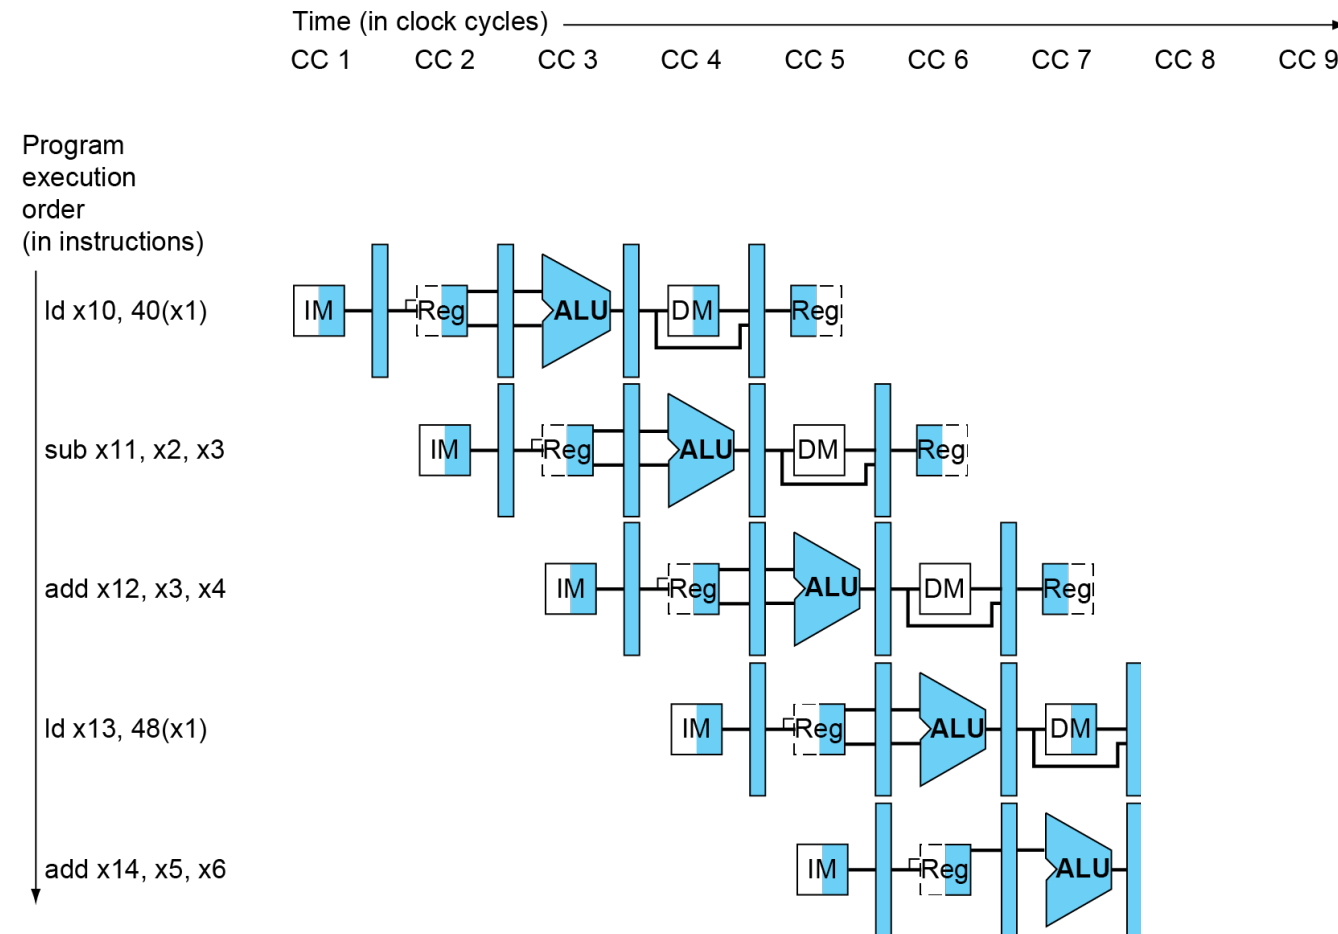
Multi-Cycle Pipeline Diagram

- Form showing resource usage



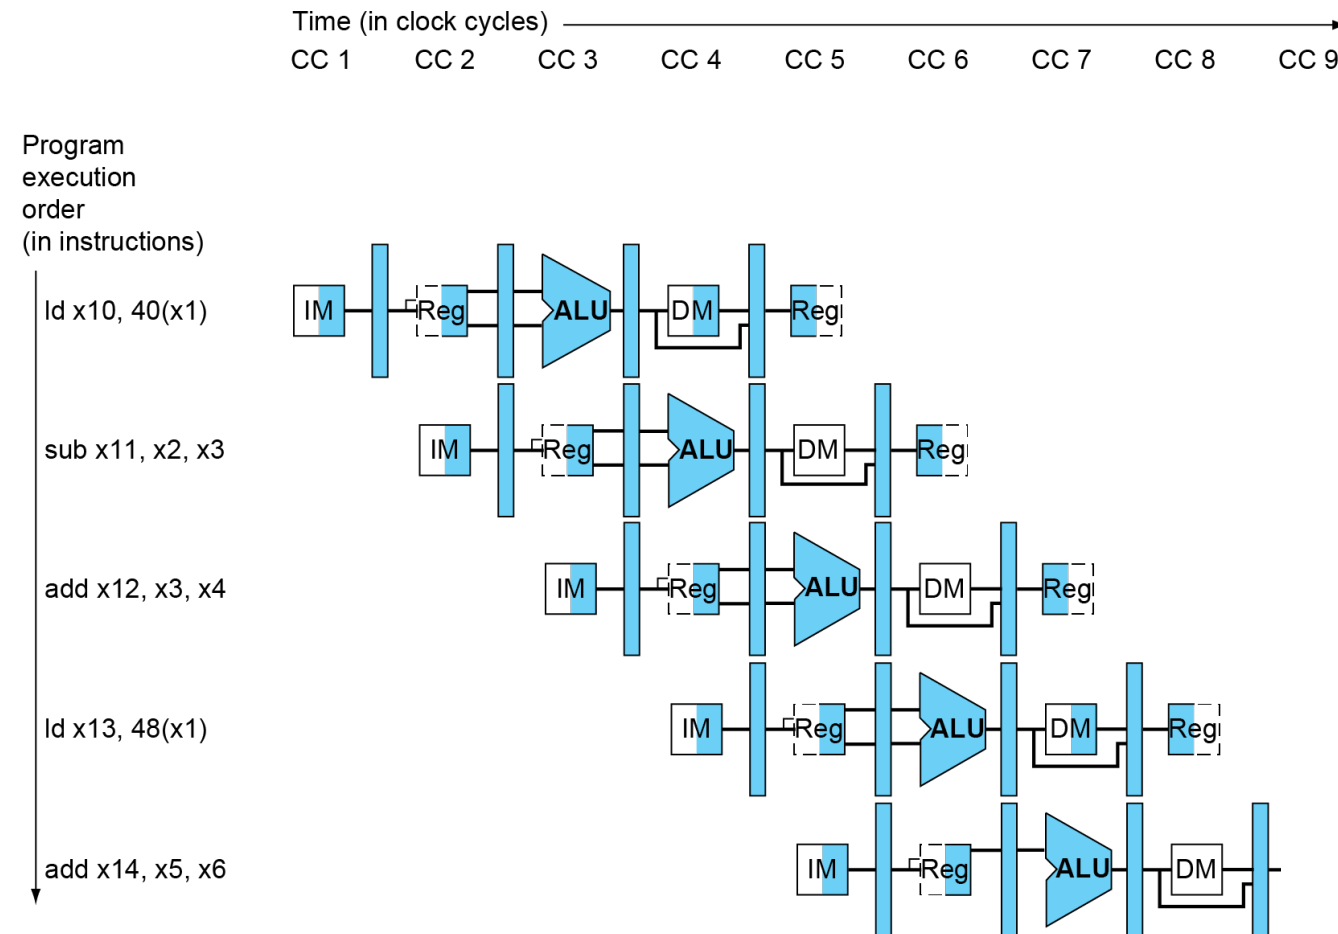
Multi-Cycle Pipeline Diagram

- Form showing resource usage



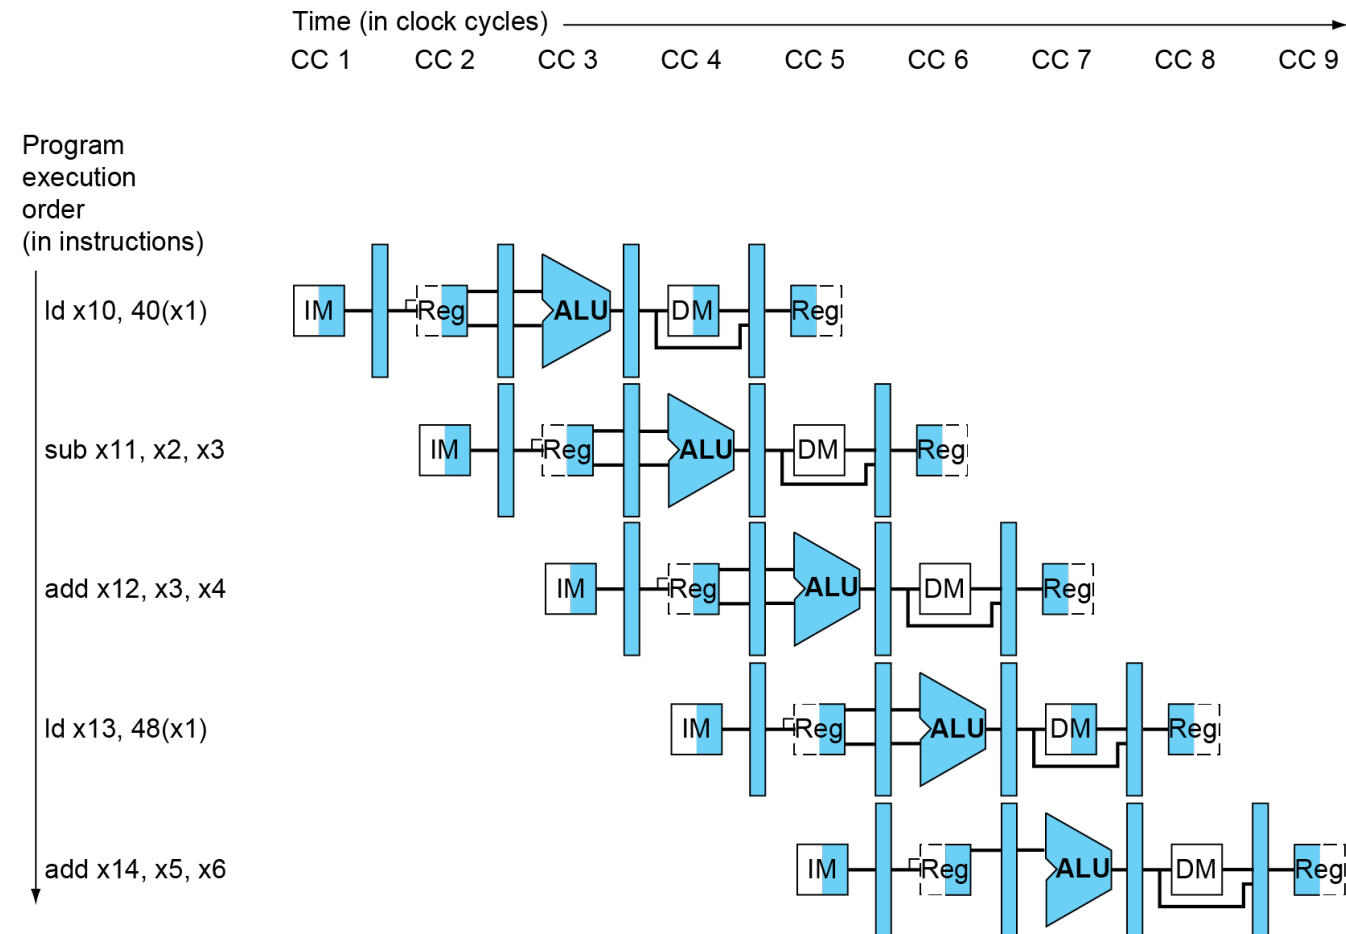
Multi-Cycle Pipeline Diagram

- Form showing resource usage



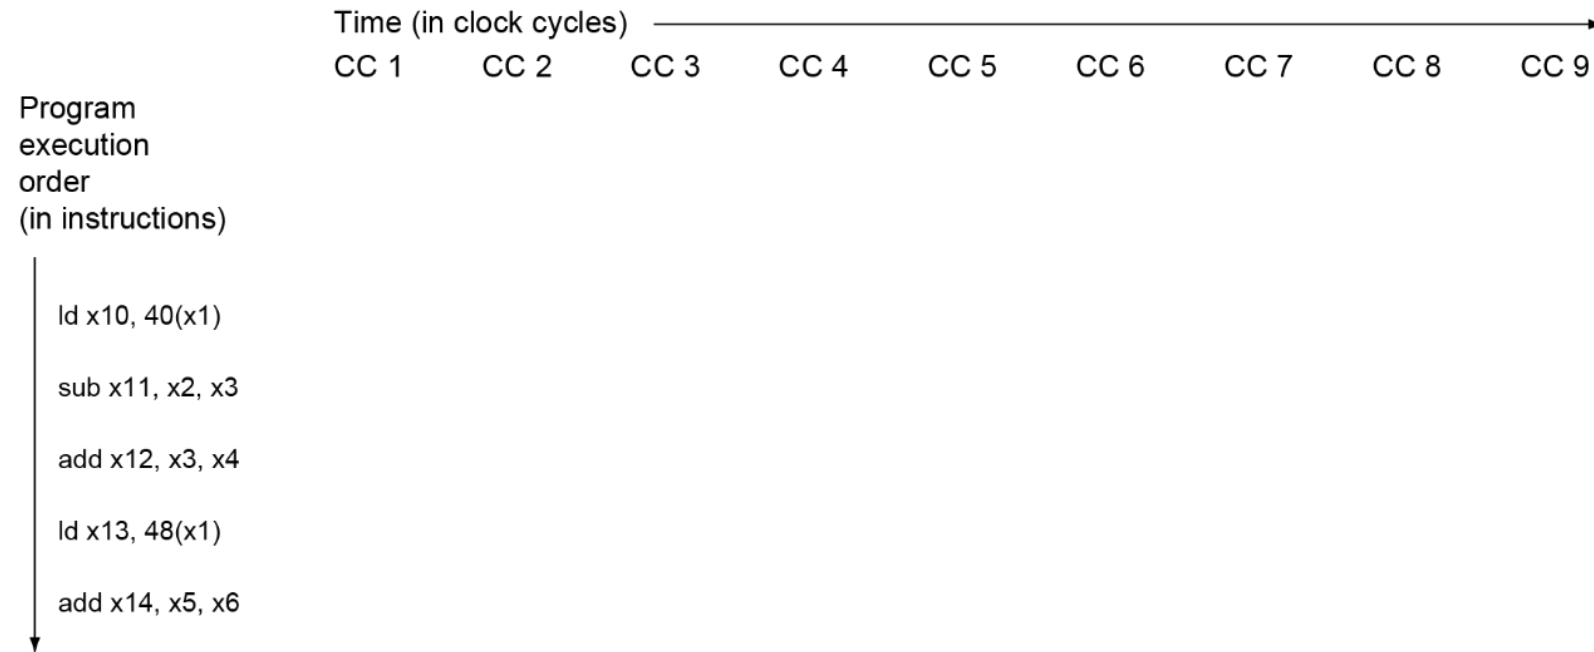
Multi-Cycle Pipeline Diagram

- Form showing resource usage



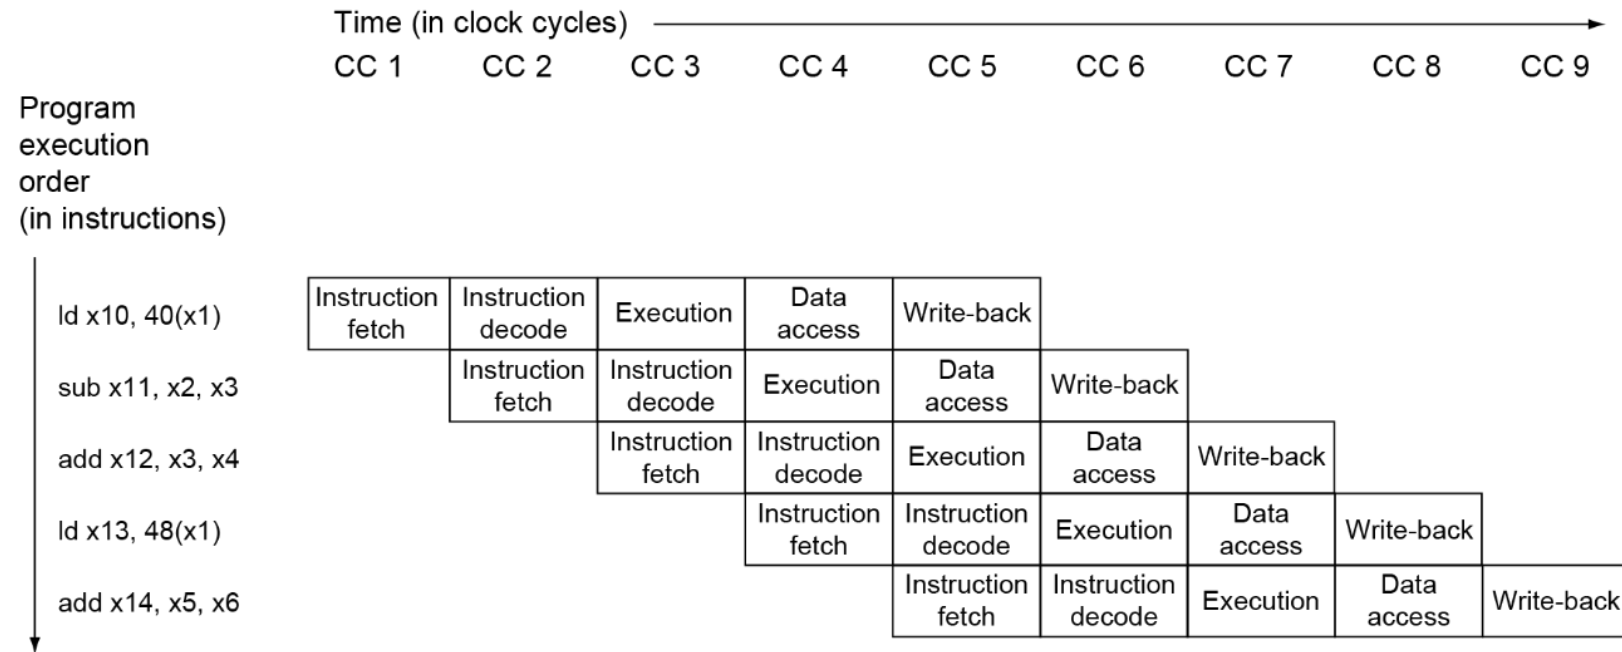
Multi-Cycle Pipeline Diagram

- Traditional form



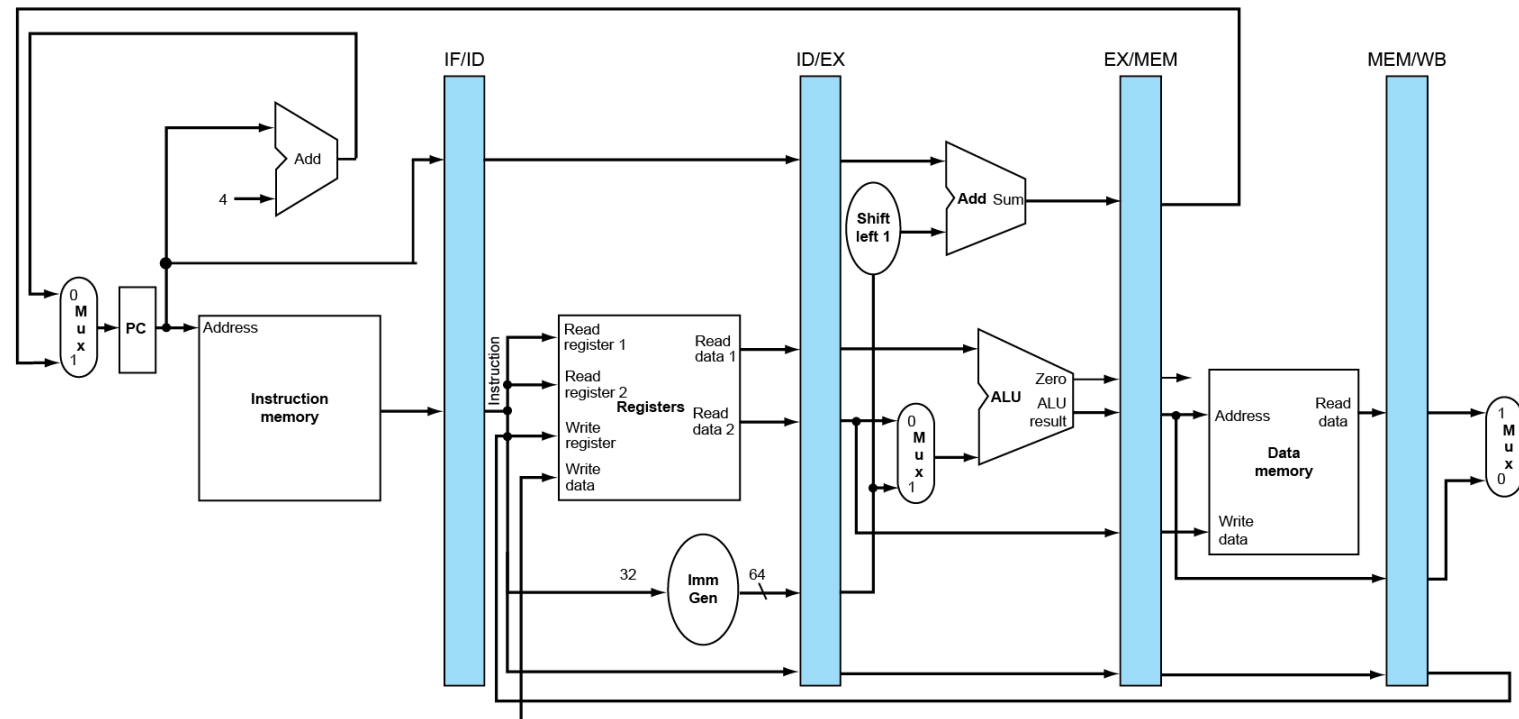
Multi-Cycle Pipeline Diagram

- Traditional form



Single-Cycle Pipeline Diagram

- State of pipeline in a given cycle



Hazards/Alee

- Situations that prevent starting the next instruction in the next cycle
- Structure hazards
 - A required resource is busy
- Data hazard
 - Need to wait for previous instruction to complete its data read/write
- Control hazard
 - Deciding on control action depends on previous instruction

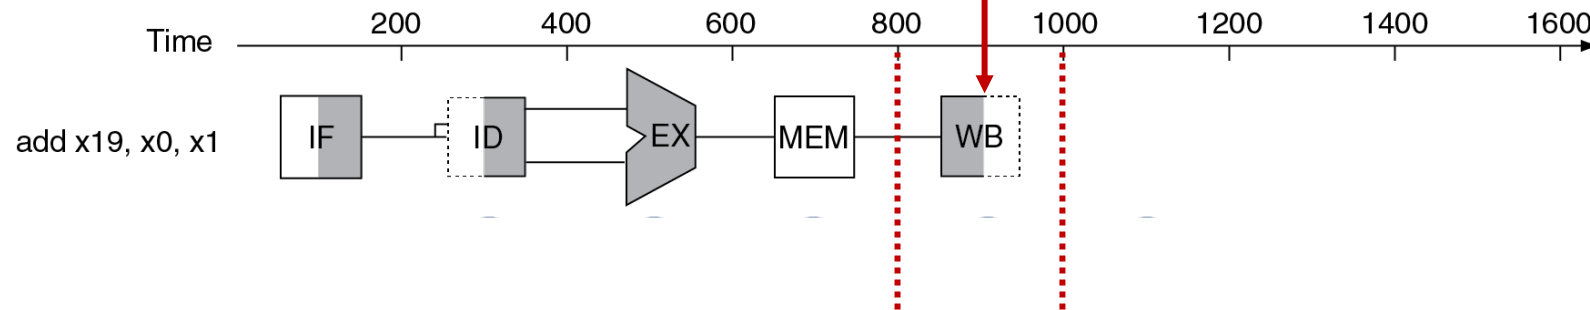
Structural Hazards / Alee Strutturali

- Conflict for use of a resource
- In RISC-V pipeline with a single memory
 - Load/store requires data access
 - Instruction fetch would have to *stall* for that cycle
 - Would cause a pipeline “bubble”
 - Hence, pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches
 - Register file accessed in ID and WB.
 - Write the Register File on clock's negative edges (write on clock's first semi-period)
 - Read the Register File on clock's positive edges (read on clock's second semi-period)

Data Hazards

- An instruction depends on completion of data access by a previous instruction

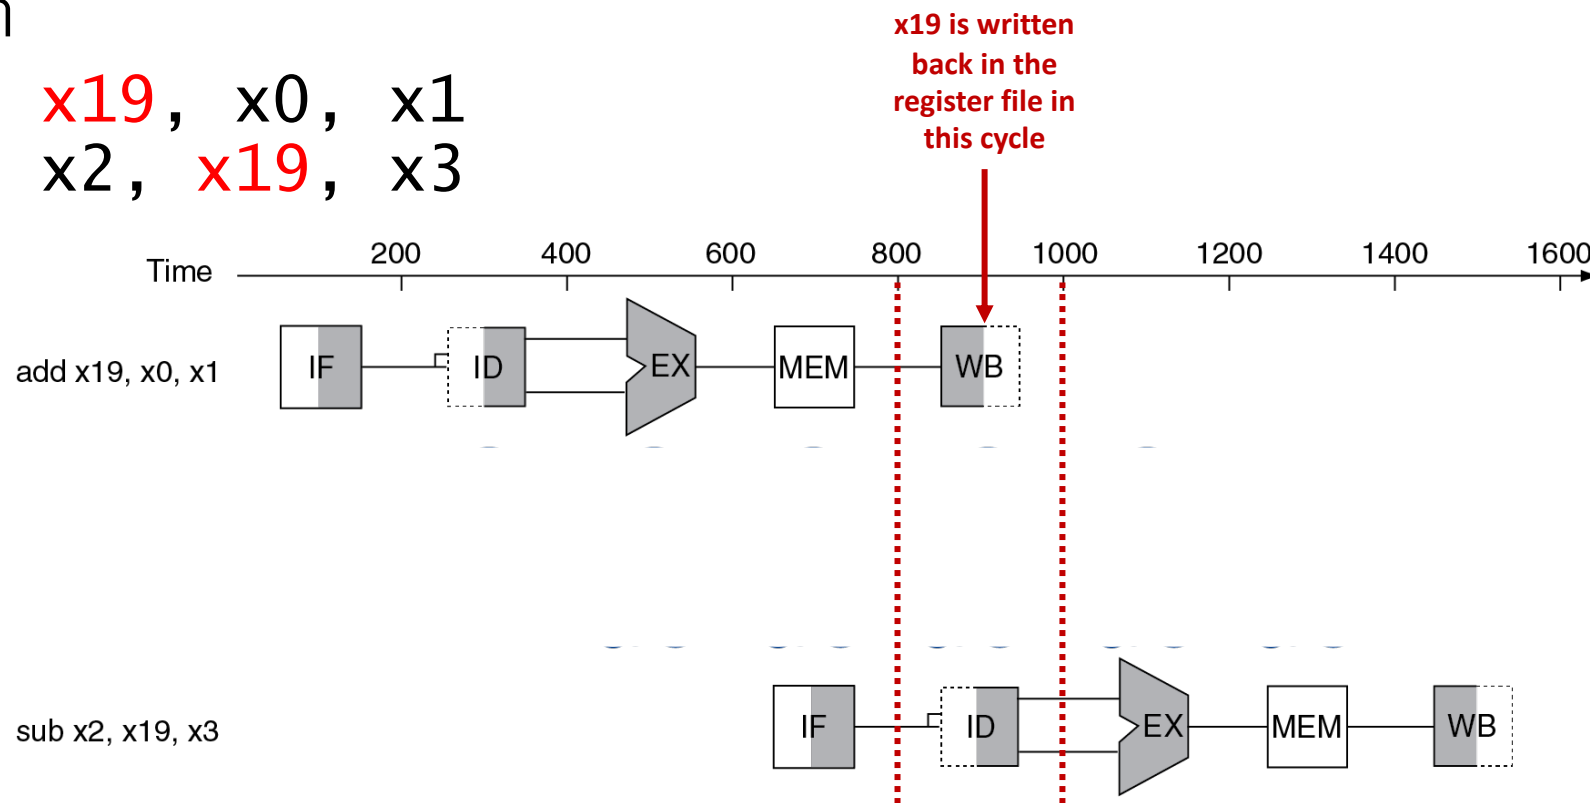
- add **x19**, x0, x1
sub x2, **x19**, x3



Data Hazards

- An instruction depends on completion of data access by a previous instruction

- add **x19**, x0, x1
- sub x2, **x19**, x3



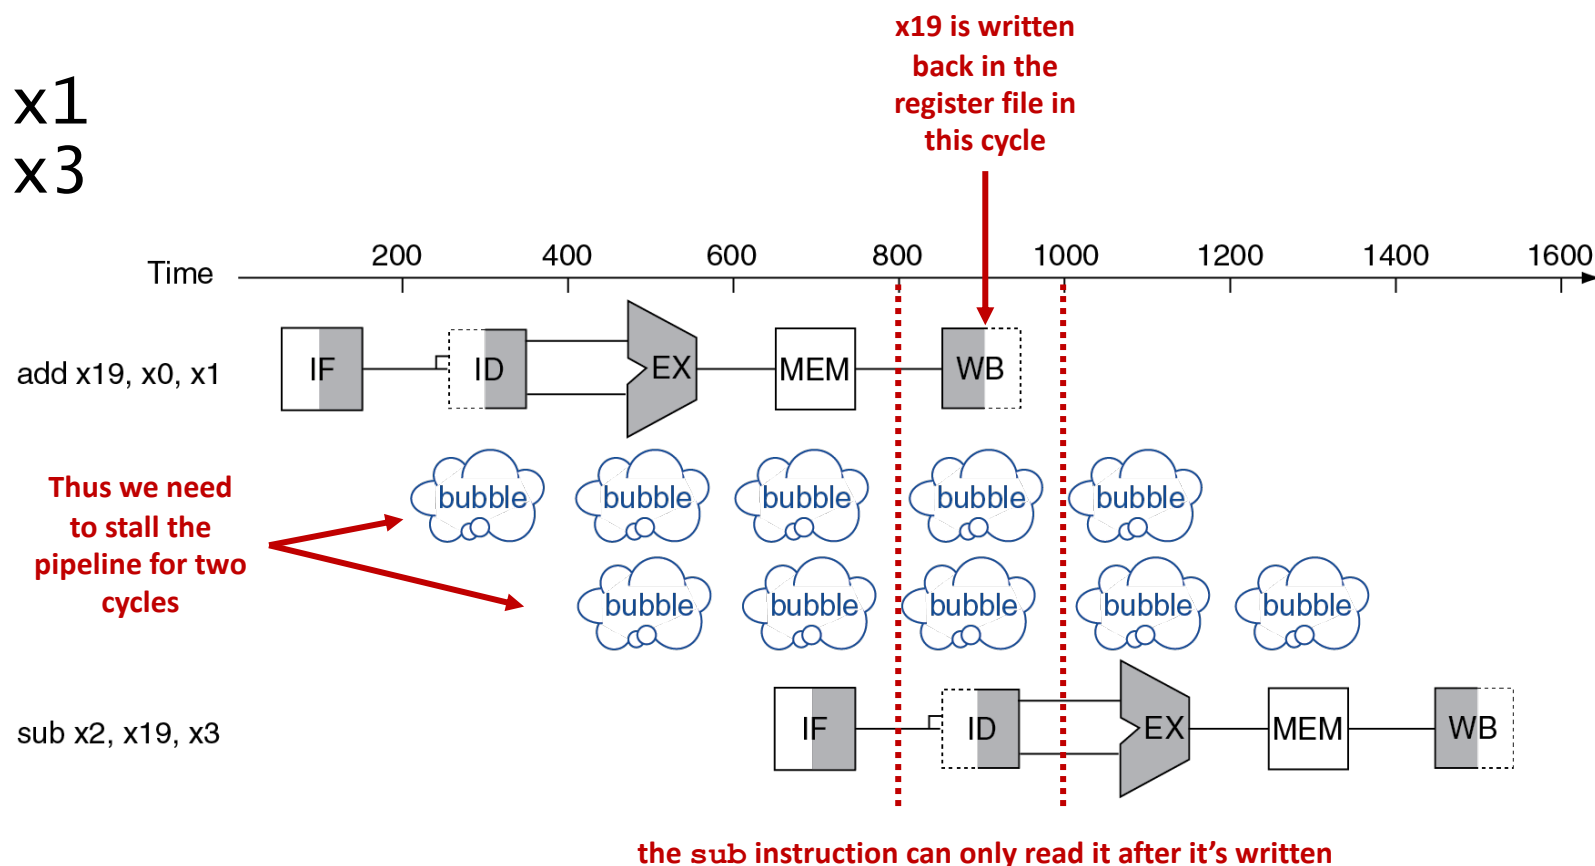
the sub instruction can only read it after it's written

Data Hazards

- An instruction depends on completion of data access by a previous instruction
 - add **x19**, x0, x1
 - sub x2, **x19**, x3

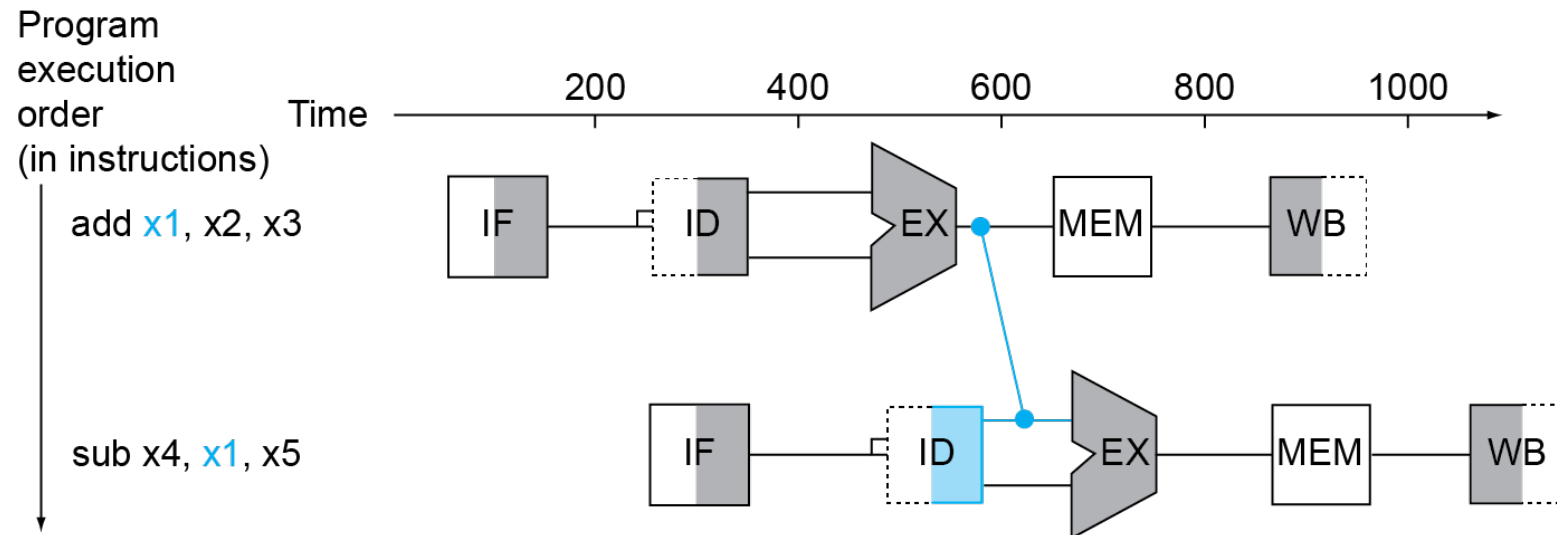
Pipeline interlock:

Il componente *pipeline interlock* rileva le alee e stalla la pipeline finchè l'alea non è risolta. Il CPI dell'istruzione cresce con la lunghezza degli stalli.



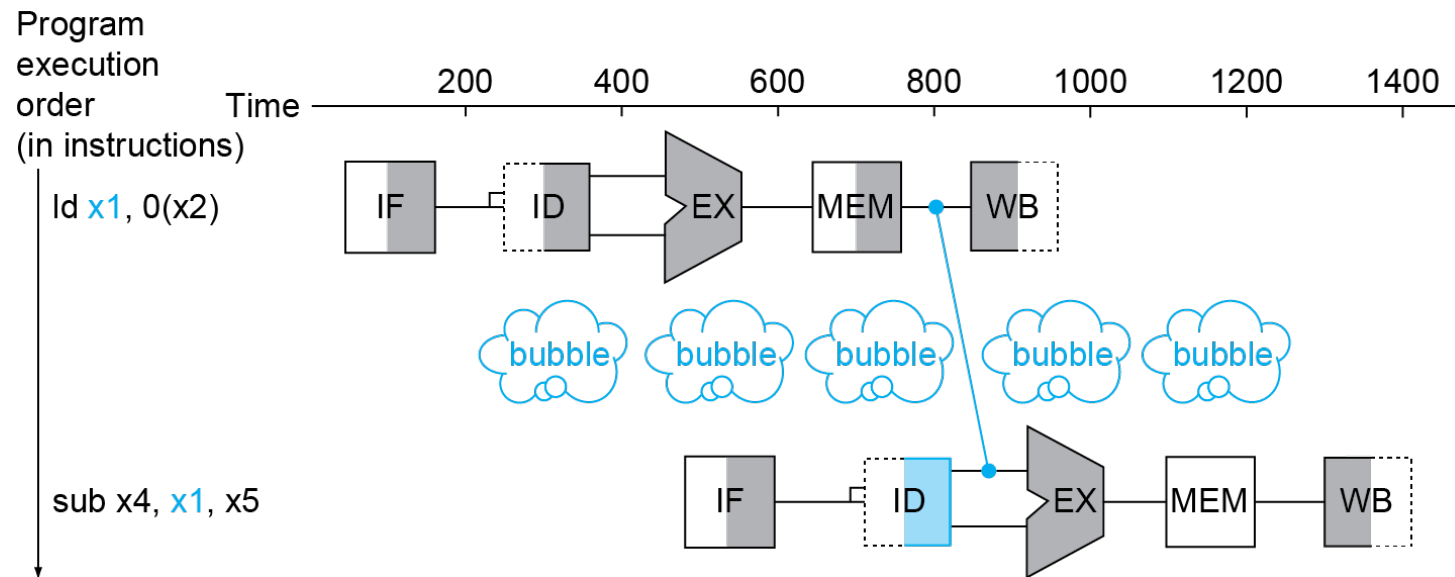
Forwarding (aka Bypassing)

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



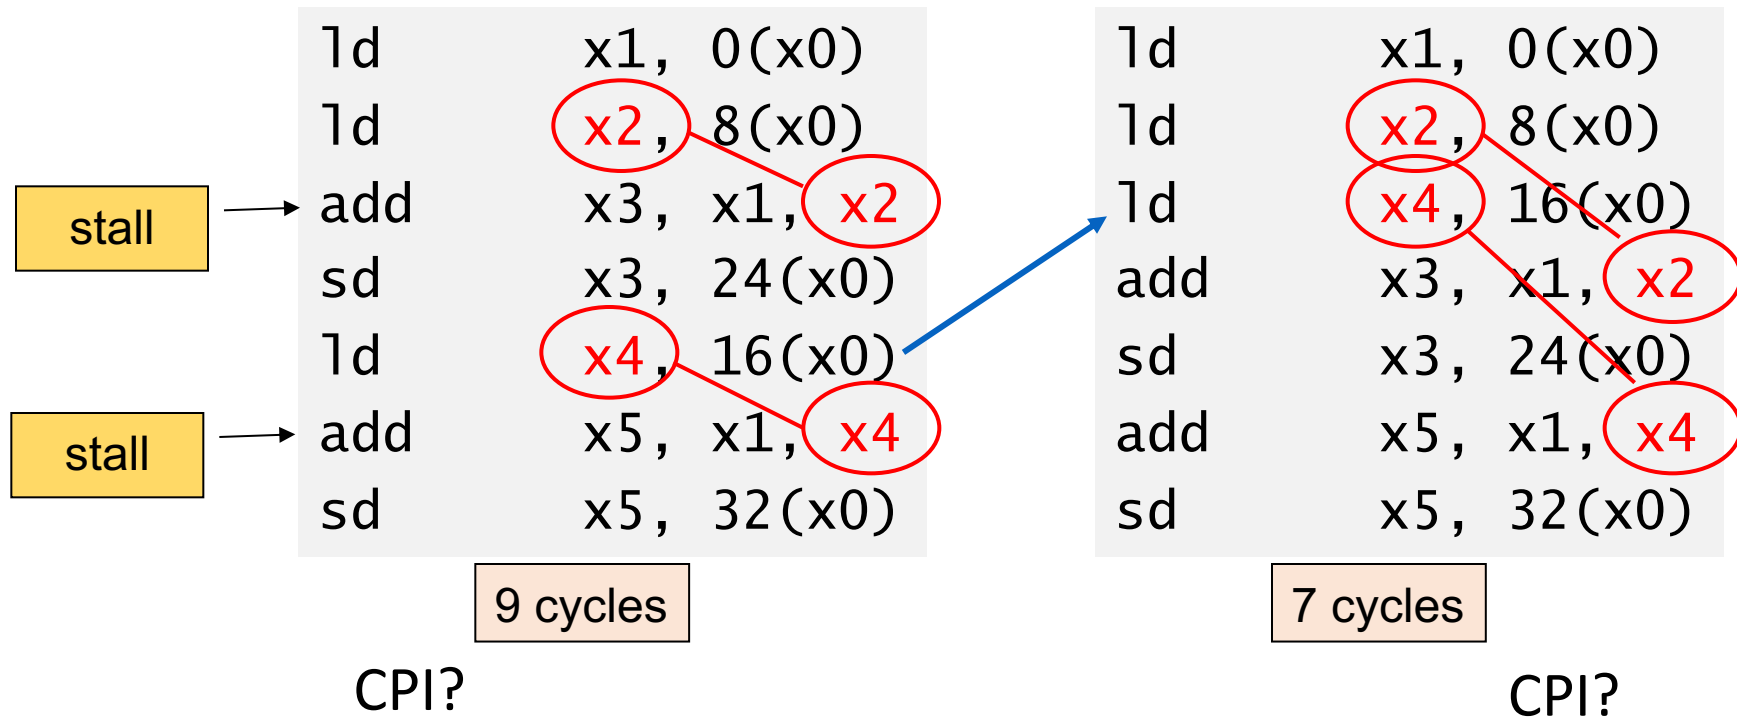
Load-Use Data Hazard

- Can't always avoid stalls by forwarding
 - If value not computed when needed
 - Can't forward backward in time!



Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for **a = b + e; c = b + f;**

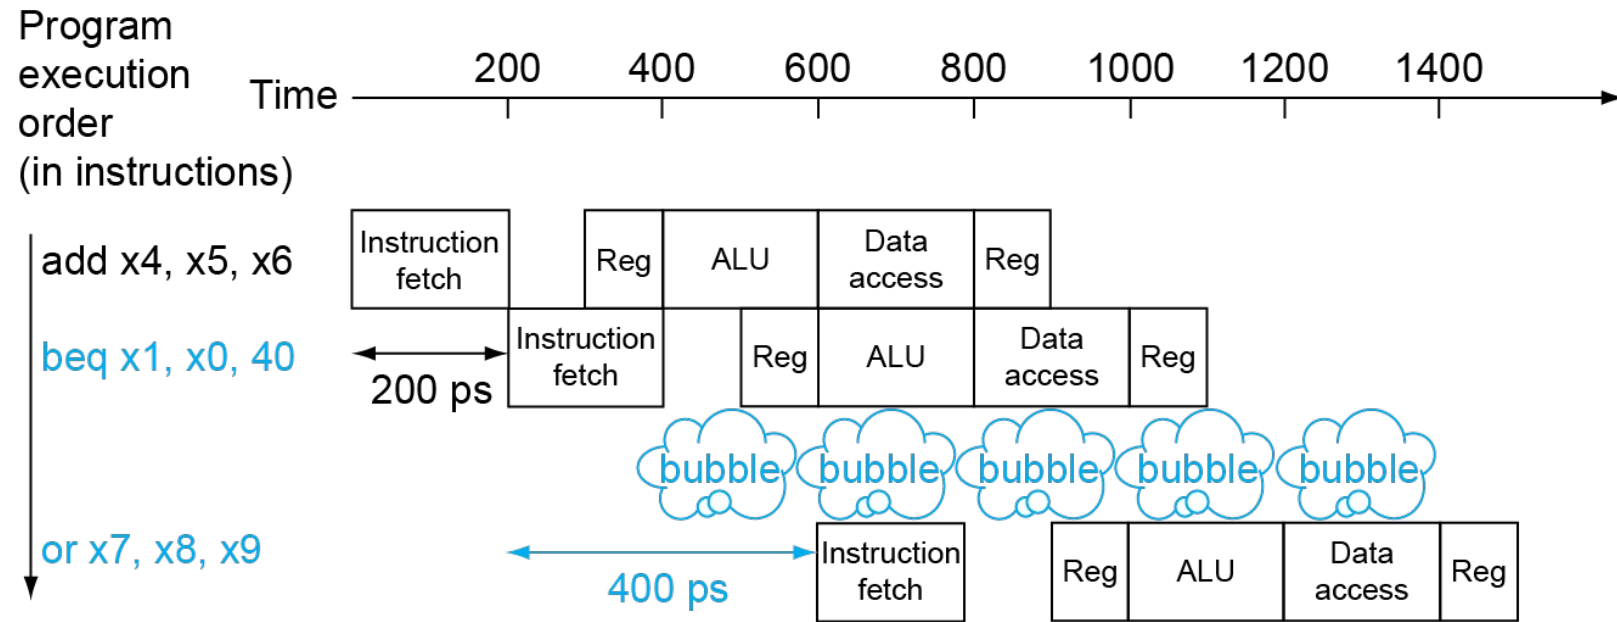


Control Hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch
- In RISC-V pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage
- Countermeasure (assuming branch target address calculation in ID):
 - Static :
 1. Stall on every Branch
 2. Always predict taken or predict untaken
 3. Branch delay slot – not used anymore.
 - Dynamic branch prediction.

1.Static: Stall on Branch

- Wait until branch outcome determined before fetching next instruction



2. Static: Always predict “taken” or “untaken”

- If prediction wrong need to turn fetched instruction in a no-op to avoid changing the state of the processor.
- Compiler can write code to maximize branch prediction correctness.

Untaken branch instruction	IF	ID	EX	MEM	WB				
Instruction $i + 1$		IF	ID	EX	MEM	WB			
Instruction $i + 2$			IF	ID	EX	MEM	WB		
Instruction $i + 3$				IF	ID	EX	MEM	WB	
Instruction $i + 4$					IF	ID	EX	MEM	WB
Taken branch instruction	IF	ID	EX	MEM	WB				
Instruction $i + 1$		IF	idle	idle	idle	idle			
Branch target			IF	ID	EX	MEM	WB		
Branch target+1				IF	ID	EX	MEM	WB	
Branch target+2					IF	ID	EX	MEM	WB

Processor State:

Insieme di registri ed elementi di memoria che determina in modo univoco il comportamento della CPU (i.e. RegisterFile, Mem).

Ogni modifica allo stato del processore è irreversibile.

Branch Prediction

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong
- In RISC-V pipeline
 - Can predict branches not taken
 - Fetch instruction after branch, with no delay

More-Realistic Branch Prediction

- Dynamic branch prediction
 - Hardware measures actual branch behavior
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history

Pipeline Summary

- Pipelining improves performance by increasing instruction throughput
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- Subject to hazards
 - Structure, data, control
- Instruction set design affects complexity of pipeline implementation

Stalls and Performance

- Stalls reduce performance
 - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

$\text{Clock Cycles} = \text{Instruction Count} \times \text{Cycles per Instruction}$

$\text{CPU Time}_{NO\ STALLS} = \text{Instruction Count} \times \text{CPI}_{min} \times \text{Cycle Time}$

$\text{CPU Time}_{WITH\ STALLS} = \text{Instruction Count} \times (\text{CPI}_{min} + \#Stalls) \times \text{Clock Cycle Time}$

- Given an ISA and a Program the number of Stalls ($\#Stalls$) depends on the micro-architecture.
 - Branch resolved at MEM or ID?
 - Forwarding unit or interlock unit?

Schema per ridurre il costo di branch

