

– algoritmo di tommasulo

Scheduling dinamico con esecuzione speculativa

Scheduling Dinamico (es. con algoritmo di TOMASULO):

Obiettivo: stallare la pipeline solo in caso di alee strutturali e di controllo.

- | | |
|-------------------------------|-------------------------|
| – Attivazione sequenziale: | IN-ORDER DISPATCHING |
| – Esecuzione fuori ordine: | OUT-OF-ORDER EXECUTION |
| – Completamento fuori ordine: | OUT-OF-ORDER COMPLETION |

Scheduling Dinamico con esecuzione speculativa:

▪ **Obiettivo:** stallare la pipeline solo in caso di alee strutturali

▪ **Soluzione:** predizione, evitando azioni “irrevocabili” prima di essere certi che tale azione deve essere effettuata

▪ **Conseguenza:** si amplia la finestra su cui cercare il parallelismo a livello di istruzione

- | | |
|-------------------------------|---------------------------------------|
| – Attivazione sequenziale: | PREDICTION BASED IN-ORDER DISPATCHING |
| – Esecuzione fuori ordine | OUT-OF-ORDER EXECUTION |
| – Completamento fuori ordine: | OUT-OF-ORDER COMPLETION |
| – Avanzamento in ordine: | IN-ORDER COMMIT |

Soluzione

- È necessario separare la fase di esecuzione data driven dalla fase in cui si rendono impegnativi i risultati (COMMIT)
- Si utilizza uno schema simile a quello di Tomasulo ma con aggiunto il Reorder Buffer

Hardware-Based Speculation

- Execute instructions along predicted execution paths but only commit the results if prediction was correct
- Instruction commit: allowing an instruction to update the register file when instruction is no longer speculative
- Need an additional piece of hardware to prevent any irrevocable action until an instruction commits
 - I.e. updating state or taking an execution

Reorder Buffer

- Reorder buffer – holds the result of instruction between completion and commit
- Four fields:
 - Instruction type: branch/store/register
 - Destination field: register number
 - Value field: output value
 - Ready field: completed execution?
- Modify reservation stations:
 - Operand source is now reorder buffer instead of functional unit


Reorder Buffer

- Issue:
 - Allocate RS and ROB, read available operands
- Execute:
 - Begin execution when operand values are available
- Write result:
 - Write result and ROB tag on CDB
- Commit:
 - When a predicted branch reaches head of ROB, update register
 - When a mispredicted branch reaches head of ROB, discard all entries

Reorder Buffer

- Register values and memory values are not written until an instruction commits
- On misprediction:
 - Speculated entries in ROB are cleared
- Exceptions:
 - Not recognized until it is ready to commit

Reorder Buffer



Reorder Buffer

Status	Wait until	Action or bookkeeping
Issue all instructions		<pre> if (RegisterStat[rs].Busy) /*in-flight instr. writes rs*/ { h ← RegisterStat[rs].Reorder; if (ROB[h].Ready) /* Instr completed already */ { RS[r].Vj ← ROB[h].Value; RS[r].Qj ← 0; } else { RS[r].Qj ← h; } /* wait for instruction */ } else { RS[r].Vj ← Regs[rs]; RS[r].Qj ← 0; }; RS[r].Busy ← yes; RS[r].Dest ← b; ROB[b].Instruction ← opcode; ROB[b].Dest ← rd; ROB[b].Ready ← no; </pre>
FP operations and stores	Reservation station (r) and ROB (b) both available	<pre> if (RegisterStat[rt].Busy) /*in-flight instr writes rt*/ { h ← RegisterStat[rt].Reorder; if (ROB[h].Ready) /* Instr completed already */ { RS[r].Vk ← ROB[h].Value; RS[r].Qk ← 0; } else { RS[r].Qk ← h; } /* wait for instruction */ } else { RS[r].Vk ← Regs[rt]; RS[r].Qk ← 0; }; </pre>
FP operations		<pre> RegisterStat[rd].Reorder ← b; RegisterStat[rd].Busy ← yes; ROB[b].Dest ← rd; </pre>
Loads		<pre> RS[r].A ← imm; RegisterStat[rt].Reorder ← b; RegisterStat[rt].Busy ← yes; ROB[b].Dest ← rt; </pre>
Stores		<pre> RS[r].A ← imm; </pre>
Execute FP op	(RS[r].Qj = 0) and (RS[r].Qk = 0)	Compute results—operands are in Vj and Vk
Load step 1	(RS[r].Qj = 0) and there are no stores earlier in the queue	RS[r].A ← RS[r].Vj + RS[r].A;
Load step 2	Load step 1 done and all stores earlier in ROB have different address	Read from Mem[RS[r].A]
Store	(RS[r].Qj = 0) and store at queue head	ROB[h].Address ← RS[r].Vj + RS[r].A;
Write result all but store	Execution done at r and CDB available	<pre> b ← RS[r].Dest; RS[r].Busy ← no; ∀x (if (RS[x].Qj==b) { RS[x].Vj ← result; RS[x].Qj ← 0; }); ∀x (if (RS[x].Qk==b) { RS[x].Vk ← result; RS[x].Qk ← 0; }); ROB[b].Value ← result; ROB[b].Ready ← yes; </pre>
Store	Execution done at r and (RS[r].Qk = 0)	ROB[h].Value ← RS[r].Vk;
Commit	Instruction is at the head of the ROB (entry h) and ROB[h].ready = yes	<pre> d ← ROB[h].Dest; /* register dest, if exists */ if (ROB[h].Instruction==Branch) { if (branch is mispredicted) { clear ROB[h], RegisterStat; fetch branch dest; }; } else if (ROB[h].Instruction==Store) { Mem[ROB[h].Destination] ← ROB[h].Value; } else /* put the result in the register destination */ { Regs[d] ← ROB[h].Value; }; ROB[h].Busy ← no; /* free up ROB entry */ /* free up dest register if no one else writing it */ if (RegisterStat[d].Reorder==h) { RegisterStat[d].Busy ← no; }; </pre>

Reorder Buffer

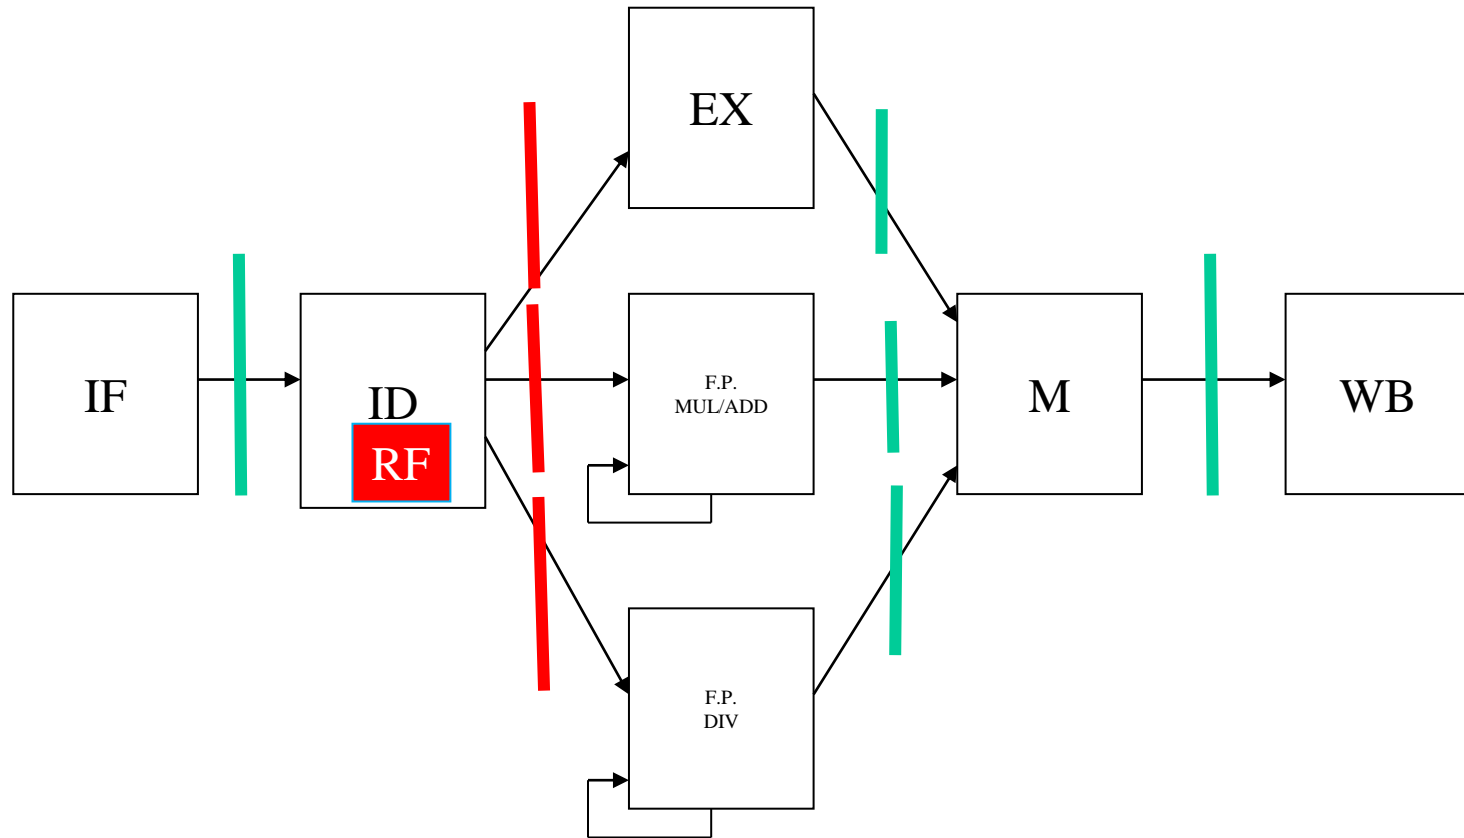
Reorder buffer						
Entry	Busy	Instruction		State	Destination	Value
1	No	f1d	f6,32(x2)	Commit	f6	Mem[32 + Regs[x2]]
2	No	f1d	f2,44(x3)	Commit	f2	Mem[44 + Regs[x3]]
3	Yes	fmul.d	f0,f2,f4	Write result	f0	#2 × Regs[f4]
4	Yes	fsub.d	f8,f2,f6	Write result	f8	#2 − #1
5	Yes	fdiv.d	f0,f0,f6	Execute	f0	
6	Yes	fadd.d	f6,f8,f2	Write result	f6	#4 + #2

Reservation stations								
Name	Busy	Op	Vj	Vk	Qj	Qk	Dest	A
Load1	No							
Load2	No							
Add1	No							
Add2	No							
Add3	No							
Mult1	No	fmul.d	Mem[44 + Regs[x3]]	Regs[f4]			#3	
Mult2	Yes	fdiv.d		Mem[32 + Regs[x2]]	#3		#5	

FP register status										
Field	f0	f1	f2	f3	f4	f5	f6	f7	f8	f10
Reorder #	3						6		4	5
Busy	Yes	No	No	No	No	No	Yes	...	Yes	Yes

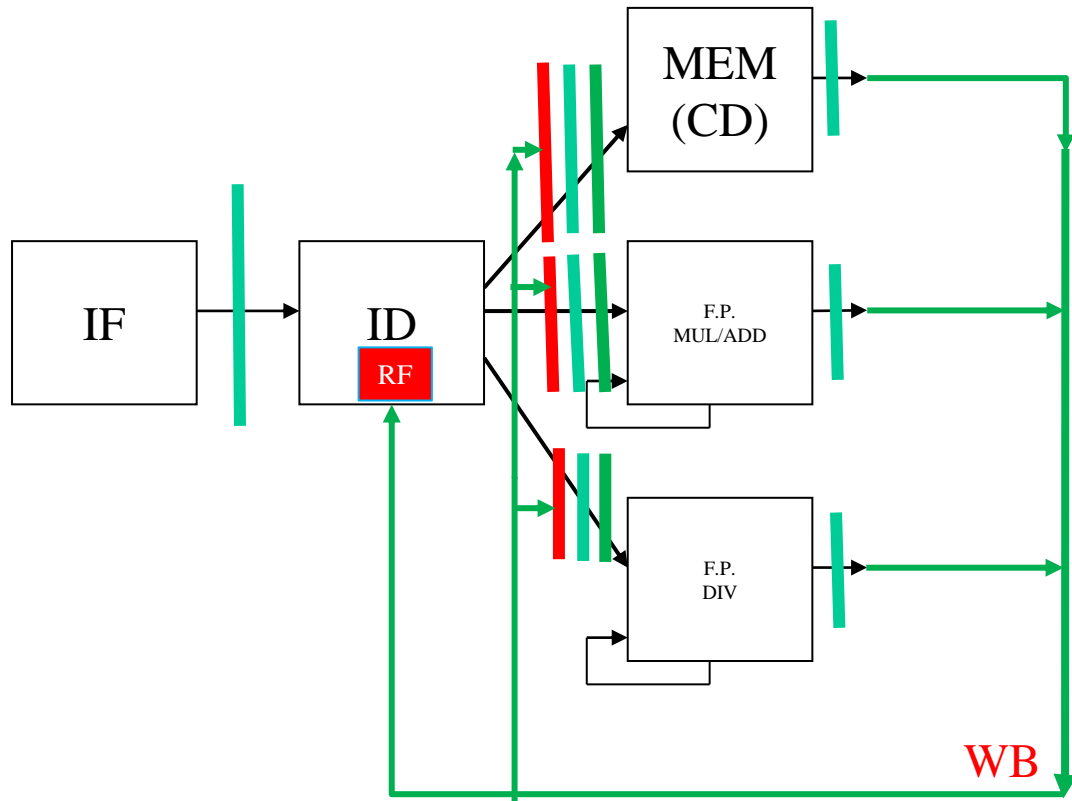
Pipeline bloccante

Stalla con alee RAW, di controllo e strutturali



Pipeline non bloccante senza rob

Stalla con alee di controllo e strutturali



B/F	tag	rd	OP	valore	e	c
B	1	1	+	5	1	1
B	2		BR		0	0
B	3	1	-	-2	1	0
B	4	1	*	1	0	0
B	5	1	*	1	0	0
F						
F						

Reorder Buffer

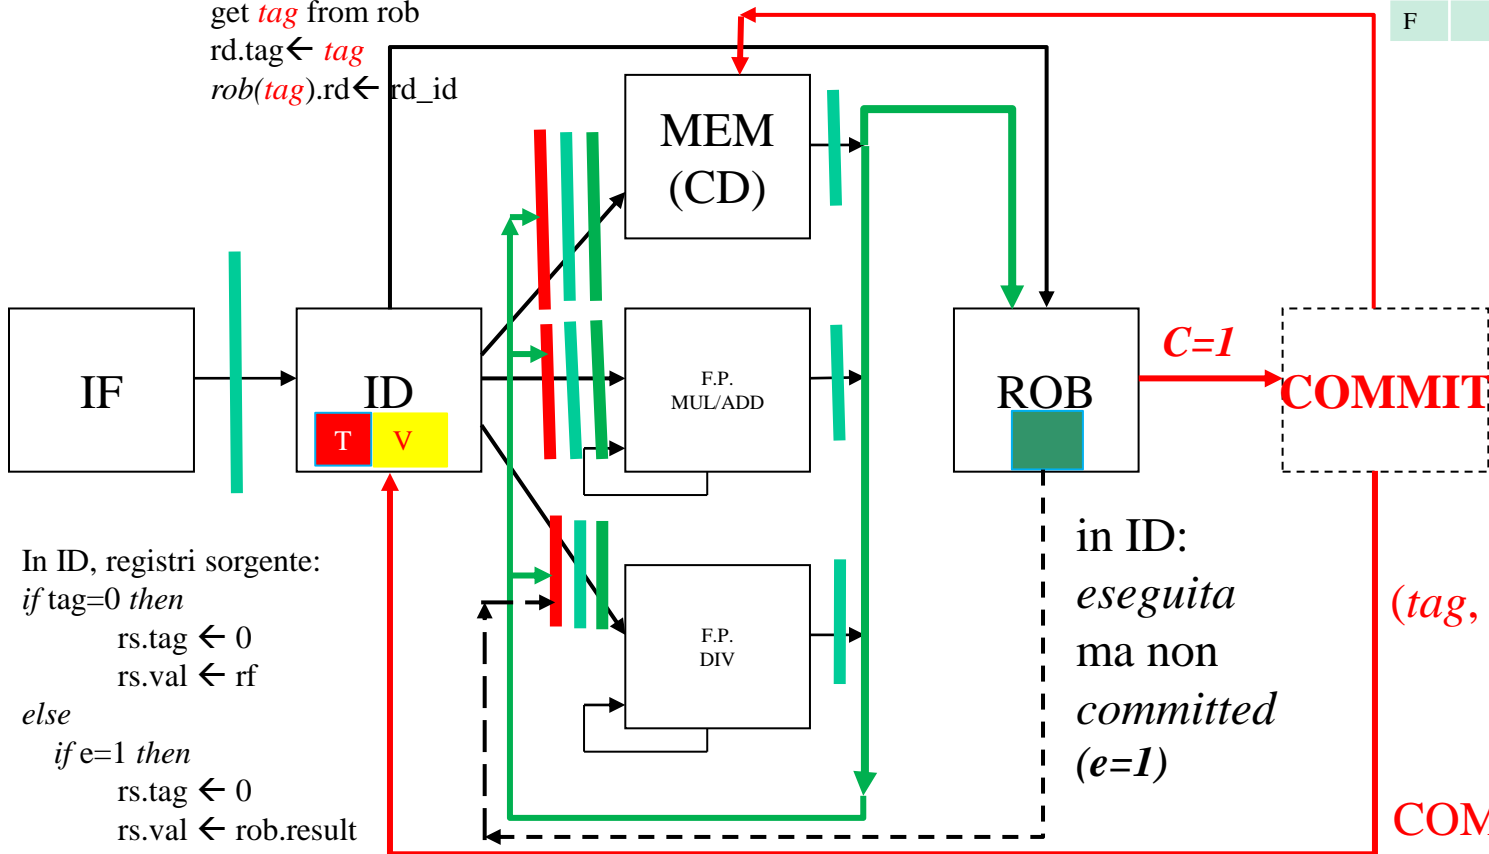
Quando c'è il ROB, si ha che: TAG = (Indice nel ROB)

In ID, registro destinazione

get *tag* from rob

rd.tag ← *tag*

rob(*tag*).rd ← rd_id



COMMIT

Il reorder buffer (ROB)

- Viene caricato in ordine dallo stadio ID contestualmente con il caricamento della Reservation Station
- Nell'intervallo tra il completamento dell'esecuzione e la fase di commit, il ROB e non il register file mette a disposizione i risultati
- Pertanto:
- Il TAG: (Unit_ID, RS_ID) diventa: TAG: (Indice nel ROB)
- Il risultato di una istruzione diventa “impegnativo”, cioè passa ad aggiornare lo stato del “sistema” se:
 - L'istruzione è stata eseguita
 - I risultati di tutte le istruzioni precedenti sono “impegnativi”

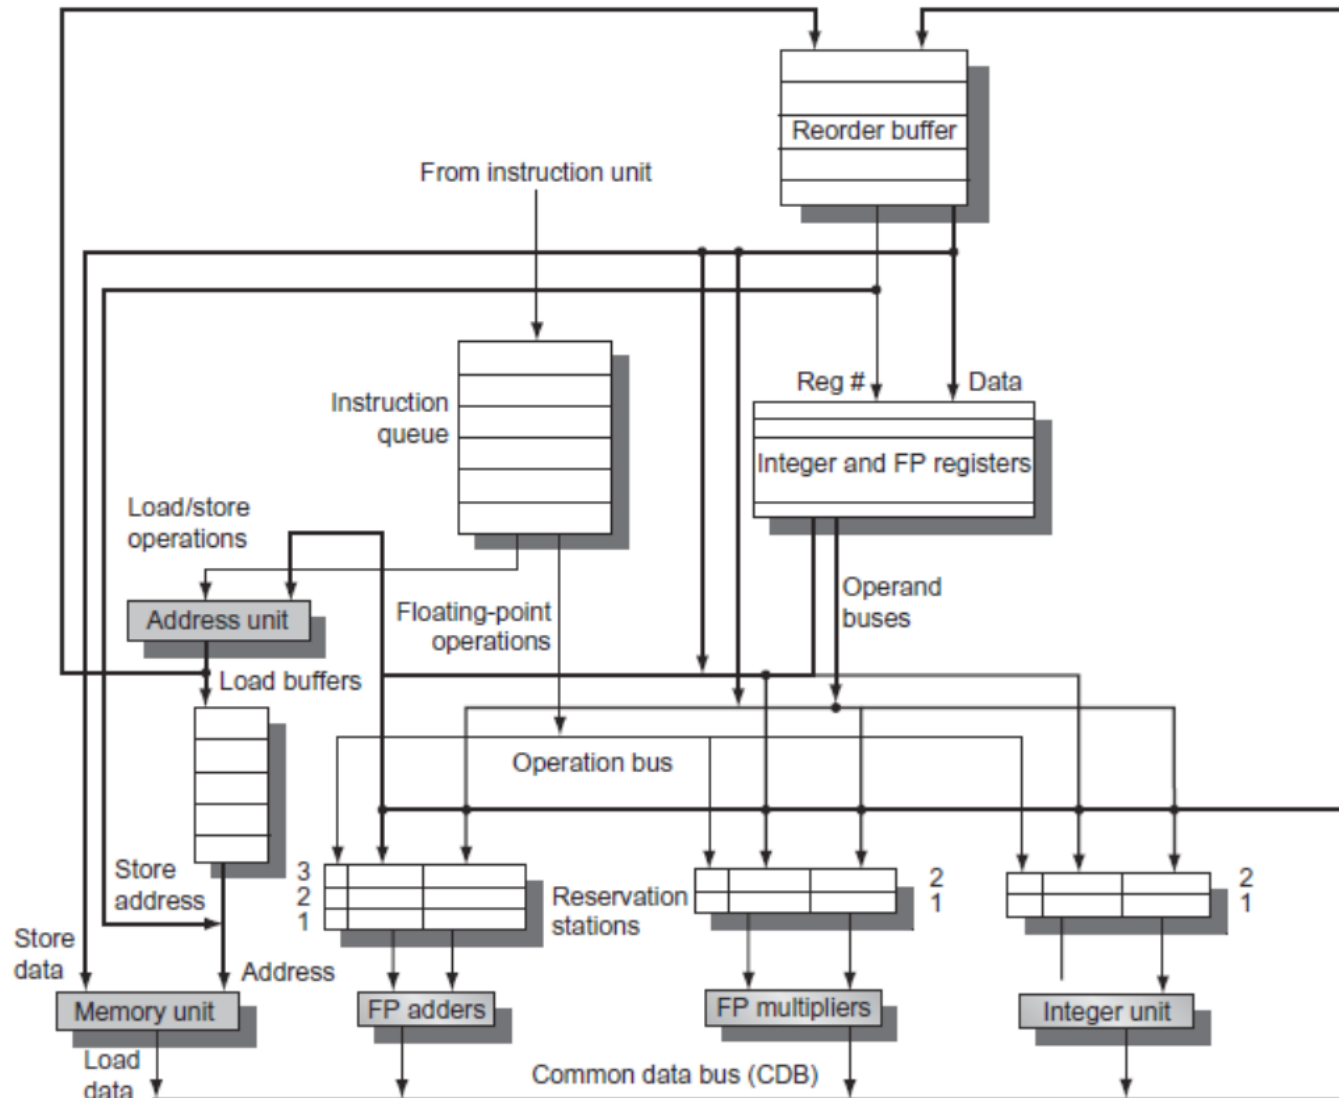
Multiple Issue and Static Scheduling

- To achieve $CPI < 1$, need to complete multiple instructions per clock
- Solutions:
 - Statically scheduled superscalar processors
 - Dynamically scheduled superscalar processors

Dynamic Scheduling, Multiple Issue, and Speculation

- Modern microarchitectures:
 - Dynamic scheduling + multiple issue + speculation
- Two approaches:
 - Assign reservation stations and update pipeline control table in half clock cycles
 - Only supports 2 instructions/clock
 - Design logic to handle any possible dependencies between the instructions
- Issue logic is the bottleneck in dynamically scheduled superscalars

Overview of Design



Multiple Issue

- Examine all the dependencies among the instructions in the bundle
- If dependencies exist in bundle, encode them in reservation stations
- Also need multiple completion/commit
- To simplify RS allocation:
 - Limit the number of instructions of a given class that can be issued in a “bundle”, i.e. on FP, one integer, one load, one store

Example

Loop: ld x2,0(x1)	//x2=array element
addi x2,x2,1	//increment x2
sd x2,0(x1)	//store result
addi x1,x1,8	//increment pointer
bne x2,x3,Loop	//branch if not last

Example (No Speculation)

Iteration number	Instructions	Issues at clock cycle number	Executes at clock cycle number	Memory access at clock cycle number	Write CDB at clock cycle number	Comment
1	ld x2,0(x1)	1	2	3	4	First issue
1	addi x2,x2,1	1	5		6	Wait for ld
1	sd x2,0(x1)	2	3	7		Wait for addi
1	addi x1,x1,8	2	3		4	Execute directly
1	bne x2,x3,Loop	3	7			Wait for addi
2	ld x2,0(x1)	4	8	9	10	Wait for bne
2	addi x2,x2,1	4	11		12	Wait for ld
2	sd x2,0(x1)	5	9	13		Wait for addi
2	addi x1,x1,8	5	8		9	Wait for bne
2	bne x2,x3,Loop	6	13			Wait for addi
3	ld x2,0(x1)	7	14	15	16	Wait for bne
3	addi x2,x2,1	7	17		18	Wait for ld
3	sd x2,0(x1)	8	15	19		Wait for addi
3	addi x1,x1,8	8	14		15	Wait for bne
3	bne x2,x3,Loop	9	19			Wait for addi

Example (Multiple Issue with Speculation)

Iteration number	Instructions	Issues at clock number	Executes at clock number	Read access at clock number	Write CDB at clock number	Commits at clock number	Comment
1	ld x2,0(x1)	1	2	3	4	5	First issue
1	addi x2,x2,1	1	5		6	7	Wait for ld
1	sd x2,0(x1)	2	3			7	Wait for addi
1	addi x1,x1,8	2	3		4	8	Commit in order
1	bne x2,x3,Loop	3	7			8	Wait for addi
2	ld x2,0(x1)	4	5	6	7	9	No execute delay
2	addi x2,x2,1	4	8		9	10	Wait for ld
2	sd x2,0(x1)	5	6			10	Wait for addi
2	addi x1,x1,8	5	6		7	11	Commit in order
2	bne x2,x3,Loop	6	10			11	Wait for addi
3	ld x2,0(x1)	7	8	9	10	12	Earliest possible
3	addi x2,x2,1	7	11		12	13	Wait for ld
3	sd x2,0(x1)	8	9			13	Wait for addi
3	addi x1,x1,8	8	9		10	14	Executes earlier
3	bne x2,x3,Loop	9	13			14	Wait for addi