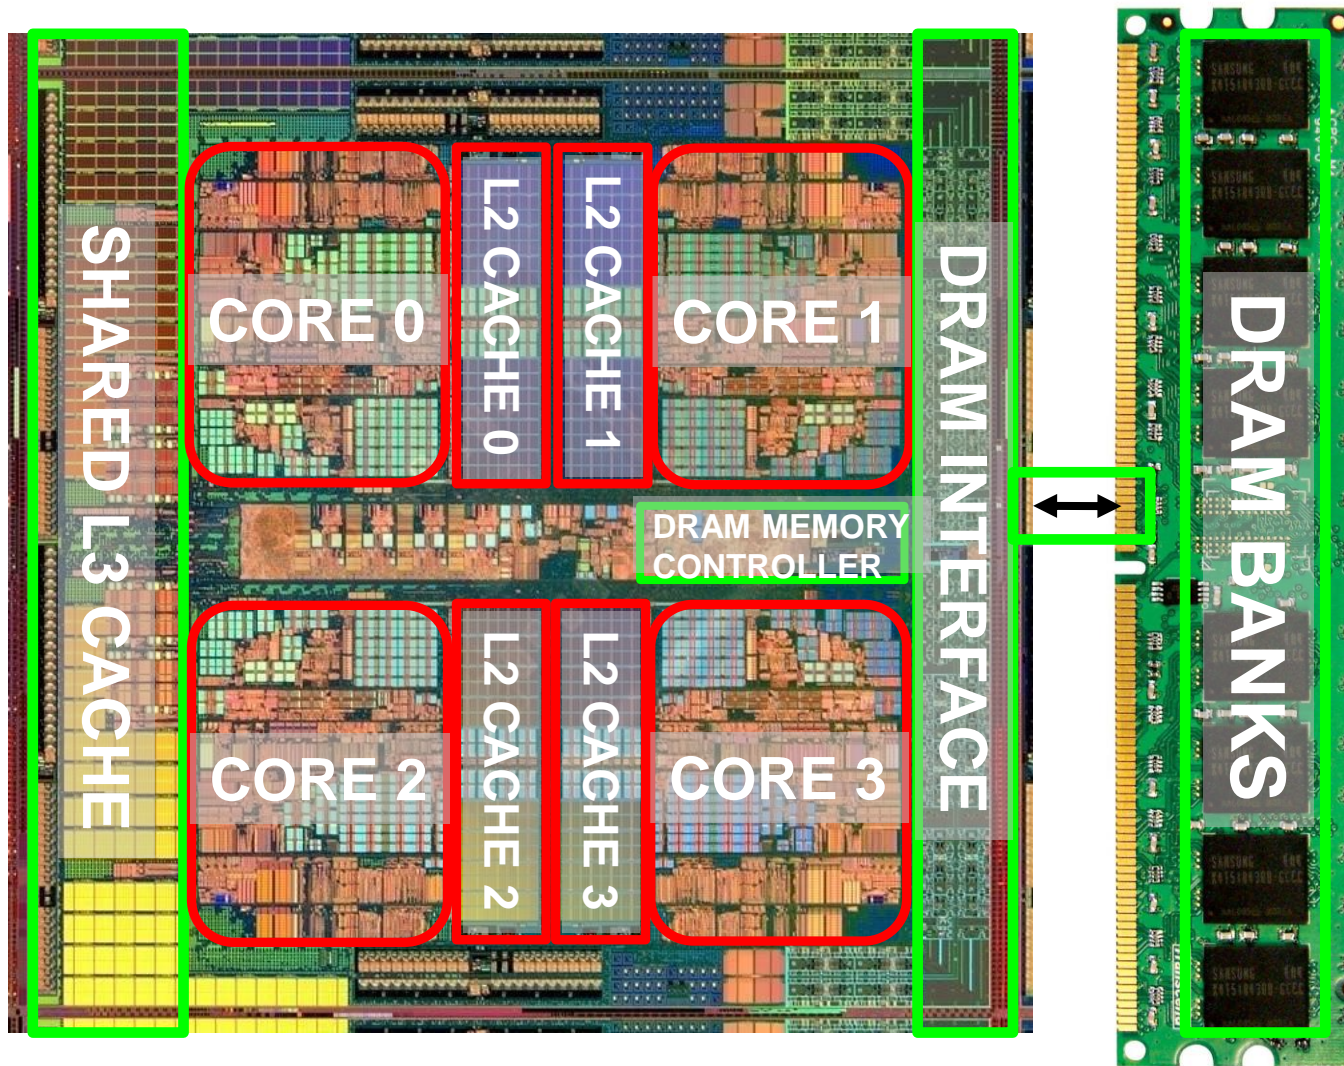


# Gerarchia di Memoria e Cache

Andrea Bartolini – [a.bartolini@unibo.it](mailto:a.bartolini@unibo.it)

# Multi-Core Issues in Caching

# Caches in a Multi-Core System





# Caches in Multi-Core Systems

---

- L'efficienza della cache diventa ancora più importante in un sistema multi-core/multithread
  - Memory bandwidth è prioritaria
  - Lo spazio di cache è una risorsa limitata tra core/thread
- Come progettiamo le cache in un sistema multi-core?

# Private vs. Shared Caches

- **Private** cache: la cache appartiene ad un core (un blocco condiviso può essere in più cache)
- **Shared** cache: La cache è condivisa da più core



# Shared Caches Between Cores

---

## ■ Advantages:

- ❑ High effective capacity
- ❑ **Dynamic partitioning** of available cache space
  - No fragmentation due to static partitioning
  - If one core does not utilize some space, another core can
- ❑ **Easier to maintain coherence (a cache block is in a single location)**

## ■ Disadvantages

- ❑ Slower access (cache not tightly coupled with the core)
- ❑ Cores incur **conflict misses due to other cores' accesses**
  - Misses due to inter-core interference
  - Some cores can destroy the hit rate of other cores
- ❑ Guaranteeing a minimum level of service (or fairness) to each core is harder (how much space, how much bandwidth?)

# Caching in Multiprocessors

---

- Caching not only complicates ordering of **all operations**...
  - A memory location can be present in multiple caches
  - Prevents the effect of a store or load to be seen by other processors → **makes it difficult for all processors to see the same global order of (all) *memory operations***
- ... but it also complicates ordering of **operations on a single memory location**
  - A single memory location can be present in multiple caches
  - **Makes it difficult for processors that have cached the same location to have the correct value of that location (in the presence of updates to that location)**

# Memory Consistency vs. Cache Coherence

---

- **Consistency** is about ordering of **all memory operations** from different processors (i.e., to different memory locations)
  - **Global ordering** of accesses to *all* memory *locations*
- **Coherence** is about ordering of **operations** from different processors **to the same memory location**
  - **Local ordering** of accesses to *each* cache *block*




# Cache Coherence

# Shared Memory Model

---

- Many parallel programs communicate through *shared memory*
- Proc 0 writes to an address, followed by Proc 1 reading
  - This implies communication between the two




- Each read should receive the value last written by anyone
  - This requires synchronization (what does last written mean?)
- What if Mem[A] is cached (at either end)?

# Cache Coherence

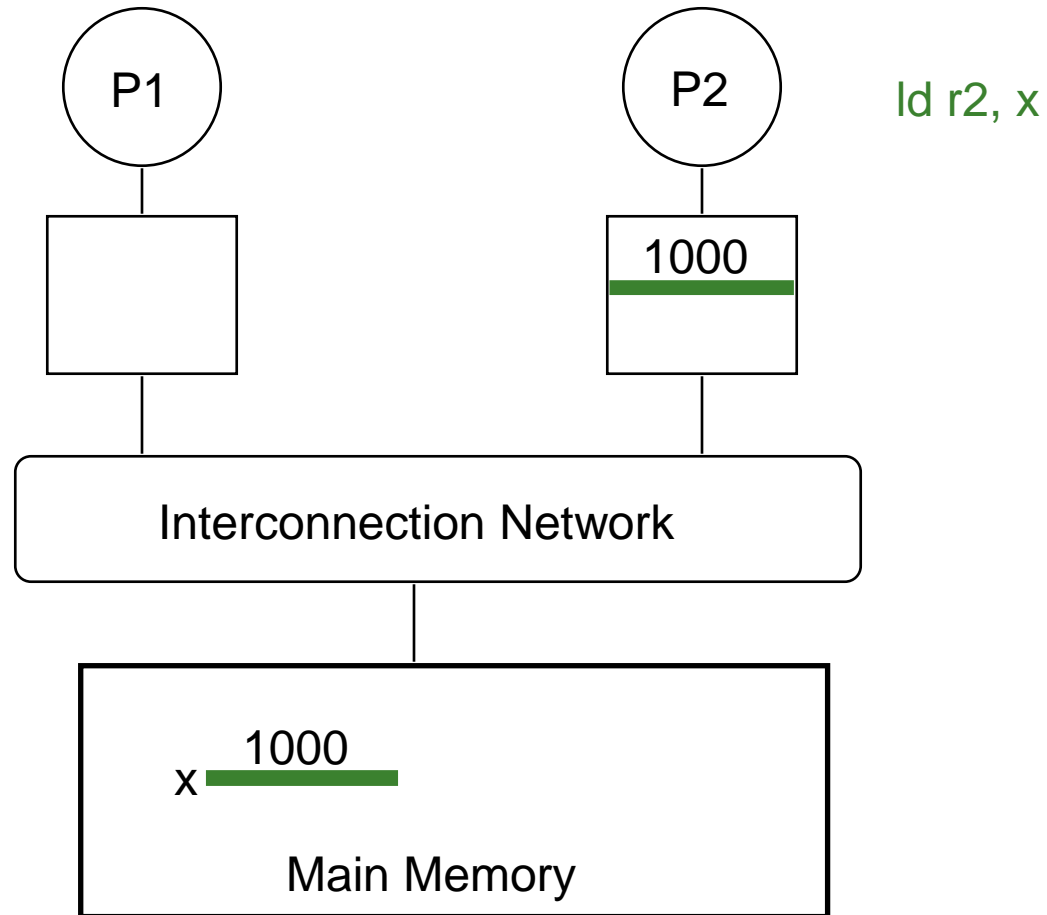
---

- Se più processori memorizzano nella cache lo stesso blocco, come fanno a garantire che tutti vedano uno stato coerente?



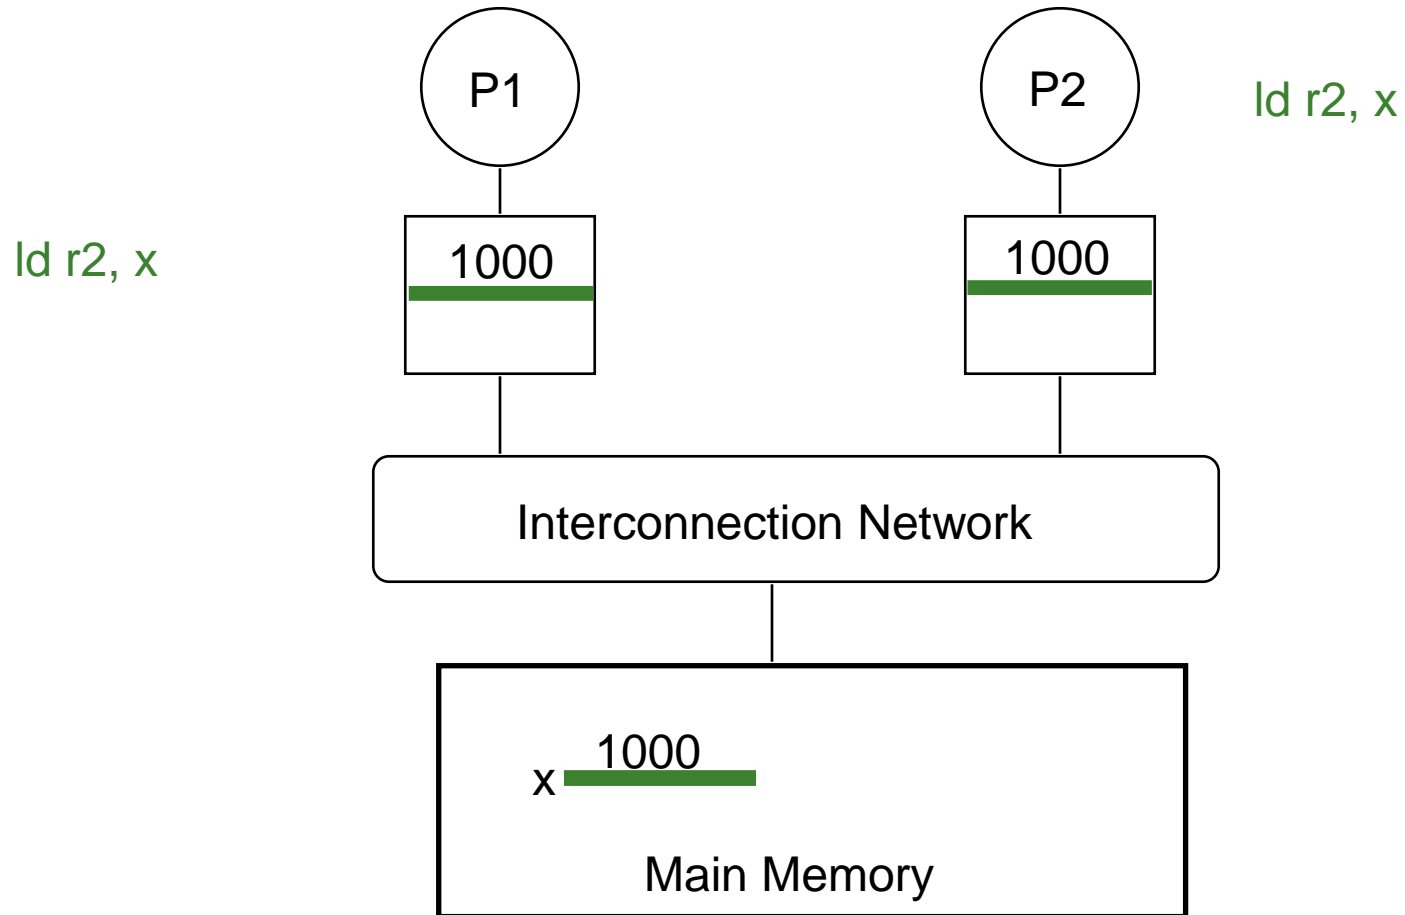
# The Cache Coherence Problem

---



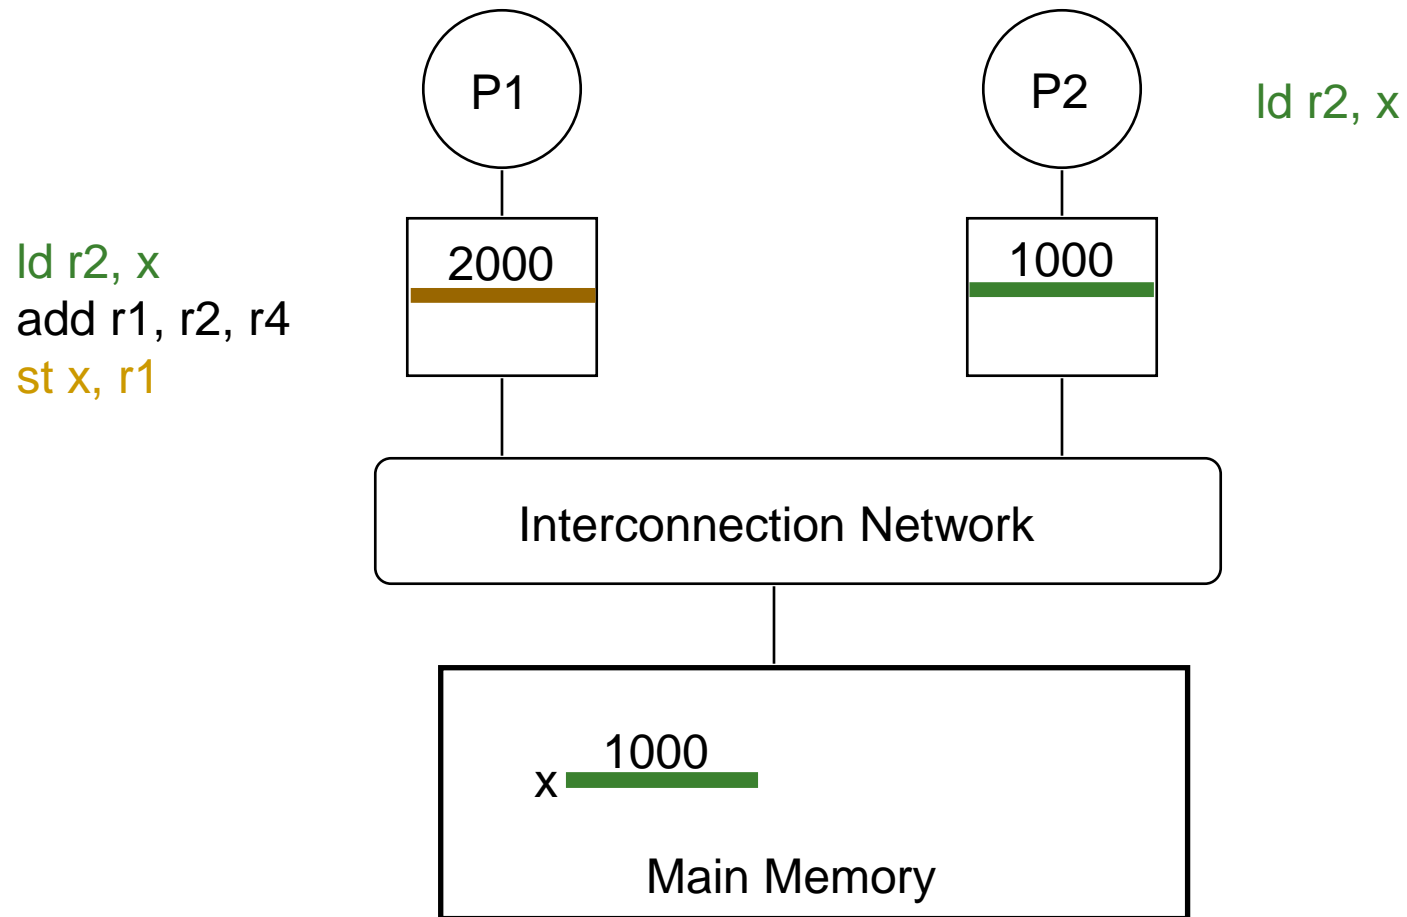
# The Cache Coherence Problem

---



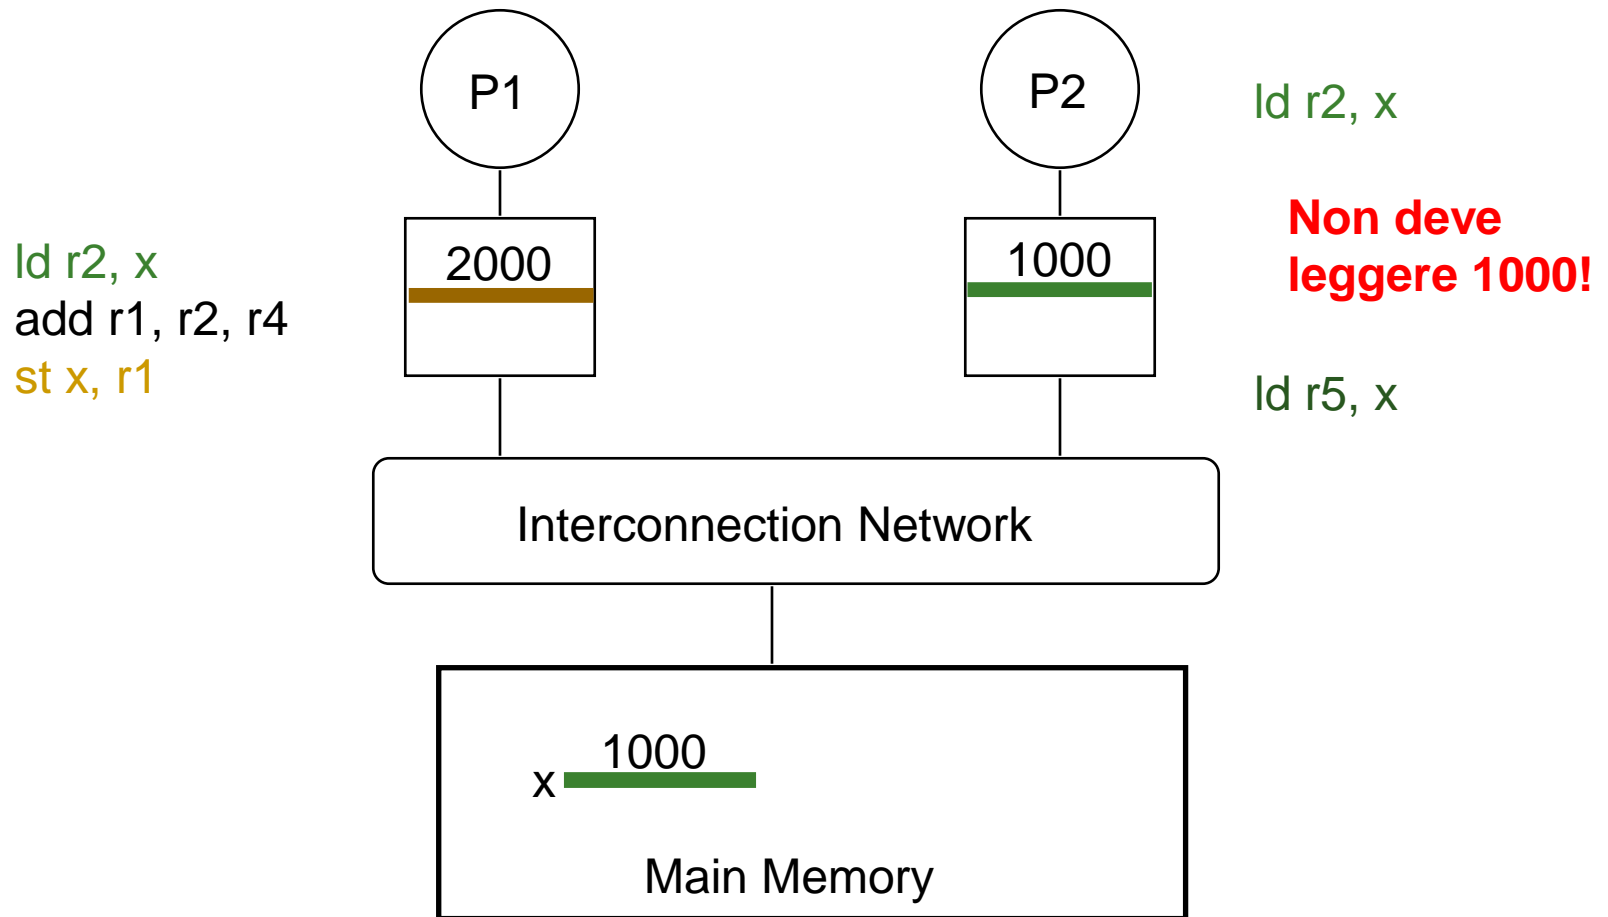
# The Cache Coherence Problem

---



# The Cache Coherence Problem

---



# Cache Coherence: Whose Responsibility?

---

## ■ Software

- ❑ Can programmer ensure coherence if caches invisible to software?
- ❑ **Coarse-grained:** Page-level coherence has overheads
- ❑ Non-solution: Make shared locks/data non-cacheable
- ❑ **A combination of non-cacheable and coarse-grained is doable**
- ❑ **Fine-grained:** What if the ISA provided a cache flush instruction?
  - FLUSH-LOCAL A: Flushes/invalidates the cache block containing address A from a processor's local cache.
  - FLUSH-GLOBAL A: Flushes/invalidates the cache block containing address A from all other processors' caches.
  - FLUSH-CACHE X: Flushes/invalidates all blocks in cache X.

## ■ Hardware

- ❑ **Greatly simplifies software's job**
- ❑ One idea: Invalidate all other copies of block A when a core writes to A



# (Non-)Solutions to Cache Coherence

---

## ■ No hardware based coherence

- Keeping caches coherent is software's responsibility
- + Makes microarchitect's life easier
- Makes average programmer's life much harder
  - need to worry about hardware caches to maintain program correctness?
- Overhead in ensuring coherence in software (e.g., page protection, page-based software coherence, non-cacheable)

## ■ All caches are shared between all processors

- + No need for coherence
- Shared cache becomes the bandwidth bottleneck
- Very hard to design a scalable system with low-latency cache access this way

# Maintaining Coherence


---

- Need to guarantee that all processors see a consistent value (i.e., consistent updates) for the same memory location
- Writes to location A by P0 should be seen by P1 (eventually), and all writes to A should appear in some order
- Coherence needs to provide:
  - **Write propagation:** guarantee that updates will propagate
  - **Write serialization:** provide a consistent order seen by all processors for the same memory location
- Need a global point of serialization for this store ordering

# Cache Coherence

---


- Se più processori memorizzano nella cache lo stesso blocco, come fanno a garantire che tutti vedano uno stato coerente?



# Cache Coherence

---

- Se più processori memorizzano nella cache lo stesso blocco, come fanno a garantire che tutti vedano uno stato coerente?



# Hardware Cache Coherence

---

- Basic idea:
  - A processor/cache broadcasts its write/update to a memory location to all other processors
  - Another cache that has the location either updates or invalidates its local copy

# Coherence: Update vs. Invalidate

---

- How can we *safely update replicated data*?
  - Option 1 (Update protocol): push an update to all copies
  - Option 2 (Invalidate protocol): ensure there is only one copy (local), update it
- **On a Read:**
  - If local copy is Invalid, put out request
  - (If another node has a copy, it returns it, otherwise memory does)

# Coherence: Update vs. Invalidate (II)

---

## ■ **On a Write:**

- ❑ Read block into cache as before

## **Update Protocol:**

- ❑ Write to block, and simultaneously broadcast written data and address to sharers
- ❑ (Other nodes update the data in their caches if block is present)

## **Invalidate Protocol:**

- ❑ Write to block, and simultaneously broadcast invalidation of address to sharers
- ❑ (Other nodes invalidate block in their caches if block is present)

# Update vs. Invalidate Tradeoffs

---

- Which do we want?
  - Write frequency and sharing behavior are critical
- **Update**
  - + If sharer set is constant and updates are infrequent, avoids the cost of invalidate-reacquire (broadcast update pattern)
  - If data is rewritten without intervening reads by other cores, updates would be useless
  - Write-through cache policy → bus becomes bottleneck
- **Invalidate**
  - + After invalidation broadcast, core has exclusive access rights
  - + Only cores that keep reading after each write retain a copy
  - If write contention is high, leads to ping-ponging (rapid invalidation-reacquire traffic from different processors)



# Two Cache Coherence Methods

---

- ❑ How do we ensure that the proper caches are updated?
- ❑ **Snoopy Bus** [Goodman ISCA 1983, Papamarcos+ ISCA 1984]
  - Bus-based, *single point of serialization for all memory requests*
  - Processors observe other processors' actions
    - ❑ E.g.: P1 makes “read-exclusive” request for A on bus, P0 sees this and invalidates its own copy of A
- ❑ **Directory** [Censier and Feautrier, IEEE ToC 1978]
  - *Single point of serialization per block*, distributed among nodes
  - Processors make explicit requests for blocks
  - Directory tracks which caches have each block
  - Directory coordinates invalidation and updates
    - ❑ E.g.: P1 asks directory for exclusive copy, directory asks P0 to invalidate, waits for ACK, then responds to P1

# Snoopy Cache Coherence

---

- Idea:
  - ❑ All caches “snoop” all other caches’ read/write requests and keep the cache block coherent
  - ❑ Each cache block has “coherence metadata” associated with it in the tag store of each cache
- Easy to implement if all caches share a common bus
  - ❑ Each cache broadcasts its read/write operations on the bus
  - ❑ Good for small-scale multiprocessors
  - ❑ What if you would like to have a 10,000-node multiprocessor?

# Directory Based Coherence


---

- Idea: A logically-central directory keeps track of where the copies of each cache block reside. Caches consult this directory to ensure coherence.
  
- An example mechanism:
  - For each cache block in memory, store  $P+1$  bits in directory
    - One bit for each cache, indicating whether the block is in cache
    - Exclusive bit: indicates that a cache has the only copy of the block and can update it without notifying others
  - On a read: set the cache's bit and arrange the supply of data
  - On a write: invalidate all caches that have the block and reset their bits
  - Have an "exclusive bit" associated with each block in each cache (so that the cache can update the exclusive block silently)

# A Very Simple Coherence Scheme (VI)




---

- Caches “snoop” (observe) each other’s write/read operations. If a processor writes to a block, all others invalidate the block.
- A simple protocol:





- Write-through, no-write-allocate cache
- Actions of the local processor on the cache block: PrRd, PrWr,
- Actions that are broadcast on the bus for the block: BusRd, BusWr


# (VI) - Esempio




# (VI) - Esempio




P1



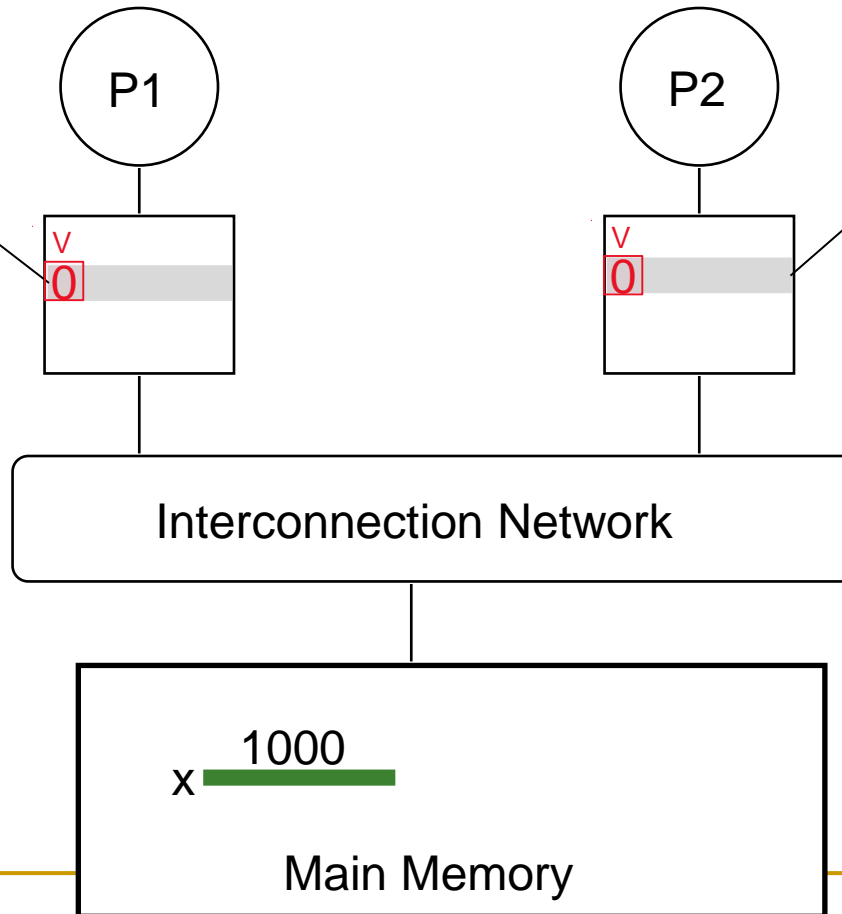


P2






Interconnection Network



# (VI) - Esempio



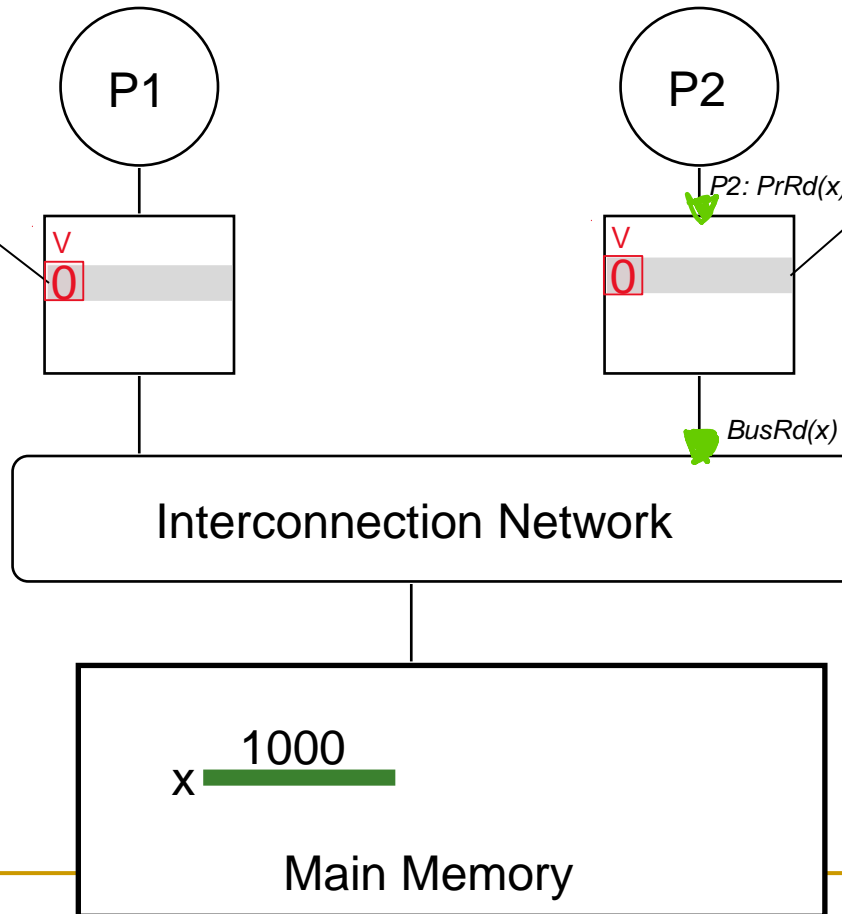
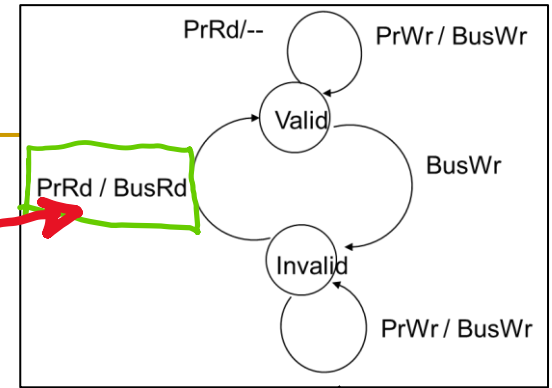

# (VI) - Esempio



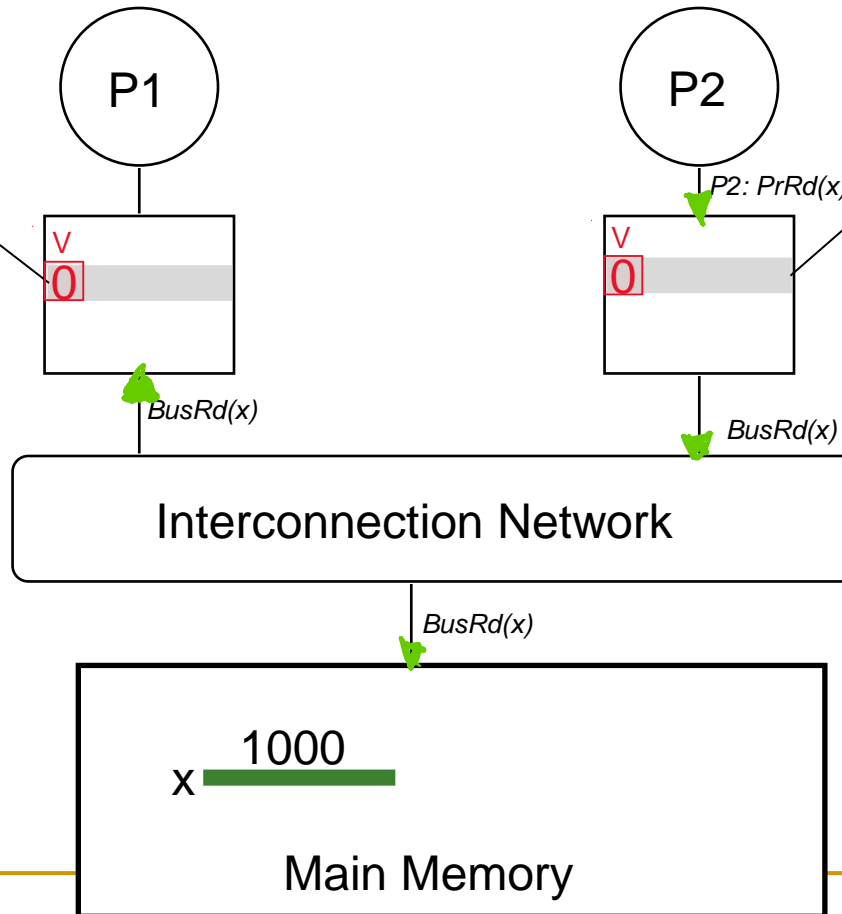


ld r2, x



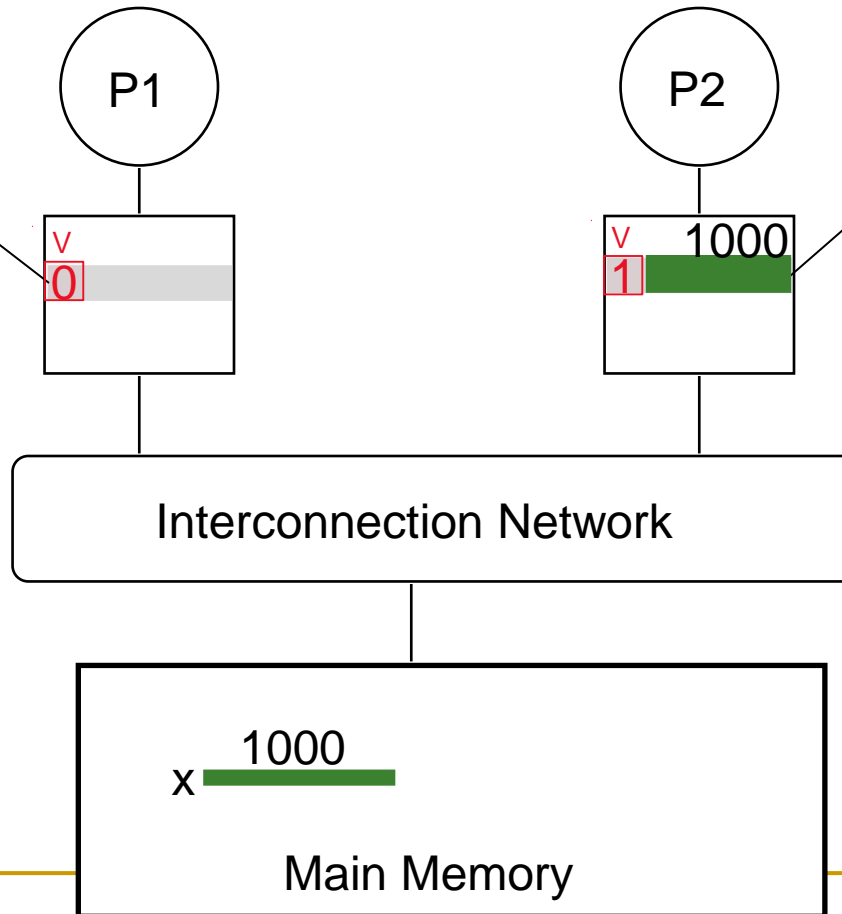


# (VI) - Esempio



# (VI) - Esempio





# (VI) - Esempio




ld r2, x

# (VI) - Esempio




ld r2, x



P1





P2





ld r2, x






# (VI) - Esempio




ld r2, x





BusRd(x)




ld r2, x





# (VI) - Esempio




ld r2, x





BusRd(x)





ld r2, x





# (VI) - Esempio



ld r2, x





BusRd(x)





BusRd(x)

ld r2, x




# (VI) - Esempio




ld r2, x

P1




P2





ld r2, x

Interconnection Network








# (VI) - Esempio





ld r2, x  
add r1, r2, r4  
st x, r1





ld r2, x





# (VI) - Esempio





ld r2, x  
add r1, r2, r4  
st x, r1



ld r2, x




# (VI) - Esempio




ld r2, x  
add r1, r2, r4  
st x, r1

P1




P2





ld r2, x

Interconnection Network




# (VI) - Esempio




ld r2, x  
add r1, r2, r4  
st x, r1

P1




P2





ld r2, x  
ld r5, x

Interconnection Network




# (VI) - Esempio




ld r2, x  
add r1, r2, r4  
st x, r1

P1



BusRd(x)

P2




BusRd(x)



ld r2, x  
ld r5, x



BusRd(x)




# (VI) - Esempio




ld r2, x  
add r1, r2, r4  
st x, r1

P1



BusRd(x)

P2




BusRd(x)



ld r2, x  
ld r5, x



BusRd(x)




# (VI) - Esempio




ld r2, x  
add r1, r2, r4  
st x, r1

P1




P2



ld r2, x  
ld r5, x

Interconnection Network



# Extending the Protocol

---

- What if you want write-back caches?
  - We want a “modified” state

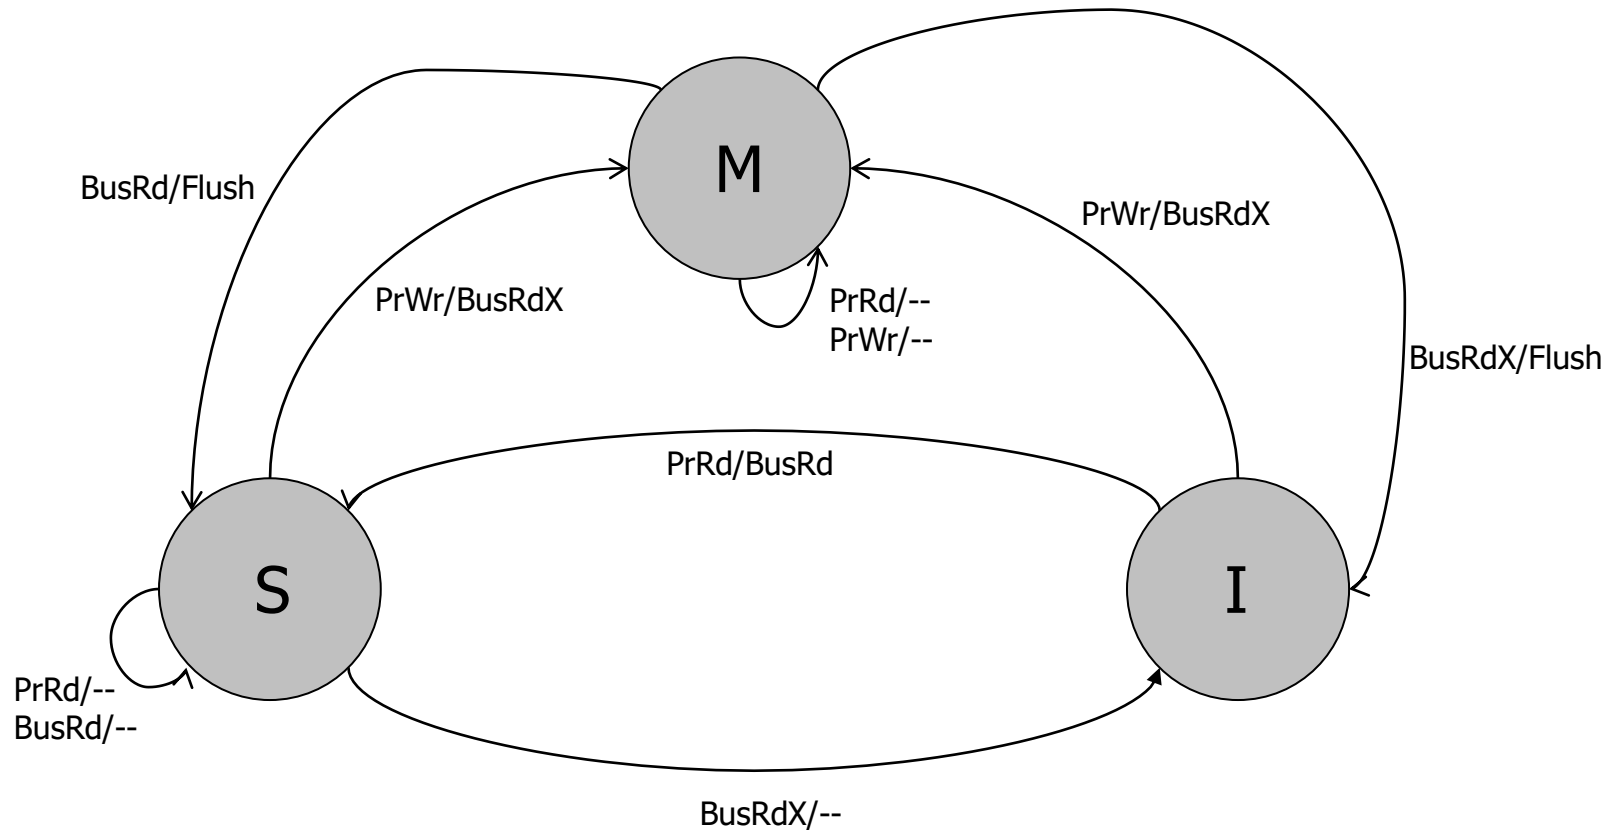


# A More Sophisticated Protocol: MSI

---

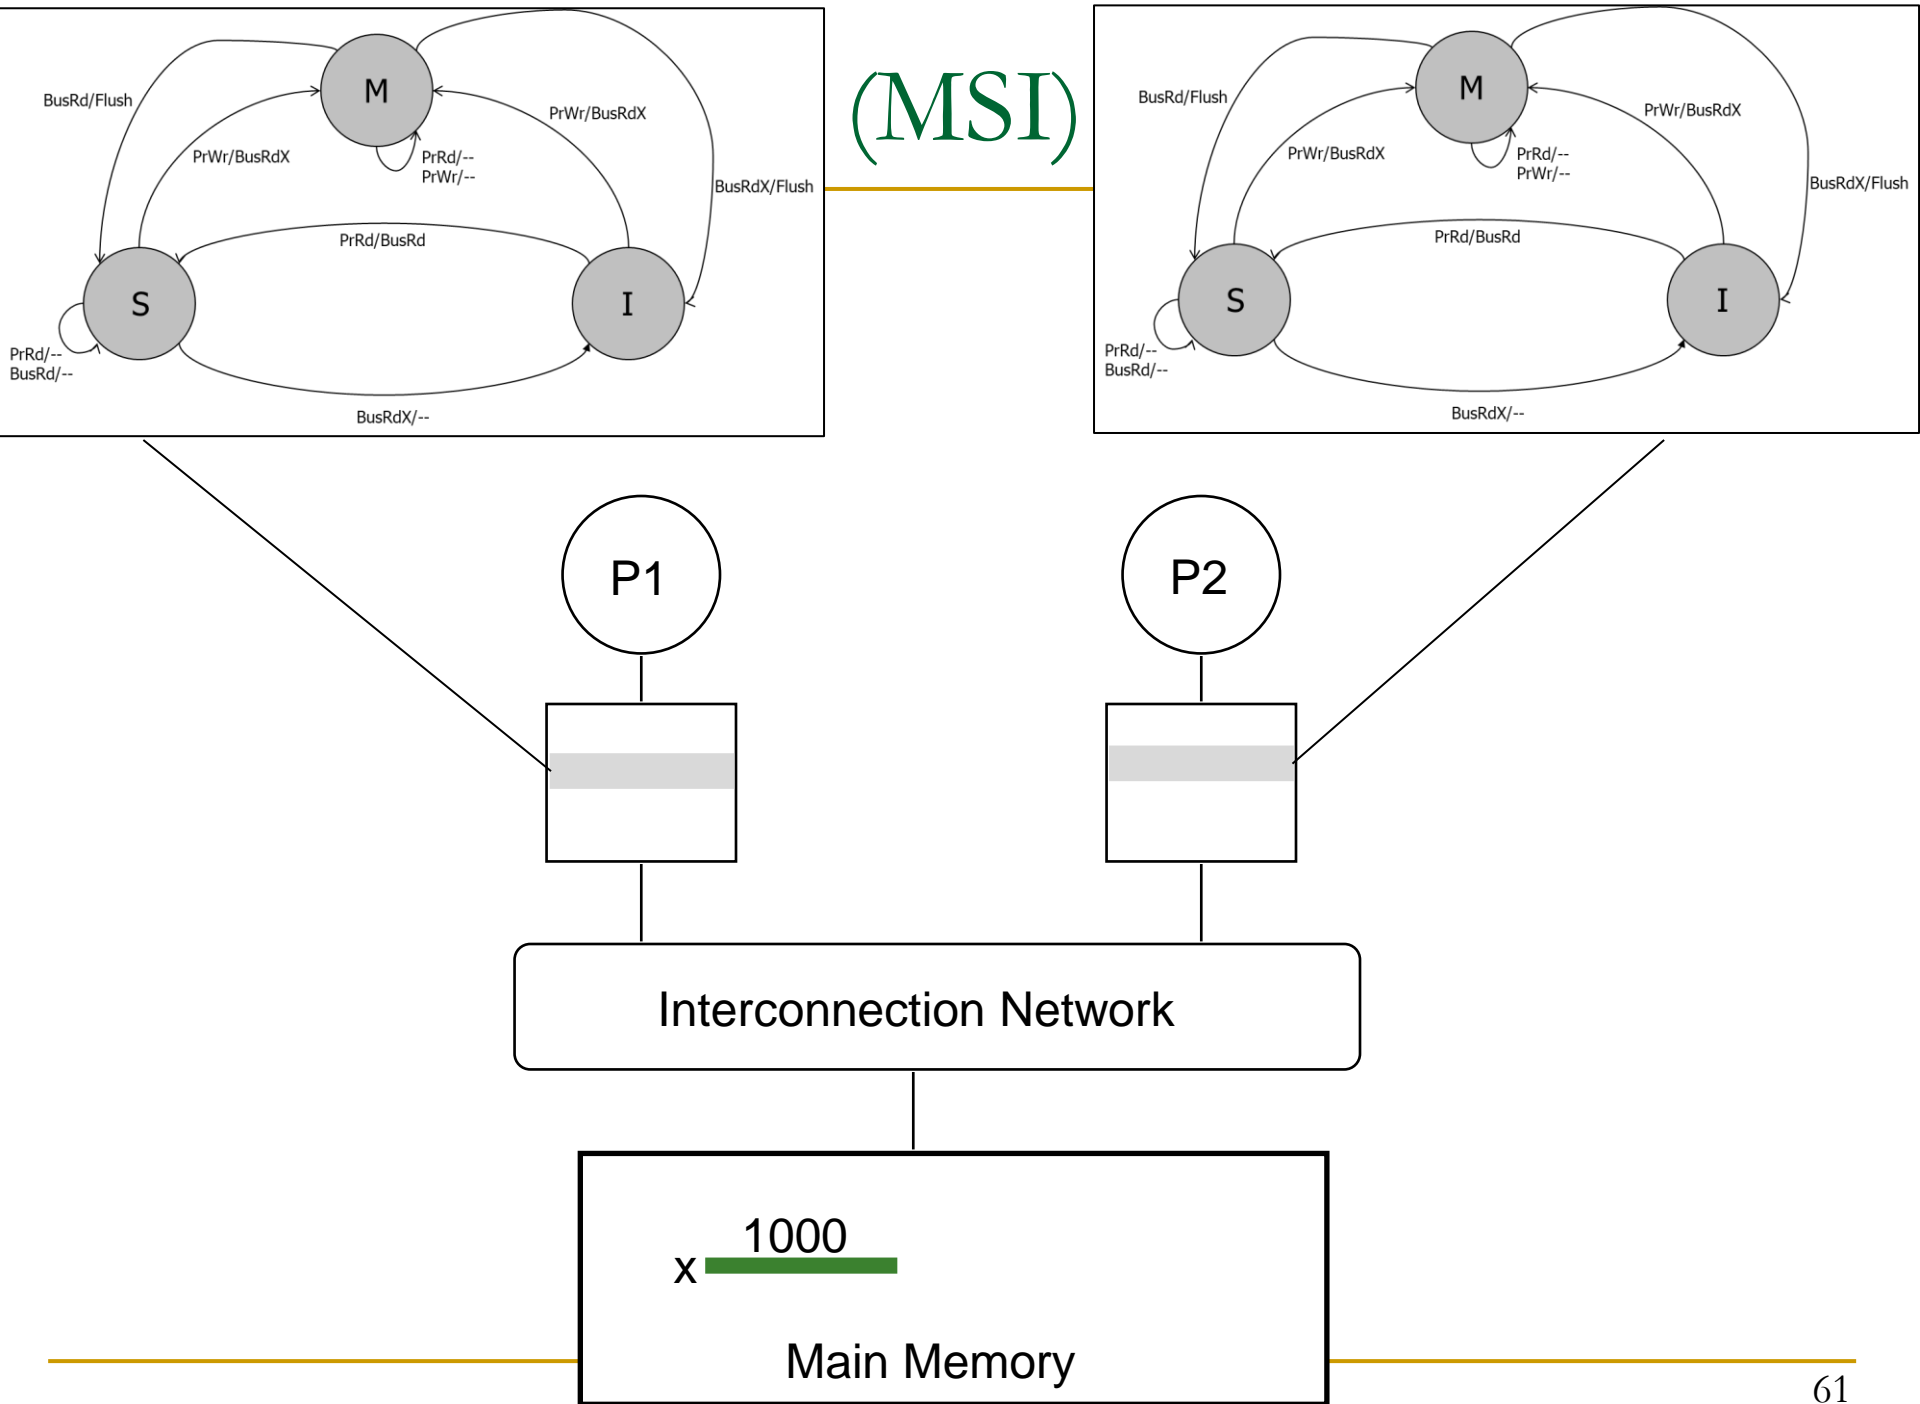
- Extend metadata per block to encode three states:
  - **M**(odified): cache line is the only cached copy and is dirty
  - **S**(hared): cache line is one of potentially several cached copies and it is clean (i.e., at least one clean cached copy)
  - **I**(nvalid): cache line is not present in this cache
- Read miss makes a *Read* request on bus, transitions to **S**
- Write miss makes a *ReadEx* request, transitions to **M** state
- When a processor snoops *ReadEx* from another writer, it must invalidate its own copy (if any)
- $S \rightarrow M$  *upgrade* can be made without re-reading data from memory (via *Invalidations*)

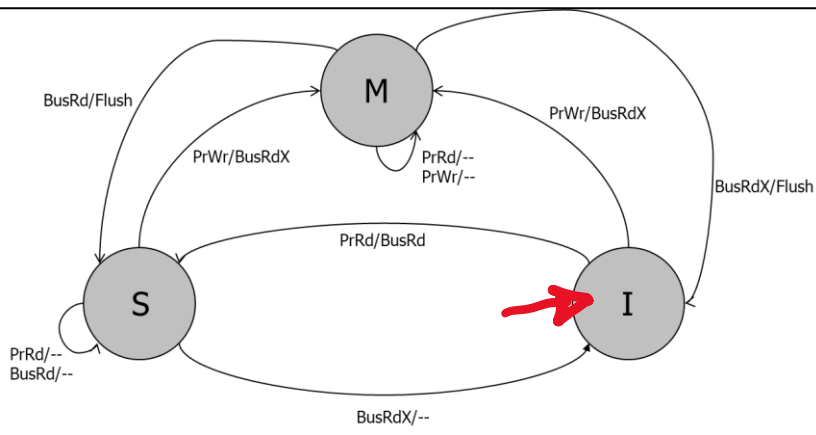
# MSI State Machine




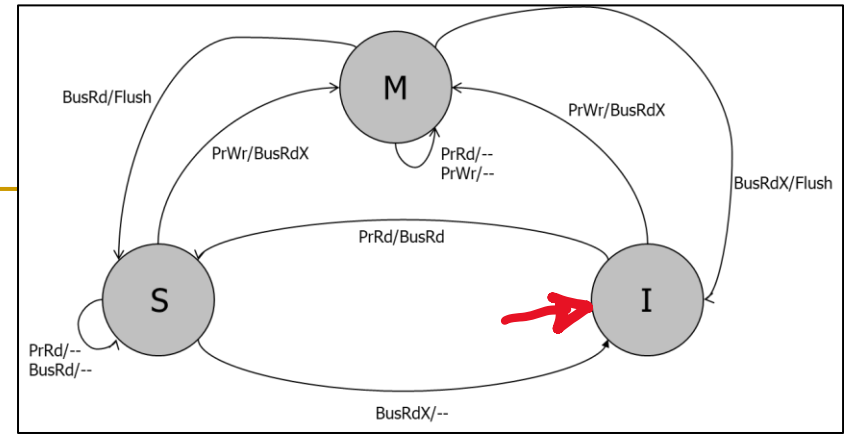
ObservedEvent/Action


[Culler/Singh96]







(MSI)






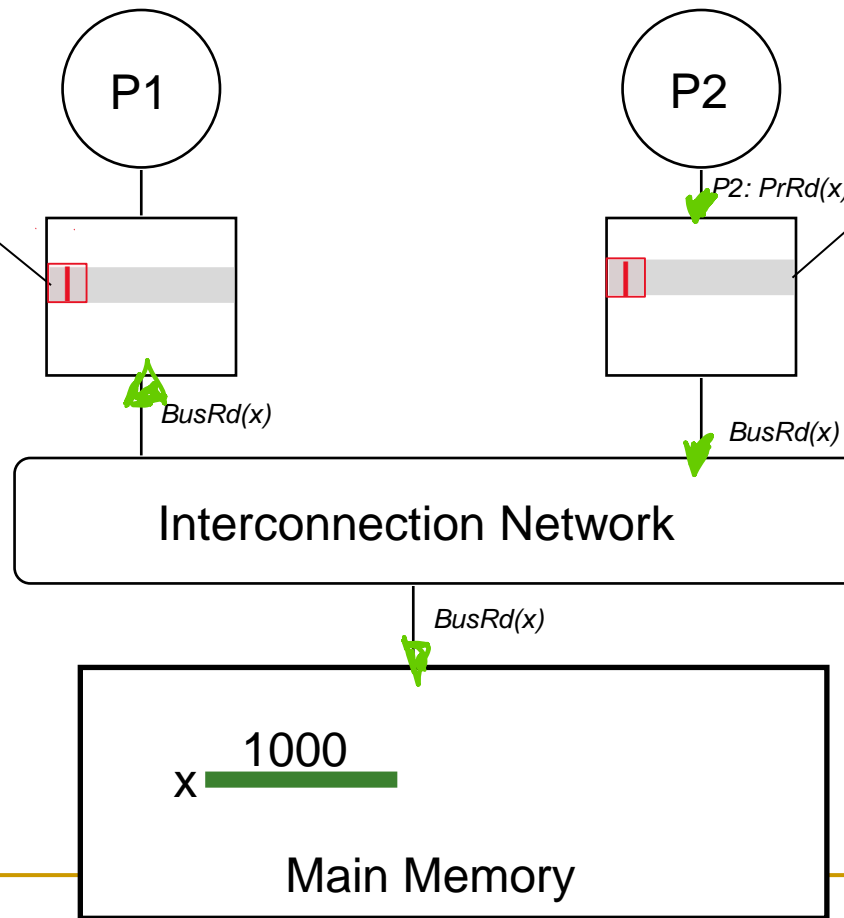

(MSI)



ld r2, x

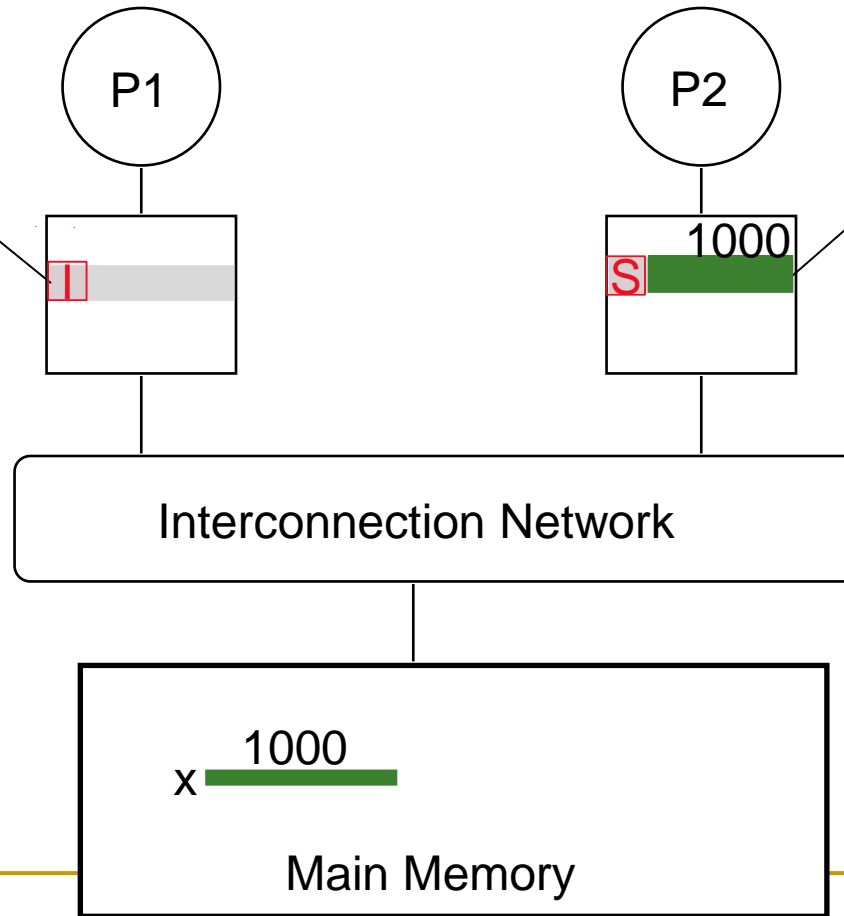
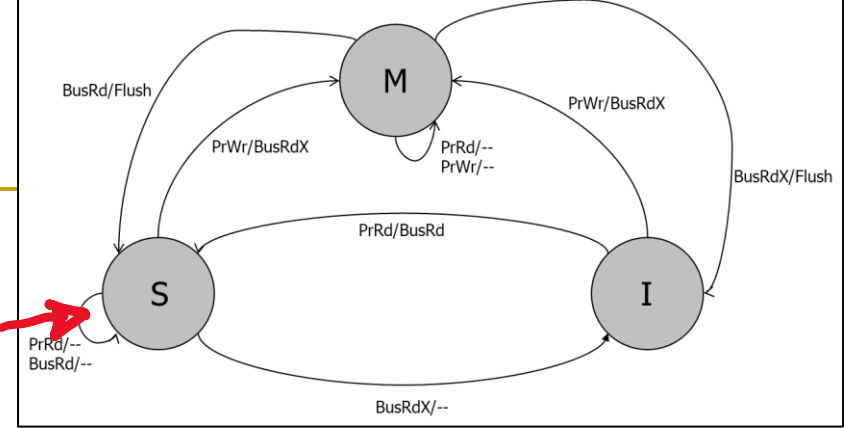
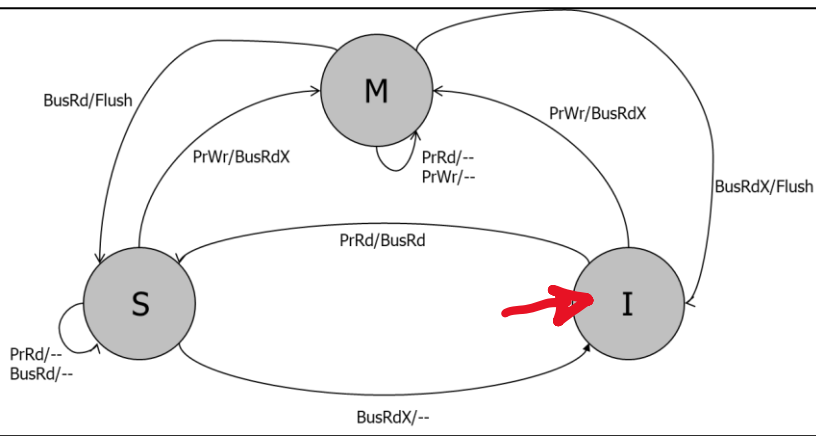


(MSI)



ld r2, x

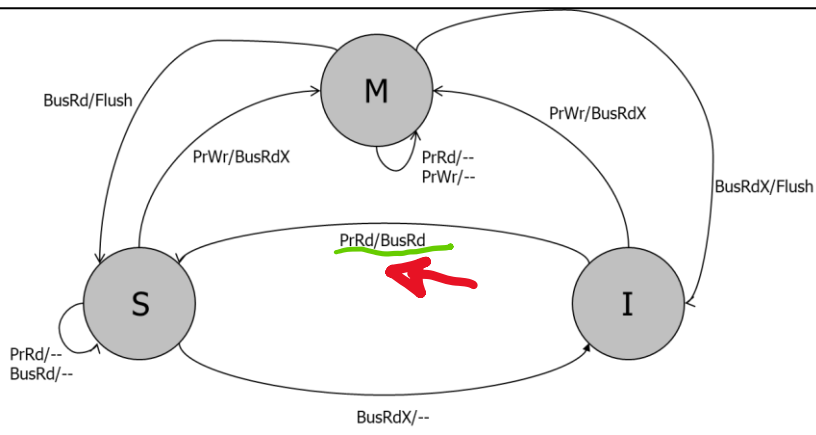
(MSI)



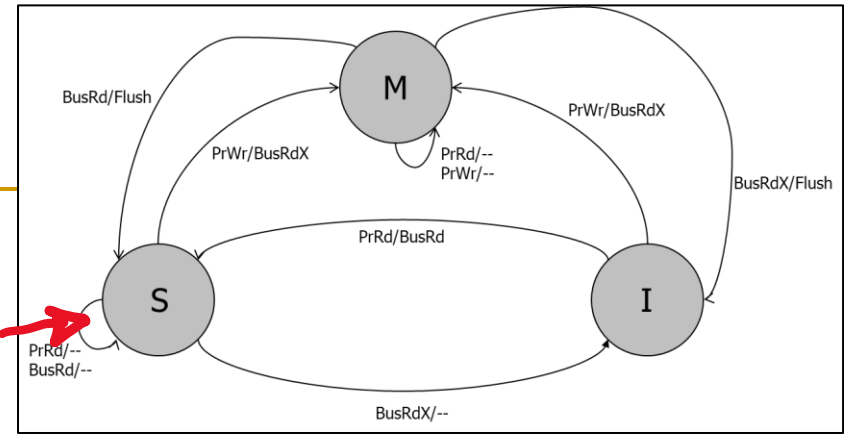
ld r2, x








(MSI)




ld r2, x




BusRd(x)



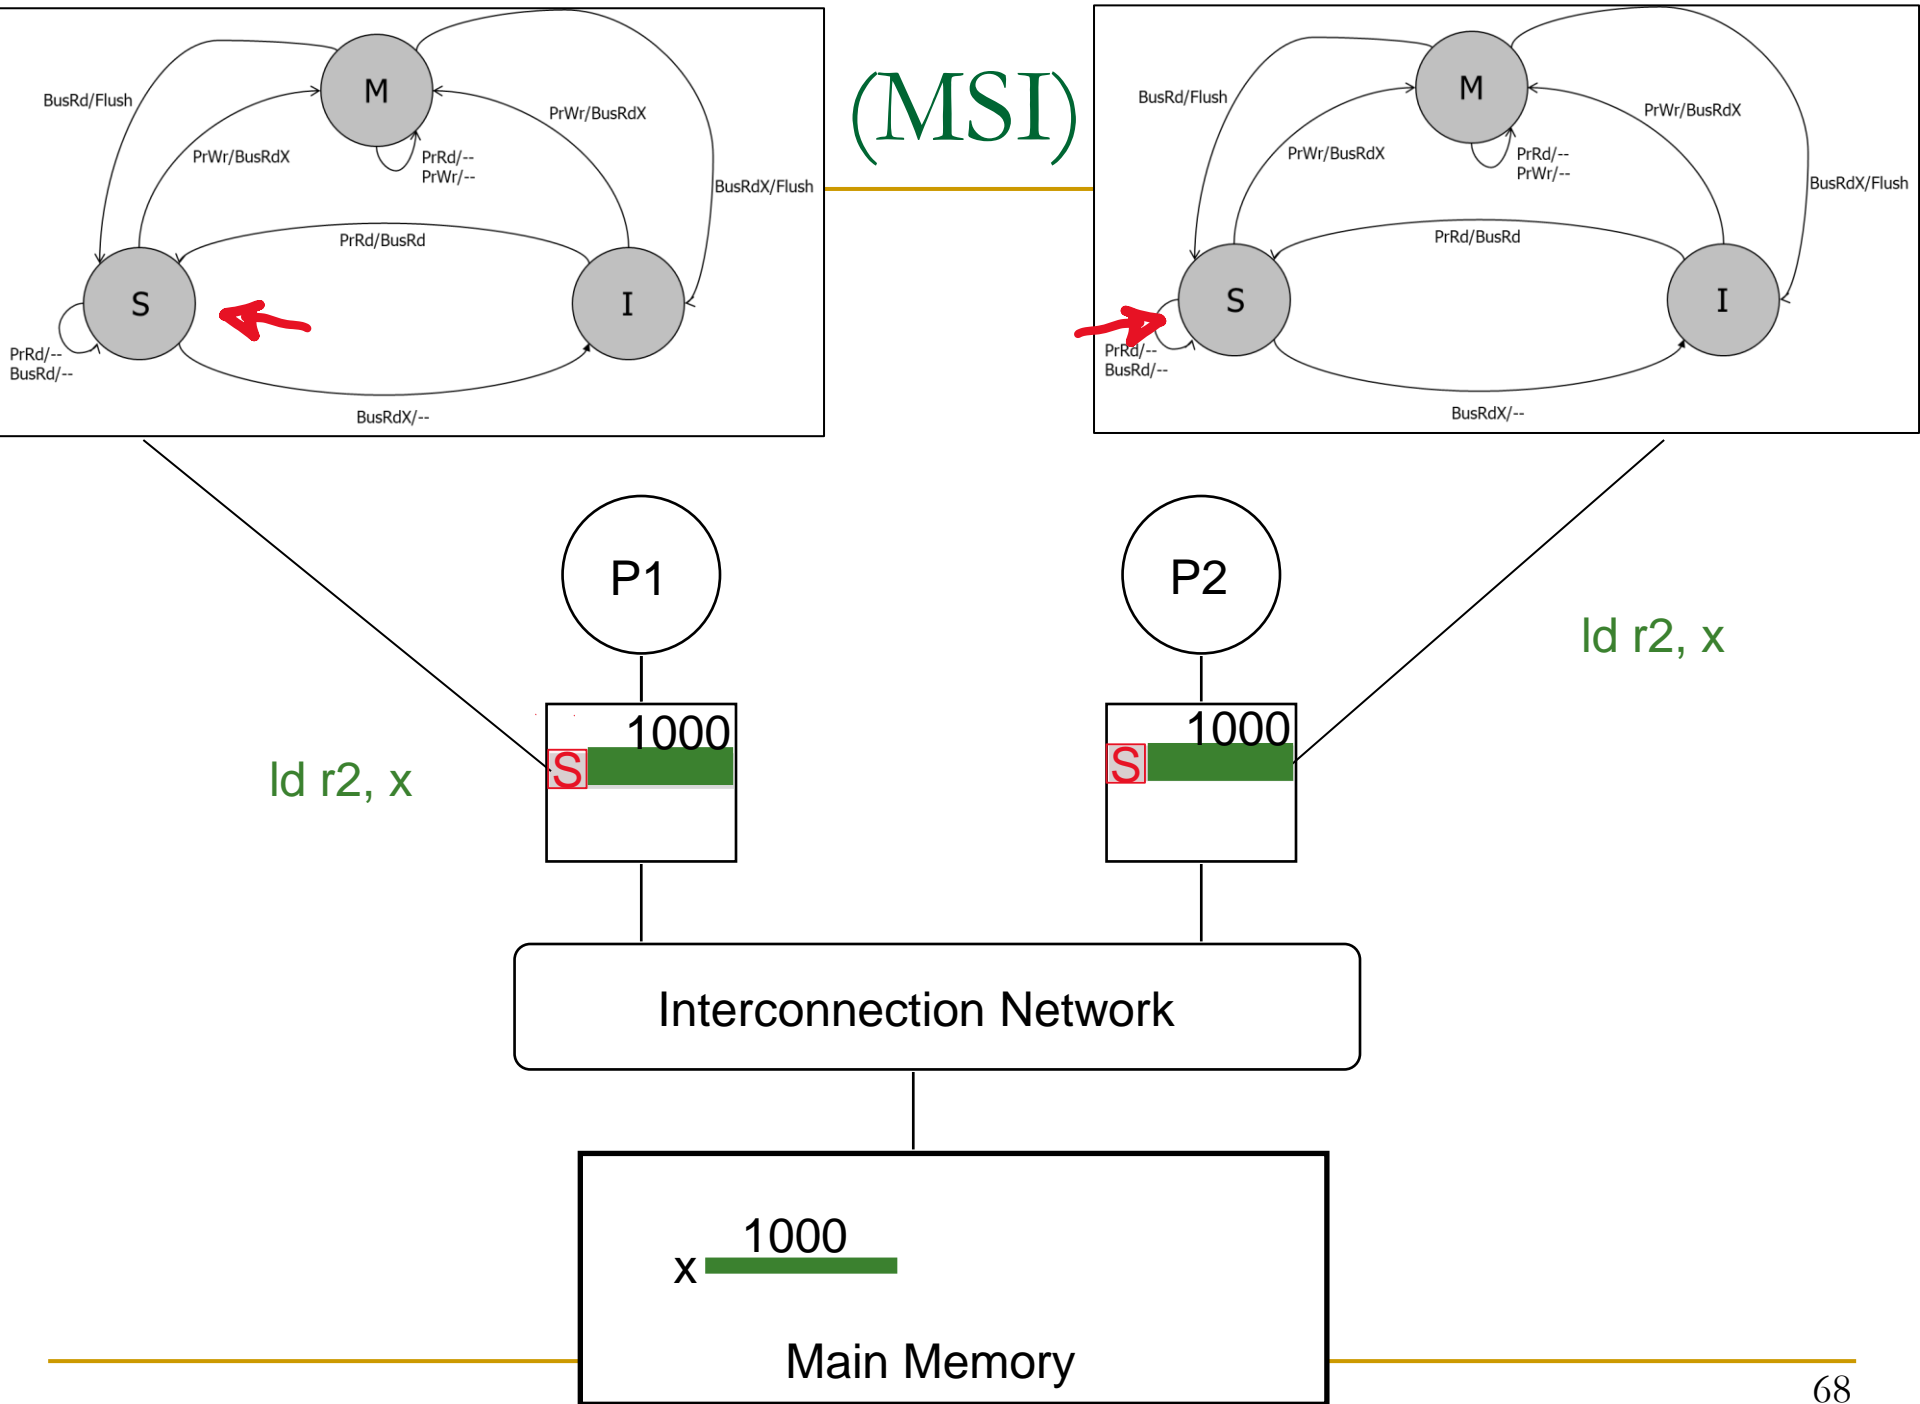
BusRd(x)




ld r2, x

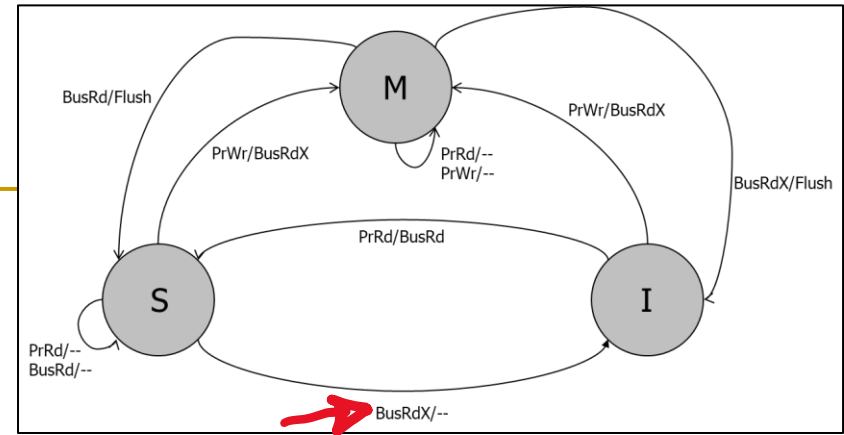


BusRd(x)







(MSI)



ld r2, x  
add r1, r2, r4  
st x, r1




BusRdX(x)





BusRdX(x)

ld r2, x

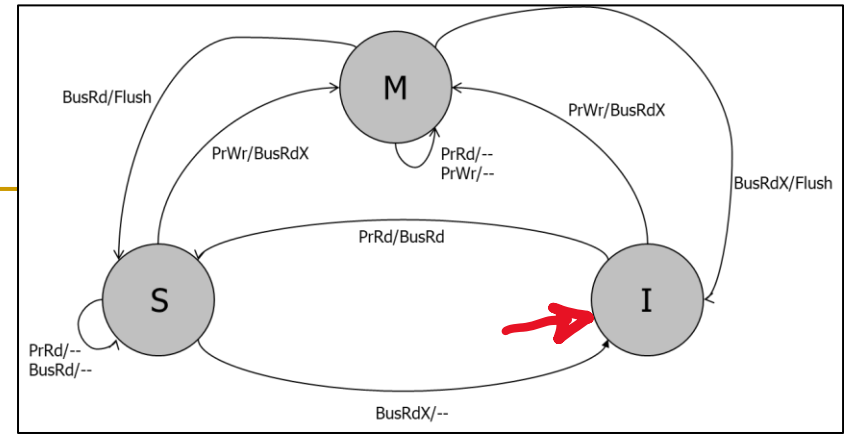


BusRdX(x)







(MSI)



ld r2, x  
add r1, r2, r4  
st x, r1




BusRdX(x)

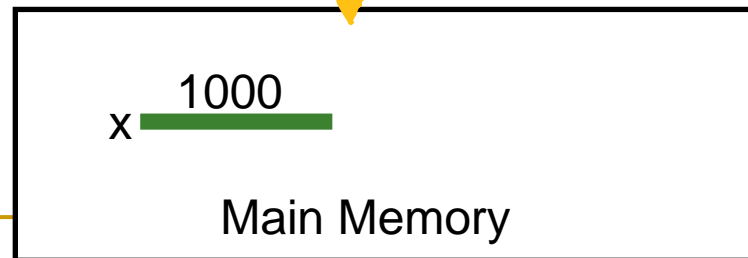


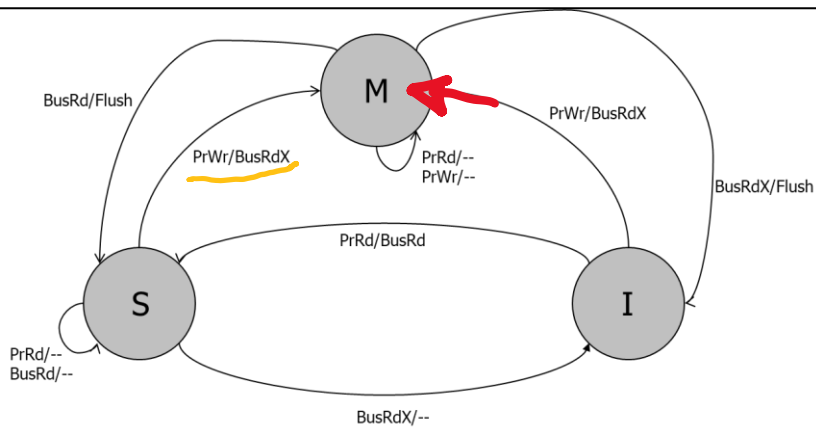
BusRdX(x)

ld r2, x




BusRdX(x)







(MSI)





ld r2, x  
add r1, r2, r4  
st x, r1




P1: PrWr(x)




ld r2, x






(MSI)




ld r2, x  
add r1, r2, r4  
st x, r1

P1




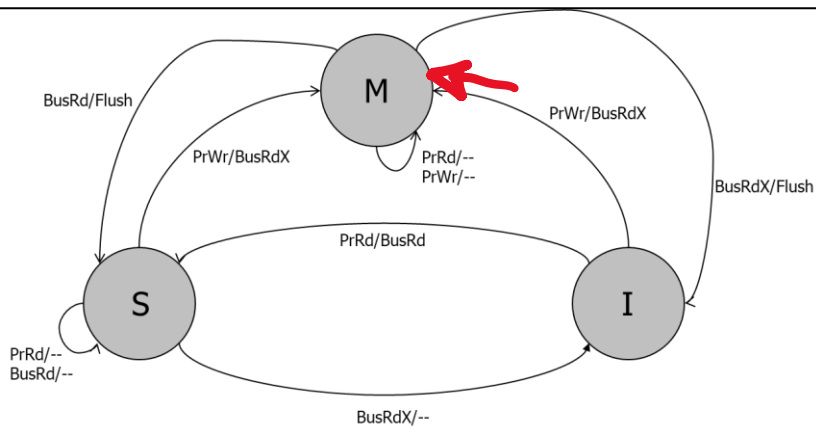
P2




ld r2, x

Interconnection Network








(MSI)

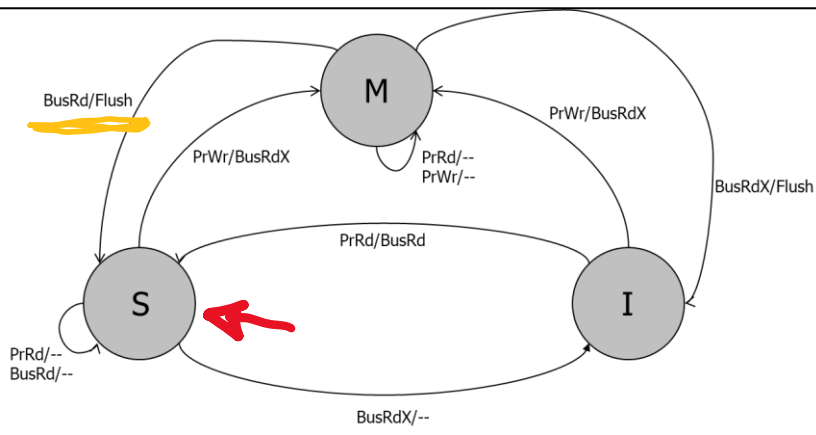


ld r2, x  
add r1, r2, r4  
st x, r1




ld r2, x  
ld r5, x







(MSI)



ld r2, x  
add r1, r2, r4  
st x, r1





Flush(x)

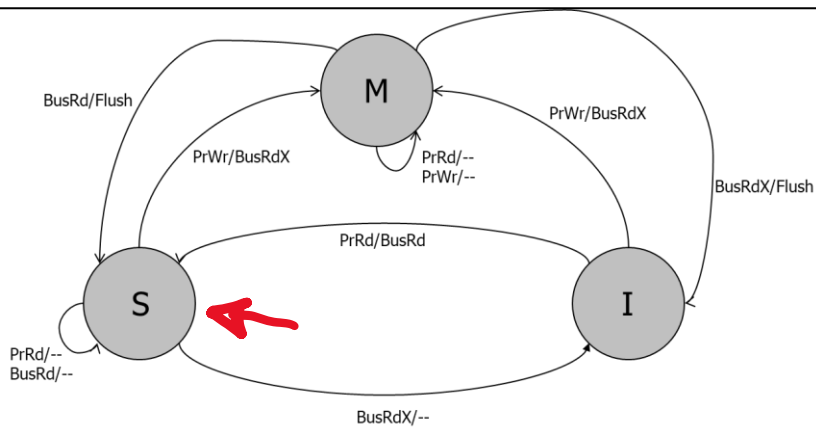


BusRd(x)


ld r2, x  
ld r5, x










(MSI)

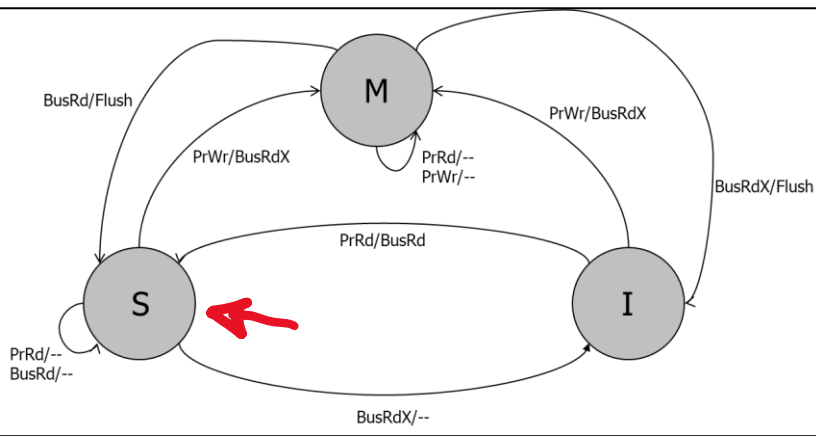


ld r2, x  
add r1, r2, r4  
st x, r1




ld r2, x  
ld r5, x






(MSI)




ld r2, x  
add r1, r2, r4  
st x, r1

P1



P2



ld r2, x  
ld r5, x

Interconnection Network

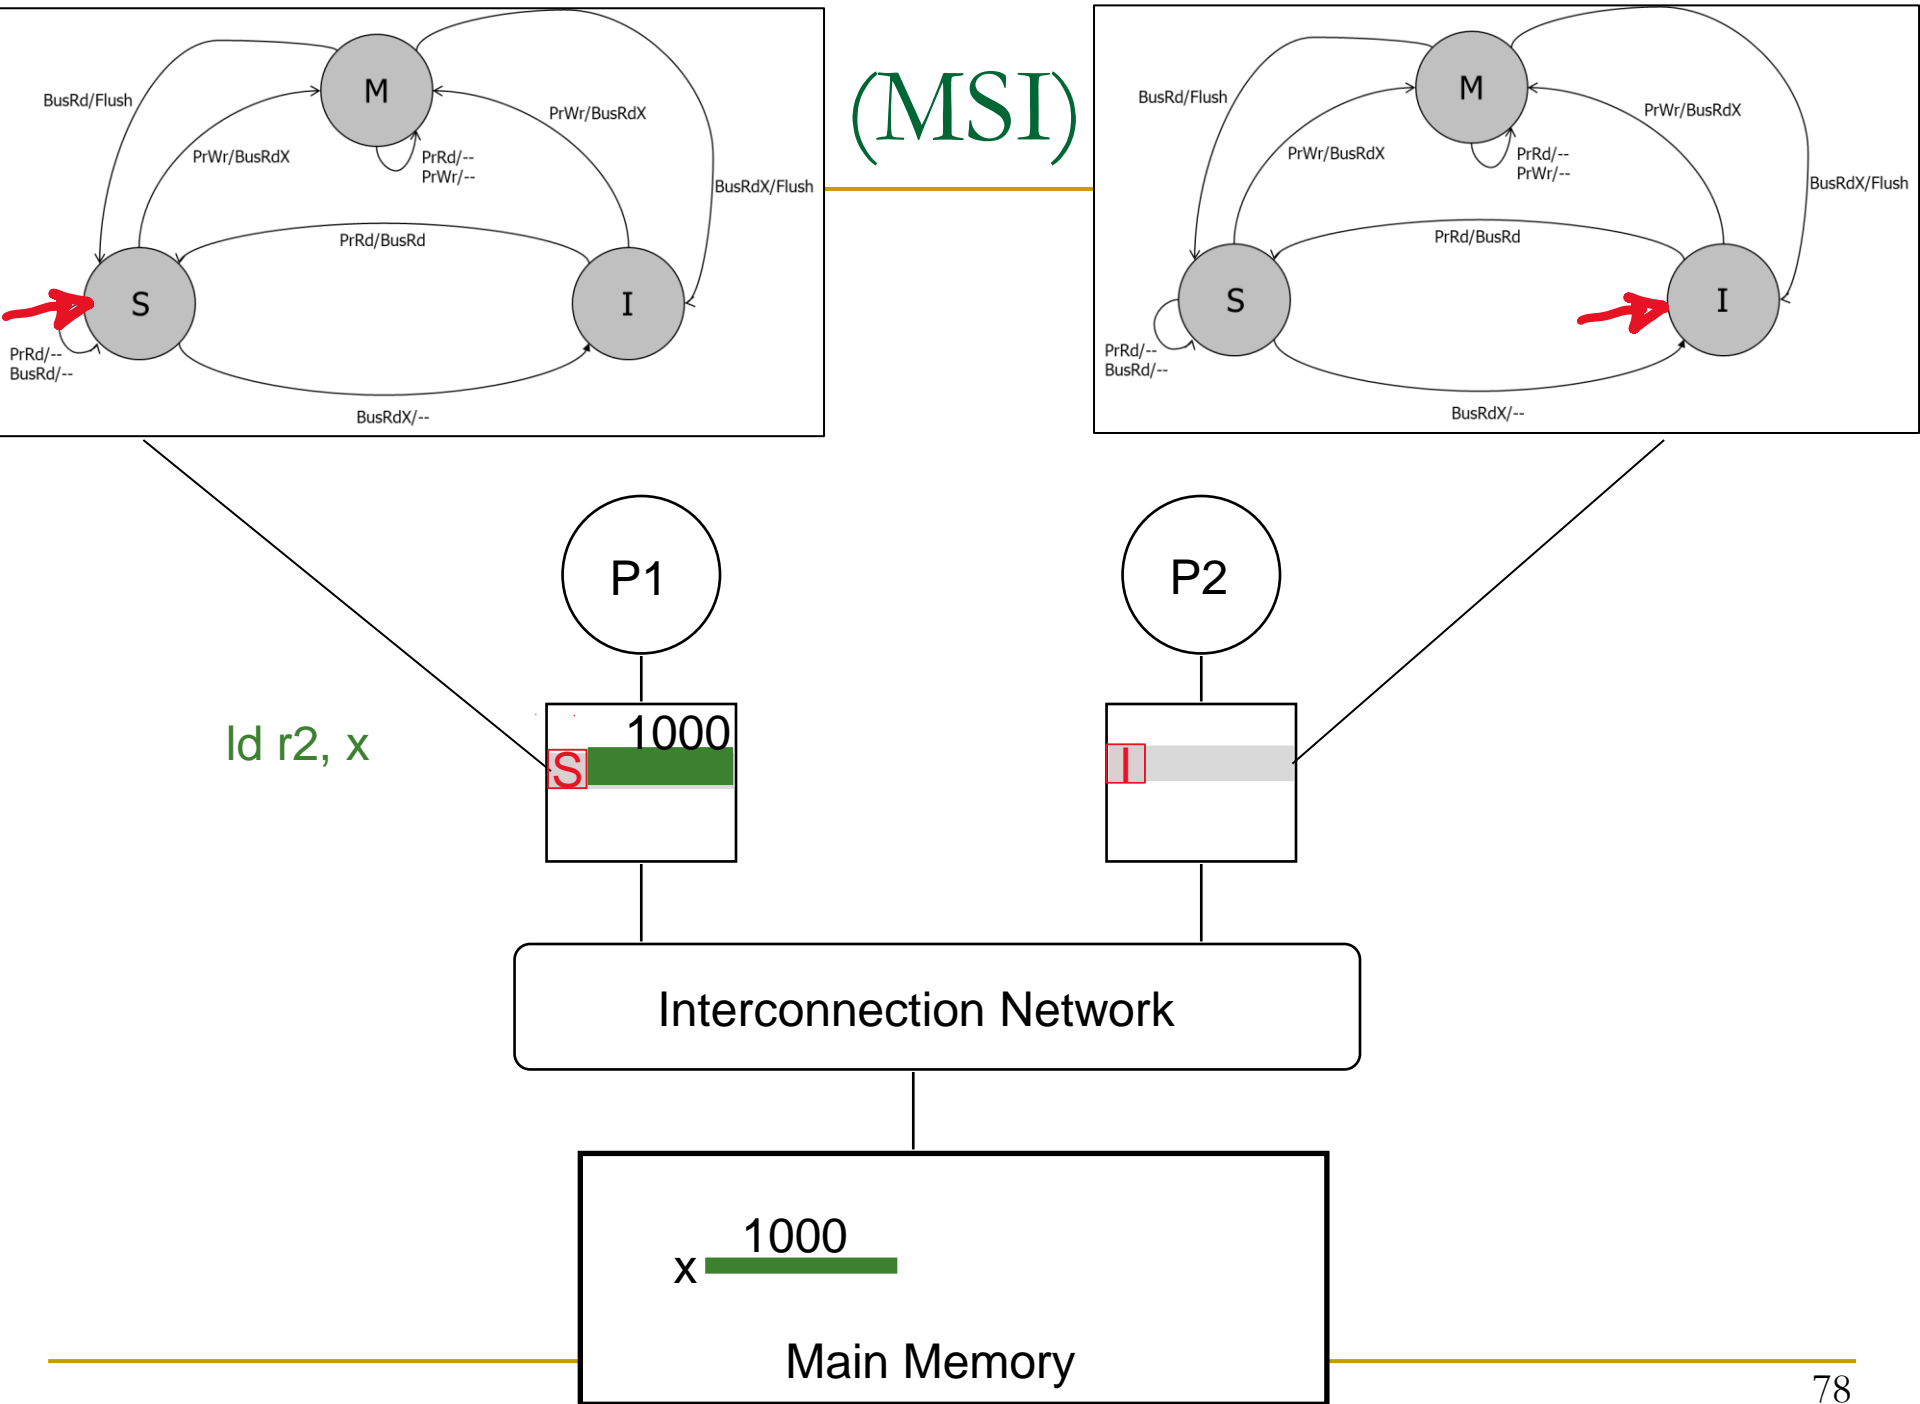
x 2000

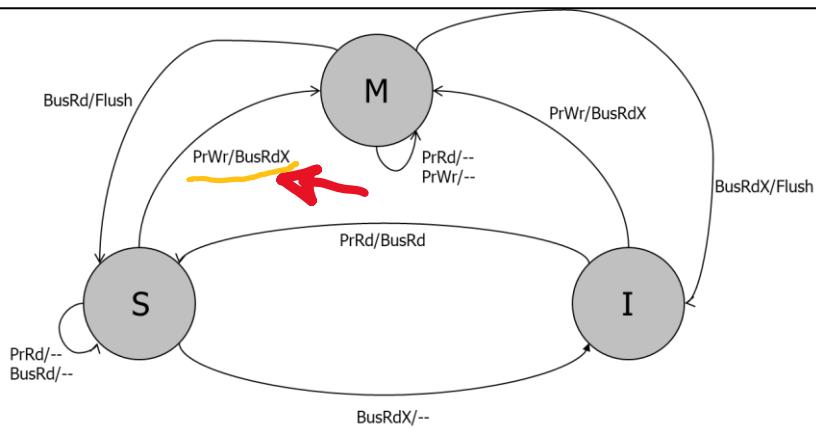
Main Memory

# The Problem with MSI


---

- A block is in no cache to begin with
- Problem: On a read, the block immediately goes to "Shared" state although it may be the only copy to be cached (i.e., no other processor will cache it)
- Why is this a problem?
  - Suppose the cache that reads the block wants to write to it at some point
  - It needs to broadcast "invalidate" even though it has the only cached copy!
  - *If the cache knew it had the only cached copy in the system, it could have written to the block without notifying any other cache → saves unnecessary broadcasts of invalidations*

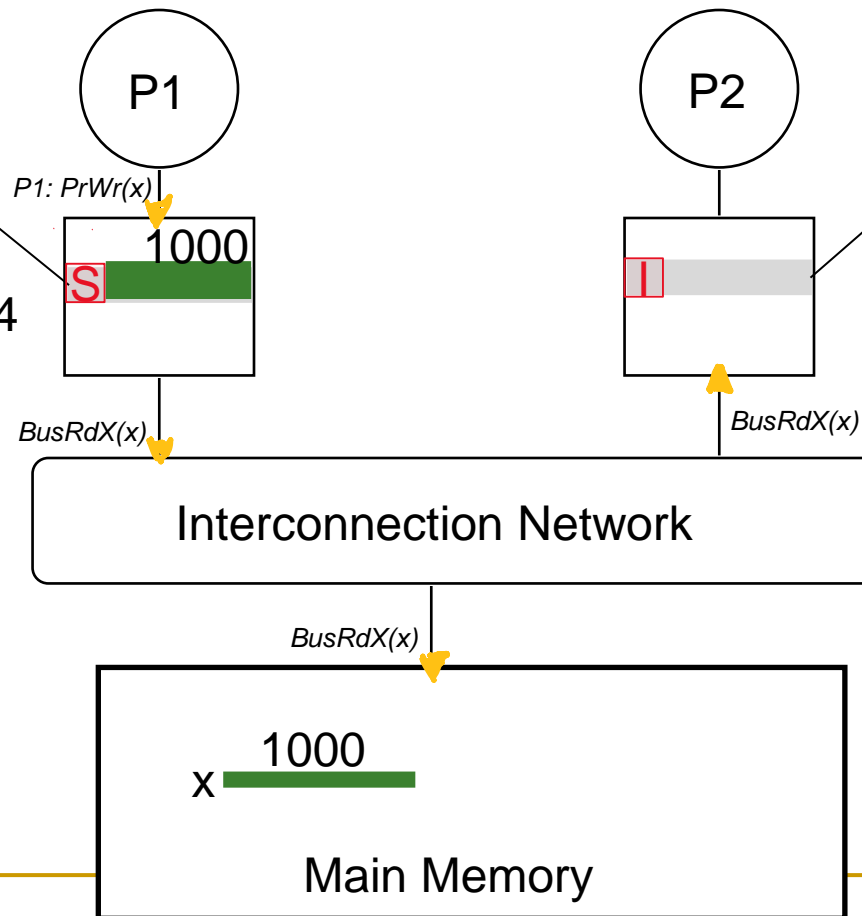





(MSI)




ld r2, x  
add r1, r2, r4  
st x, r1







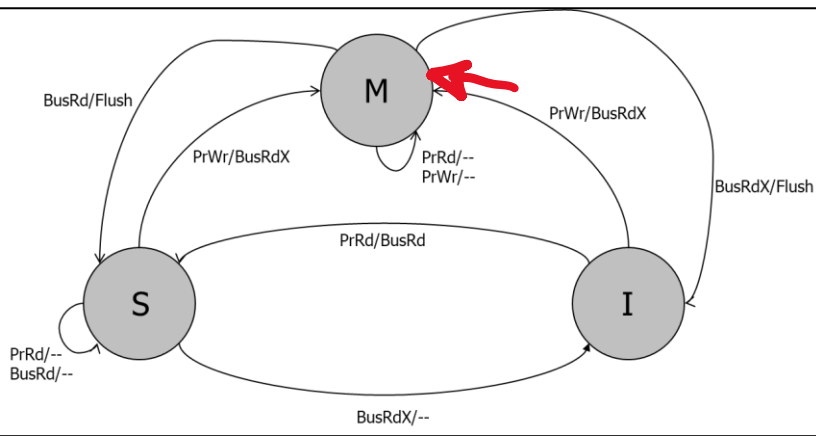


(MSI)




ld r2, x  
add r1, r2, r4  
st x, r1

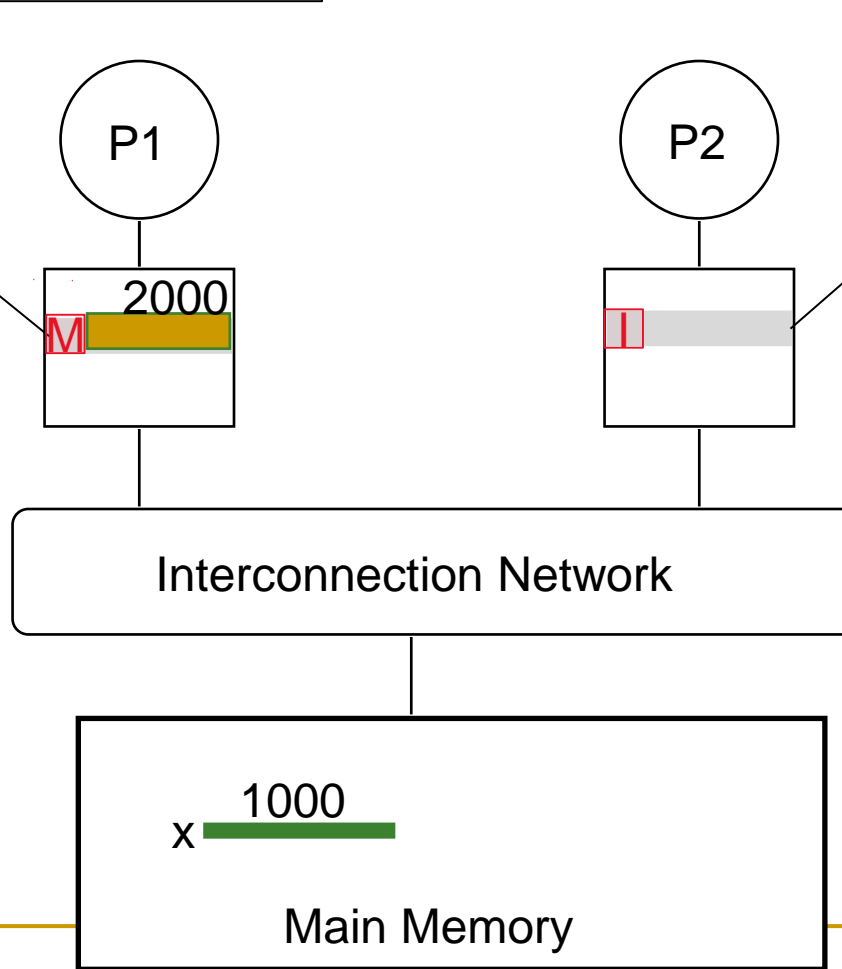




(MSI)



ld r2, x  
add r1, r2, r4  
st x, r1



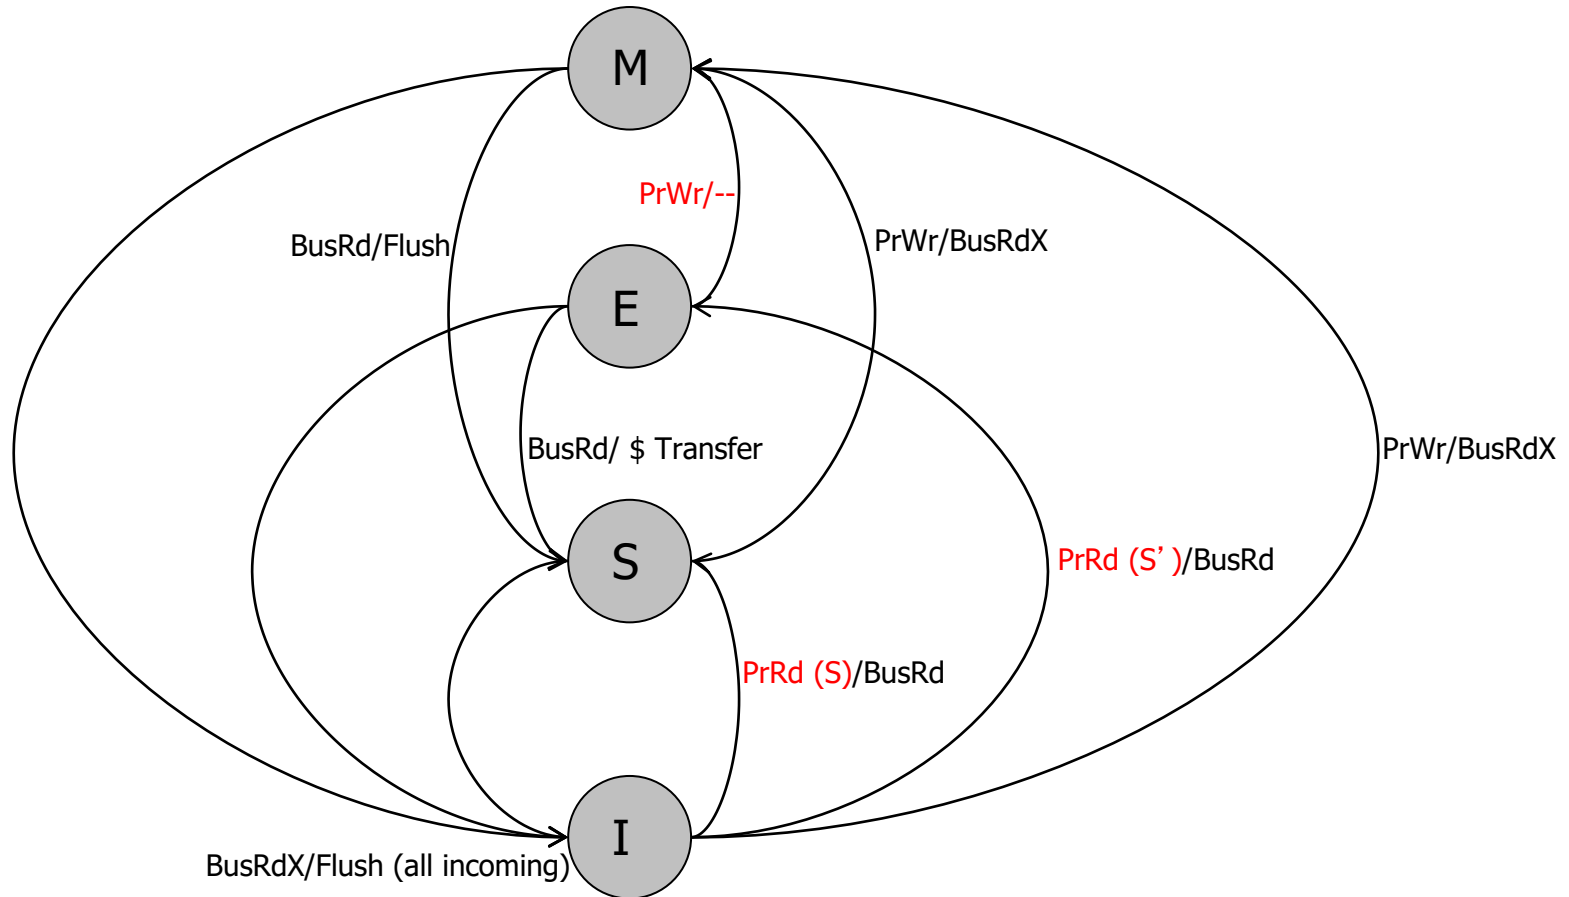
# The Solution: MESI

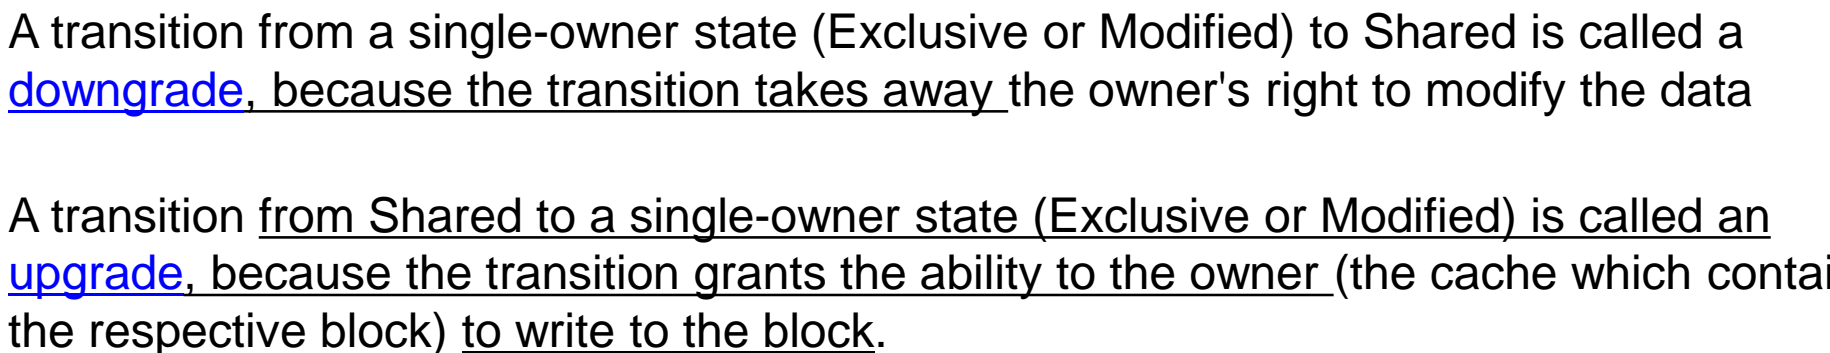
---

- Idea: Add another state indicating that this is the only cached copy and it is clean.
  - *Exclusive* state
- Block is placed into the *exclusive* state if, during *BusRd*, no other cache had it
  - Wired-OR “shared” signal on bus can determine this:  
snooping caches assert the signal if they also have a copy
- Silent transition *Exclusive* → *Modified* is possible on write!
- MESI is also called the *Illinois protocol*
  - Papamarcos and Patel, “A low-overhead coherence solution for multiprocessors with private cache memories,” ISCA 1984.




# MESI State Machine






# MESI State Machine



# Illinois Protocol



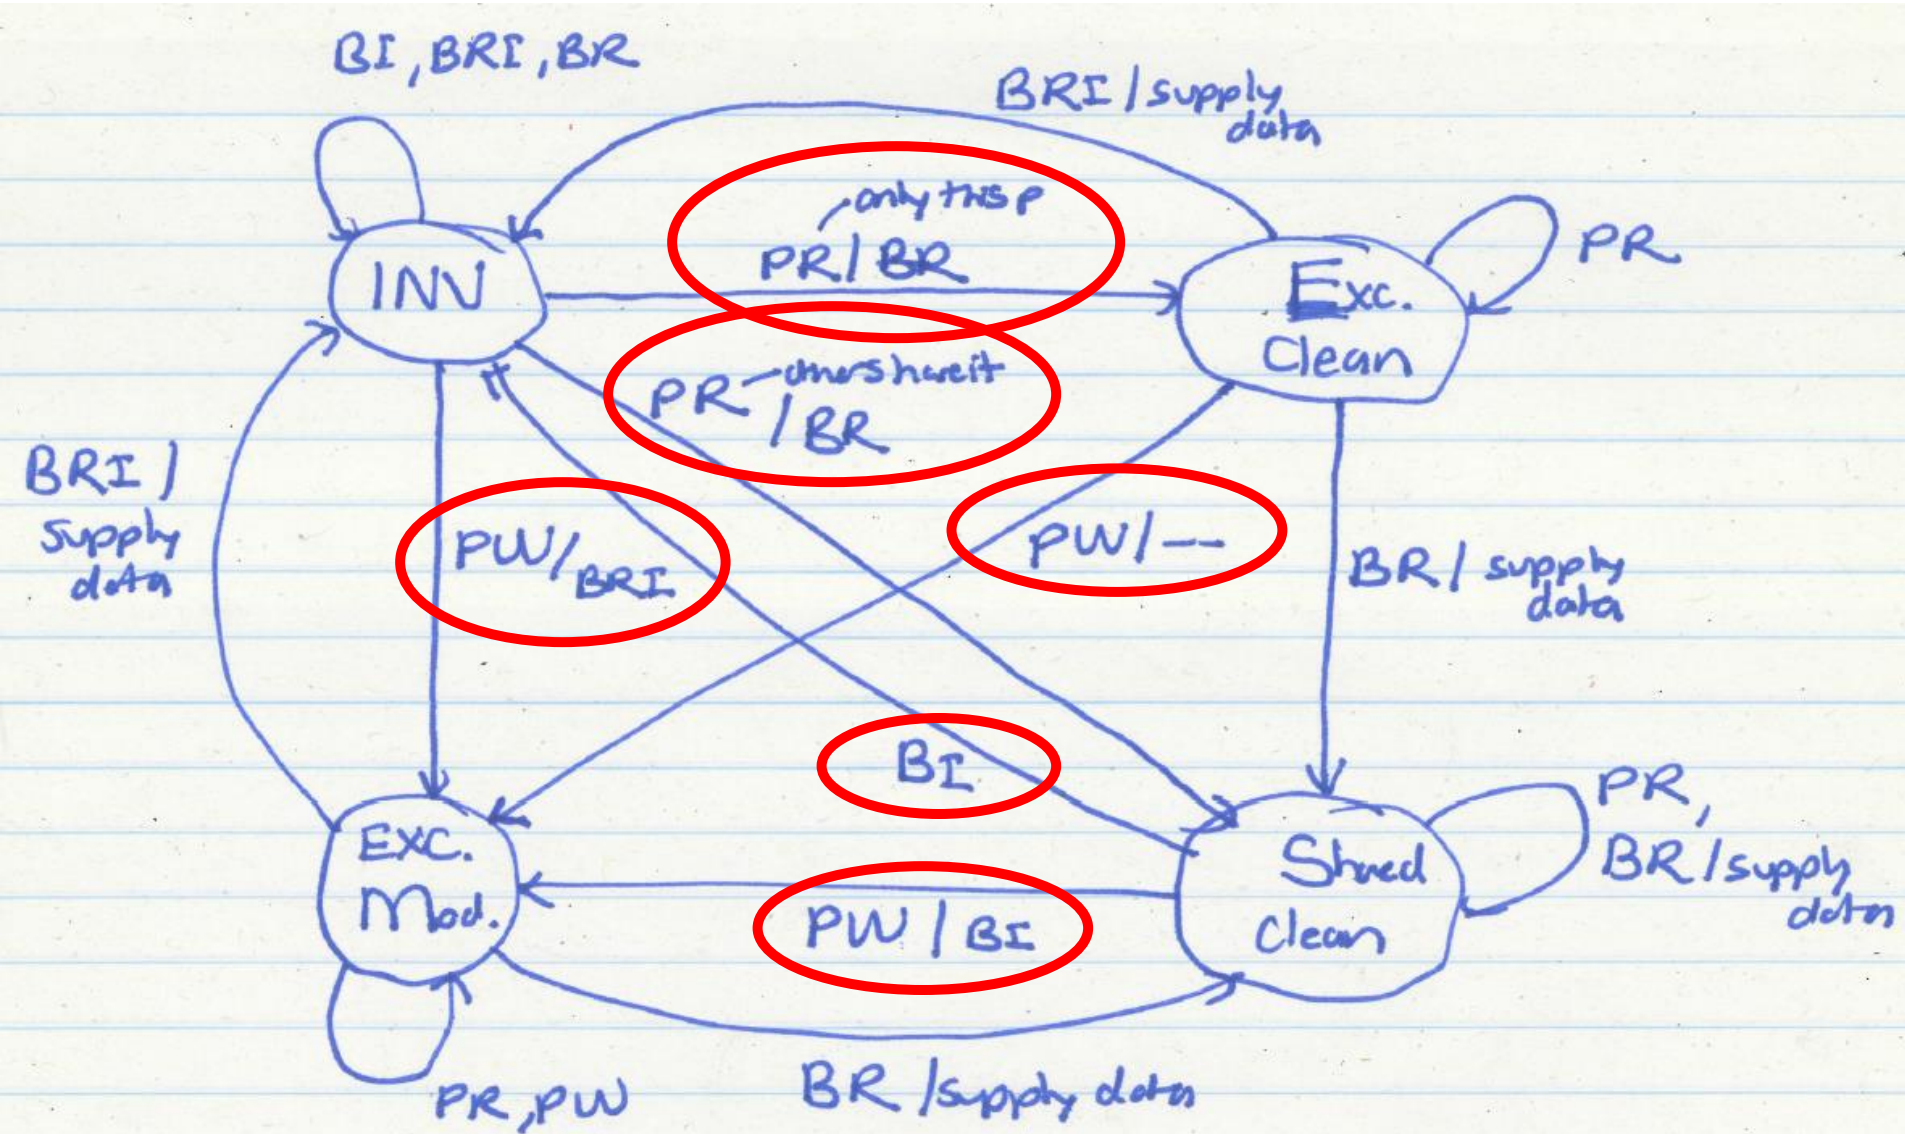
**BI:** Invalidate, but already have the data (do not supply it)

**BRI:** Invalidate, but also need the data (supply it)

## 4 States


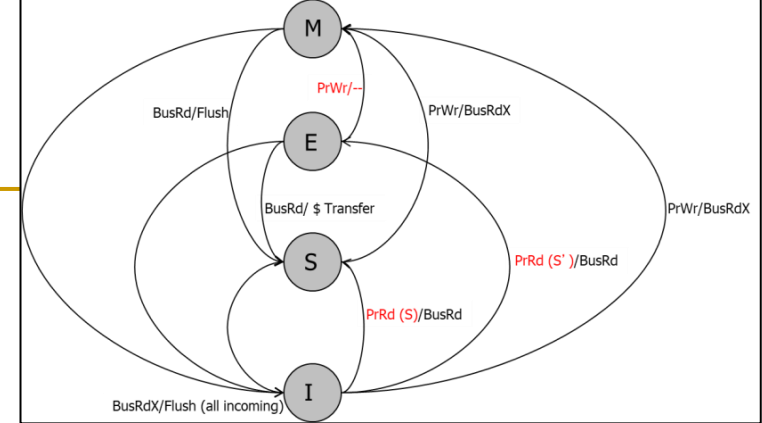
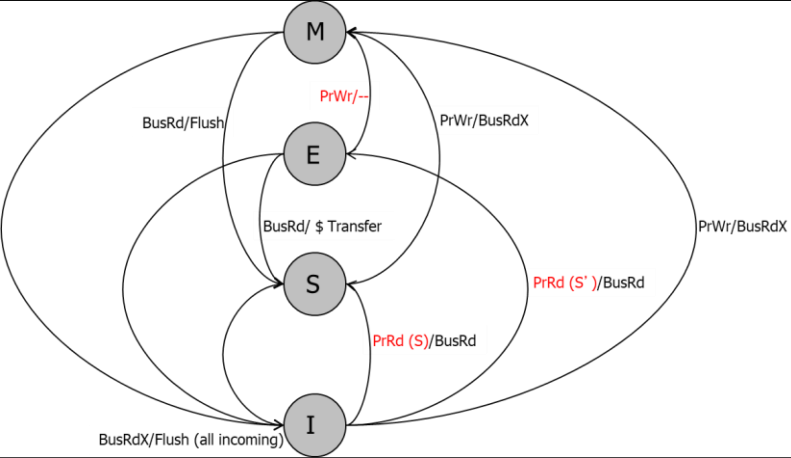
- M: Modified (Exclusive copy, modified)
- E: Exclusive (" " , clean)
- S: Shared (Shared copy, clean)
- I: Invalid

# MESI State Machine









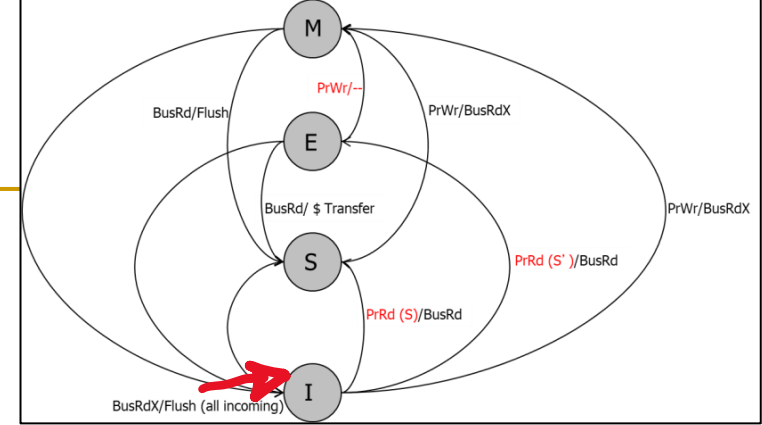

# (MESI)



# (MESI)



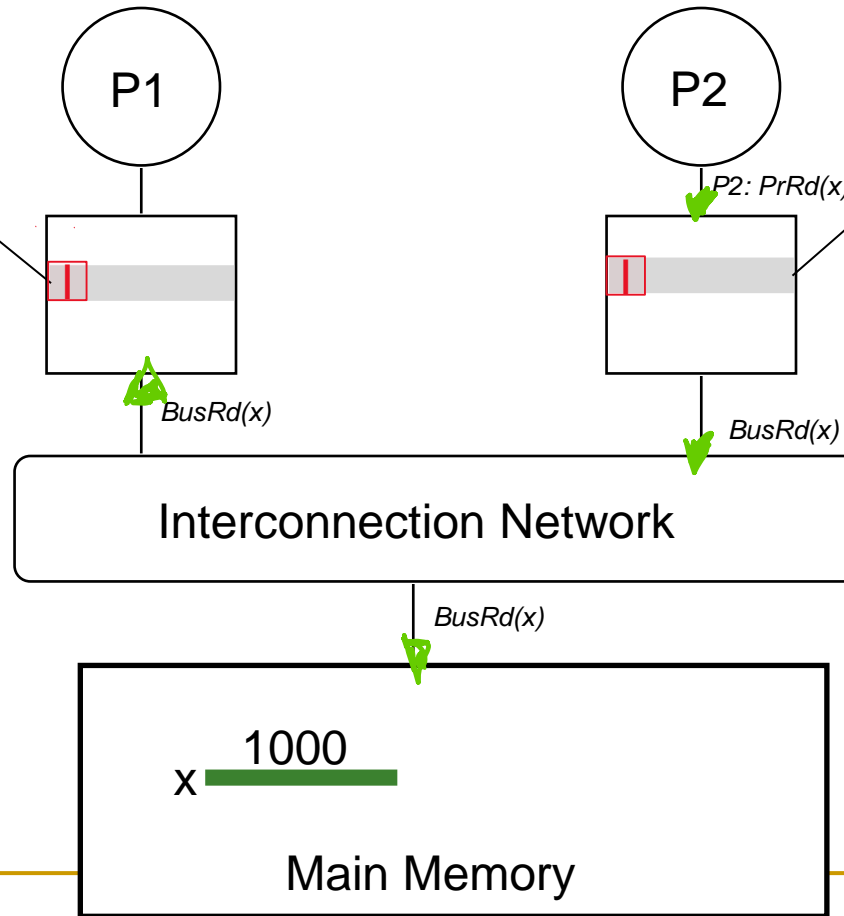
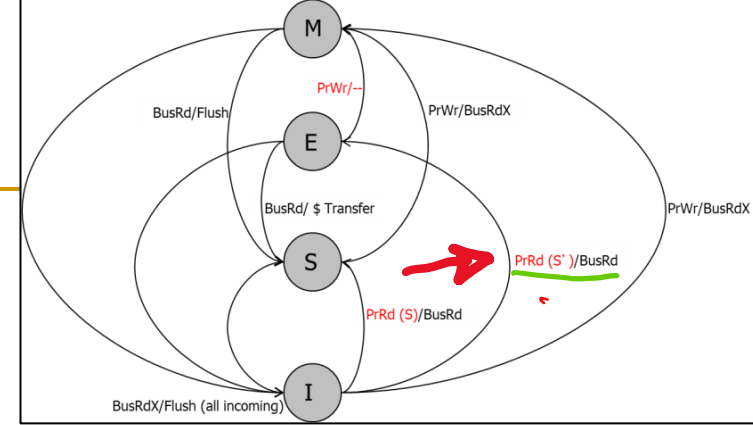
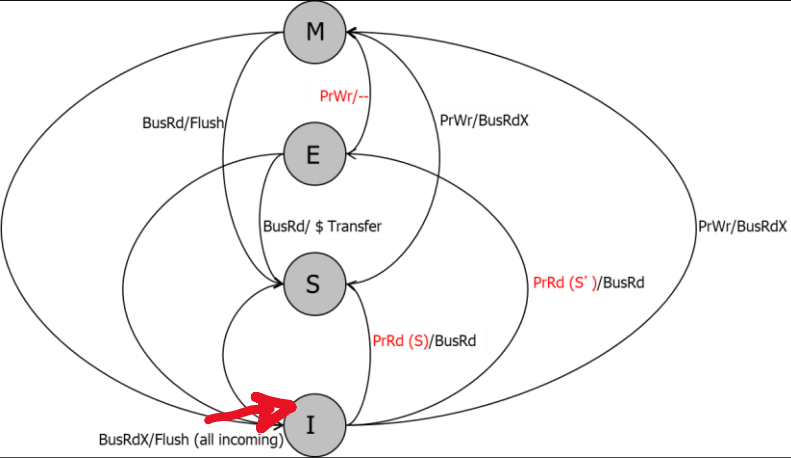
# (MESI)



ld r2, x

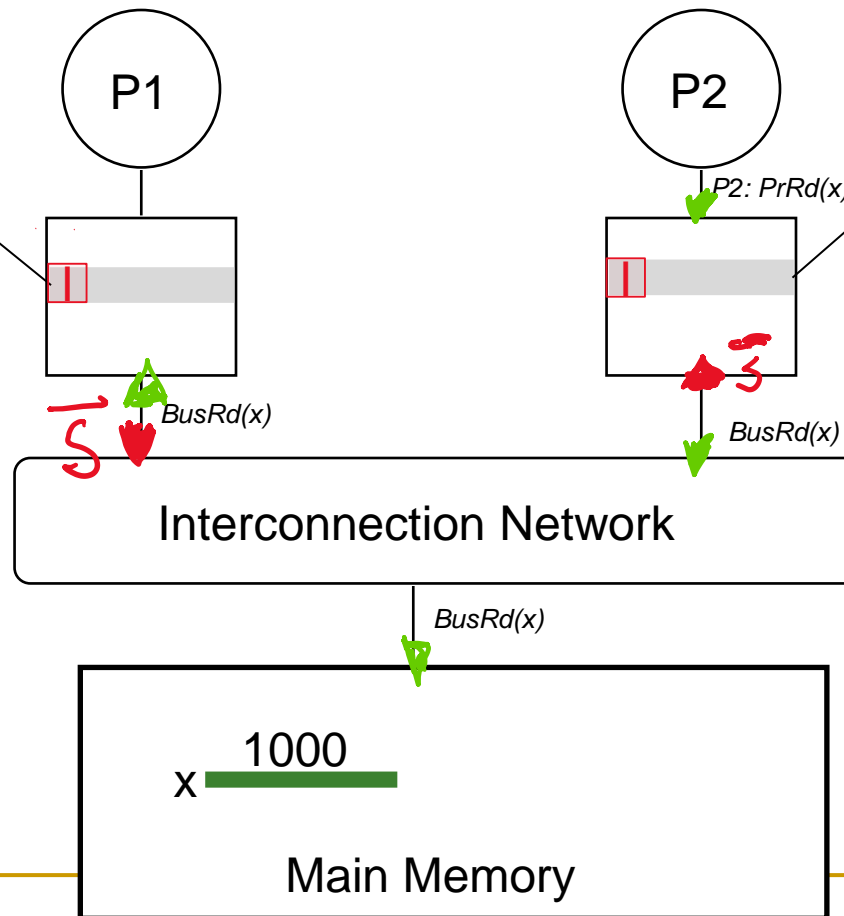
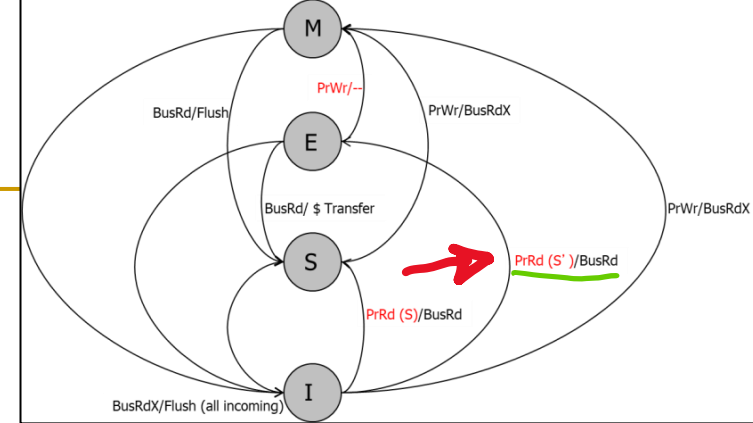
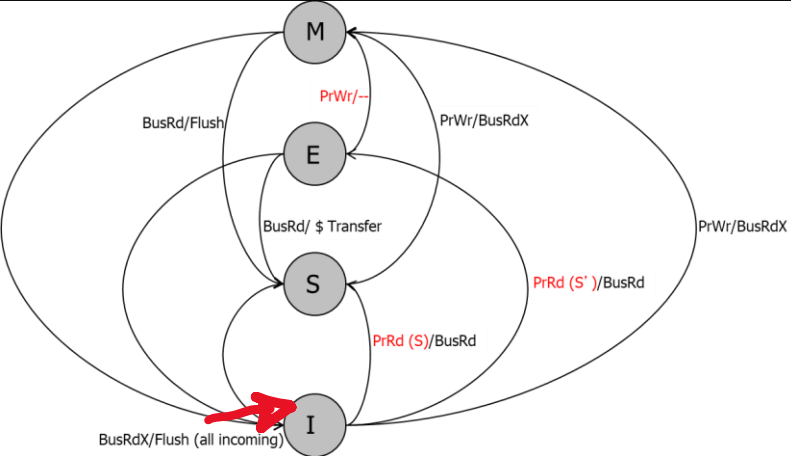


# (MESI)



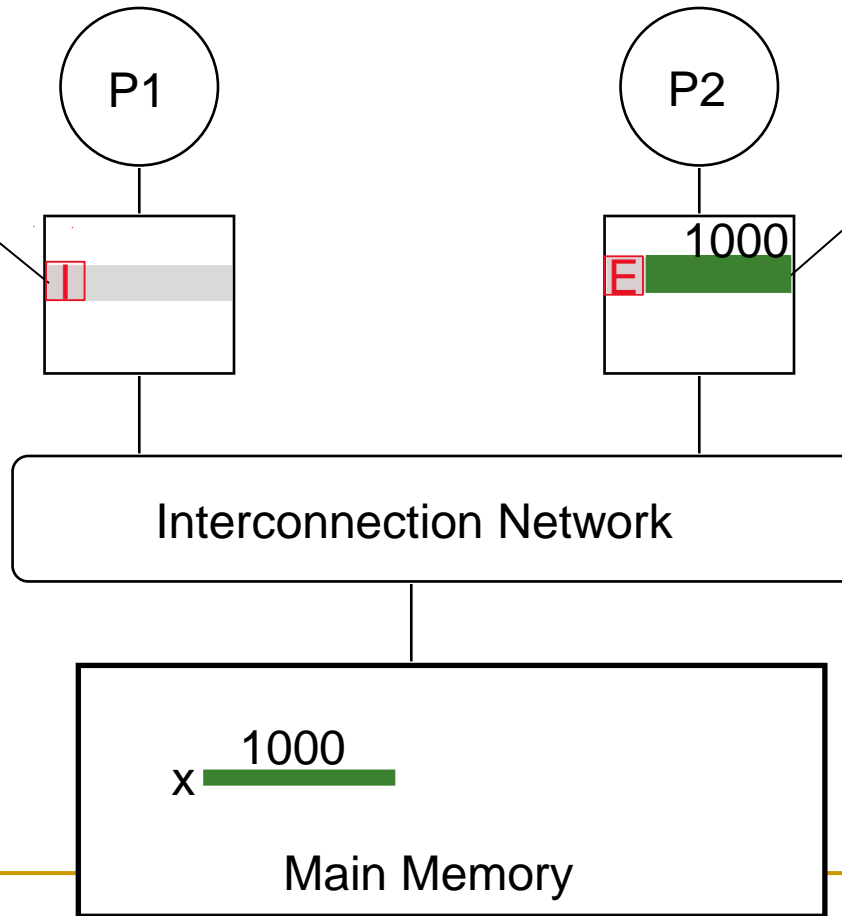
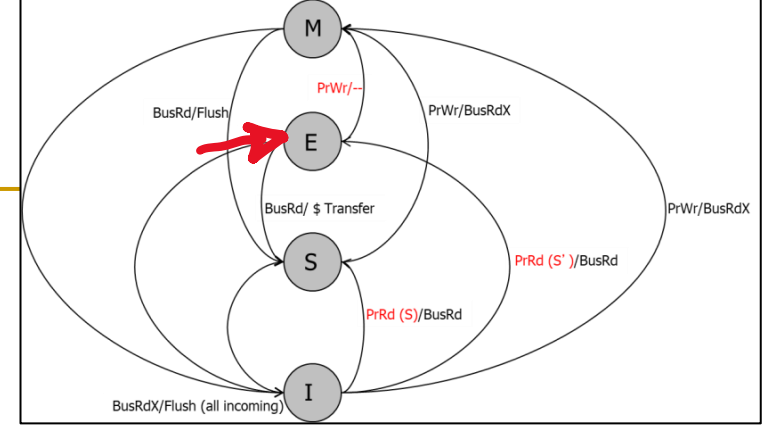

ld r2, x

# (MESI)



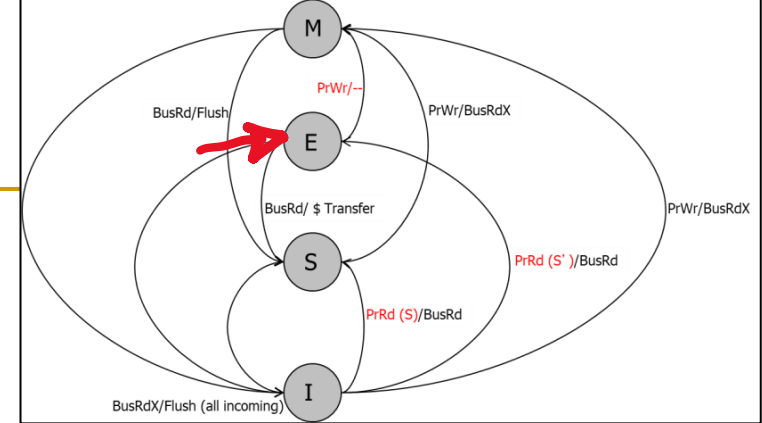

ld r2, x

# (MESI)




ld r2, x


# (MESI)



ld r2, x



BusRd(x)

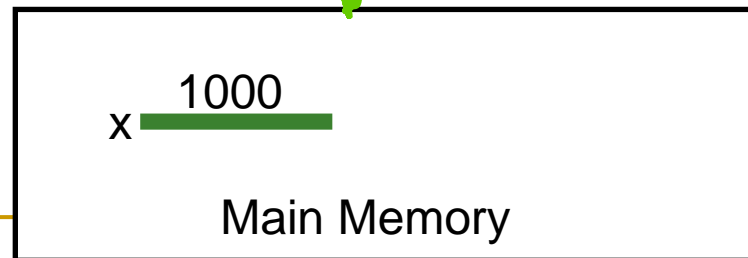


BusRd(x)

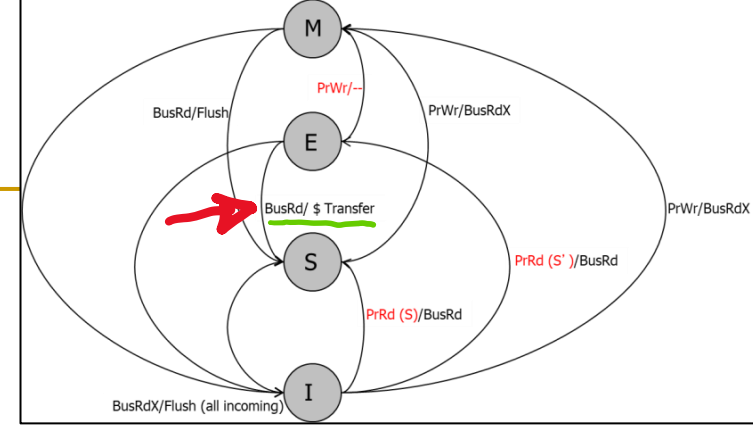
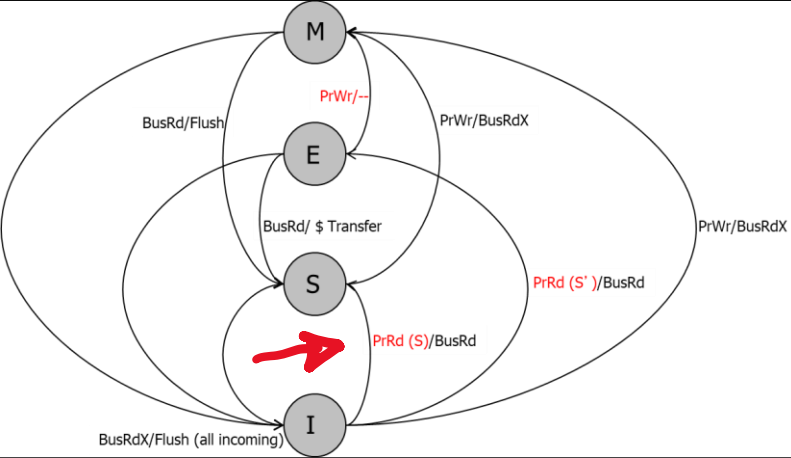
ld r2, x




BusRd(x)




# (MESI)




ld r2, x




BusRd(x)



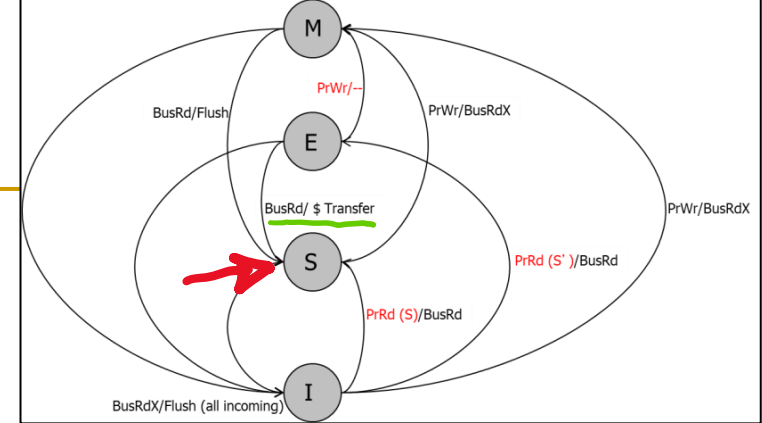
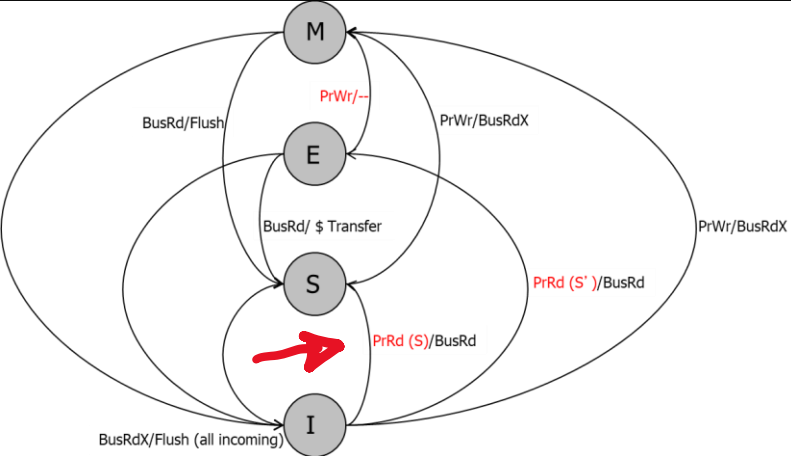
BusRd(x)




BusRd(x)



# (MESI)



ld r2, x

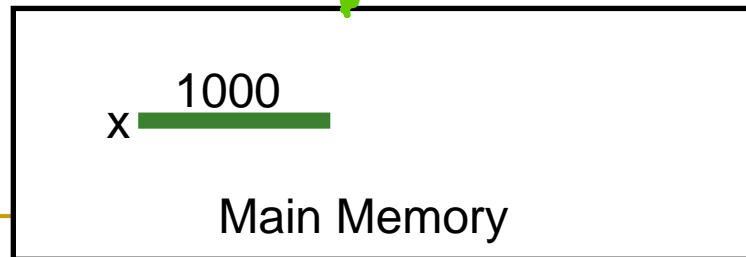


BusRd(x)


Interconnection Network

BusRd(x)

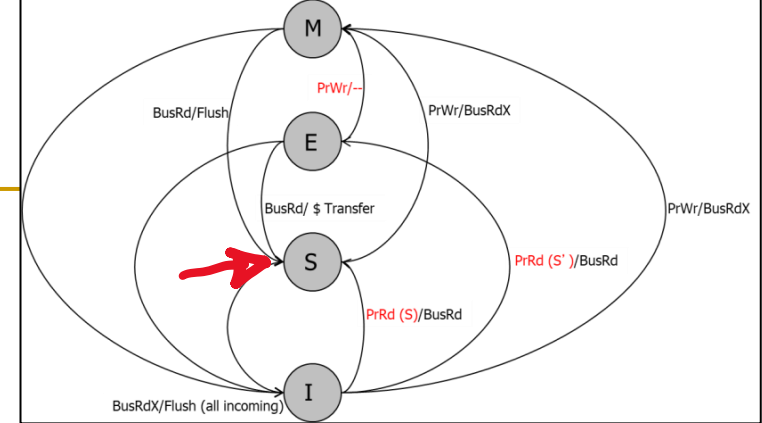
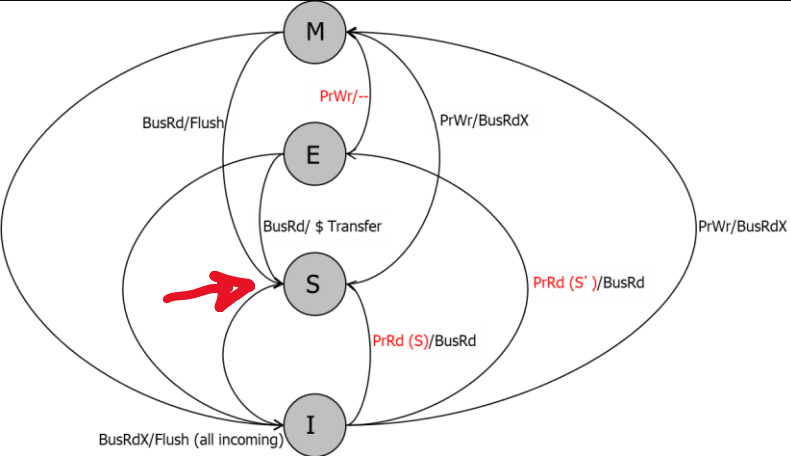
BusRd(x)







ld r2, x




# (MESI)



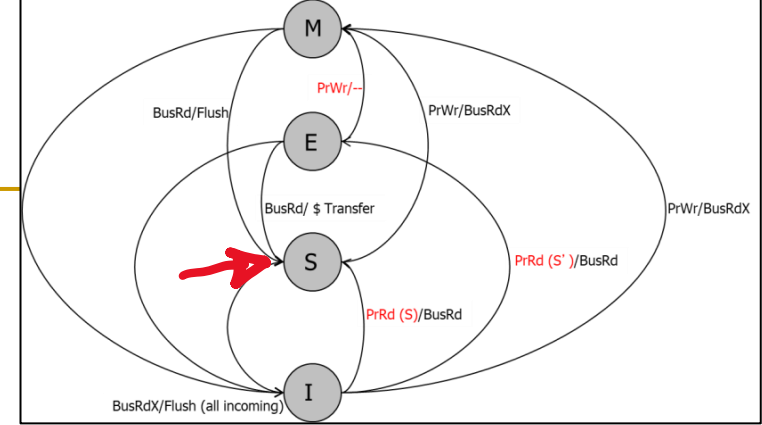
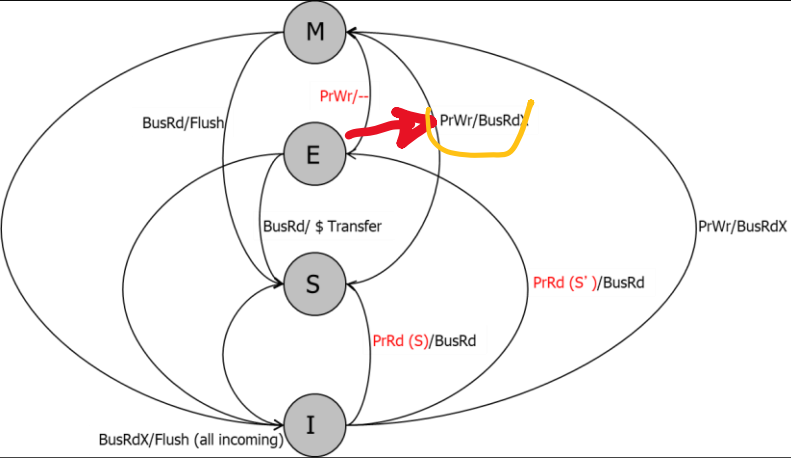
ld r2, x




ld r2, x




# (MESI)



ld r2, x  
add r1, r2, r4  
st x, r1




BusRdX(x)




BusRdX(x)

ld r2, x

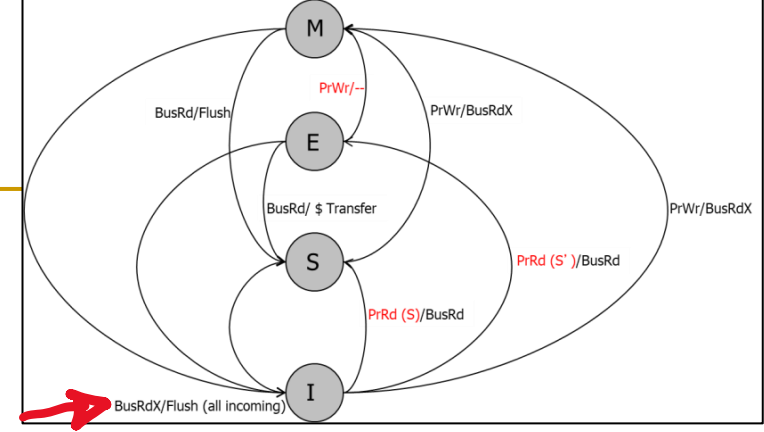
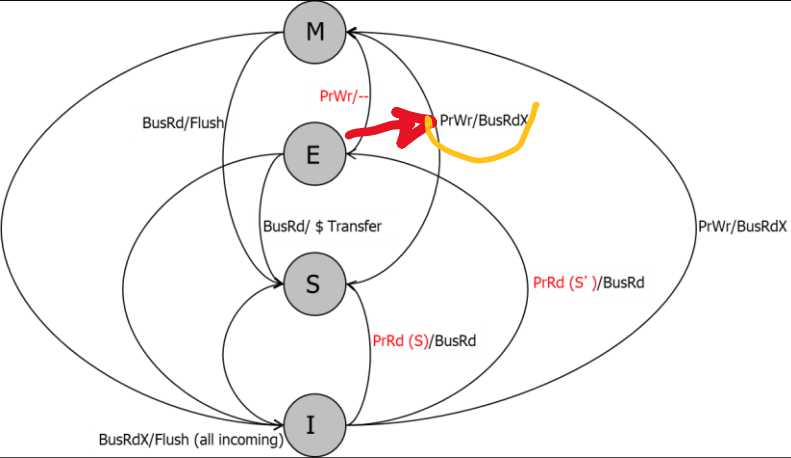


BusRdX(x)







# (MESI)



ld r2, x  
add r1, r2, r4  
st x, r1




BusRdX(x)




BusRdX(x)


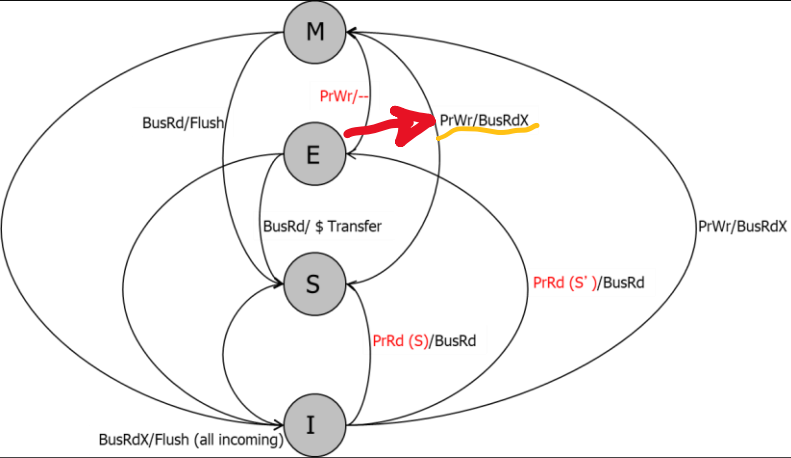
ld r2, x





BusRdX(x)





# (MESI)



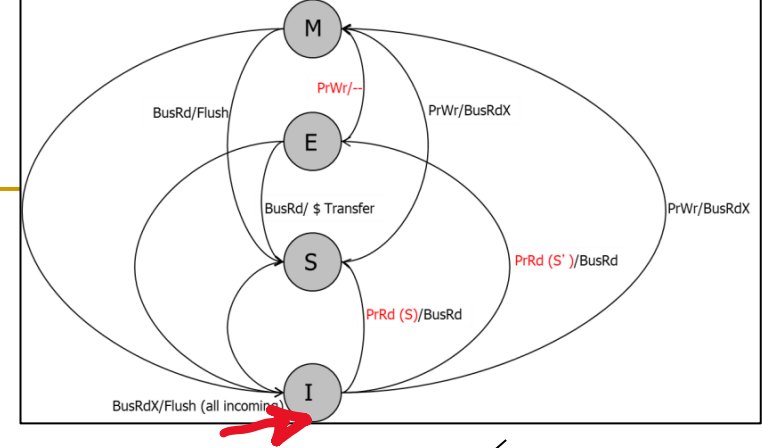
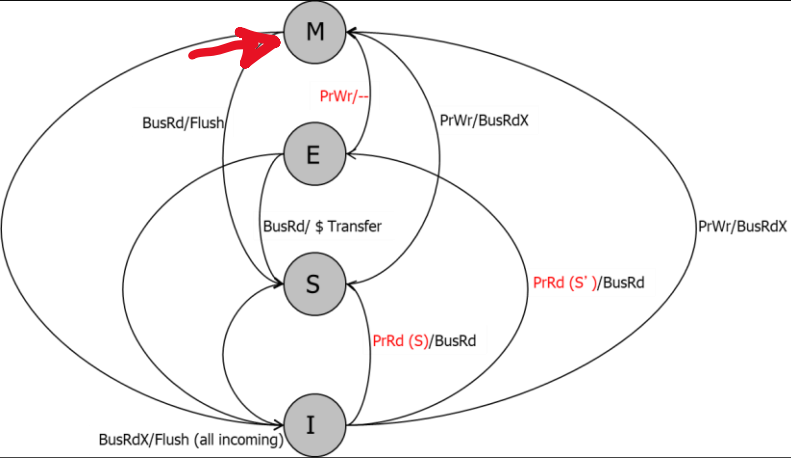
ld r2, x  
add r1, r2, r4  
st x, r1







ld r2, x




# (MESI)



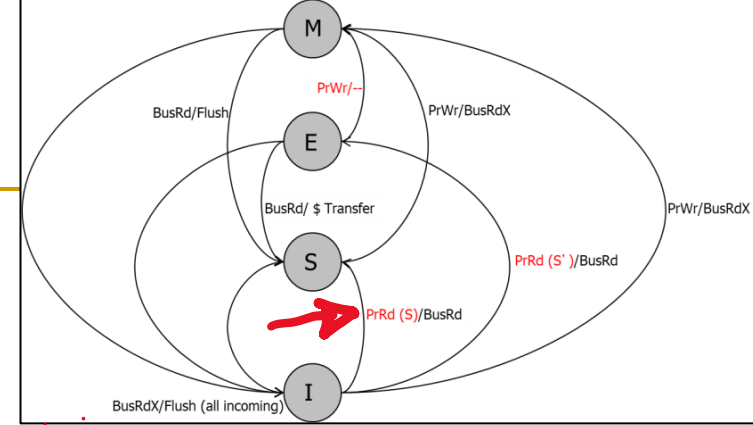

ld r2, x  
add r1, r2, r4  
st x, r1



ld r2, x




# (MESI)




ld r2, x  
add r1, r2, r4  
st x, r1

P1



BusRd(x)

P2




BusRd(x)

ld r2, x  
ld r5, x

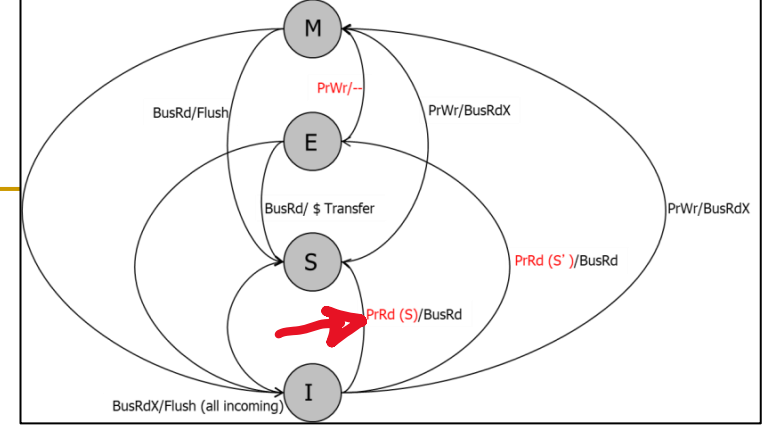
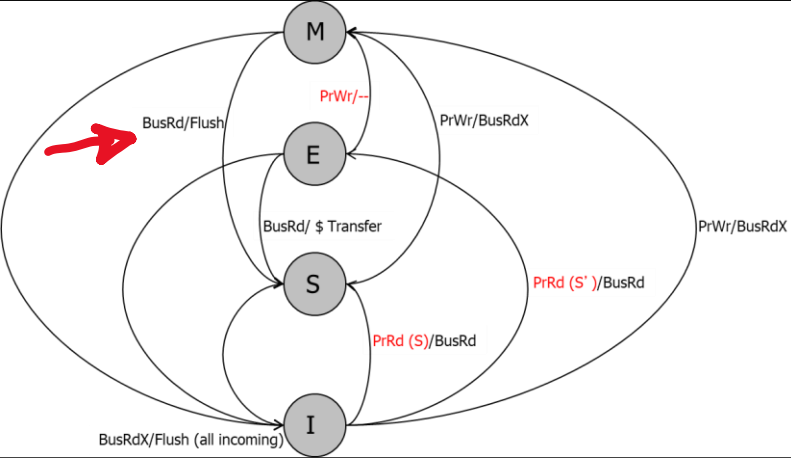


BusRd(x)




Main Memory

# (MESI)




ld r2, x  
add r1, r2, r4  
st x, r1

P1




Interconnection Network



P2



ld r2, x  
ld r5, x





# (MESI)





ld r2, x  
add r1, r2, r4  
st x, r1



P1





P2



ld r2, x  
ld r5, x




# (MESI)




ld r2, x  
add r1, r2, r4  
st x, r1


P1



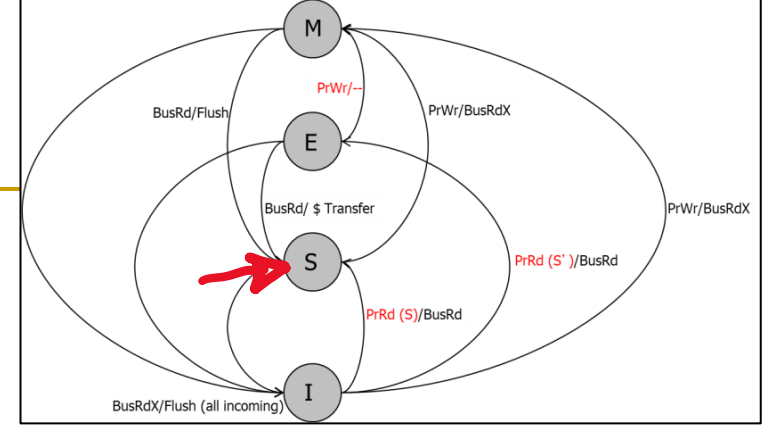
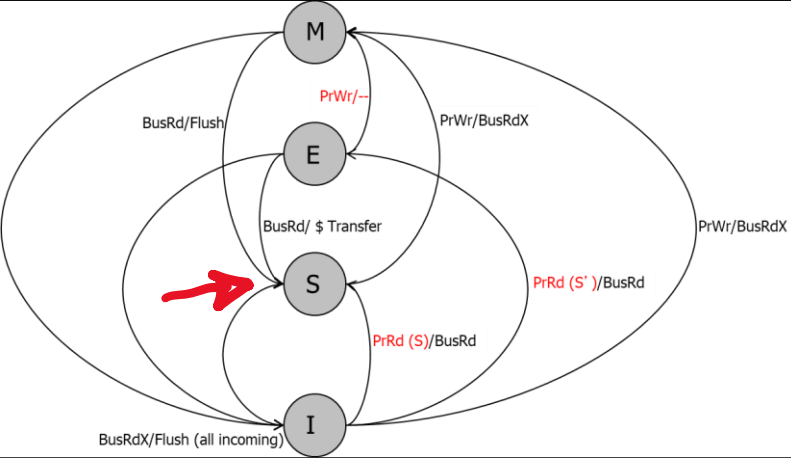
P2







ld r2, x  
ld r5, x




# (MESI)



ld r2, x  
add r1, r2, r4  
st x, r1


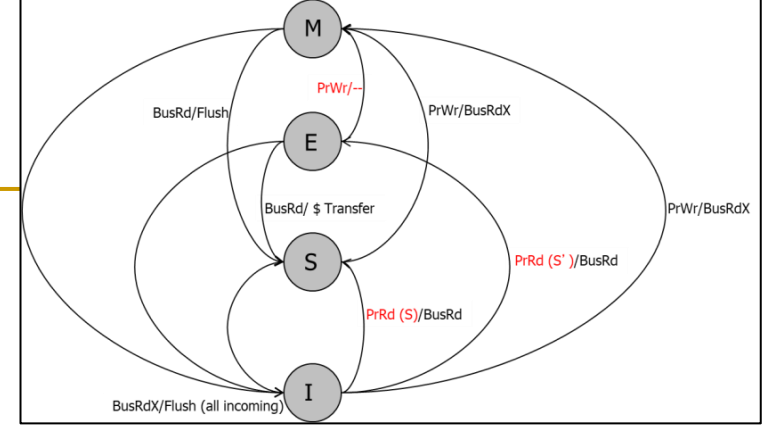
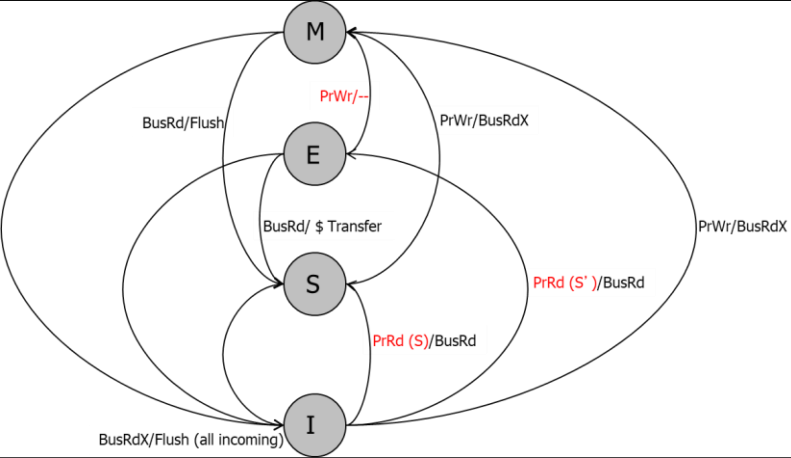


ld r2, x  
ld r5, x

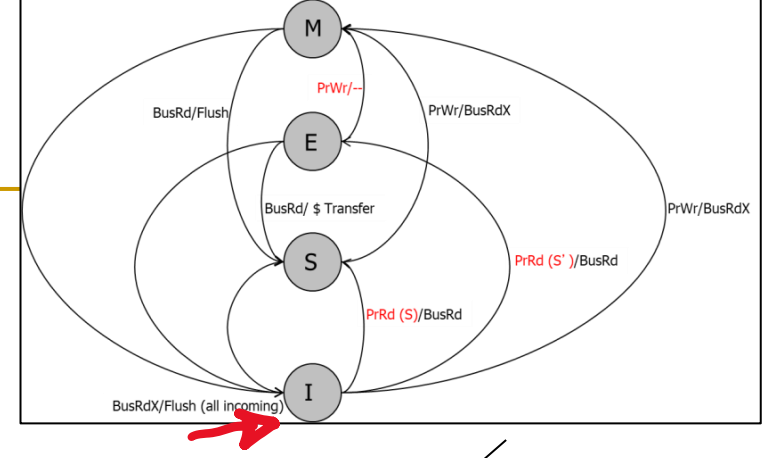
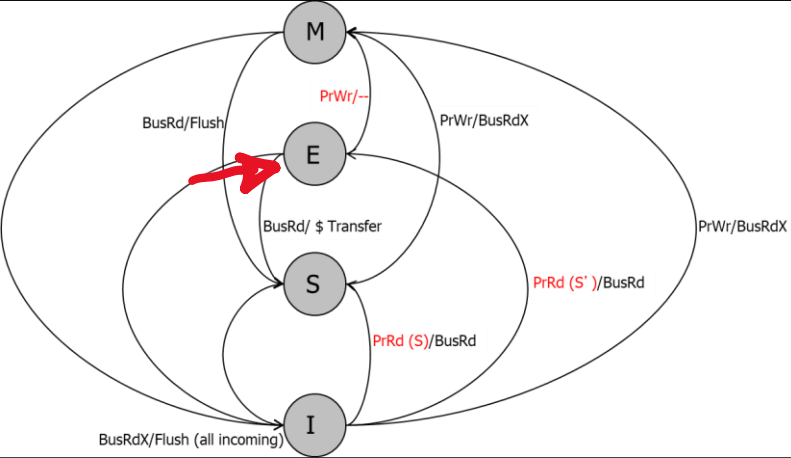




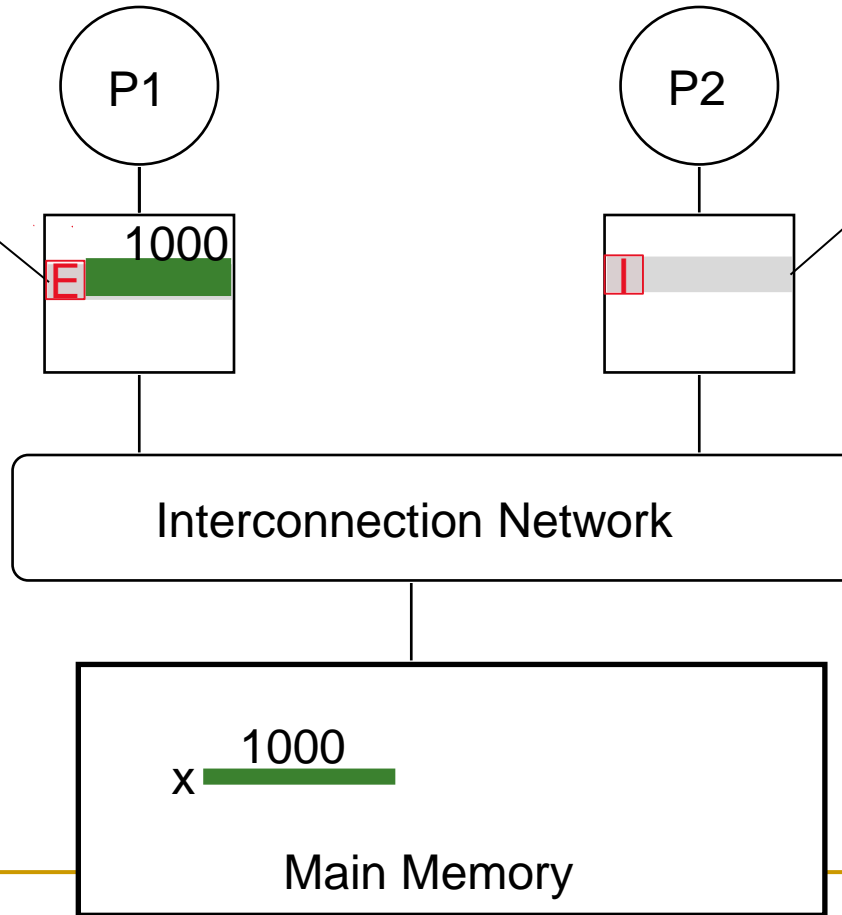
# (MESI)



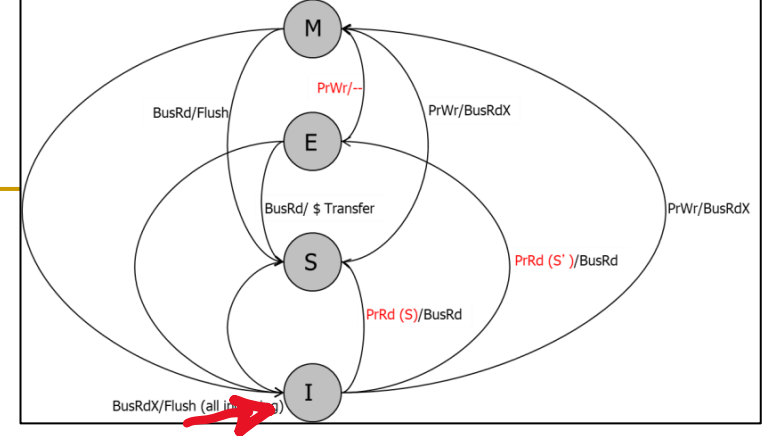
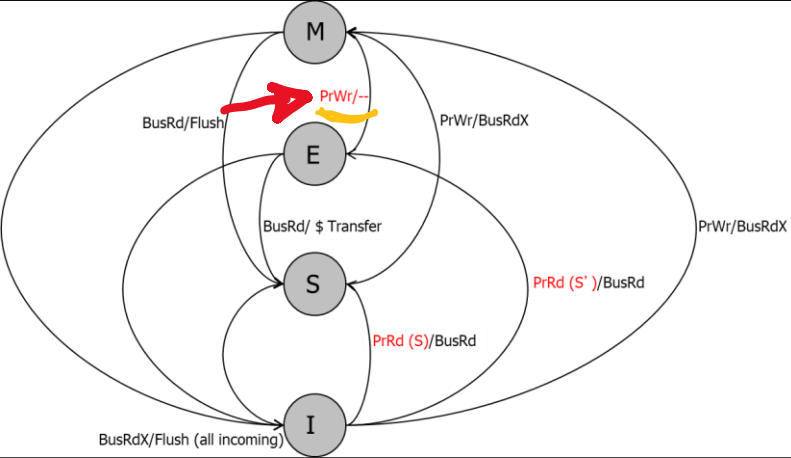
# (MESI)







ld r2, x



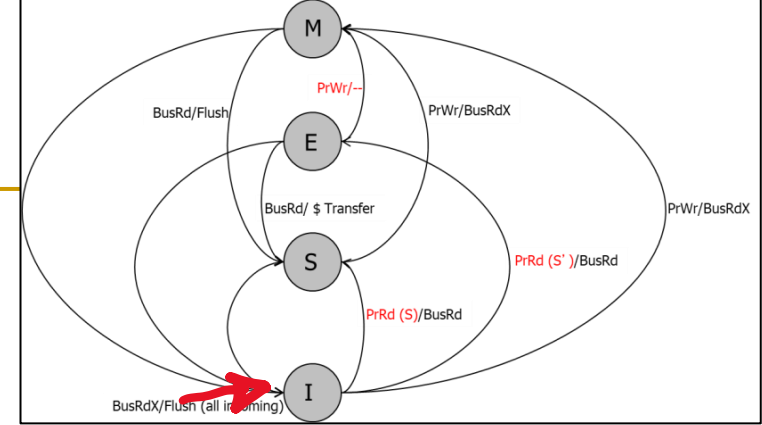

# (MESI)





ld r2, x  
add r1, r2, r4  
st x, r1



# (MESI)



ld r2, x  
add r1, r2, r4  
st x, r1


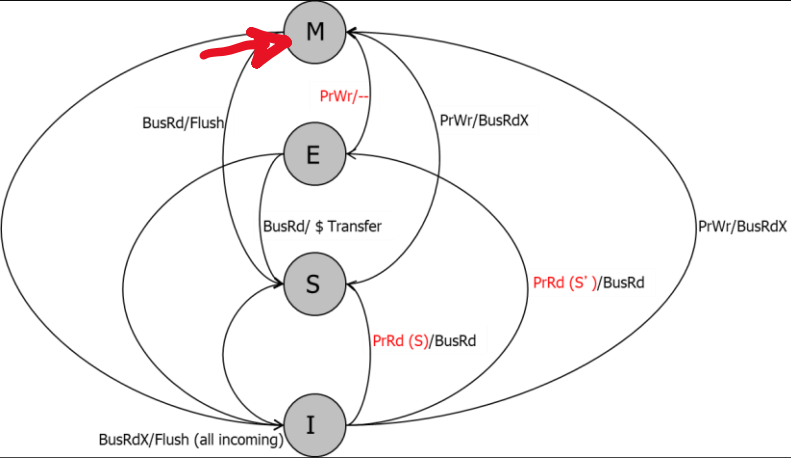


Interconnection Network





x 1000

Main Memory

# (MESI)



ld r2, x  
add r1, r2, r4  
st x, r1



# Snoopy Invalidation Tradeoffs

---

- Should a downgrade from M go to S or I?
  - S: if data is likely to be reused (before it is written to by another processor)
  - I: if data is likely to be not reused (before it is written to by another)
- Cache-to-cache transfer
  - On a BusRd, should data come from another cache or memory?
    - Another cache
      - May be faster, if memory is slow or highly contended
    - Memory
      - Simpler: no need to wait to see if another cache has the data first
      - Less contention at the other caches
      - Requires writeback on M downgrade
- Writeback on Modified->Shared: necessary?
  - One possibility: **Owner** (O) state (MOESI protocol)
    - One cache owns the latest data (memory is not updated)
    - Memory writeback happens when all caches evict copies

# The Problem with MESI

---

- Observation: Shared state requires the data to be clean
  - i.e., all caches that have the block have the up-to-date copy and so does the memory
- Problem: Need to write the block to memory when BusRd happens when the block is in Modified state
- Why is this a problem?
  - Memory can be updated unnecessarily → some other processor may want to write to the block again

# Improving on MESI

---

- Idea 1: Do not transition from  $M \rightarrow S$  on a BusRd. Invalidate the copy and supply the modified block to the requesting processor directly without updating memory
- Idea 2: Transition from  $M \rightarrow S$ , but designate one cache as the owner (O), who will write the block back when it is evicted
  - Now “Shared” means “Shared and potentially dirty”
  - This is a version of the MOESI protocol



# Tradeoffs in Sophisticated Cache Coherence Protocols

---

- The protocol can be optimized with more states and prediction mechanisms to
  - + Reduce unnecessary invalidates and transfers of blocks
- However, more states and optimizations
  - Are more difficult to design and verify (lead to more cases to take care of, race conditions)
  - Provide diminishing returns