

University of Bologna

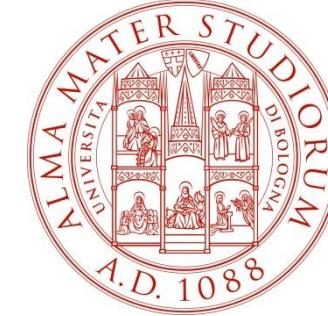
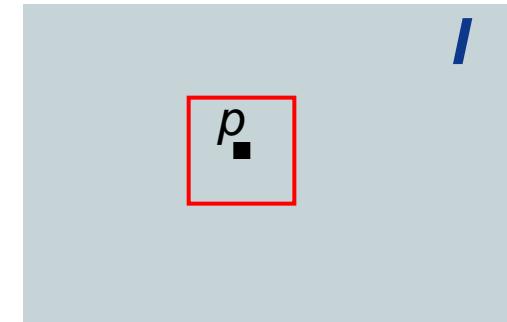


Image Filtering

Luigi Di Stefano (luigi.distefano@unibo.it)

Introduction

- *Image Filters* are image processing operators that compute the new intensity (colour) of a pixel, p , based on the intensities (colours) of those belonging to a neighbourhood of p .
- They accomplish a variety of useful image processing functions, such as e.g. *denoising* and *sharpening* (edge enhancement).
- An important sub-class of filters is given by *Linear* and *Translation-Equivariant* (LTE) operators, which we will consider first.
- Signal theory dictates their application in image processing to consist in a *2D convolution* between the input image and the *impulse response function (point spread function or kernel)* of the LTE operator.
- LTE operators are used as feature extractors in CNNs (Convolutional Neural Networks).



LTE Operators and Convolution

- Given an input 2D signal $i(x,y)$, a 2D operator, $T\{\cdot\}$: $o(x, y) = T\{i(x, y)\}$, is said to be **Linear iff:**

$$T\{ai_1(x, y) + bi_2(x, y)\} = ao_1(x, y) + bo_2(x, y), \text{ with } o_1(\cdot) = T\{i_1(\cdot)\}, o_2(\cdot) = T\{i_2(\cdot)\}$$

with a, b any two constants.

- The operator is said to be **Translation-Equivariant** iff:

$$T\{i(x - x_0, y - y_0)\} = o(x - x_0, y - y_0)$$

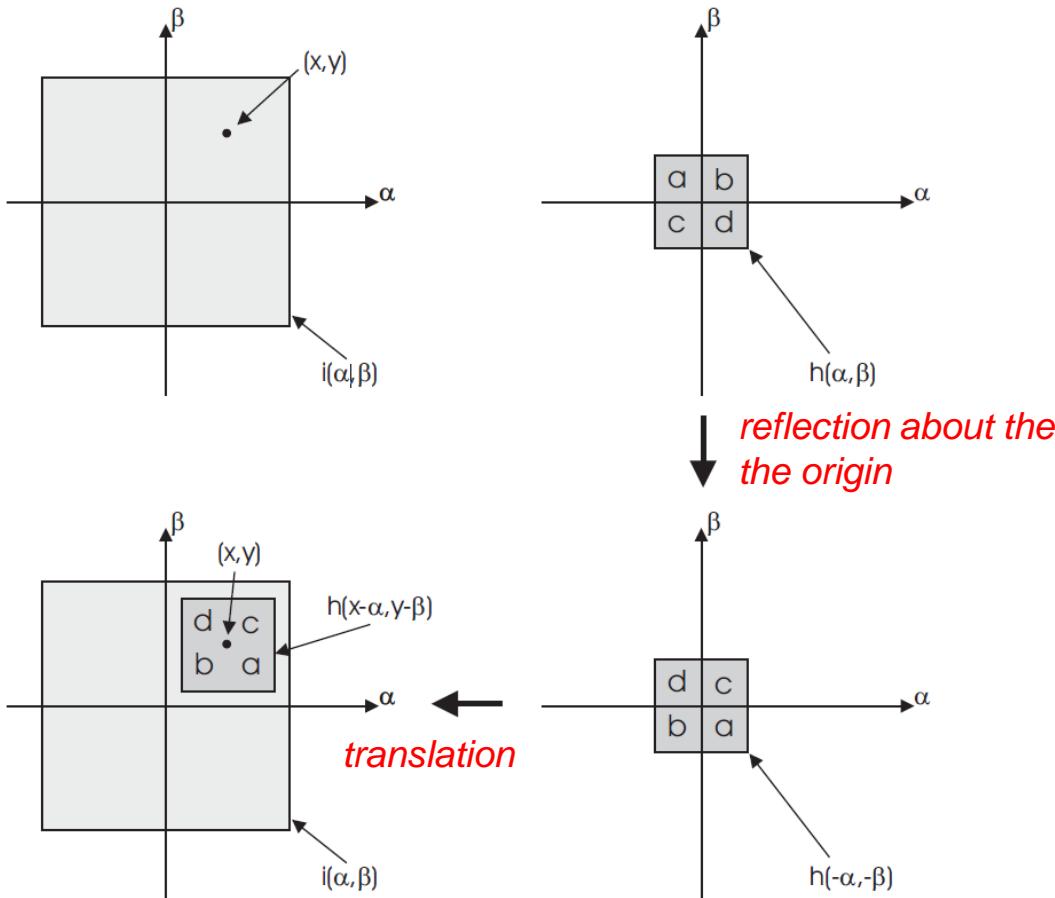
- If the operator is **LTE**, the output signal is given by the **convolution** between the **impulse response (point spread function)**, $h(x, y) = T\{\delta(x, y)\}$, of the operator and the input signal:

$$o(x, y) = T\{i(x, y)\} = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x - \alpha, y - \beta) d\alpha d\beta$$

Unit impulse
(Dirac delta
function)

A Graphical View of Convolution

$$o(x, y) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x - \alpha, y - \beta) d\alpha d\beta$$



Properties of Convolution



- We will often denote the convolution operation by the symbol “ * ”, e.g.

$$o(x, y) = i(x, y) * h(x, y)$$

- Some useful properties of convolution are as follows:

1. $f * (g * h) = (f * g) * h$ *(Associative Property)*
2. $f * g = g * f$ *(Commutative Property)*
3. $f * (g + h) = f * g + f * h$ *(Distributive Property wrt the Sum)*
4. $(f * g)' = f' * g = f * g'$ *(Convolution Commutes with Differentiation)*

Correlation

- The correlation of signal $i(x,y)$ wrt signal $h(x,y)$ is defined as:

$$i(x,y) \circ h(x,y) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x + \alpha, y + \beta) d\alpha d\beta$$

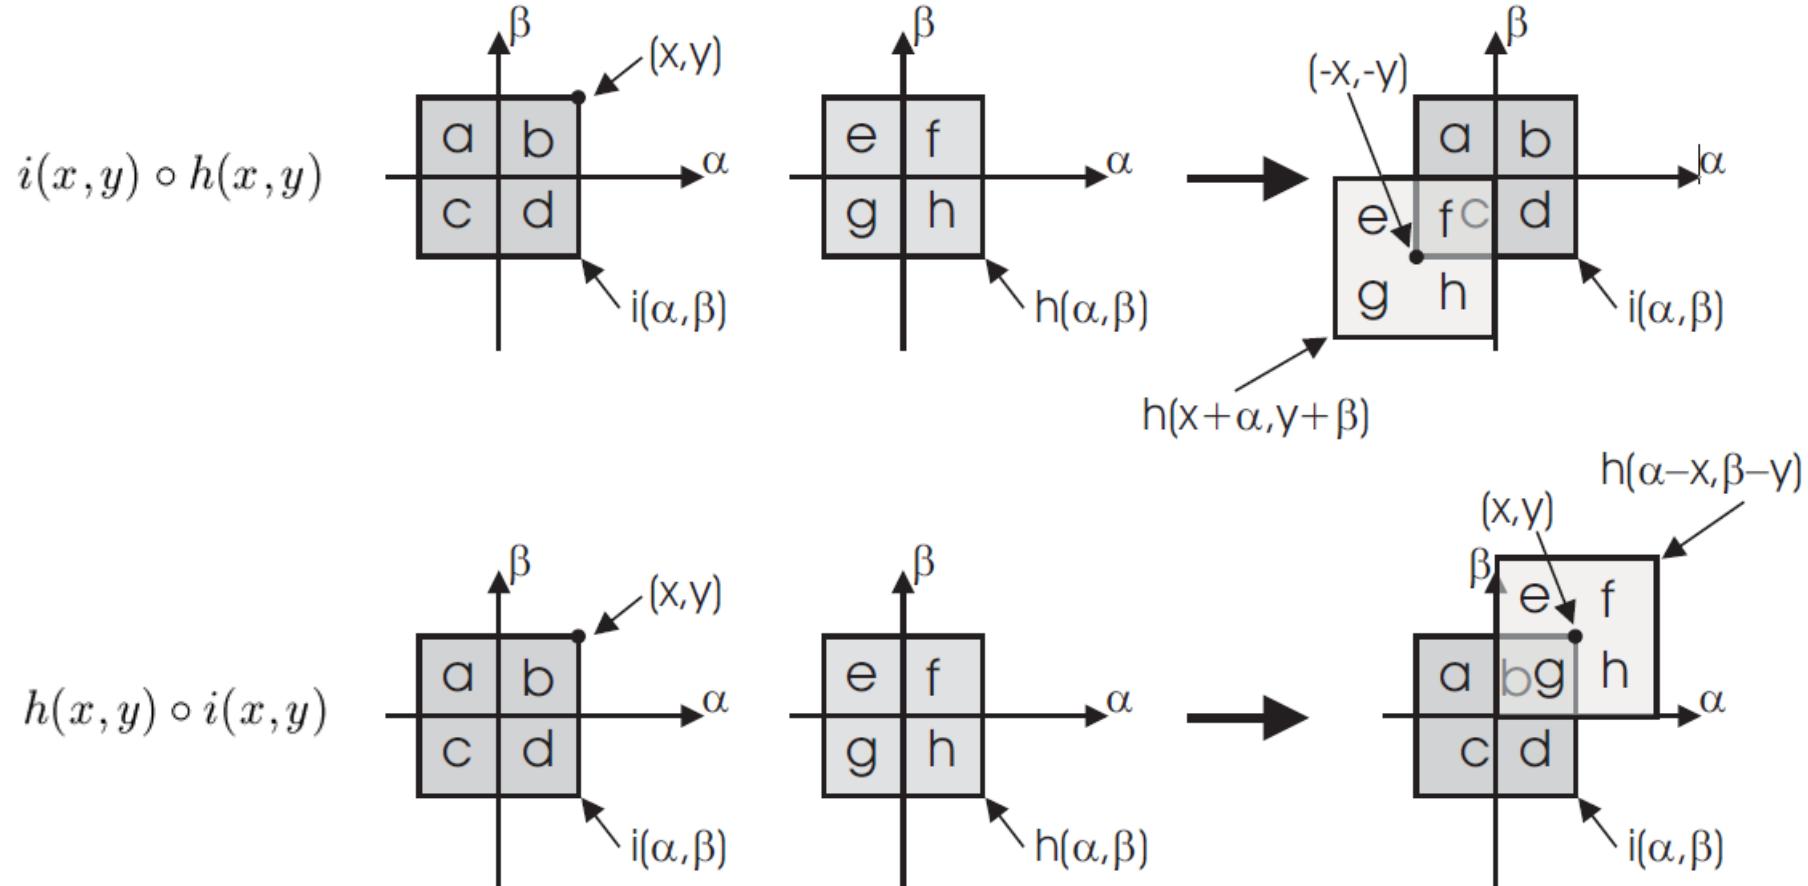
- Accordingly, the correlation of $h(x,y)$ wrt $i(x,y)$ is given by:

$$h(x,y) \circ i(x,y) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} h(\alpha, \beta) i(x + \alpha, y + \beta) d\alpha d\beta$$

- Unlike convolution, correlation is not commutative:

$$\begin{aligned} h(x,y) \circ i(x,y) &= \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} h(\alpha, \beta) i(x + \alpha, y + \beta) d\alpha d\beta \\ &= \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\xi, \eta) h(\xi - x, \eta - y) d\xi d\eta \\ &= \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(\alpha - x, \beta - y) d\alpha d\beta \\ &\neq i(x,y) \circ h(x,y) \end{aligned}$$

A Graphical View of Correlation



Convolution and Correlation



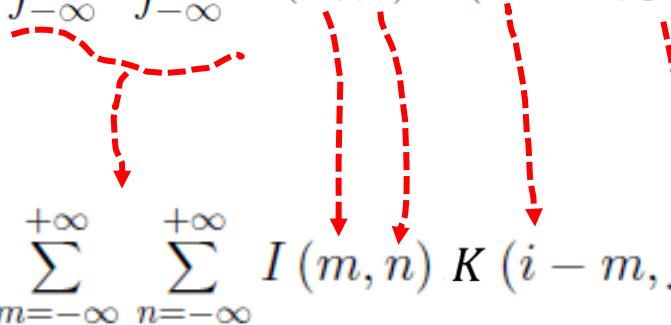
- The correlation of h wrt i is similar to convolution: the product of the two signals is integrated after translating h without reflection. Hence, if h is an even function ($h(x,y)=h(-x,-y)$), the convolution between i and h ($i^*h=h^*i$) is the same as the correlation of h wrt i :

$$\begin{aligned} i(x, y) * h(x, y) &= h(x, y) * i(x, y) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x - \alpha, y - \beta) d\alpha d\beta \\ &= \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(\alpha - x, \beta - y) d\alpha d\beta \\ &= h(x, y) \circ i(x, y) \end{aligned}$$

- It is worth observing that correlation is never commutative, even if h is an even function. To recap:
 - $i * h = h * i$ *(convolution is commutative)*
 - $i \circ h \neq h \circ i$ *(correlation is not commutative)*
 - $i * h = h * i = h \circ i$ *(if h is an even function)*

Discrete Convolution

$$o(x, y) = T\{i(x, y)\} = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x - \alpha, y - \beta) d\alpha d\beta$$


 $O(i, j) = T\{I(i, j)\} = \sum_{m=-\infty}^{+\infty} \sum_{n=-\infty}^{+\infty} I(m, n) K(i - m, j - n)$

where $I(i, j)$ and $O(i, j)$ are the discrete 2D input and output signals, respectively, and $K(i, j) = T\{\delta(i, j)\}$ is the **Kernel** of the discrete LTE operator, i.e. the response to the 2D discrete unit impulse (**Kronecker delta function**), $\delta(i, j)$.

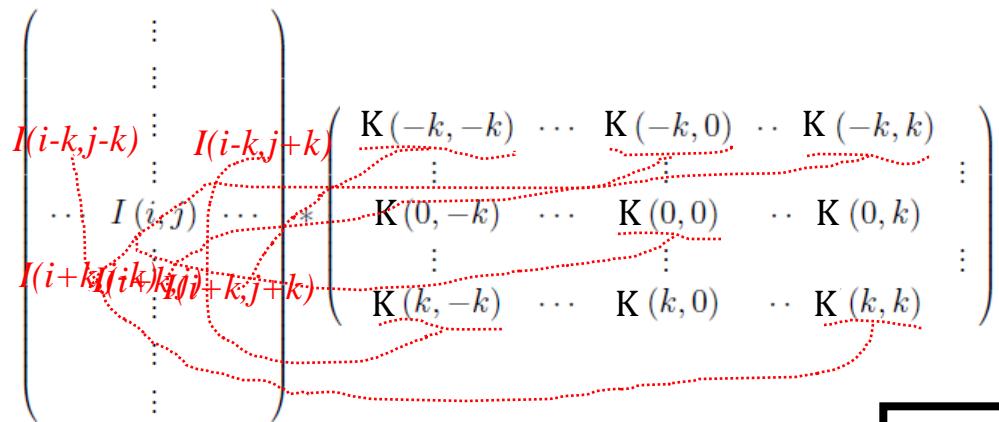
Akin to continuous signals, discrete convolution consists in summing the product of the two signals where one has been reflected about the origin and translated. The previously highlighted four major convolution properties hold for discrete convolution too.

Practical Implementation

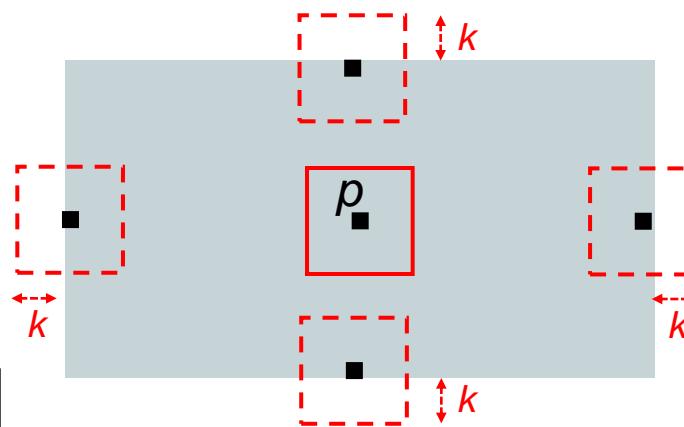
In image processing both the input image and the kernel are stored into matrixes of given *finite sizes*, with the image being much larger than the kernel. One would cycle through the kernel, thus:

$$O(i, j) = \sum \sum K(m, n)I(i - m, j - n)$$

Conceptually, to obtain the output image we need to slide the kernel across the whole input image and compute the convolution at each pixel (do not overwrite the input matrix !)



OpenCV: cv2.filter2D



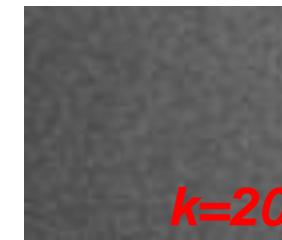
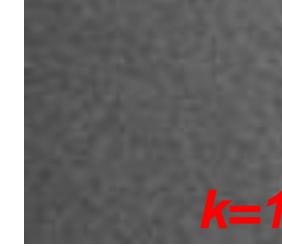
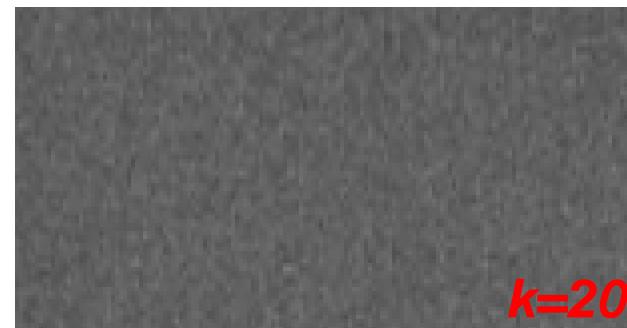
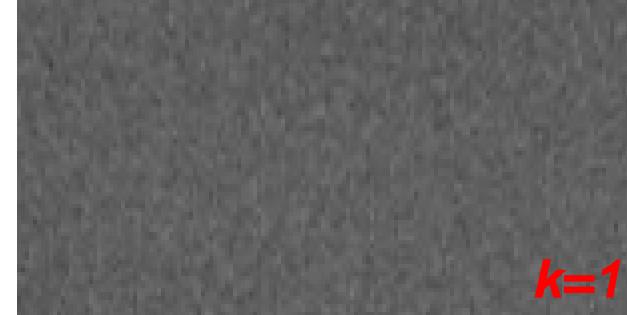
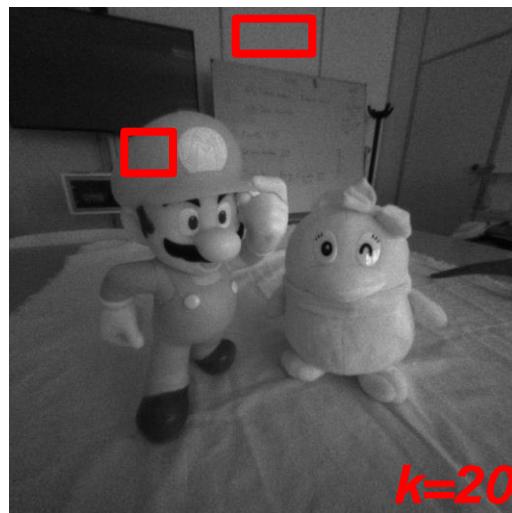
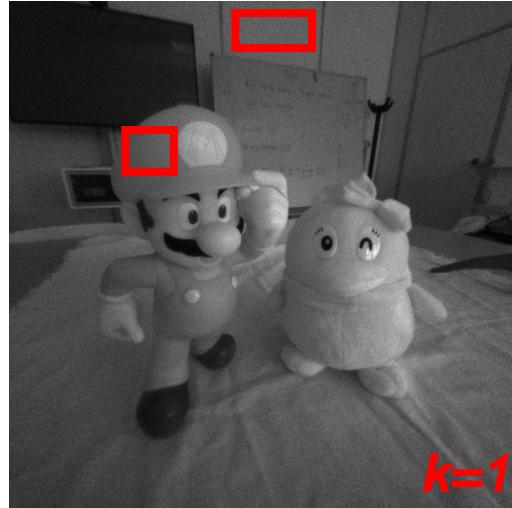
Border Issue. Two main options: CROP (common in image processing), PAD (preferred in CNNs). How to pad: zero-padding, replicate (aaala.....ddd), reflect (cbalabc.....dfglgfd), reflect_101 (dcblabcd....efghlgfe),

Visualizing and Understanding Noise



t_k

↓



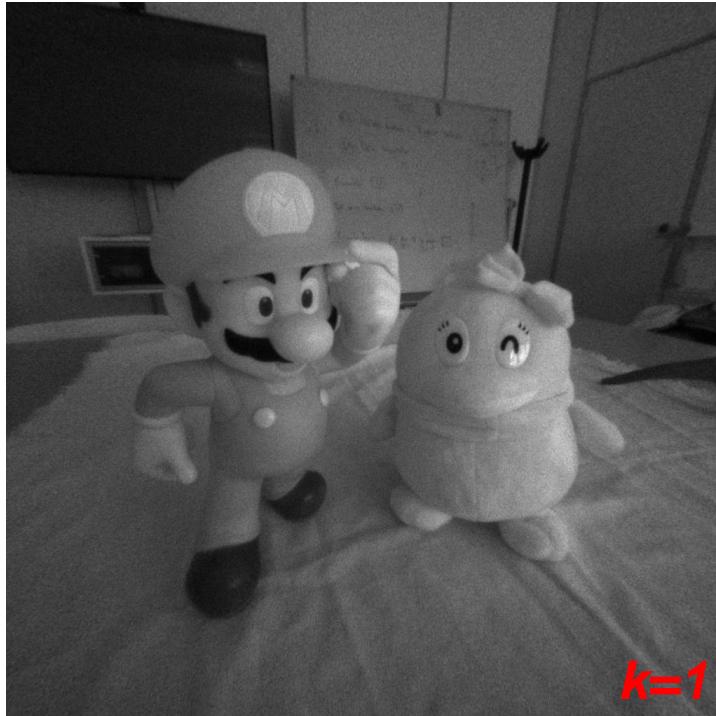
$$I_k(p) = \tilde{I}(p) + n_k(p)$$

with $n_k(p)$ i.i.d.
and $n_k(p) \sim N(0, \sigma^2)$

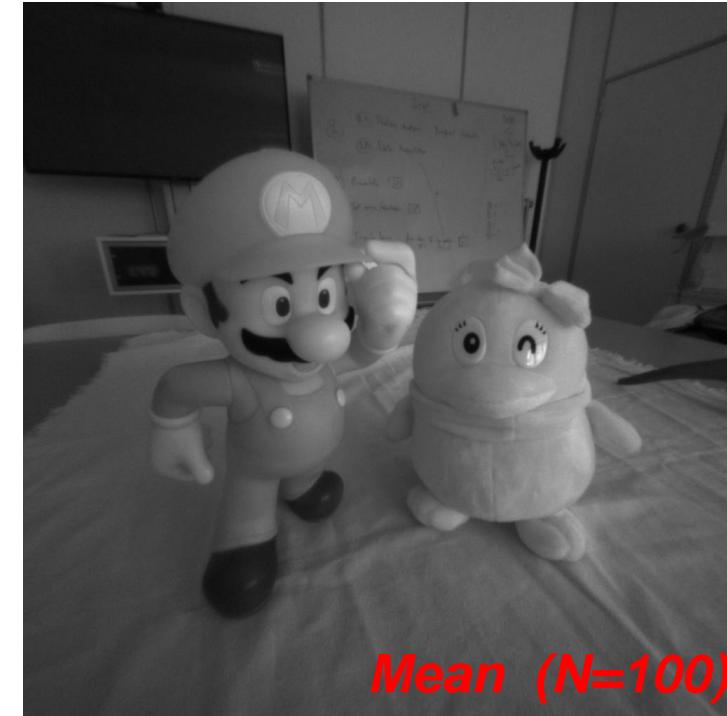
Denoising by taking a mean across time



$$O(p) = \frac{1}{N} \sum_{k=1}^N I_k(p) = \frac{1}{N} \sum_{k=1}^N (\tilde{I}(p) + n_k(p)) = \frac{1}{N} \sum_{k=1}^N \tilde{I}(p) + \frac{1}{N} \sum_{k=1}^N n_k(p) \Rightarrow N \nearrow : \cong \tilde{I}(p)$$



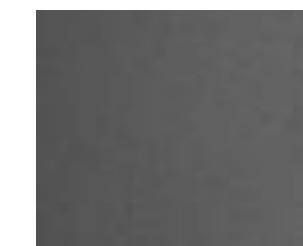
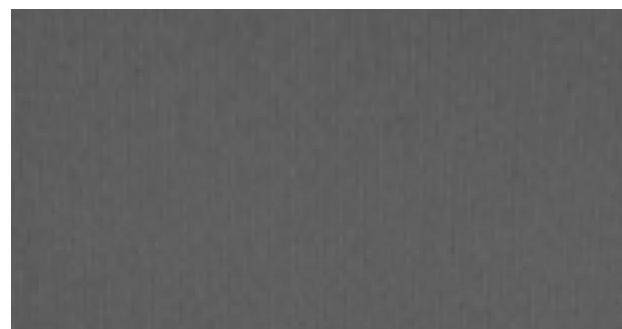
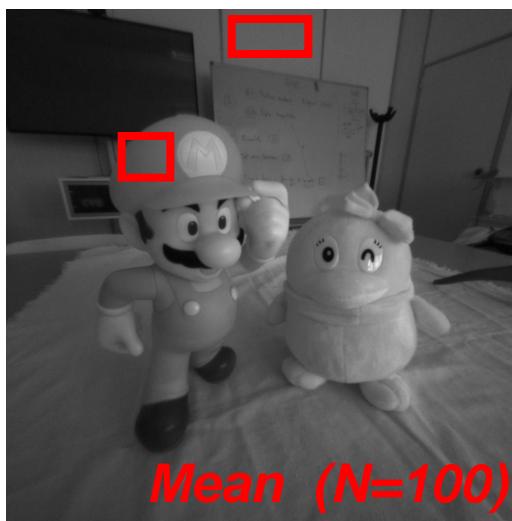
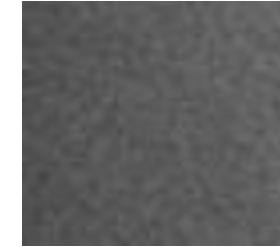
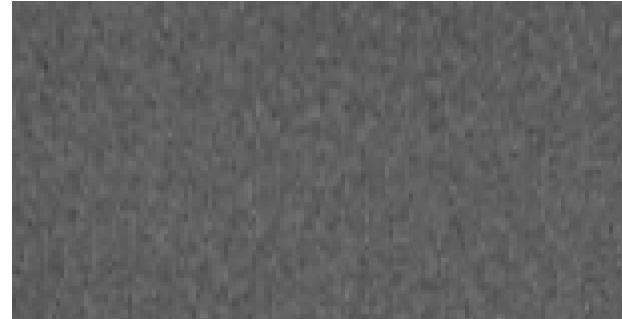
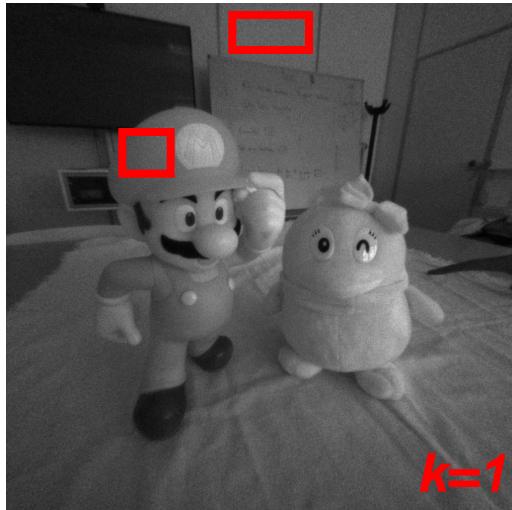
$k=1$



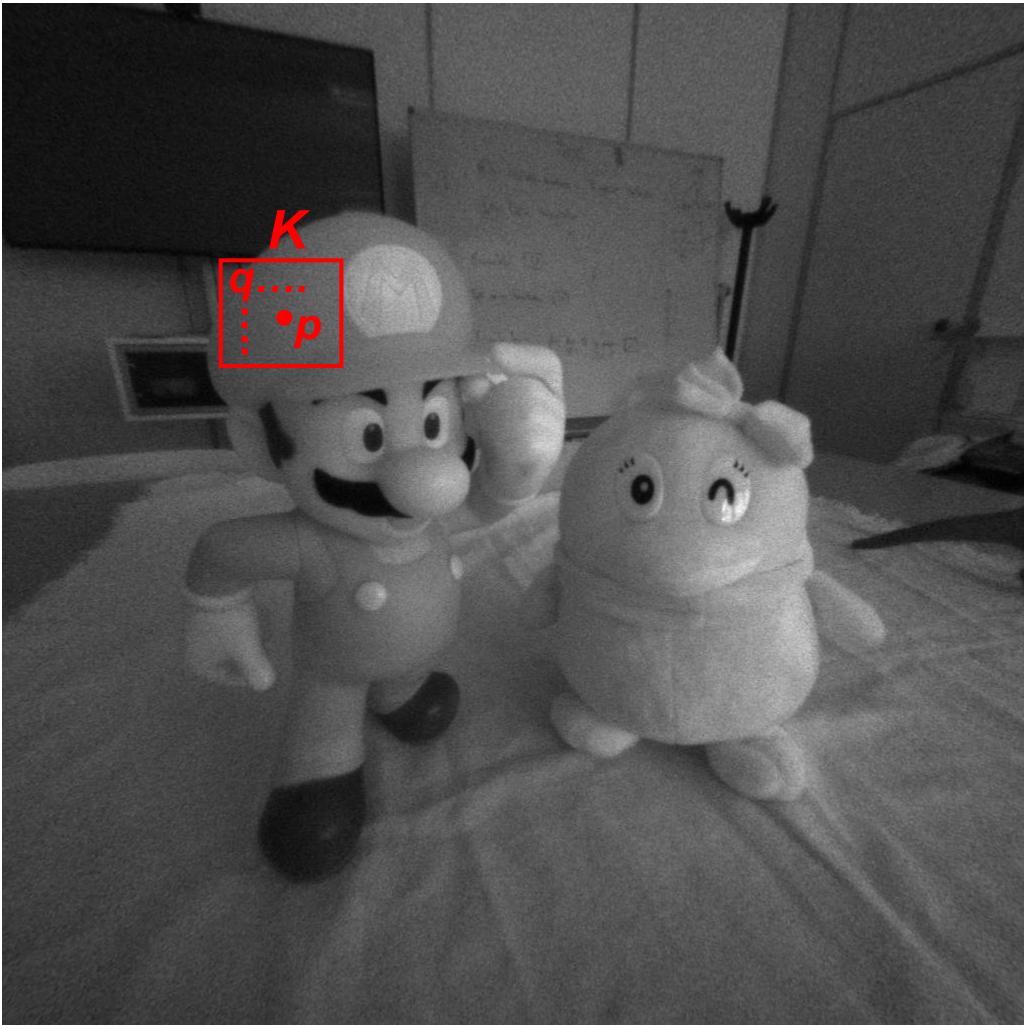
Mean (N=100)

$$\sigma_M^2 = \frac{\sigma^2}{N}$$

The mean image is far less noisy



What if we are given a single image ?



$$\begin{aligned}O(p) &= \frac{1}{|K|} \sum_{q \in K} I(q) \\&= \frac{1}{|K|} \sum_{q \in K} (\tilde{I}(q) + n(q)) \\&= \frac{1}{|K|} \sum_{q \in K} \tilde{I}(q) + \frac{1}{|K|} \sum_{q \in K} n(q) \\&\Rightarrow |K| \nearrow \wedge (\tilde{I}(q) = \tilde{I}(p) \forall q \in K) : \cong \tilde{I}(p)\end{aligned}$$

We may compute a mean across neighbouring pixels, i.e. a spatial rather than temporal mean.

↓

Denoising Filters

Mean Filter

- Mean filtering is the simplest (and fastest) way to denoise an image. It consists in replacing each pixel intensity by the average intensity over a chosen neighbourhood (e.g. 3x3, 5x5, 7x7...).
- The Mean Filter is an LTE operator as it can be expressed as a convolution with a kernel. Below, the kernels for a 3x3 and 5x5 mean filter:

$$\begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix} = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \\ 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \\ 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \\ 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \\ 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \end{bmatrix} = \frac{1}{25} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

OpenCV: cv2.blur

- According to signal processing theory, the Mean Filter carries out a ***low-pass filtering*** operation, which in image processing is also referred to as ***image smoothing***.
- Smoothing is often aimed at image denoising, though sometimes the purpose is to cancel out small-size unwanted details that might hinder the image analysis task. For example, in feature-based computer vision algorithms smoothing is key to create the so called ***scale-space***, which endows these approaches with ***scale invariance***.
- Mean filtering is inherently fast because multiplications are not needed. Moreover, it can be implemented very efficiently by incremental calculation schemes (***box-filtering***).

Denoising by a Mean Filter

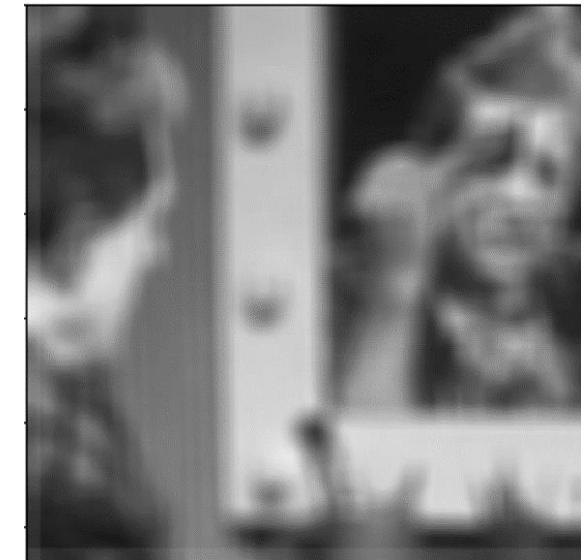
Original Image



Image corrupted by Noise



Smoothing by a Mean Filter



$$n_k(p) \sim \mathcal{N}(\mu = 0, \sigma = 10)$$

$$K = 7 \times 7$$

$$K = 13 \times 13$$

Mean filtering can reduce noise but blurs the image

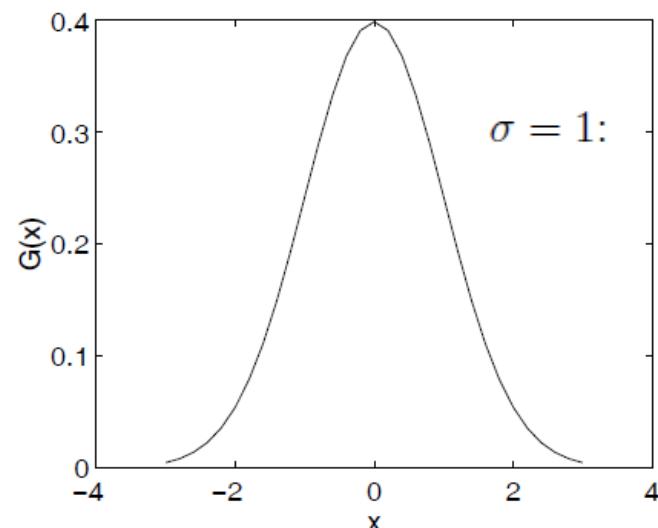
Gaussian Filter (1)



- LTE operator whose impulse response is a 2D Gaussian function (with zero mean and constant diagonal covariance matrix).

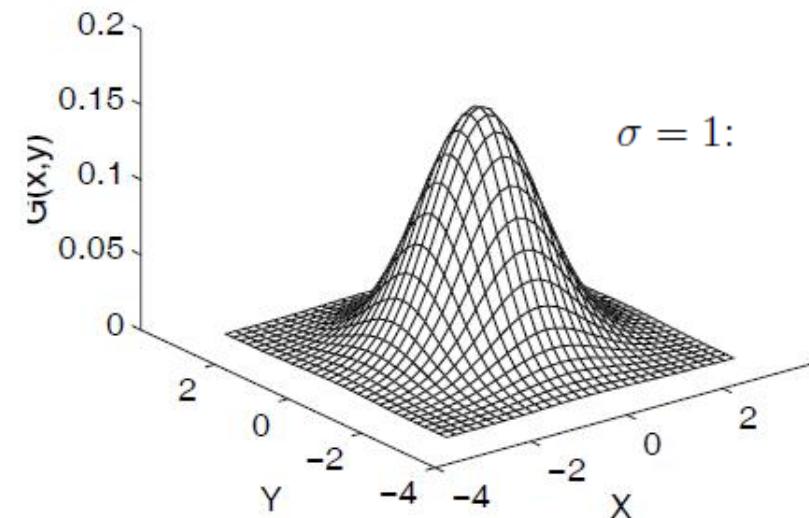
1D Gaussian

$$G(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}}$$



2D Gaussian

$$G(x, y) = G(x)G(y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



OpenCV: cv2.GaussianBlur

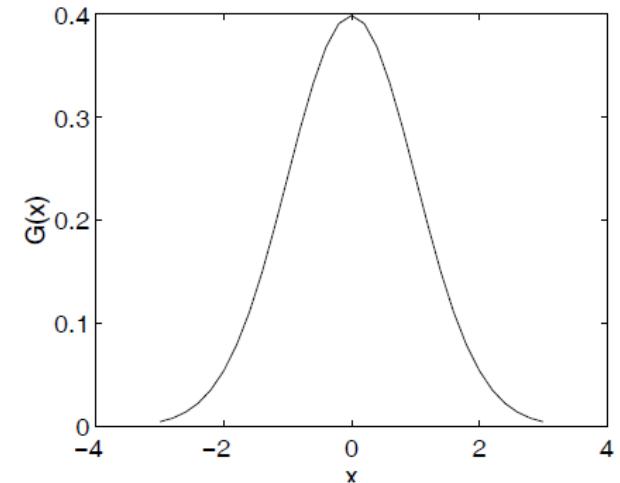
Circularly Symmetric

Practical Implementation



- The discrete Gaussian kernel can be obtained by sampling the corresponding continuous function, which is however of infinite extent. A finite size must therefore be properly chosen.

- To this purpose, we can observe that:
 - ✓ The larger is the size, the more accurate turns out the discrete approximation of the ideal continuous filter.
 - ✓ The computational cost grows with filter size.
 - ✓ The Gaussian gets smaller and smaller as we move away from the origin.



- Therefore, we should use larger sizes for filters with high σ , smaller sizes whenever σ is smaller. A *rule-of-thumb* to chose the size of the filter given σ would then be quite useful.

- As the interval $[-3\sigma, +3\sigma]$ captures 99% of the area ("energy") of the Gaussian function, a typical *rule-of-thumb* dictates taking a $(2k+1) \times (2k+1)$ kernel with:

$$k = \lceil 3\sigma \rceil$$

$\sigma = 1$	\Rightarrow	7×7
$\sigma = 1.5$	\Rightarrow	11×11
$\sigma = 2$	\Rightarrow	13×13
$\sigma = 3$	\Rightarrow	19×19

Mean vs. Gaussian

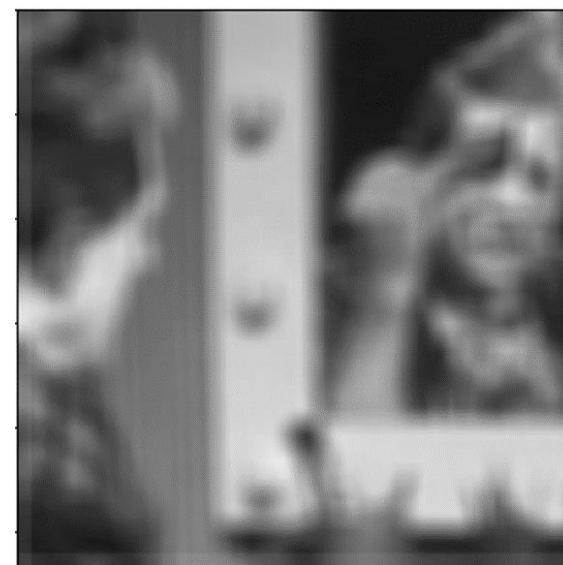


Image corrupted by Noise



$$n_k(p) \sim \mathcal{N}(\mu = 0, \sigma = 10)$$

Smoothing by a Mean Filter



$$K = 13 \times 13$$

Smoothing by a Gaussian Filter



$$K = 13 \times 13$$

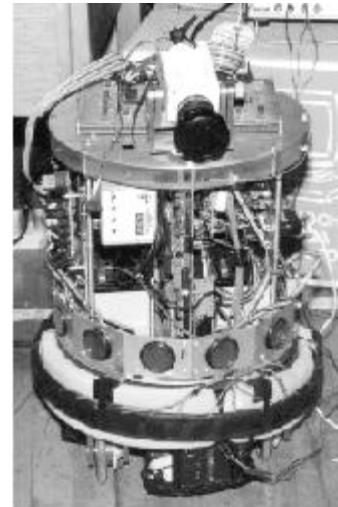
Linear filtering can reduce noise but blurs the image.

Given the same kernel size, Gaussian Filtering yields less blur than Mean Filtering.

Parameter σ sets the amount of smoothing



Original Image

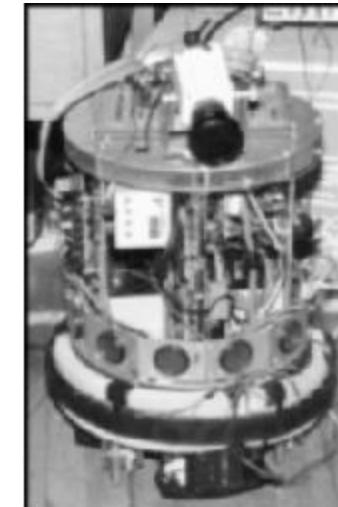


The higher σ , the stronger the smoothing caused by the filter. This can be understood, e.g., by observing that as σ increases, the weights of closer points get smaller while those of farther points get larger.

Smoothing by
a Gaussian
Filter with $\sigma=2$



Smoothing by
a Gaussian Filter
with $\sigma = 1$



As σ gets larger, small details disappear and the image content deals with larger size structures. Thus, filtering with a chosen σ can be thought of as setting the “scale” of interest to analyse image content.



Smoothing by
a Gaussian Filter
with $\sigma = 4$

Impulse Noise



Original Image



Smoothing by
a 3x3 Mean



Image corrupted
by Impulse Noise
(aka Salt-and-
Pepper Noise)



Linear filtering is ineffective toward impulse noise (and blurs the image)



Smoothing by
a 5x5 Mean

Median Filter

- ***Non-linear filter*** whereby each pixel intensity is replaced by the *median* over a given neighbourhood, the *median* being the value falling half-way in the sorted set of intensities.

123	125	126	130	140
122	124	126	127	135
118	120	150	125	134
119	115	119	123	133
111	116	110	120	130

Neighbourhood values:

115, 119, 120, 123, 124,
125, 126, 127, 150

Median value: 124

OpenCV: cv2.Medianblur

$$\text{median } [A(x) + B(x)] \neq \text{median } [A(x)] + \text{median } [B(x)]$$

- Median filtering counteracts impulse noise effectively, as *outliers* (i.e. noisy pixels) tend to fall at either the top or bottom end of the sorted intensities.
- Median filtering tends to keep sharper edges than linear filters such as the Mean or Gaussian:

$$\dots 10 \ 10 \ 40 \ 40 \ \dots \Rightarrow \dots 10 \ 20 \ 30 \ 40 \ \dots \ (\text{Mean})$$

$$\Rightarrow \dots 10 \ 10 \ 40 \ 40 \ \dots \ (\text{Median})$$

Example

Original Image



Image Corrupted by Impulse Noise (0.05)



Filtering by a 3x3 Median

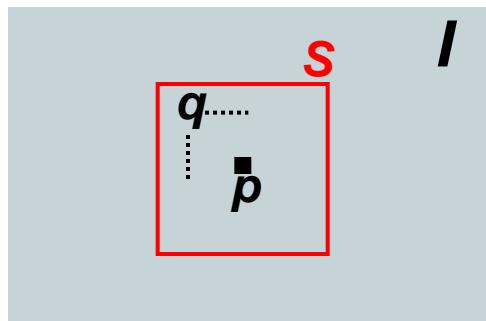


When dealing with impulse noise, the *Median Filter* can effectively denoise the image without introducing significant blur.

Yet, Gaussian-like noise, such as sensor noise, cannot be dealt with by the Median, as this would require computing new noiseless intensities. Purposely, the *Median* may be followed by *linear filtering*.

Bilateral Filter (1)

- Advanced *non-linear filter* to accomplish denoising of Gaussian-like noise without blurring the image (aka *edge preserving smoothing*).



$$O(p) = \sum_{q \in S} H(p, q) \cdot I_q$$

$H(p, q) = \frac{1}{W(p)} G_{\sigma_s}(d_s(p, q)) G_{\sigma_r}(d_r(I_p, I_q))$

OpenCV: cv2.bilateralFilter

$$d_s(p, q) = \|p - q\|_2 = \sqrt{(u_p - u_q)^2 + (v_p - v_q)^2} \longrightarrow \text{Spatial Distance}$$

$$d_r(I_p, I_q) = |I_p - I_q| \longrightarrow \text{Range (Intensity) Distance}$$

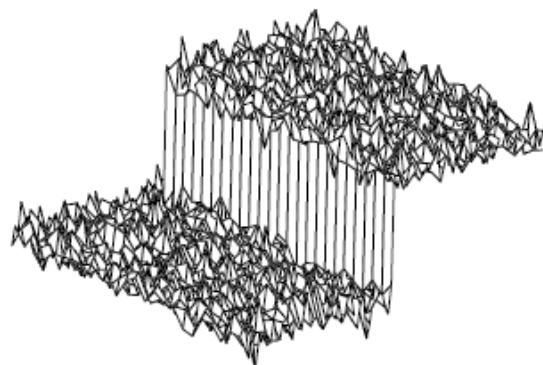
$$W(p) = \sum_{q \in S} G_{\sigma_s}(d_s(p, q)) G_{\sigma_r}(d_r(I_p, I_q)) \longrightarrow \text{Normalization Factor (Unity Gain)}$$

Bilateral Filter (2)

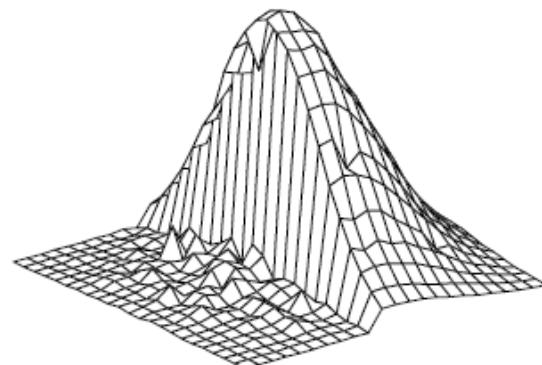
**Step-edge as wide
as 100 gray-levels**

**$H(p,q)$ at a pixel just across
the edge
in the brighter region**

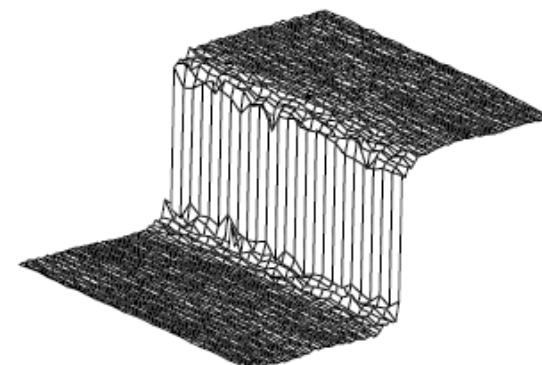
Output provided by the filter
 $\sigma_s = 5, \sigma_r = 50$



(a)



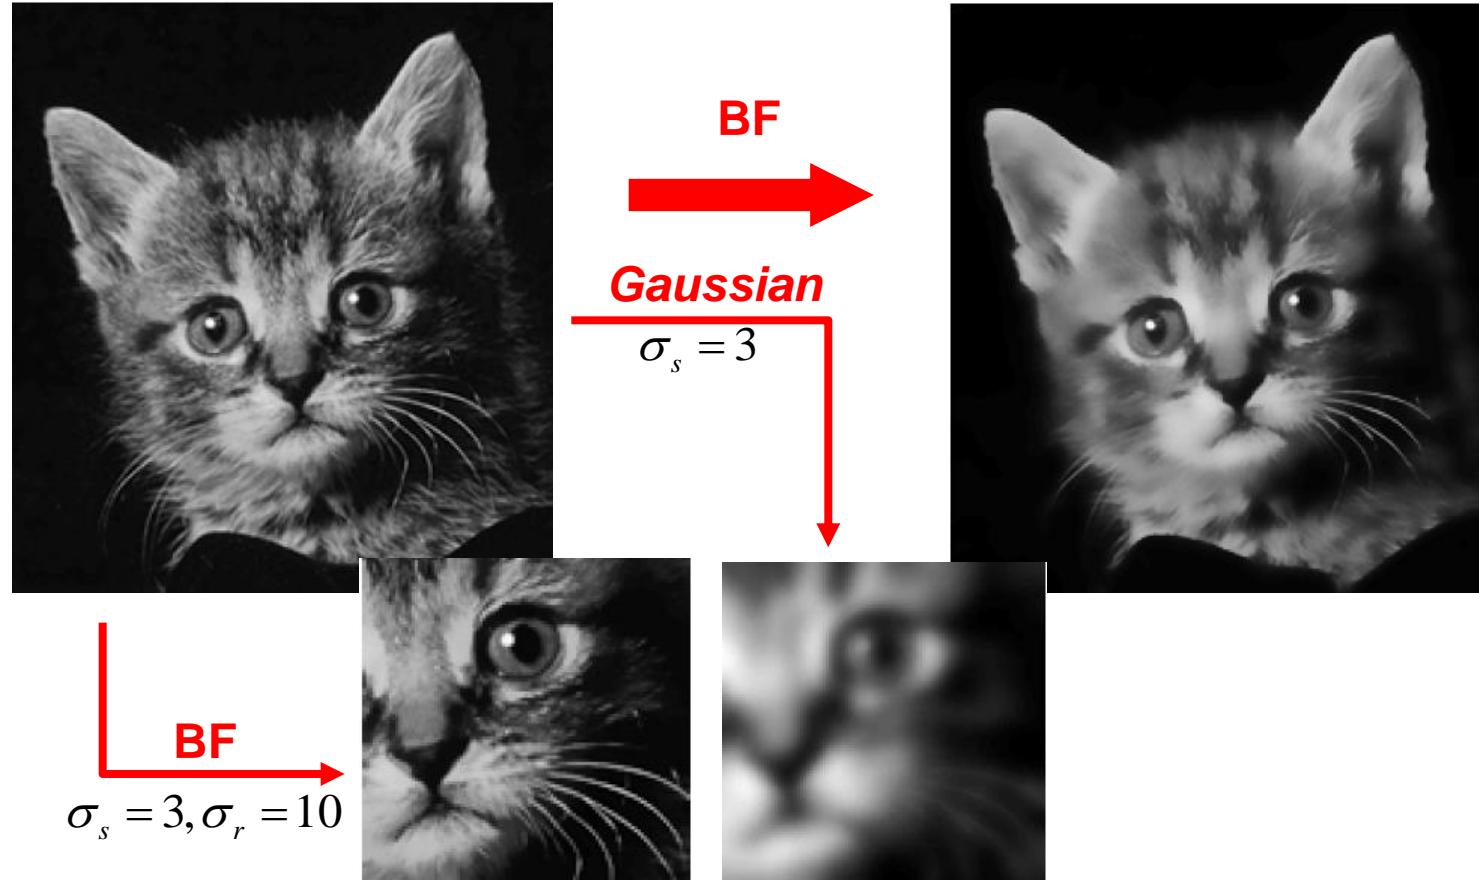
(b)



(c)

- Given the supporting neighbourhood, neighbouring pixels take a larger weight as they are both closer and more similar to the central pixel.
- At a pixel nearby an edge, the neighbours falling on the other side of the edge look quite different and thus cannot contribute significantly to the output value due to their weights being small.

Bilateral Filter (3)



More examples at:

http://people.csail.mit.edu/sparis/siggraph07_course/

Guided Filter

<https://kaiminghe.github.io/eccv10/index.html>

Mean vs. Gaussian vs. Bilateral

Image corrupted by Noise



$n_k(p) \sim \mathcal{N}(\mu = 0, \sigma = 10)$

Smoothing by a Mean Filter



$K = 13 \times 13$

Smoothing by a Gaussian Filter



$K = 13 \times 13$

Smoothing by a Bilateral Filter



$K = 13 \times 13$

Thanks to its weight function, the Bilateral Filter can accomplish denoising without blurring the image (aka *edge preserving smoothing*).

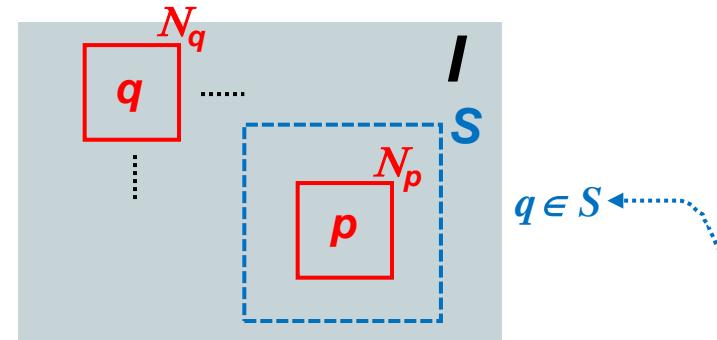
Non-local Means Filter (1)



- Another well-known *non-linear edge preserving smoothing filter*. The key idea is that the similarity among patches spread over the image can be deployed to achieve denoising.



OpenCV: cv2.fastNIMeansDenoising



$$O(p) = \sum_{q \in I} w(p, q)I(q)$$

$$w(p, q) = \frac{1}{Z(p)} e^{-\frac{\|N_p - N_q\|_2^2}{h^2}}$$

$$Z(p) = \sum_{q \in I} e^{-\frac{\|N_p - N_q\|_2^2}{h^2}}$$

Typical choices: $N=7 \times 7$,
 $S=21 \times 21$, $h = 10 \cdot \sigma$

Non-local Means vs. Bilateral

Image corrupted by Noise



Smoothing by a Bilteral Filter



$K = 13 \times 13$

Smoothing by a NLM Filter



$K = 13 \times 13$



Main References

- 1) V. S. Nalwa, "A Guided Tour of Computer Vision", Addison-Wesley Publishing Company, 1993.
- 2) R. Fisher, S. Perkins, A. Walker, E. Wolfart, "Hypermedia Image Processing Reference", Wiley, 1996 (<http://homepages.inf.ed.ac.uk/rbf/HIPR2/>)
- 3) R. Schalkoff, "Digital Image Processing And Computer Vision, Wiley, 1989.
- 4) C. Tomasi, R. Manduchi "Bilateral Filtering for Gray and Color Images", ICCV 1998.
- 5) S.Paris, P. Kornprobst, J. Tumblin, F. Durand "A Gentle Introduction to Bilateral Filtering and its Applications" (SIGGRAPH 2008,CVPR 2008, SIGGRAPH 2007)
http://people.csail.mit.edu/sparis/siggraph07_course/
- 6) A. Buades, B. Coll, J.M. Morel, "A non-local algorithm for image denoising", CVPR 2005.
- 7) Kaiming He, Jian Sun, and Xiaou Tang, "Guided Image Filtering", IEEE Trans. On PAMI, 2013