

University of Bologna

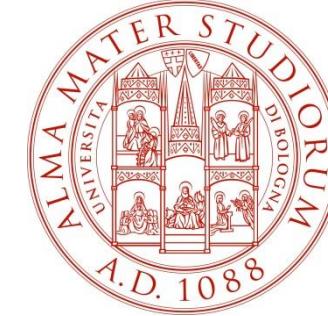


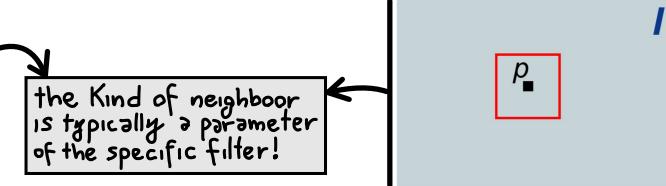
Image Filtering

Luigi Di Stefano (luigi.distefano@unibo.it)

Introduction



- **Image Filters** are image processing operators that compute the new intensity (colour) of a pixel, p , based on the intensities (colours) of those belonging to a neighbourhood of p .



- They accomplish a variety of useful image processing functions, such as e.g. denoising and sharpening (edge enhancement).

In the image processing context filters are typically chosen by the designer, meanwhile in computer vision and in particular in CNN they are present (and their presence is fundamental) BUT they're automatically learnt and NOT pre-set; nevertheless the way they behave and are implemented is of course always the same!

- An important sub-class of filters is given by **Linear and Translation-Equivariant (LTE)** operators, which we will consider first.

What are filters used for? A filter (in particular in image processing) has the purpose to "improve" an image in its quality accordingly to a certain way and purpose! Also, for example, in CV they are used to alter images into new ones that are more prone to be extracted information from!

- Signal theory dictates their application in image processing to consist in a **2D convolution** between the input image and the *impulse response function (point spread function or kernel)* of the LTE operator.

- LTE operators are used as feature extractors in CNNs (Convolutional Neural Networks).

It is known from signal theory that when a 2D input signal (as an image is) is given to an LTE filter (aka a linear and translation-equivariant operator) THEN its behavior is implemented through the usage of the convolution operator AND the impulse response of the filter itself!

LTE Operators and Convolution

- Given an **input 2D signal $i(x,y)$** , a **2D operator, $T\{\cdot\}$** : $\widehat{o(x,y)} = T\{i(x,y)\}$, is said to be **Linear iff:**

$$T\{ai_1(x,y) + bi_2(x,y)\} = ao_1(x,y) + bo_2(x,y), \text{ with } o_1(\cdot) = T\{i_1(\cdot)\}, o_2(\cdot) = T\{i_2(\cdot)\}$$

INDEED, LINEARITY MEANS SUPERPOSITION OF EFFECTS DO HOLD!!

"o" stands for "output" and $o(x,y)$ is the function obtained applying the T operator to the "input" function $i(x,y)$ (BOTH of the input and the output are 2D funct)

with a,b any two constants.

- The operator is said to be **Translation-Equivariant** iff:

$$T\{i(x - x_0, y - y_0)\} = o(x - x_0, y - y_0)$$

- If the operator is **LTE**, the output signal is given by the **convolution** between the **impulse response (point spread function)**, $h(x,y) = T\{\delta(x,y)\}$, of the operator and the input signal:

ANOTHER NAME FOR THE
"IMPULSE RESPONSE"

$$o(x,y) = T\{i(x,y)\} = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x - \alpha, y - \beta) d\alpha d\beta$$

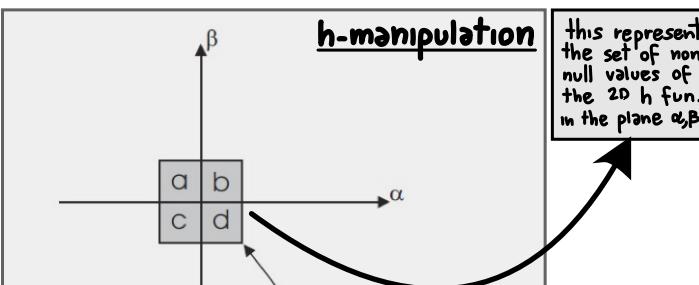
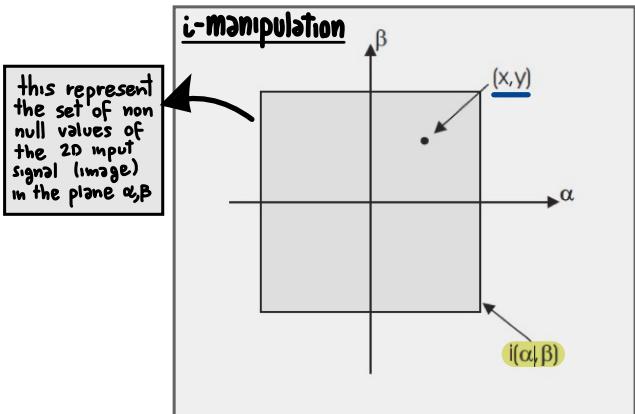
**Unit impulse
(Dirac delta
function)**

REMARK: this is the one peculiar function $\delta(x,y)$ which is non-null only in the origin AND for which it holds

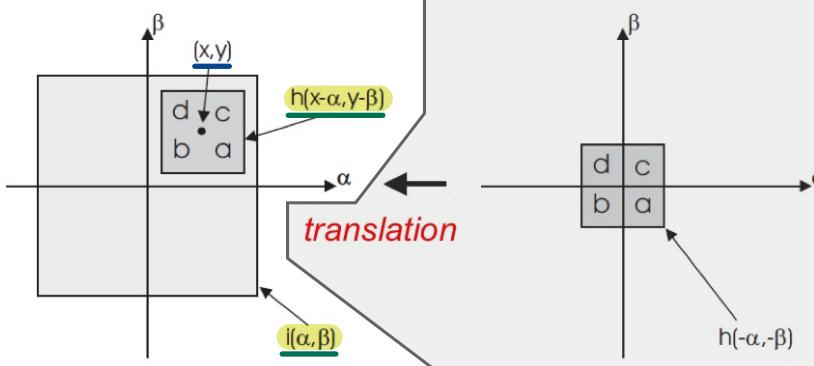
$$\int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} \delta(\alpha, \beta) d\alpha d\beta = 1$$

A Graphical View of Convolution

$$o(x, y) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x - \alpha, y - \beta) d\alpha d\beta$$



In practice, we'll see cases in which h is defined as way smaller in its cardinality w.r.t. i (notice that also h is a 2D functional typically not-null only in a specific region around the origin).



Indeed, convolution is about **MAD multiply and add**, also called **MAC multiply and cumulate**: given i and h , the convolution of i w.r.t. h (2D signals) in x, y takes each value of i and multiplies it by the corresponding value of h' ; then cumulating in a sum all of these contributes, where h' is obtained by translating to (x, y) the reflection of h about the origin!
 For that reason it's typically said that **MAD\MAC** is the basic operation behind the convolutional product!

Properties of Convolution



- We will often denote the convolution operation by the symbol “ * ”, e.g.

$$o(x, y) = i(x, y) * h(x, y)$$

- Some useful properties of convolution are as follows:

1. $f * (g * h) = (f * g) * h$ (Associative Property)
2. $f * g = g * f$ (Commutative Property)
3. $f * (g + h) = f * g + f * h$ (Distributive Property wrt the Sum)
4. $(f * g)' = f' * g = f * g'$ (Convolution Commutes with Differentiation)
to be intended as a derivative of any order

Correlation

- The correlation of signal $i(x,y)$ wrt signal $h(x,y)$ is defined as:

$$i(x,y) \circ h(x,y) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x + \alpha, y + \beta) d\alpha d\beta$$

- Accordingly, the correlation of $h(x,y)$ wrt $i(x,y)$ is given by:

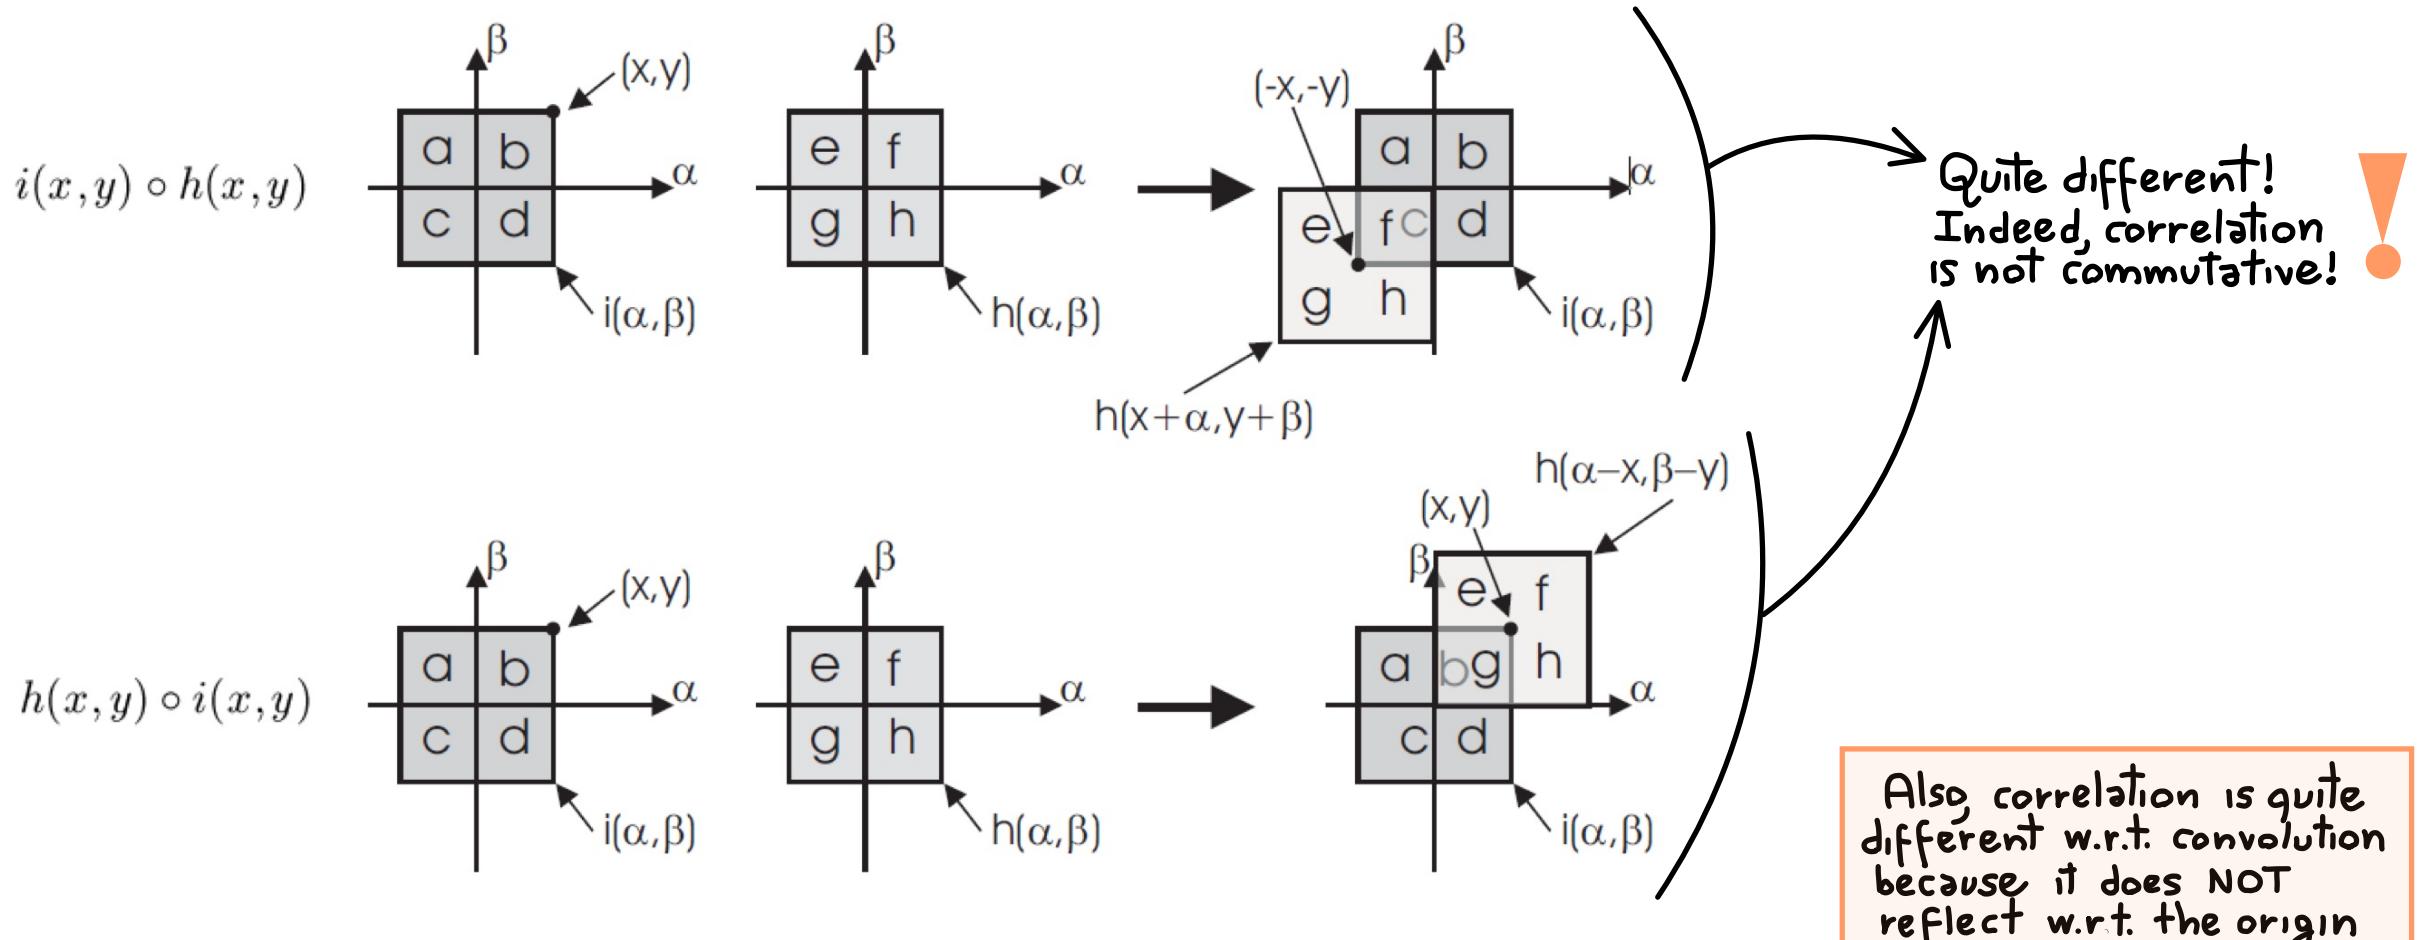
$$h(x,y) \circ i(x,y) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} h(\alpha, \beta) i(x + \alpha, y + \beta) d\alpha d\beta$$

- Unlike convolution, correlation is not commutative:

$$\begin{aligned}
 h(x,y) \circ i(x,y) &= \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} h(\alpha, \beta) i(x + \alpha, y + \beta) d\alpha d\beta \\
 &\stackrel{x+\alpha=\xi}{=} \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\xi, \eta) h(\xi - x, \eta - y) d\xi d\eta \\
 &\stackrel{\text{red dashed box}}{=} \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(\alpha - x, \beta - y) d\alpha d\beta \\
 &\neq i(x,y) \circ h(x,y)
 \end{aligned}$$

$x+\alpha=\xi$; $y+\beta=\eta$

A Graphical View of Correlation



Convolution and Correlation



- The correlation of h wrt i is similar to convolution: the product of the two signals is integrated after translating h without reflection. Hence, if h is an even function ($h(x,y)=h(-x,-y)$), the convolution between i and h ($i*h=h*i$) is the same as the correlation of h wrt i :

$$\begin{aligned} i(x, y) * h(x, y) &= h(x, y) * i(x, y) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x - \alpha, y - \beta) d\alpha d\beta \\ &= \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(\alpha - x, \beta - y) d\alpha d\beta \\ &= h(x, y) \circ i(x, y) \end{aligned}$$

→ If the first term is an even function, then convolution and correlation are equivalent!
 $\gamma * \pi \text{ eq. } \gamma \circ \pi \Leftrightarrow \gamma \text{ even}$!

- It is worth observing that correlation is never commutative, even if h is an even function. To recap:

1. $i * h = h * i$ (convolution is commutative)

2. $i \circ h \neq h \circ i$ (correlation is not commutative)

3. $i * h = h * i = h \circ i = i \circ h$ (if h is an even function)

Discrete Convolution

notice: in a discrete setting, the delta function $\delta(x,y)$ is zero everywhere except for the origin in which it takes value exactly one, cause they must hold:

$$\sum_{K_1=-\infty}^{+\infty} \sum_{K_2=-\infty}^{+\infty} \delta(K_1, K_2) = 1$$

$$\delta(K_1, K_2) = 0 \quad \forall (K_1, K_2) \neq (0,0)$$

$$o(x, y) = T \{ i(x, y) \} = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} i(\alpha, \beta) h(x - \alpha, y - \beta) d\alpha d\beta$$

$$O(i, j) = T \{ I(i, j) \} = \sum_{m=-\infty}^{+\infty} \sum_{n=-\infty}^{+\infty} I(m, n) K(i - m, j - n)$$

Why "K"? Because in the DISCRETE SETTINGS typically the input response function is actually called KERNEL!

Notice that by its own definition the Kernel is composed of constant predetermined values!

- where $I(i,j)$ and $O(i,j)$ are the discrete 2D input and output signals, respectively, and $K(i,j) = T\{\delta(i,j)\}$ is the **Kernel** of the discrete LTE operator, i.e. the response to the 2D discrete unit impulse (Kronecker delta function), $\delta(i,j)$.

- Akin to continuous signals, discrete convolution consists in summing the product of the two signals where one has been reflected about the origin and translated. The previously highlighted four major convolution properties hold for discrete convolution too.

Practical Implementation

- In image processing both the **input image** and the **kernel** are stored into matrixes of given **finite sizes**, with the **image being much larger than the kernel**. One would cycle through the kernel, thus:

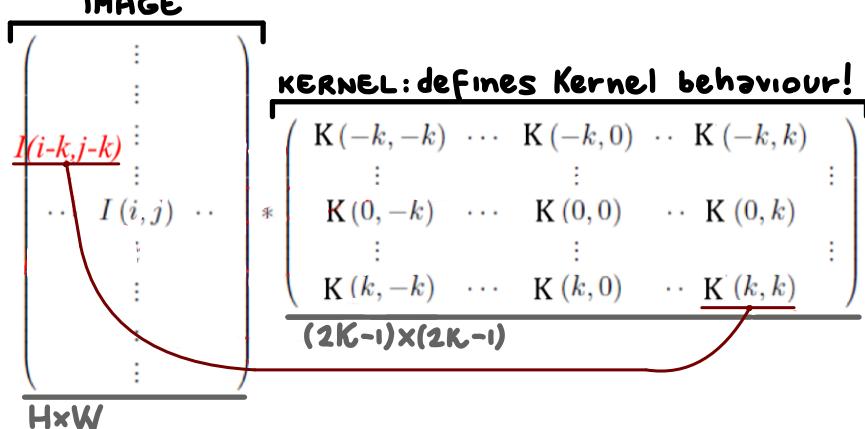
$$|I| = HW$$

$$|K| = (2K+1)^2$$

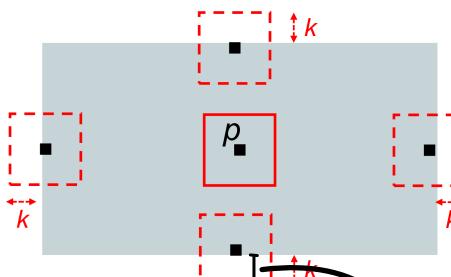
$$|I| \gg |K|$$

RECALL THAT CONVOLUTION IS COMMUTATIVE

$$O(i, j) = \sum_m \sum_n K(m, n) I(i - m, j - n)$$



Conceptually, to obtain the output image we need to slide the kernel across the whole input image and compute the convolution at each pixel (do not overwrite the input matrix !)



Practically, given a pixel $(i, j) \in I$, we simply iterate through ALL the Kernel with indexes m, n in ranges $[K, K]$ and for each Kernel element $K(m, n)$ we go pick the related correct element in the image which is $I(i - m, j - n)$. VISUALLY WE HAVE THE KERNEL OF THE FILTER SLIDING ACCROSS THE WHOLE IMAGE: this is called **SLIDING WINDOW APPROACH** and it's how typically the filter behaviour (aka its convolutional product) is conceptualized!

BE AWARE: for image pixels that are in the first\last K rows\columns, the Kernel is in part outside the image, so we CANNOT immediately compute the convolution for them! Various solution for this problem can be adopted: in CROPPING we ignore these pixels and we produce an output image which is smaller w.r.t. the input one, losing its peripheral part; in PADDING we properly ENLARGE the input image BEFORE filtering accordingly to a predetermined rule with the proper amount of pixels!

Border Issue. Two main options: **CROP** (common in image processing), **PAD** (preferred in CNNs). How to pad: zero-padding, replicate (aaala.....ddd), reflect (cbaabc.....dfgfgfd), reflect_101 (dcblabcd.....efghfgfe),

OpenCV: cv2.filter2D

ATTENTION: openCV.filter2D uses correlation, NOT convolution; to use it to implement an LTE filter, just pass to it the ORIGIN-REFLECTION of the kernel!

A cool remark: the importance of the associative property

Let's suppose to have f, w 2D signals where f is $H \times W$ and w is $K \times K$ (where these dimensions indicate the dimension of the portion of the plane in which these signal are supposed to be not null), where f cardinality is way bigger w.r.t. w cardinality ($|f| = HW \gg |w| = K^2$). If we consider the convolutional product $f * w$ then we are in a typical image processing situation of filtering image f through Kernel w .

If we want to compute $f * w$, the amount of DAC to be done is $n_1 = HWK^2$ (indeed we do K^2 MACs for each one of the HW f pixels) and the output (under PADDING hypothesis for borders management) is $H \times W$ (still a 2D signal of course)

Let's also suppose that w is such that it can be written as $w = g * h$, where g is $K \times 1$ and h is $1 \times K$.

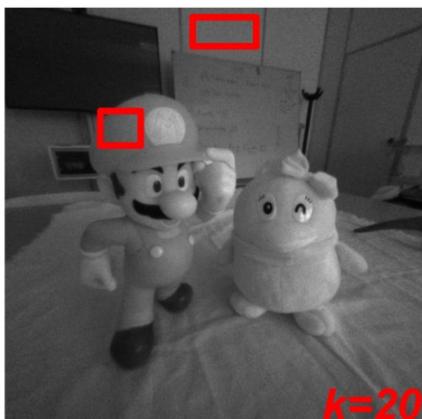
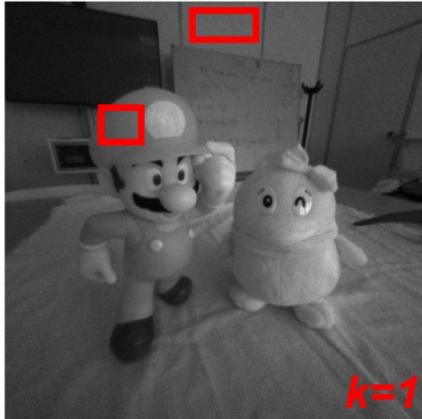
THEN, it holds $f * w = f * (g * h)$ and by virtue of the associative property we can write $f * w = (f * g) * h$. It's worth noticing that $f * g$ requires HWK MACs and outputs a $H \times W$, THEN the convolutional product between the already known $f * g$ and h again requires HWK MACs and outputs a $H \times W$.

Indeed, w.r.t. the previous case, the amount of MAC done to get to the same result is lower, because it's $n_2 = 2HWK$ where $n_1 > n_2$.

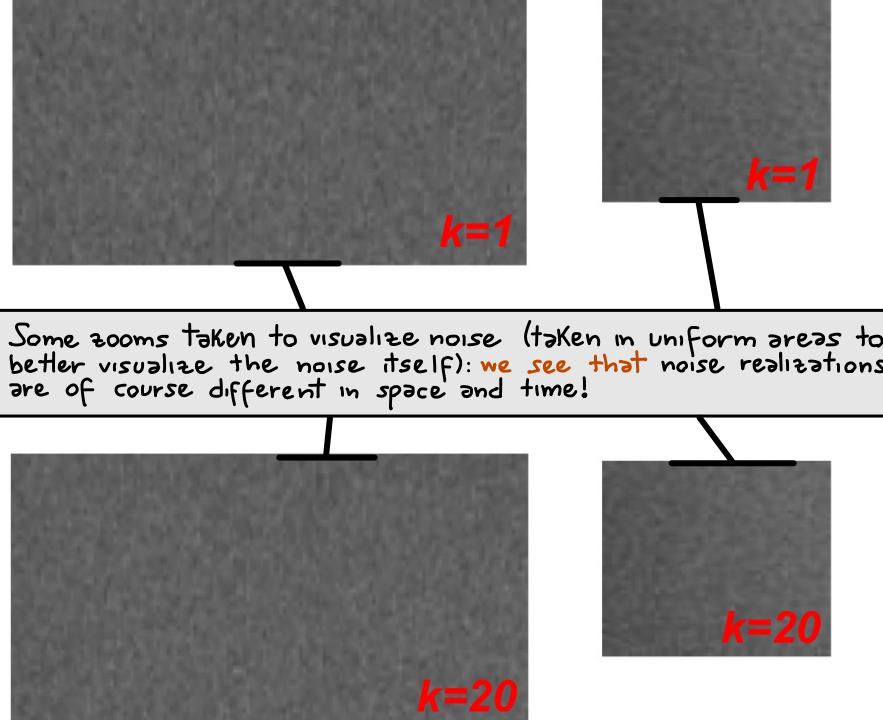
Indeed, when the proper conditions hold, big filter can be thought of convolutions\series of smaller ones, cause it's more efficient to apply the series of convolutions given by the smaller filters w.r.t. applying directly the big one, ending up indeed with LESS MAC operations!

Visualizing and Understanding Noise

t_k
IN THIS EXAMPLE k INDICATES
DISCRETE INSTANTS OF TIME



PIXEL in
euclidian coords



$$I_k(p) = \tilde{I}(p) + n_k(p)$$

ideal deterministic denoised intensity value

with $n_k(p)$ i. i. d.
and $n_k(p) \sim N(0, \sigma^2)$

→ independent and identically distributed (across pixels and time)

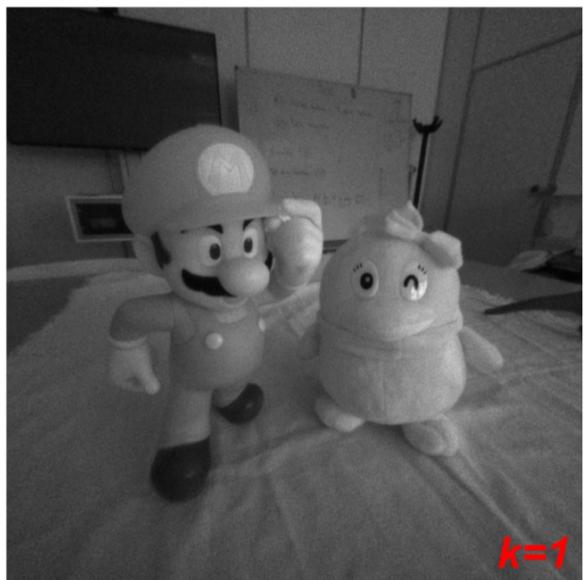
- The **experiment** we will now consider: we take through time multiple images of the same static scene in perfectly identical conditions!
- Also, we'll model noise in the classic way: IID Gaussian distr with zero mean value!

Denoising by taking a mean across time

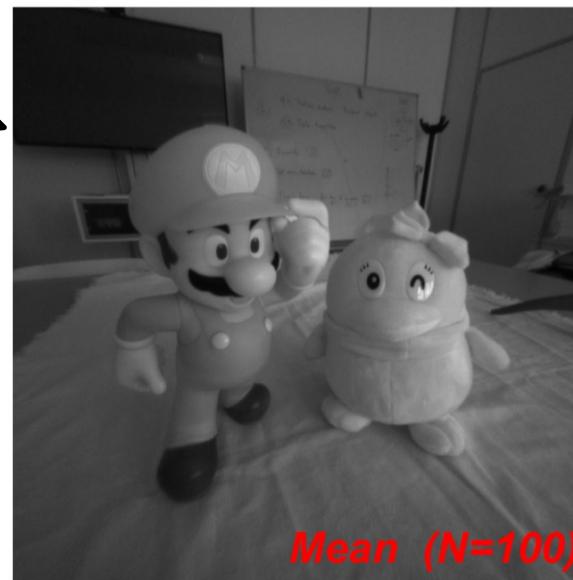


$$O(p) = \frac{1}{N} \sum_{k=1}^N I_k(p) = \frac{1}{N} \sum_{k=1}^N (\tilde{I}(p) + n_k(p)) = \frac{1}{N} \sum_{k=1}^N \underbrace{\tilde{I}(p)}_{\text{CONST}} + \frac{1}{N} \sum_{k=1}^N n_k(p) \Rightarrow N \nearrow: \cong \tilde{I}(p)$$

THIS IS THE SAMPLE MEAN, aka the mean estimated through a set of N samples!
Ideally, for $N \rightarrow +\infty$ this is the EXACT mean of $n_k(p)$, which is zero ($n_k(p) \sim \mathcal{N}(0, \sigma^2)$)



$k=1$



Mean ($N=100$)

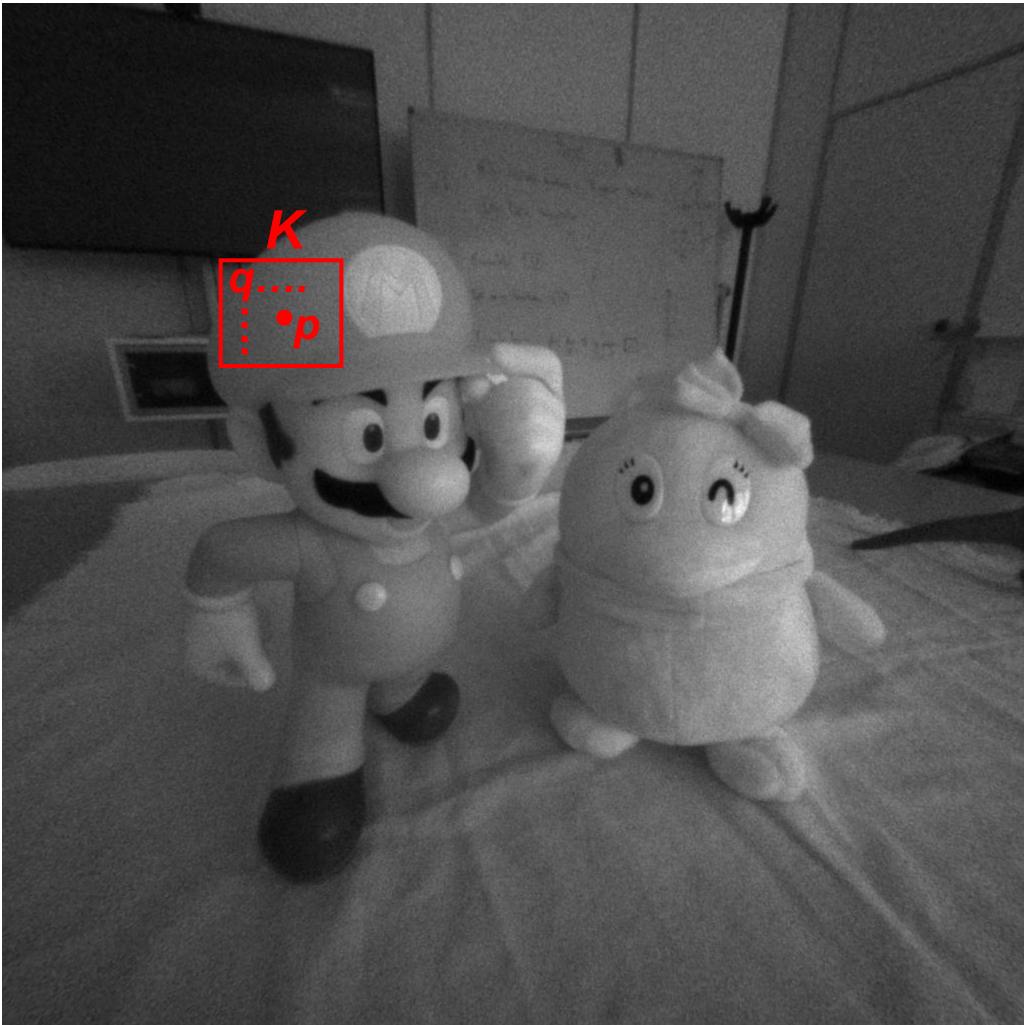
$$\sigma_M^2 = \frac{\sigma^2}{N}$$

IT CAN BE PROVEN
THAT ALSO THE
NOISE VARIANCE IS
REDUCED, ACCORDINGLY
TO THIS FORMULA!

- This is the best possible denoiser in denoising an image (the higher N , the more the noise is mean-ed out AKA the more good is the denoiser)
BUT it is applicable ONLY for static scenes for which we take multiple images through time in the same conditions!

THIS IS THE MOST COMMON CASE

What if we are given a single image ?



$$\begin{aligned}O(p) &= \frac{1}{|K|} \sum_{q \in K} I(q) \\&= \frac{1}{|K|} \sum_{q \in K} (\tilde{I}(q) + n(q)) \\&= \frac{1}{|K|} \sum_{q \in K} \tilde{I}(q) + \frac{1}{|K|} \sum_{q \in K} n(q)\end{aligned}$$

$$\Rightarrow |K| \nearrow \wedge (\tilde{I}(q) = \tilde{I}(p) \forall q \in K) : \cong \tilde{I}(p)$$

We may compute a mean across neighbouring pixels, i.e. a spatial rather than temporal mean.

Denoising Filters

We do not take the mean accross time, but... we take it accross space!

About the condition we've obtained: strong trade off!

Computing mean accross space in a neighbour (for each pixel) K of dimension $(2K+1) \times (2K+1)$ with the aim of denoising brings us to the condition $|K| > 0$ AND $\tilde{I}(q) = \tilde{I}(p) \forall q \in K$: THIS CONDITION EMBED A **STRONG TRADE OFF COMPROMISE!**

The choice of the dimension of K (supposed to be squared) is fundamental in these kind of denoising filters (called **mean filters**): ON ONE HAND we aim to have $|K| \downarrow$ to guarantee $\tilde{I}(q) = \tilde{I}(p) \forall q \in K$ (avoiding mixing together pixels that hold different unnoised values, in particular in non uniform areas, AKA avoiding **BLUR\SMOOTH** effect), ON THE OTHER HAND we aim to have $|K| \uparrow$ to achieve an effective denoising!

In general, we have to optimize between these two objectives.

We're describing an LTE filter of Kernel $K = I / |K|$!

Notice that we've described an operation in which $\forall p \in I$ we compute its new intensity value as a weighted average of all intensities of all pixels q in a neighbour K of p (centered in p and of dimension $(2K+1) \times (2K+1)$) where all the weights are all the same (AKA $1/|K|$, in order to get the mean value of them): THIS IS exactly the definition of apply a Kernel K through a convolution, AKA applying an LTE filter of Kernel $K = I / |K|$.

\Rightarrow This is a mean filter!

Mean Filter



- Mean filtering is the simplest (and fastest) way to denoise an image. It consists in replacing each pixel intensity by the average intensity over a chosen neighbourhood (e.g. 3x3, 5x5, 7x7...).
- The Mean Filter is an LTE operator as it can be expressed as a convolution with a kernel. Below, the kernels for a 3x3 and 5x5 mean filter:

$$\begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix} = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \\ 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \\ 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \\ 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \\ 1/25 & 1/25 & 1/25 & 1/25 & 1/25 \end{bmatrix} = \frac{1}{25} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

|K| is typically way small wrt img dimensions!

OpenCV: cv2.blur

- According to signal processing theory, the Mean Filter carries out a **low-pass filtering** operation, which in image processing is also referred to as **image smoothing**.
- Smoothing is often aimed at image denoising, though sometimes the purpose is to cancel out small-size unwanted details that might hinder the image analysis task. For example, in feature-based computer vision algorithms smoothing is key to create the so called **scale-space**, which endows these approaches with **scale invariance**.
- Mean filtering is inherently fast because multiplications are not needed. Moreover, it can be implemented very efficiently by incremental calculation schemes (**box-filtering**).

The mean filter: some intrinsic properties

They are very fast: → they are ones of the faster possible filters, also because thanks of the nature of their Kernel (all elements are the same) $\nabla pE I$ applying the Kernel can be done NOT ONLY through the usual $|K|$ MACs BUT ALSO through $|K|$ simple sums (by intensity) and one final division (by $|K|$), which is a way lighter set of operations.

also, there are specific implementations in which meanwhile the Kernel is sliding across $\nabla pE I$ is possible to re-use some stored past data, saving some computations (incremental calculation schemes called **box-filtering**).

Behaving as a low pass filter: from the signal theory, it can be proved that a mean filter operates as a 2D low pass filter of frequencies on the input 2D signal image I . This is called **blurring** or **smoothing** in IP/CV context AND intrinsically implies achieving denoising BUT with blurring!

IMPORTANT: this low-pass-filtering behaviour is true for ALL filters that are implemented as convolutions of predetermined Kernels, AKA LTE filters; all LTE filters by their own nature implies BOTH denoising AND blurring/smoothing!

COOL REMARK: the mean filter Kernel is symmetric w.r.t. the origin; $i * K = i o K$

COOL REMARK: progressive blurring can be thought equivalent to see the image scene progressively from greater and greater distances!

Denoising by a Mean Filter

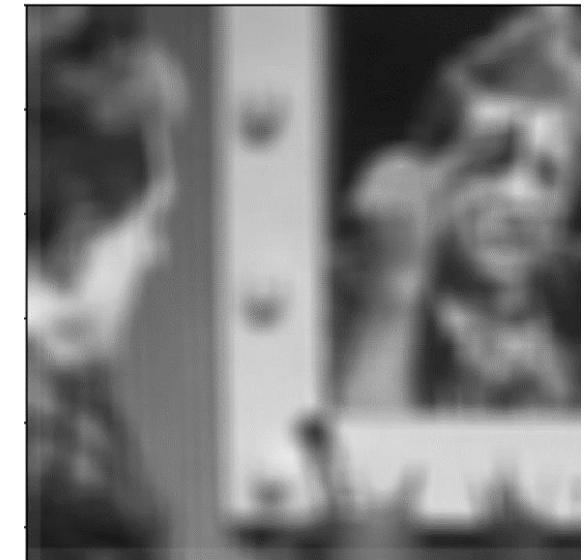
Original Image



Image corrupted by Noise



Smoothing by a Mean Filter



$$n_k(p) \sim \mathcal{N}(\mu = 0, \sigma = 10)$$

$$K = 7 \times 7$$

$$K = 13 \times 13$$

Mean filtering can reduce noise **but blurs the image**

Another LTE filter!

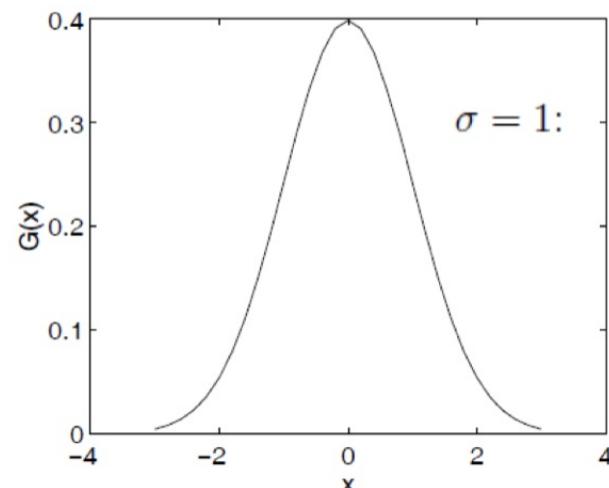
Gaussian Filter (1)



- LTE operator whose impulse response is a 2D Gaussian function (with zero mean and constant diagonal covariance matrix).

1D Gaussian

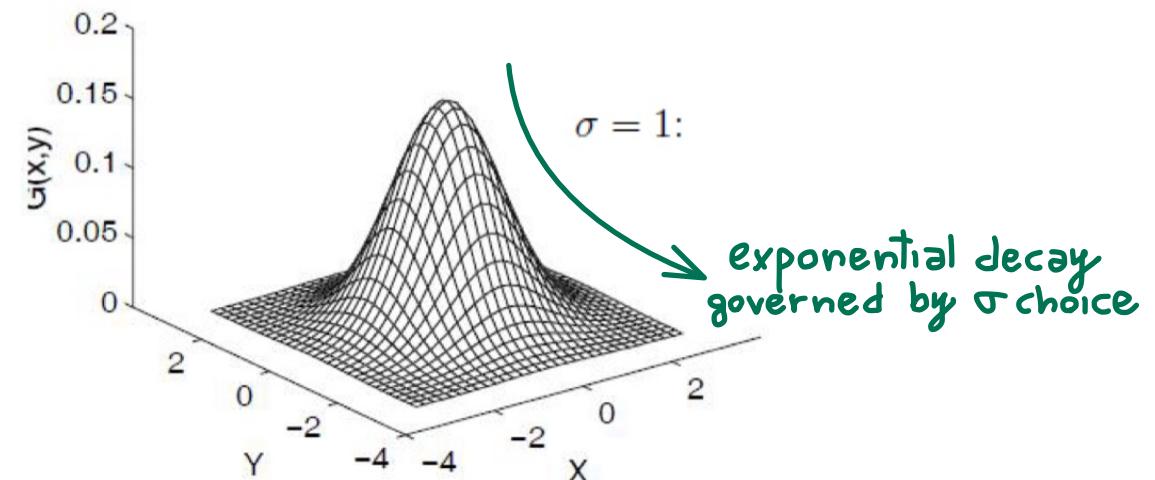
$$G(x) = \frac{1}{\sigma\sqrt{2\pi}}e^{-\frac{x^2}{2\sigma^2}}$$



OpenCV: cv2.GaussianBlur

2D Gaussian

$$G(x, y) = G(x)G(y) = \frac{1}{2\pi\sigma^2}e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Circularly Symmetric

The idea behind the Gaussian filter

If the mean filter performs a weighted sum with all the same weights, instead the Gaussian filter performs a weighted sum in which the weights exponentially decay while getting away from the Kernel center!

The idea is considering the condition of blur-avoiding $\tilde{I}(q) = \tilde{I}(p) \forall q \in K$, probably more verified NEAR the Kernel center, so we give pixels near the Kernel center more importance (w.r.t. the farest ones) still for the same cardinality of $|K|$ (recall indeed $|K| \gg 0$ is the condition for effective denoising)!

From this smart idea in optimizing between the two trade-off objectives $|K| \gg 0$ and $\tilde{I}(q) = \tilde{I}(p) \forall q \in K$, when such a weights distribution is accomplished by a 2D Gaussian function (very common choice), we get a GAUSSIAN LTE FILTER!

Some relevant notices:

- OF COURSE for the same Kardinality this LTE Gaussian filter is SLOWER w.r.t. the mean filter AND still blurs the image (although less).
- the Gaussian filter can be proven to be the BEST linear filter in compromising between denoising and not blurring too much.
- the larger σ , the larger the Gaussian function bell. Indeed, the choice of σ fix the cardinality of the Kernel.
- smaller σ , smaller Kernel, reduced blurring BUT less denoising capability!
bigger σ , bigger Kernel, more denoising capability BUT more blurring!

COOL REMARK: the Gaussian filter Kernel is symmetric w.r.t. the origin; $i * K = i \circ K$

Practical Implementation

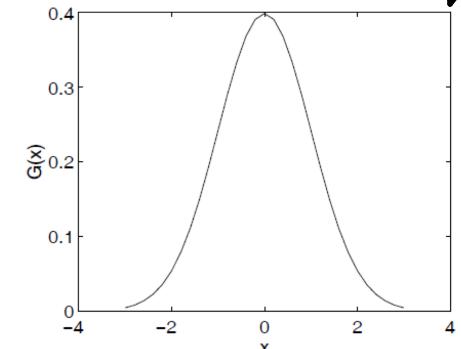
LET'S DISCRETIZE!



- The discrete Gaussian kernel can be obtained by sampling the corresponding continuous function, which is however of infinite extent. A finite size must therefore be properly chosen.
Of course we aim the discrete filter to approximate well the continuous Gaussian function bell for the choosen σ !

To this purpose, we can observe that:

- ✓ The larger is the size, the more accurate turns out the discrete approximation of the ideal continuous filter.
- ✓ The computational cost grows with filter size.
- ✓ The Gaussian gets smaller and smaller as we move away from the origin.



Therefore, we should use larger sizes for filters with high σ , smaller sizes whenever σ is smaller. A rule-of-thumb to chose the size of the filter given σ would then be quite useful.

As the interval $[-3\sigma, +3\sigma]$ captures 99% of the area ("energy") of the Gaussian function, a typical rule-of-thumb dictates taking a $(2k+1) \times (2k+1)$ kernel with:

$$k = \lceil 3\sigma \rceil$$

$\sigma = 1$	\Rightarrow	7×7
$\sigma = 1.5$	\Rightarrow	11×11
$\sigma = 2$	\Rightarrow	13×13
$\sigma = 3$	\Rightarrow	19×19

Mean vs. Gaussian

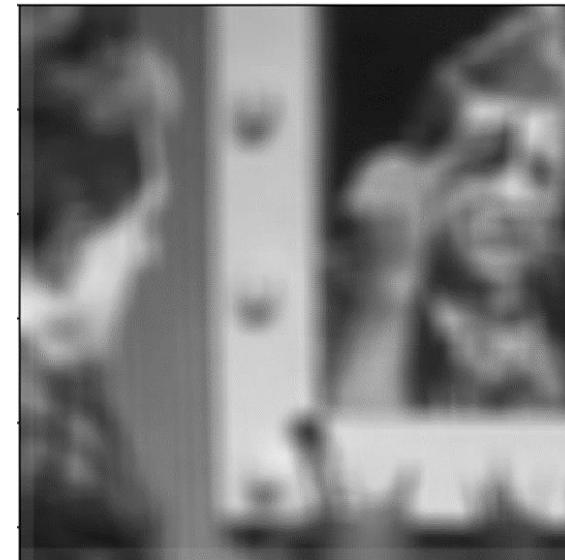
Original Image



Image corrupted by Noise



Smoothing by a Mean Filter



Smoothing by a Gaussian Filter



$$n_k(p) \sim \mathcal{N}(\mu = 0, \sigma = 10)$$

$$K = 13 \times 13$$

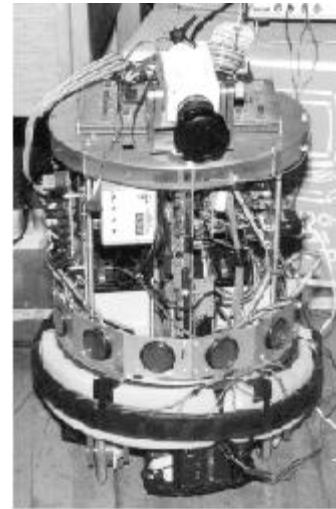
$$K = 13 \times 13$$

Linear filtering can reduce noise but blurs the image.

Given the same kernel size, Gaussian Filtering yields less blur than Mean Filtering.

Parameter σ sets the amount of smoothing

Original Image

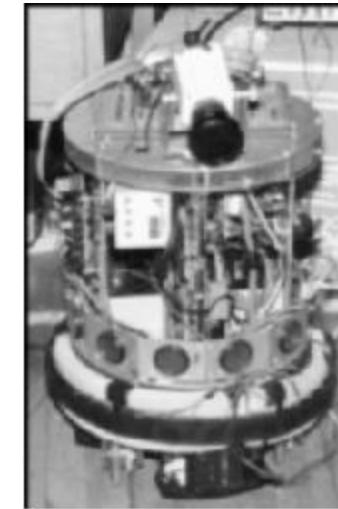


The higher σ , the stronger the smoothing caused by the filter. This can be understood, e.g., by observing that as σ increases, the weights of closer points get smaller while those of farther points get larger.

Smoothing by a Gaussian Filter with $\sigma = 2$



Smoothing by a Gaussian Filter with $\sigma = 1$



As σ gets larger, small details disappear and the image content deals with larger size structures. Thus, filtering with a chosen σ can be thought of as setting the “scale” of interest to analyse image content.



Smoothing by a Gaussian Filter with $\sigma = 4$



Blurring the image is not always bad!

When we have image processing (IP) purposes, typically the blurring\smoothing is a problem because the purpose is denoising the image maintaining it the most possible appealing to the eye: for that reason more advanced NL filters have been invented!

BUT, if we have computer visions (CV) purposes, the blurring\smoothing usually not only is not a problem but is a desired behaviour! Blurring may be a desired procedure in CV desired to be done before further analysis of the image itself!

The reason for that is that blurring makes smaller details in the image to fade away, leaving only larger structures for further image analysis!

→ it can be proved that the Gaussian filter is the best choice for this multiscale image analysis process because is the only one filter that even for sequent applications\blurs (and\or for larger and larger σ) simply cancels out (more and more) small scale structures WITHOUT introducing artifacts (AKA not existing patterns)(the more the blur, the larger the emphasized scale)!

Important remark: adapts weights to one!

In all filters (both LTE filters that are technically equivalent to 2D low pass filters BUT also NL filters) must adapt the weights to 1, meaning: $\forall p \in I$ the set of weights used $\forall q$ pixel must SUM to one; indeed, the set of weights MUST ALWAYS be normalized. The reason of that is that this sum of weights is the so called DC gain\unit gain of the filter: this is indeed the multiplicative factor that affects intensities of perfectly equal and uniform areas AND we want them to be unchanged (DC gain = 1)! For mean filter, this is satisfied by construction. Also for Gaussian filters this is ideally satisfied by construction IF almost all the Gaussian function bell is considered in defining the Kernel, otherwise, typically the Kernel itself is properly normalized in its values!

Impulse Noise

A NEW TYPE OF NOISE!



Original Image



Image corrupted
by Impulse Noise
(aka Salt-and-
Pepper Noise)

Linear filtering is ineffective toward impulse noise (and blurs the image)

Smoothing by
a 3x3 Mean



Smoothing by
a 5x5 Mean

In the Gaussian noise ALL pixels are supposed to be noisy. On the contrary, in Impulse noise almost all pixels are supposed noiseless EXCEPT for a few of them for which the intensity value is way different from the expected one (sometimes completely black or white).

Noisy pixels like that are technically called **outliers**: data coming from a distribution which IS ANOTHER w.r.t. the one we're dealing with.

For such kind of noise, not only the LTE filters are NOT effective BUT they spread the noise even more!

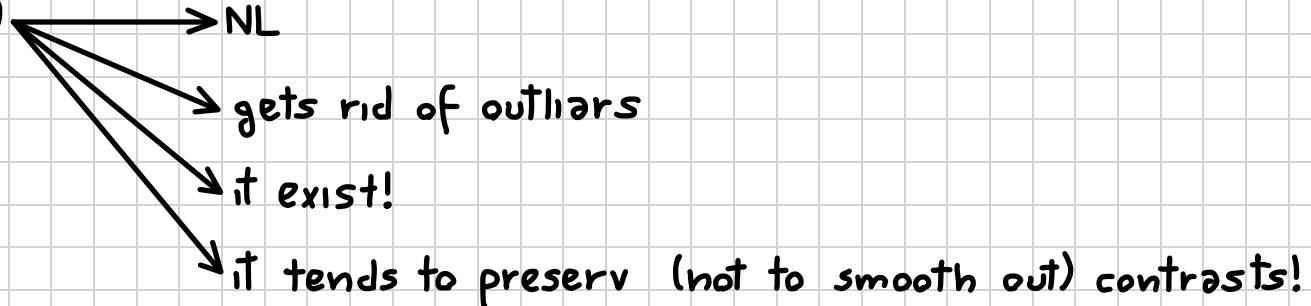
Impulse noise: filter it away with a neighborhood median!

The impulse\salt&pepper\outliers noise is NOT part of what already seen in previous chapters (alias noises that can be typically described as Gaussian).

The impulse\salt&pepper\outliers noise is typically generated when manipulating images: it may be caused by internet transmission problems OR may be present when the image is not taken directly by a sensor BUT is the output of some IP that may have done mistakes and so introducing outliers! For example, the disparities-image obtained after a correlations search onto two rectified images obtained from a stereo camera typically present outliers.

In general, and so also in image filtering, an effective way to get rid of outliers is taking the median! Given a discrete set of values, the median is PICKED AMONG THESE VALUES (no new value produced) as the value that minimizes the difference between the amounts of lower and bigger values w.r.t. it!

Median (wrt mean)



Median Filter

standard & most common
filter to address impulse noise



- Non-linear filter whereby each pixel intensity is replaced by the median over a given neighbourhood, the median being the value falling half-way in the sorted set of intensities.

123	125	126	130	140
122	124	126	127	135
118	120	150	125	134
119	115	119	123	133
111	116	110	120	130

Neighbourhood values:

115, 119, 120, 123, 124,
125, 126, 127, 150

Median value: 124

in a lot of applications we typically
use median filters with small
kernels; 3x3 or 5x5

OpenCV: cv2.Medianblur

$$\text{median} [A(x) + B(x)] \neq \text{median} [A(x)] + \text{median} [B(x)]$$

- Median filtering counteracts impulse noise effectively, as outliers (i.e. noisy pixels) tend to fall at either the top or bottom end of the sorted intensities.

- Median filtering tends to keep sharper edges than linear filters such as the Mean or Gaussian:

In place of edges and contrasts, the median existing pixel is taken, preserving and not smoothing the pixel itself!

... 10 10 40 40 ... \Rightarrow ... 10 20 30 40 ... (Mean)
 \Rightarrow ... 10 10 40 40 ... (Median)

Example

Original Image



Image Corrupted by Impulse Noise (0.05)



Filtering by a 3x3 Median



When dealing with impulse noise, the Median Filter can effectively denoise the image without introducing significant blur.

Yet, Gaussian-like noise, such as sensor noise, cannot be dealt with by the Median, as this would require computing new noiseless intensities. Purposely, the Median may be followed by linear filtering.

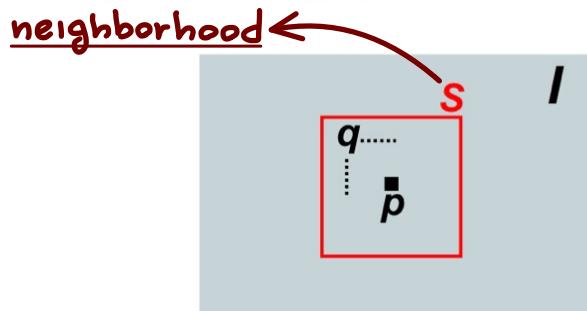
Let's introduce NL denoising filters, way more SLOW wrt linear ones BUT with the capability of almost NO BLUR!



Bilateral Filter (1)

A SIMPLE NL FILTER!

- Advanced non-linear filter to accomplish denoising of Gaussian-like noise without blurring the image (aka edge preserving smoothing).



$$O(p) = \sum_{q \in S} H(p, q) \cdot I_q$$

CLASSIC WEIGHTED SUM
WEIGHTS

OpenCV: cv2.bilateralFilter

weights not anymore predetermined BUT depends on actual values of p and q!

$$H(p, q) = \frac{1}{W(p)} G_{\sigma_s}(d_s(p, q)) G_{\sigma_r}(d_r(I_p, I_q))$$

2D GAUSSIAN FUNCTIONS

$$d_s(p, q) = \|p - q\|_2 = \sqrt{(u_p - u_q)^2 + (v_p - v_q)^2} \rightarrow \text{Spatial Distance}$$

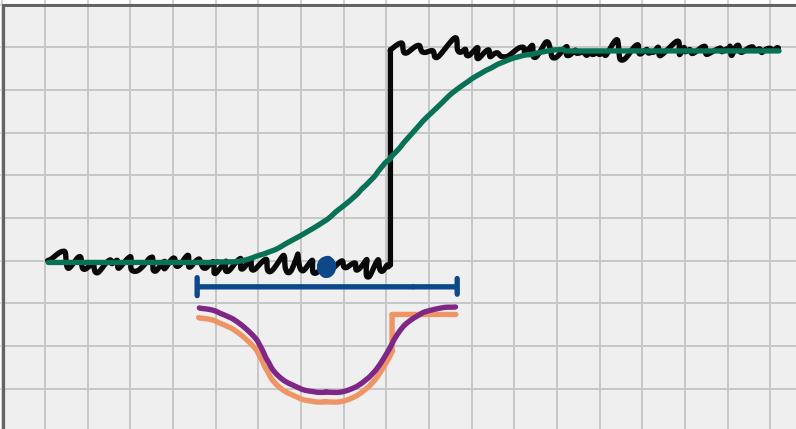
$$d_r(I_p, I_q) = |I_p - I_q| \rightarrow \text{Range (Intensity) Distance}$$

$$W(p) = \sum_{q \in S} G_{\sigma_s}(d_s(p, q)) G_{\sigma_r}(d_r(I_p, I_q)) \rightarrow \text{Normalization Factor (Unity Gain)}$$

to have a DC gain of 1!

This is called to be a SOFT implementation of the filter, that involves only simple math operations; if IF-ELSE statements where used, it would have been an HARD implemented solution!

Bilinear filters: consider also intensity neighbours (in addition to spatial ones)



- Image (1D) with an edge, noisy
- Image filtered with a gaussian filter
- an image point and its Kernel
- point weights given by the Gaussian Kernel
- point weights we aim to have with bilateral filter Kernel

→ we want to properly "cut" gaussians near edges!

With a simple Gaussian filter, the closer to the edge, the more we mix the two sides of the edge itself (getting blur).

Instead, in a bilateral filter, small weights are given NOT ONLY to neighbors that are spatially far away BUT ALSO AND to neighbors that are far away from an intensity point of view!

⇒ we consider the product (logic AND) of two gaussians, one dependent on classic **spatial distance** and the other dependent on **range\intensity** distance!

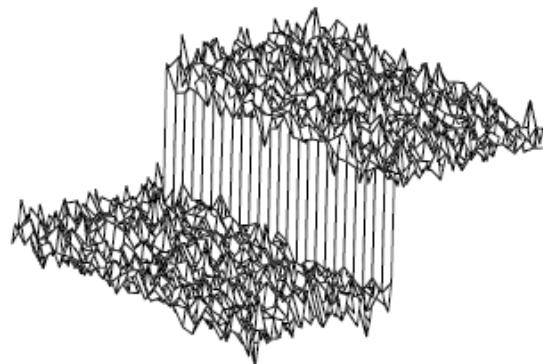
BE AWARE: of course a bilateral filter is gonna be **way slower w.r.t. LTE ones**, even if some faster variation (approximations) of the above shown bilateral filter implementation do exist!

Bilateral Filter (2)

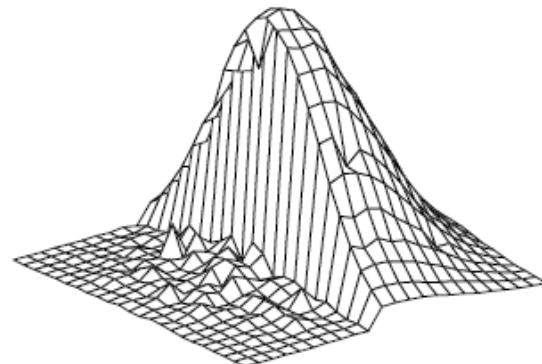
*Step-edge as wide
as 100 gray-levels*

*$H(p,q)$ at a pixel just across
the edge
in the brighter region*

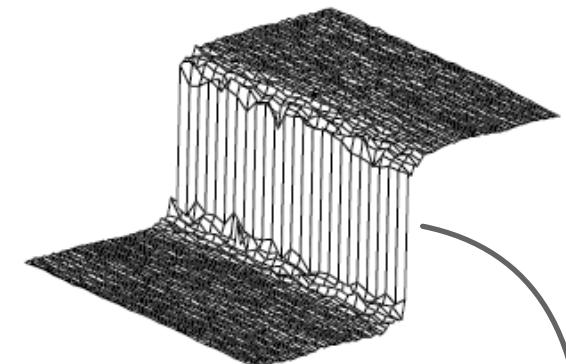
*Output provided by the filter
 $\sigma_s = 5, \sigma_r = 50$*



(a)



(b)



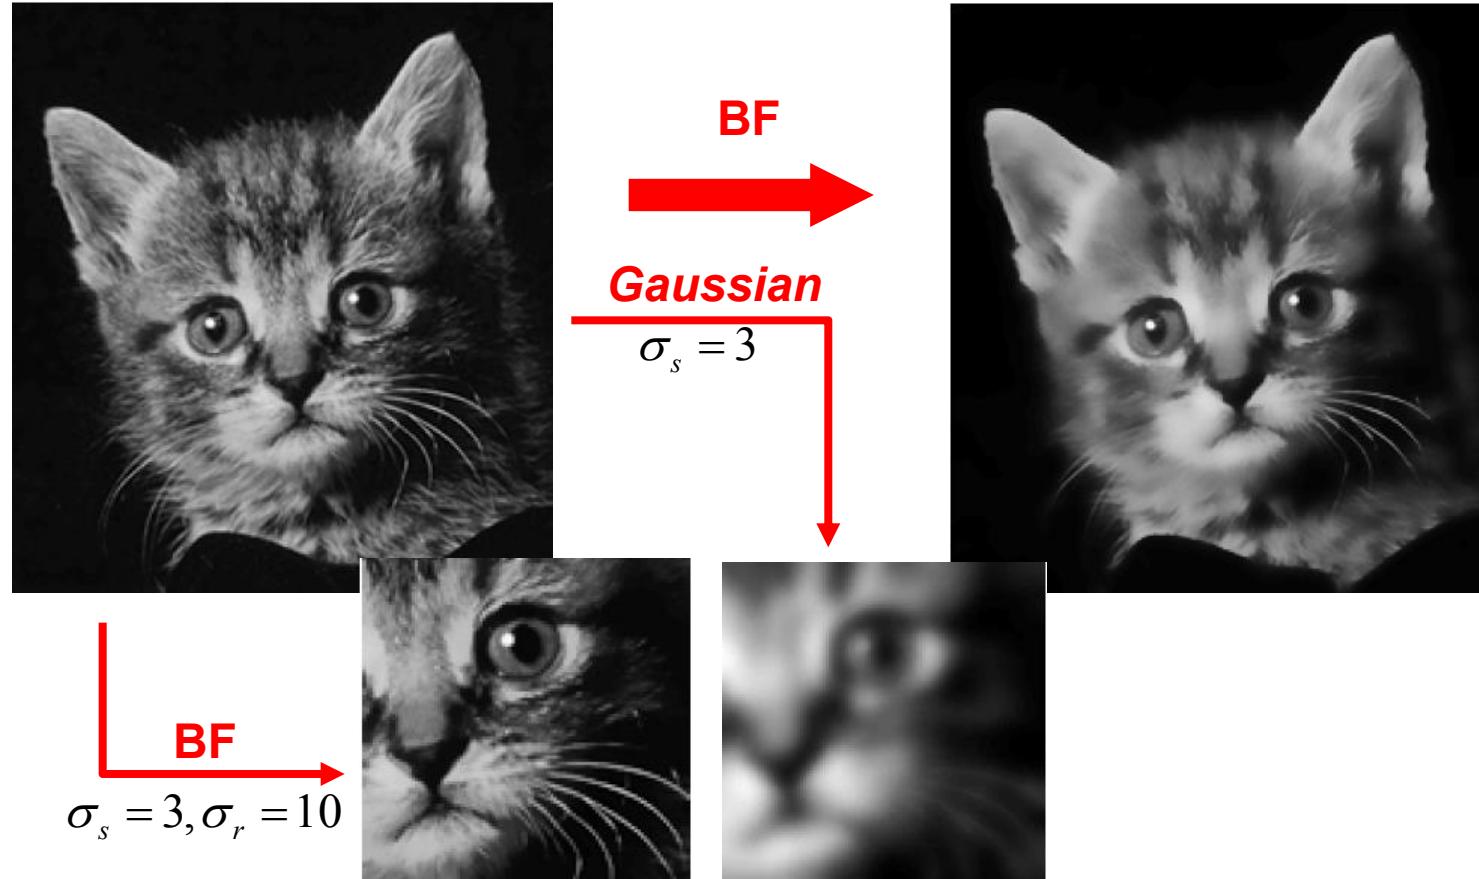
(c)

- Given the supporting neighbourhood, neighbouring pixels take a larger weight as they are both closer and more similar to the central pixel.

- At a pixel nearby an edge, the neighbours falling on the other side of the edge look quite different and thus cannot contribute significantly to the output value due to their weights being small.



Bilateral Filter (3)



More examples at:

http://people.csail.mit.edu/sparis/siggraph07_course/

Guided Filter *→ another very used NL filter!*

<https://kaiminghe.github.io/eccv10/index.html>

Mean vs. Gaussian vs. Bilateral

Image corrupted by Noise



$n_k(p) \sim \mathcal{N}(\mu = 0, \sigma = 10)$

Smoothing by a Mean Filter



$K = 13 \times 13$

Smoothing by a Gaussian Filter



$K = 13 \times 13$

Smoothing by a Bilateral Filter



$K = 13 \times 13$

Thanks to its weight function, the Bilateral Filter can accomplish denoising without blurring the image (aka edge preserving smoothing)!

EVEN SLOWER
WRT BILATERAL!

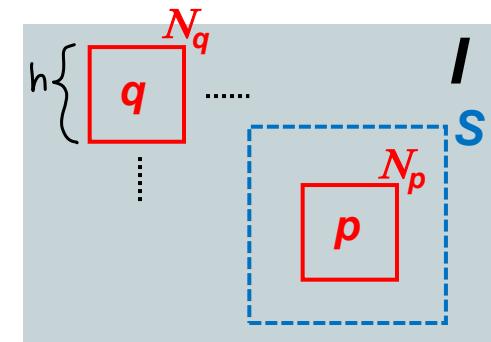
Non-local Means Filter (1)



- Another well-known *non-linear edge preserving smoothing filter*. The key idea is that the similarity among patches spread over the image can be deployed to achieve denoising.



OpenCV: cv2.fastNIMeansDenoising



$$O(p) = \sum_{q \in I} w(p, q)I(q)$$

$$w(p, q) = \frac{1}{Z(p)} e^{-\frac{\|N_p - N_q\|_2^2}{h^2}}$$

normalizing factor

$$Z(p) = \sum_{q \in I} e^{-\frac{\|N_p - N_q\|_2^2}{h^2}}$$

σ here indicates the standard deviation of the noise, typically NOT available: in practice we usually set h in a similar way w.r.t. σ in Gaussian filters!

Typical choices: $N=7 \times 7$, $S=21 \times 21$, $h = 10 \cdot \sigma$

The non local mean filter!

The basic idea behind all filters that maintain edges NOT smoothing them is to perform a mean only across pixels with little intensities differences... AND among these pixels with little intensities differences picking the highest possible amount of them!

The idea of non-local filters is take to the extreme the idea of picking high amounts of similar pixels (similar in intensity) to make the mean with: $\forall p \in I$ the considered Kernel is the WHOLE IMAGE and higher weights are given to similar-in-intensity q pixels that are supposed to belong to the same object of p itself! The belonging of a q pixel to the same object of p is comprehended through the difference among two their neighbours zones N_q and N_p both of size $h \times h$. The difference between N_q and N_p is computed by stacking them in two vectors of h^2 elements and then taking the difference among these two!

Some final notes:

- this filter is very SLOW. Faster versions use as Kernel not the whole image I but a subportion S centered in p_j ; still $|S|$ is way bigger w.r.t. cardinality of typical Kernels of local filters.
- h tunes how strong is the filtering effect in an equivalent way w.r.t. σ in Gaussian LTE filters: $h^2 \uparrow$ for collecting more pixels in N_q and N_p , filtering more, $h^2 \downarrow$ for less filtering and giving high weights only to q pixels strongly similar to p !

Non-local Means vs. Bilateral

Image corrupted by Noise



Smoothing by a Bilteral Filter



$K = 13 \times 13$

Smoothing by a NLM Filter



$K = 13 \times 13$



Main References

- 1) V. S. Nalwa, "A Guided Tour of Computer Vision", Addison-Wesley Publishing Company, 1993.
- 2) R. Fisher, S. Perkins, A. Walker, E. Wolfart, "Hypermedia Image Processing Reference", Wiley, 1996 (<http://homepages.inf.ed.ac.uk/rbf/HIPR2/>)
- 3) R. Schalkoff, "Digital Image Processing And Computer Vision, Wiley, 1989.
- 4) C. Tomasi, R. Manduchi "Bilateral Filtering for Gray and Color Images", ICCV 1998.
- 5) S.Paris, P. Kornprobst, J. Tumblin, F. Durand "A Gentle Introduction to Bilateral Filtering and its Applications" (SIGGRAPH 2008,CVPR 2008, SIGGRAPH 2007)
http://people.csail.mit.edu/sparis/siggraph07_course/
- 6) A. Buades, B. Coll, J.M. Morel, "A non-local algorithm for image denoising", CVPR 2005.
- 7) Kaiming He, Jian Sun, and Xiaou Tang, "Guided Image Filtering", IEEE Trans. On PAMI, 2013

